

Graph Neural Network- Part 2

- Definition of graph
 - Adjacency matrix , Laplacian and normalized Laplacian
 - Convolution on graphs using eigenvectors of Laplacian
 - Graph Fourier transform and inverse.
-
- Vanilla Spectral GCN,
 - Chebyshev polynomial and ChebNet

ChebNet Graph Convolution

- $x * g = U g_\theta U^T x$
- $x * g = \sum_i \theta_i U T_i(\tilde{\Lambda}) U^T x$

$$x * g_\theta = U (\sum_i \theta_i T_i(\tilde{\Lambda})) U^T x$$

- It is equivalent to:

$$x * g_\theta = \sum_{i=1}^k \theta_i T_i(\tilde{L}) x$$

We can compute \tilde{L} without the eigendecomposition of L .

The scaled Laplacian \tilde{L} is

$$\tilde{L} = \frac{2L}{\lambda_{max}} - I_n$$

Eigenvalues: [0,1,2]

Given signal $x=[1,2,3]'$

$$g_\theta \star x \approx \theta_0 T_0(L^\sim)x + \theta_1 T_1(L^\sim)x$$

$$L^\sim = 2L/\lambda_{\max} - I$$

L is normalized Laplacian

$$L^\sim = 2L/2 - I = \begin{pmatrix} 0 & -0.7071 & 0 \\ -0.7071 & 0.0000 & -0.7071 \\ 0 & -0.7071 & 0 \end{pmatrix}$$

Adjacency matrix	Degree matrix	Normalized Laplacian
$\begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}$	$\begin{pmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 1 & -\sqrt{1/2} & 0 \\ -\sqrt{1/2} & 1 & -\sqrt{1/2} \\ 0 & -\sqrt{1/2} & 1 \end{pmatrix}$

$$T_0(L^\sim)x = x \quad x = [1, 2, 3]'$$

$$\begin{aligned} T_1(L^\sim)x &= L^\sim x = -1.4142 \\ &\quad -2.8284 \\ &\quad -1.4142 \end{aligned}$$

$$\begin{aligned} y &= \theta_0 x + \theta_1 L^\sim x = \\ &0.2929 \\ &0.5858 \\ &2.2929 \end{aligned}$$

$$\theta_0 = 1 \text{ and } \theta_1 = 0.5$$

Graph Convolutional Network (GCN)

- First-order approximation of ChebNet.
- T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” in Proc. of ICLR, 2017

ChebNet and Chebyshev polynomials

- ChebNet takes the form:

$$x * g_\theta = \sum_{i=1}^k \theta_i T_i(\tilde{L}) x$$

Where T_i is Chebyshev polynomials.

- $T_0(x) = 1$
- $T_1(x) = x$
- $T_i(x) = 2xT_{i-1}(x) - T_{i-2}(x)$

First-order approximation

- Let's find the first-order approximation of

$$x * g_\theta = \sum_{i=0}^k \theta_i T_i(\tilde{L})x = \theta_0 T_0(\tilde{L})x + \theta_1 T_1(\tilde{L})x = \theta_0 x + \theta_1 \tilde{L}x$$

$$\tilde{L} = \frac{2L}{\lambda_{max}} - I_n \quad \text{Lambda_max is approximated to be 2}$$

$$\approx \theta'_0 x + \theta'_1 (L - I_N) x = \theta'_0 x - \theta'_1 D^{-\frac{1}{2}} A D^{-\frac{1}{2}} x$$

Normalized Laplacian

$$D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}} \quad I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}}$$

First-order approximation

- To restrain the number of parameters and avoid over-fitting, GCN further assume $\theta = \theta_0 = -\theta_1$

$$g_\theta \star x \approx \theta \left(I_N + D^{-\frac{1}{2}} A D^{-\frac{1}{2}} \right) x$$

- This empirically causes numerical instability to GCN.

Numerical trick

- This empirically causes numerical instability to GCN. To address this problem, GCN applies a normalization trick to replace

$$\tilde{A} = A + I$$

$$\tilde{D}_{ij} = \sum_j \tilde{A}_{ij}$$

$$\bar{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$$

a compositional layer can be defined as:

$$H' = X * g_\theta = \sigma(\bar{A} H \theta)$$

One-step GCN normalization

$$A = \begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix}, \quad X = \begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix},$$

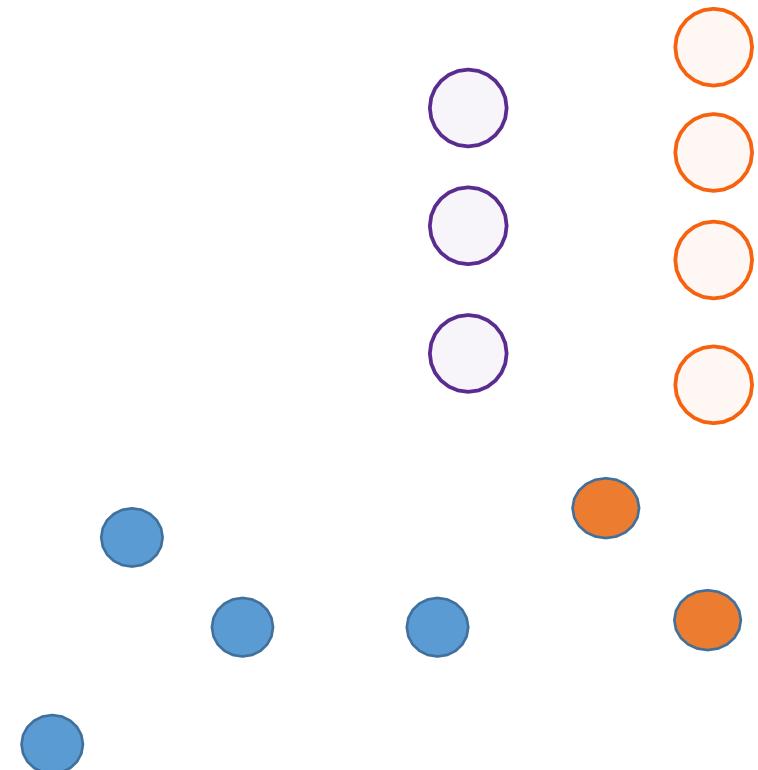
$\tilde{A} = A + I$, $\tilde{D} = \text{diag}(\text{row sums of } \tilde{A})$, and compute

$$\hat{A} = \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}, \quad \hat{X} = \hat{A}X.$$

$$\tilde{A} = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 1 \end{bmatrix}, \quad D^{\sim} = \text{diag}(2, 3, 2) \quad \hat{A} \approx \begin{bmatrix} 0.5 & 0.408 & 0 \\ 0.408 & 0.333 & 0.408 \\ 0 & 0.408 & 0.5 \end{bmatrix} \Rightarrow \hat{X} \approx \begin{bmatrix} 1.316 \\ 2.298 \\ 2.684 \end{bmatrix}$$

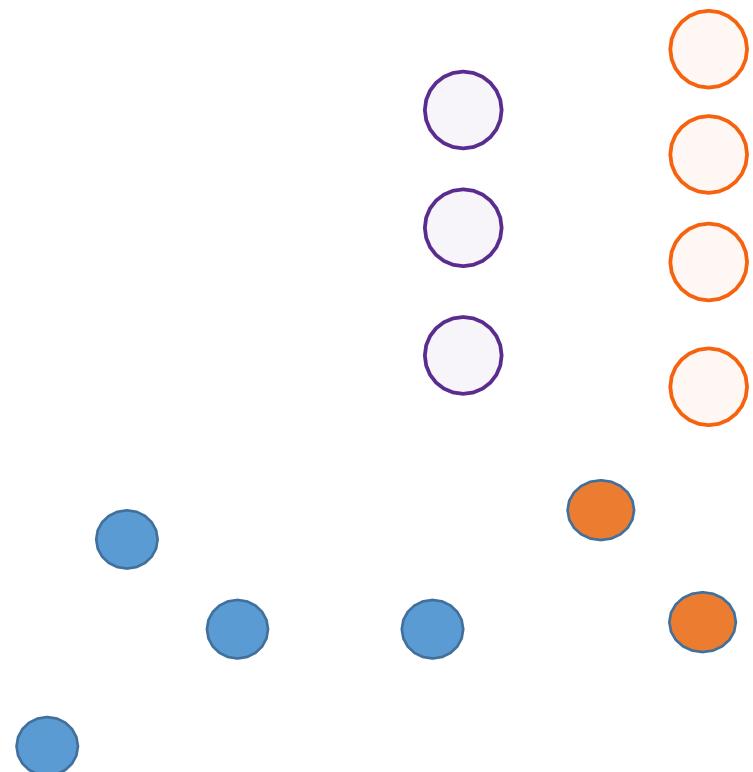
Aggregation of information

- $H' = \sigma_{\Theta}(H_i)$



$$\bullet H' = \sigma(H\Theta)$$

$$\bullet H' = \sigma(\bar{A}H\Theta)$$

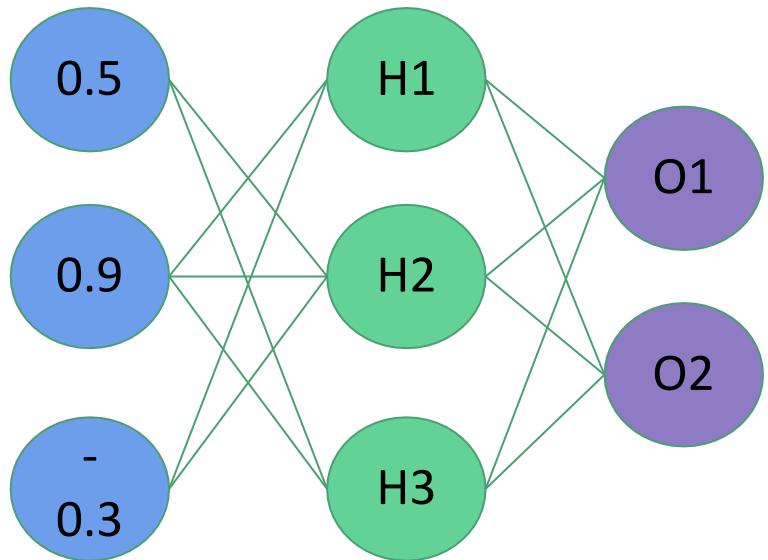


Towards more general frameworks

MLP

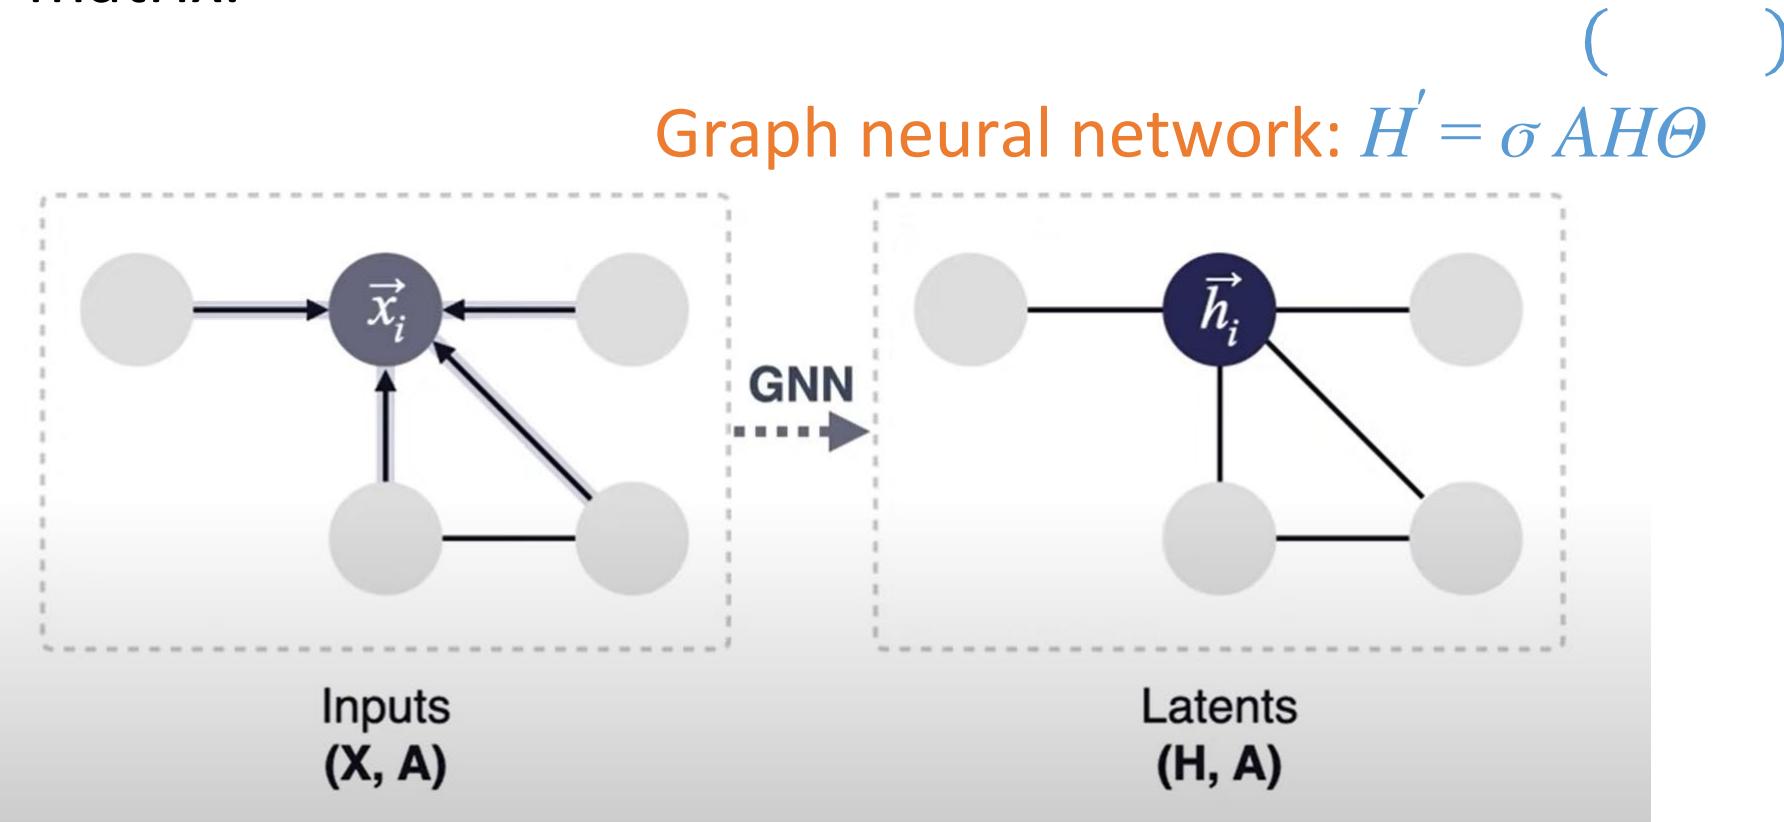
Feed forward

$$S : H' = \sigma H \Theta ()$$



GNNs

We can aggregate neighbourhoods by multiplying the adjacency matrix.



Sum-pooling

- This update rule discards the central node.

Sum-poolin g

- This update rule discards the central node.
- This can be fixed simply by

$$\tilde{A} = A + I$$

Sum-poolin g

- This update rule discards the central node.
- This can be fixed simply by

$$\tilde{A} = A + I$$

- The node-wise update rule can be written as:

$$h'_i = \sigma \left(\sum_{j \in N_i} \theta h_j \right)$$

Mean-poolin

g

- Summing the contents of the neighbouring nodes will increase the scale of the output feature.

Mean-poolin g

- Summing the contents of the neighbouring nodes will increase the scale of the output feature
- We can normalize by \tilde{D}^{-1} $\tilde{D}_{ij} = \sum_j \tilde{A}_{ij}$

$$h'_i = \sigma \left(\sum_{j \in N_i} \tilde{D}^{-1} \Theta h_j \right)$$

- The node-wise update rule can be

$$h'_i = \sigma \left(\sum_{j \in N_i} \frac{1}{|N_i|} \Theta h_j \right)$$

Graph Convolutional Networks (GCNs)

- Use symmetric normalization
- The node-wise update rule can be written as:

$$H' = \sigma(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H \Theta)$$

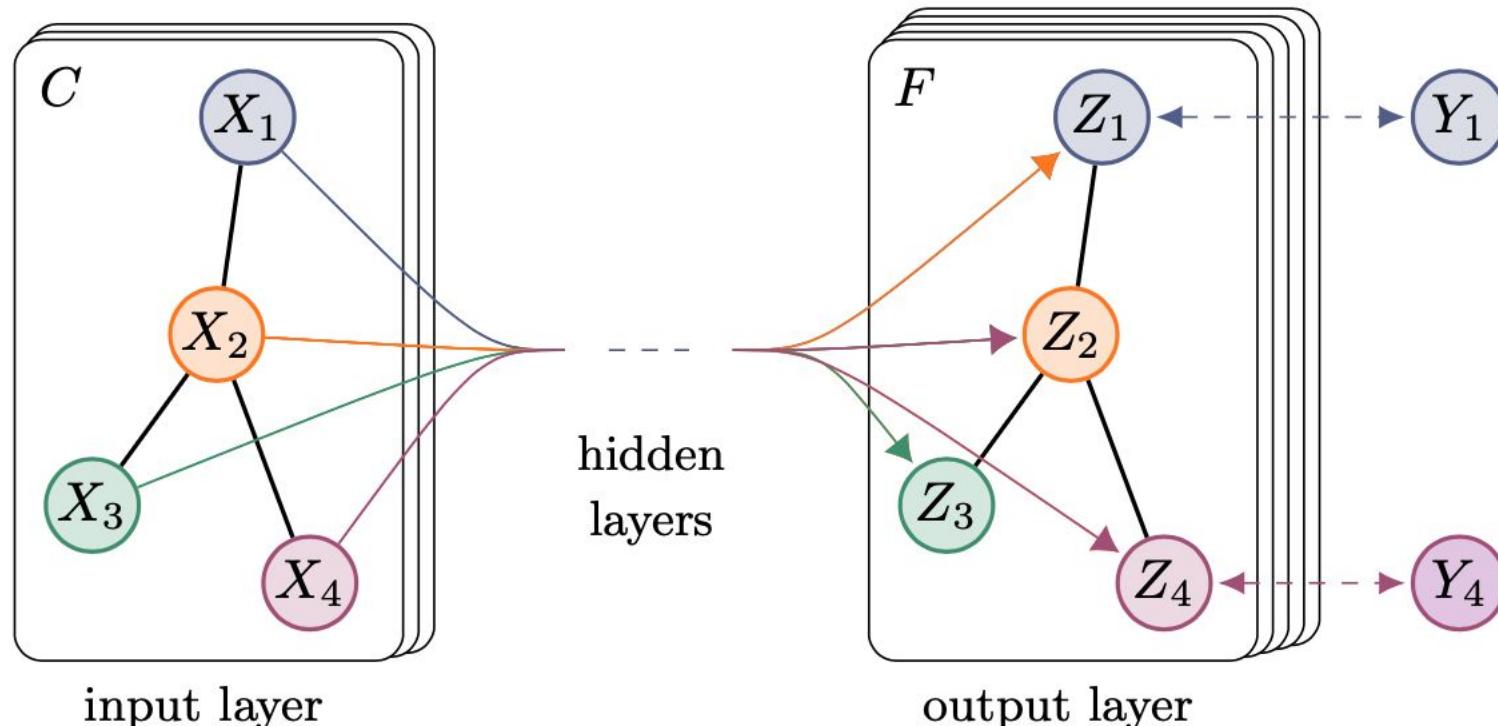
$$h'_i = \sigma \left(\sum_{j \in N_i} \frac{1}{\sqrt{|N_i| |N_j|}} \theta h_j \right)$$

GCN architecture

we consider a two-layer GCN

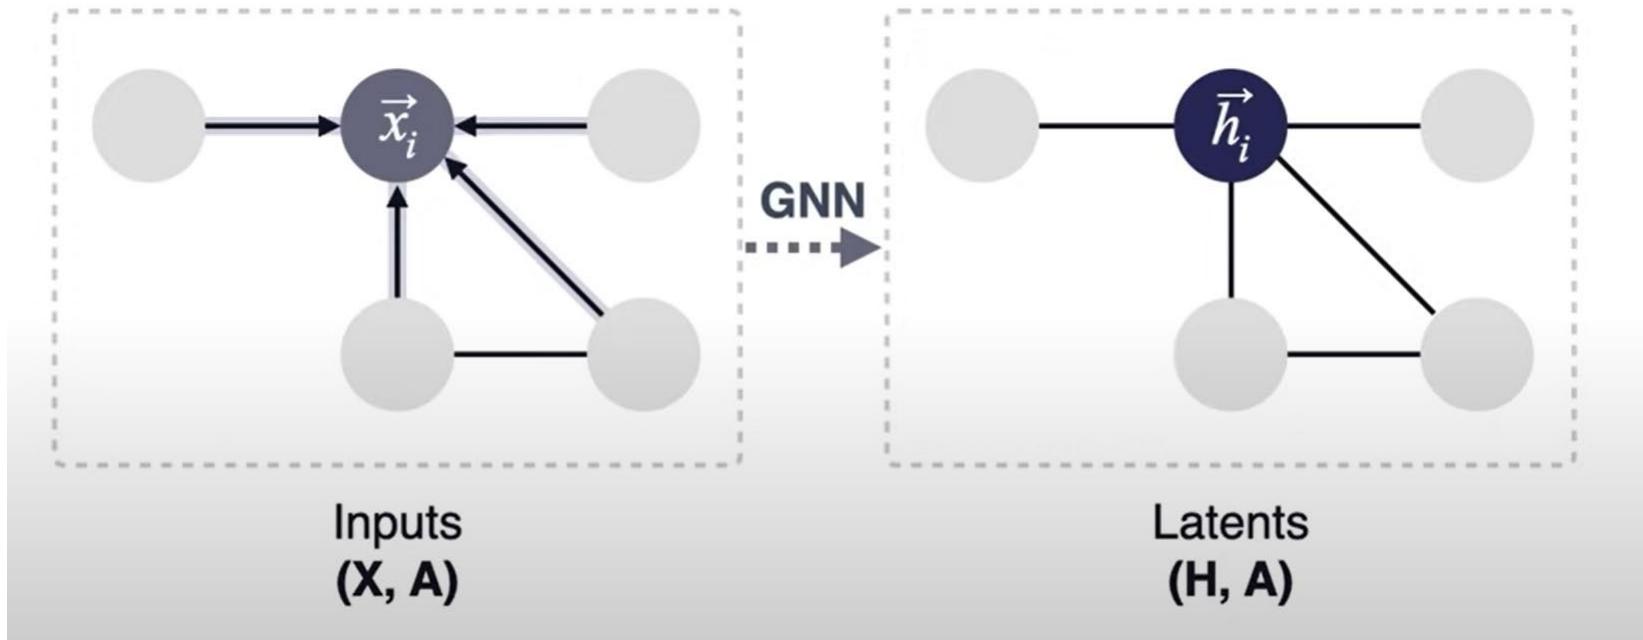
$$Z = f(X, A) = \text{sigmoid} \left(\hat{A} \text{ReLU} \left(\hat{A} X W^{(0)} \right) W^{(1)} \right)$$

$$W(0) \in \mathbb{R}^{C \times d} \quad W(1) \in \mathbb{R}^{d \times 1}$$

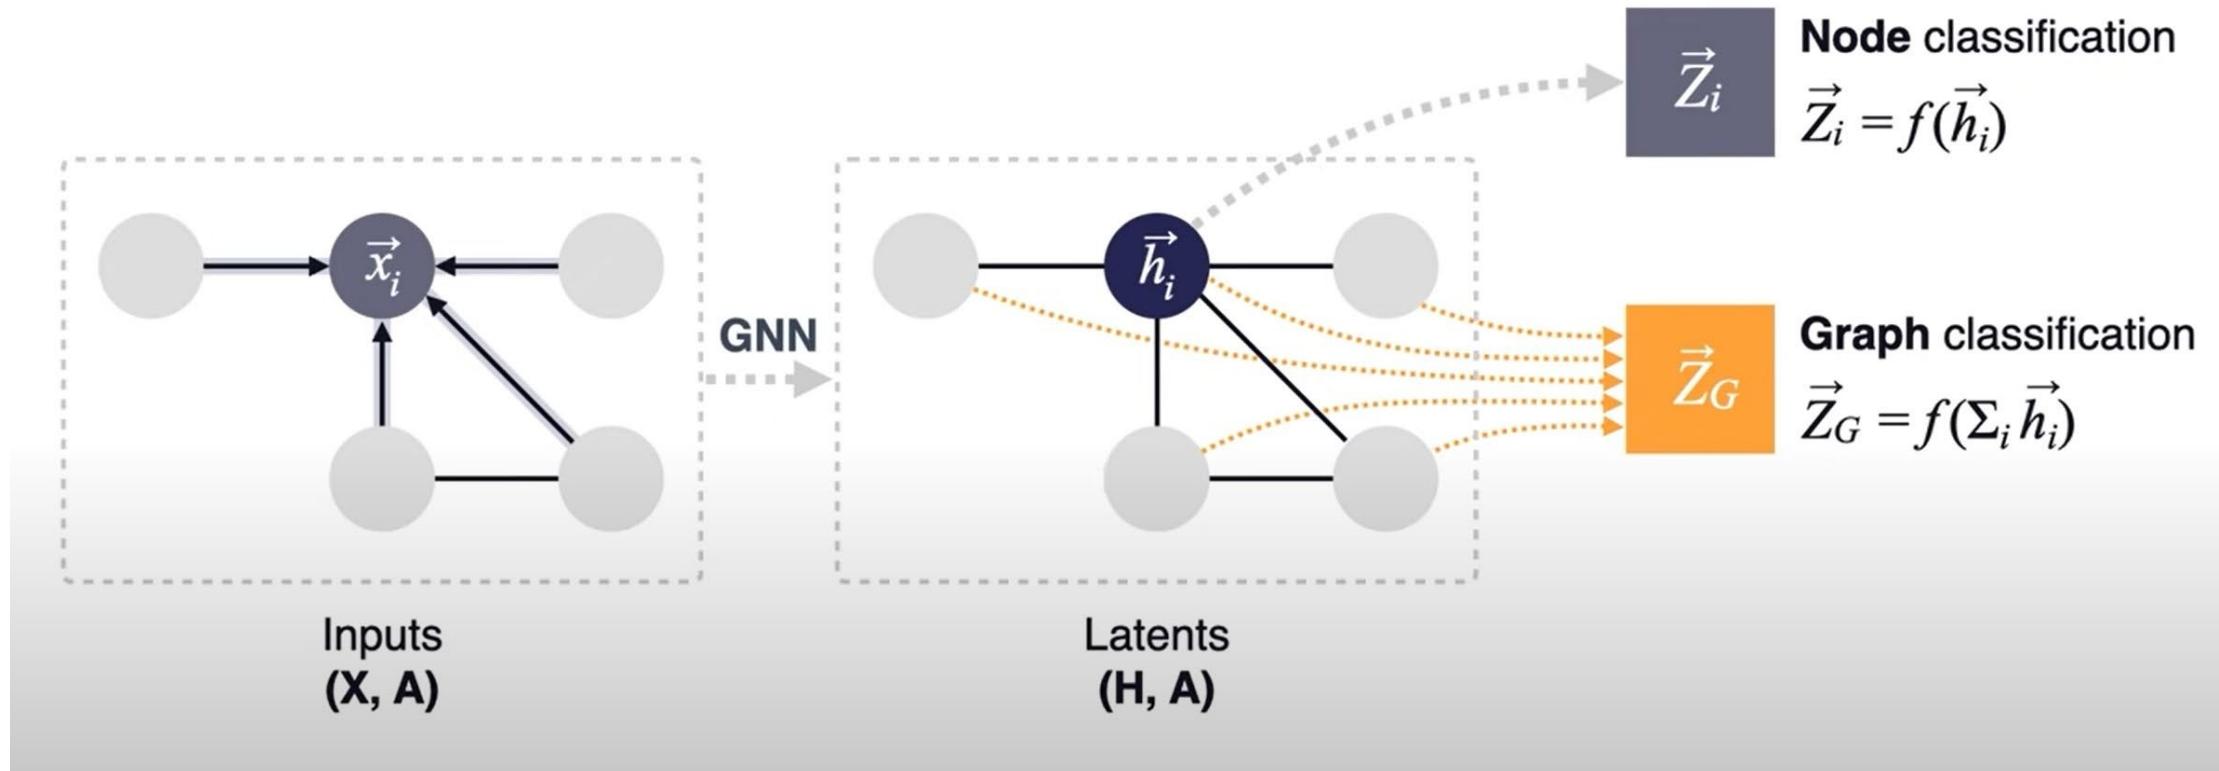


C is dimensionality of node

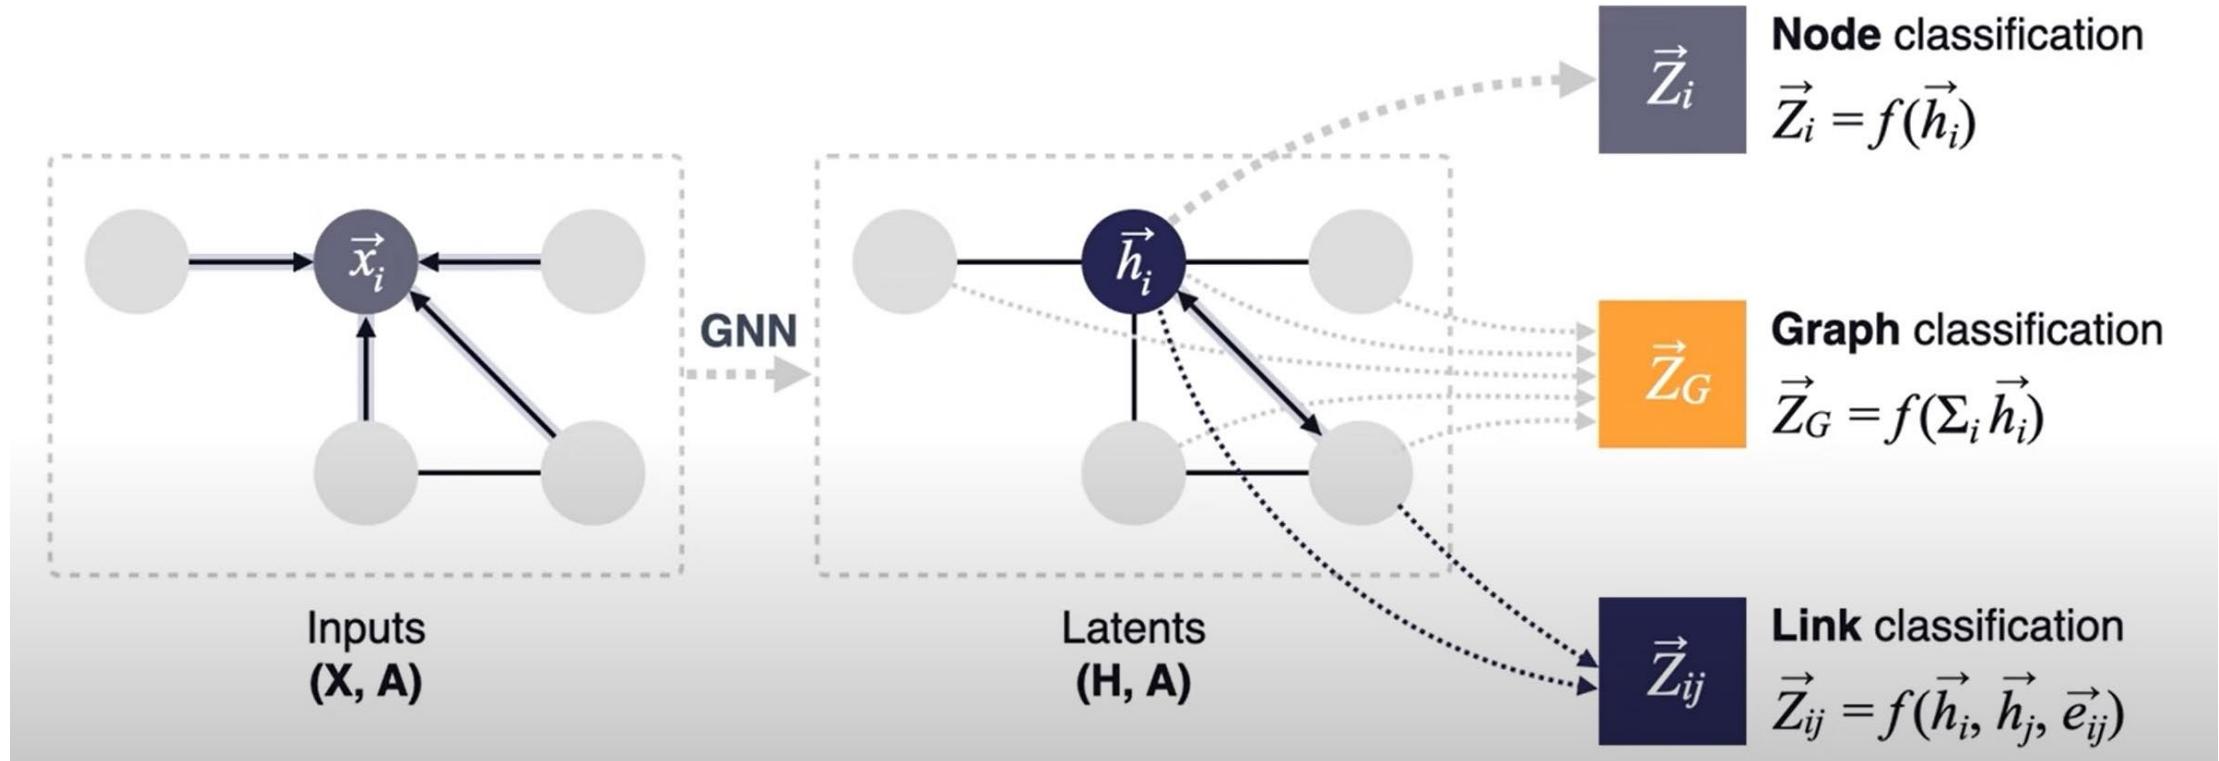
Node Classification



Graph Classification



Link Classification



- Sum pooling $h'_i = \sigma \left(\sum_{j \in N_i} \theta h_j \right)$

- Mean pooling $h'_i = \sigma \left(\sum_{j \in N_i} \frac{1}{|N_i|} \theta h_j \right)$

Can you think of alternatives?

- GCNs $h'_i = \sigma \left(\sum_{j \in N_i} \frac{1}{\sqrt{|N_i||N_j|}} \theta h_j \right)$

Graph Attention Network (GAT)

- GAT adopts attention mechanisms to learn the relative weights between two connected nodes.

$$h'_i = \sigma \left(\sum_{j \in N_i} \alpha_{ij} \Theta h_j \right)$$

The attention weight measures the influence of node j to node i

The attention weight

- The attention weight can be computed as follows:

$$a_{ij} = a(h_i, h_j)$$

or

$$a_{ij} = a(h_i, h_j, e_{ij})$$

- a is a single-layer feedforward neural network.

Attention Computation:

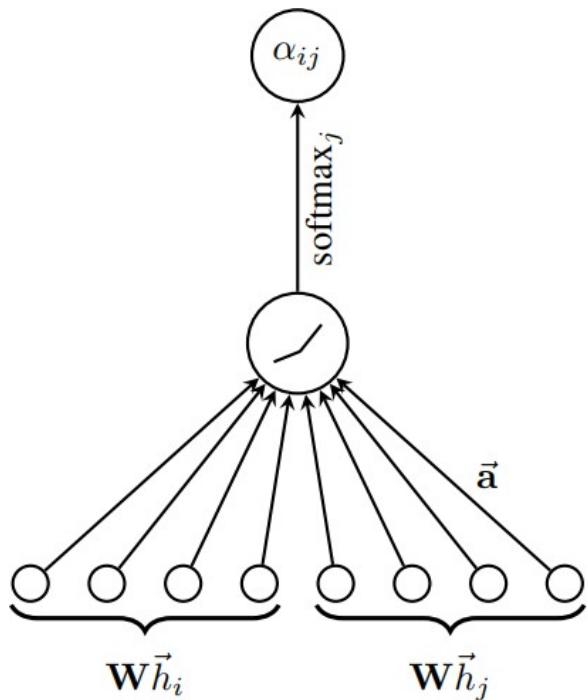
$$\blacktriangleright Z := \text{attention}(Q, K, V) = V \text{softmax} \left(\frac{1}{\sqrt{p}} Q^T K \right)$$

- a can be other functions for example a Transformer.

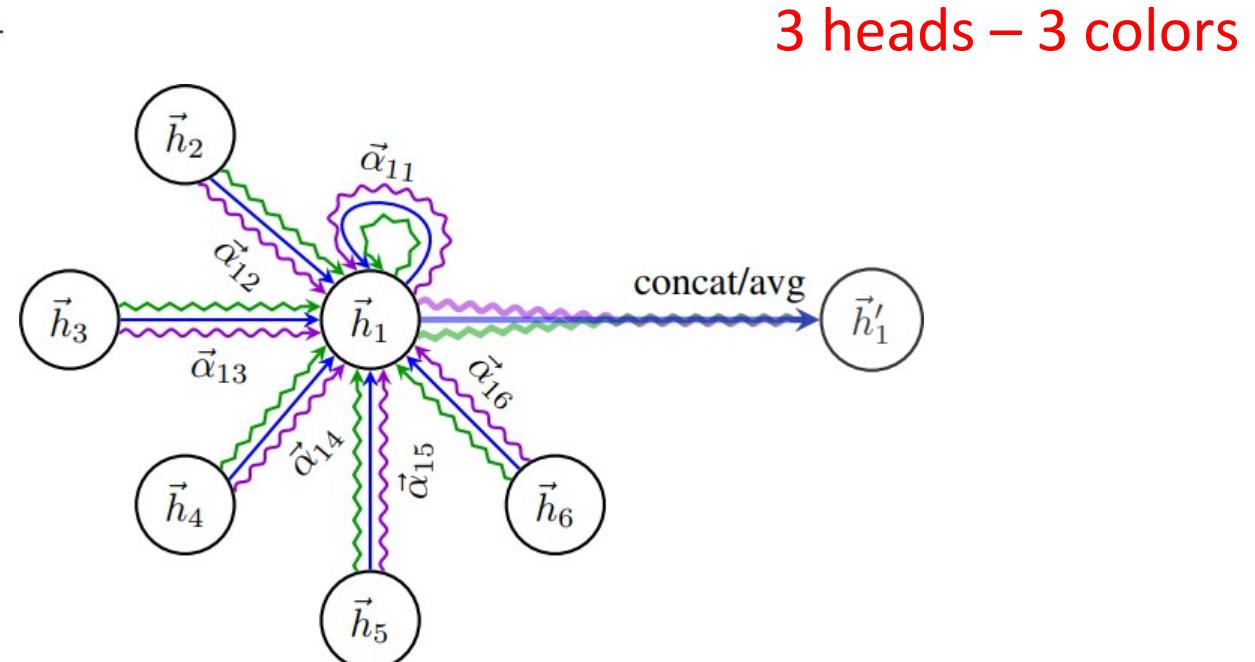
$$\alpha_{ij} = \frac{e^{a_{ij}}}{\sum_{k \in N_i} e^{a_{ik}}}$$

Multi-head attention in a single GAT

$$\alpha_{ij} = \frac{\exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left(\text{LeakyReLU} \left(\vec{\mathbf{a}}^T [\mathbf{W}\vec{h}_i \| \mathbf{W}\vec{h}_k] \right) \right)}$$



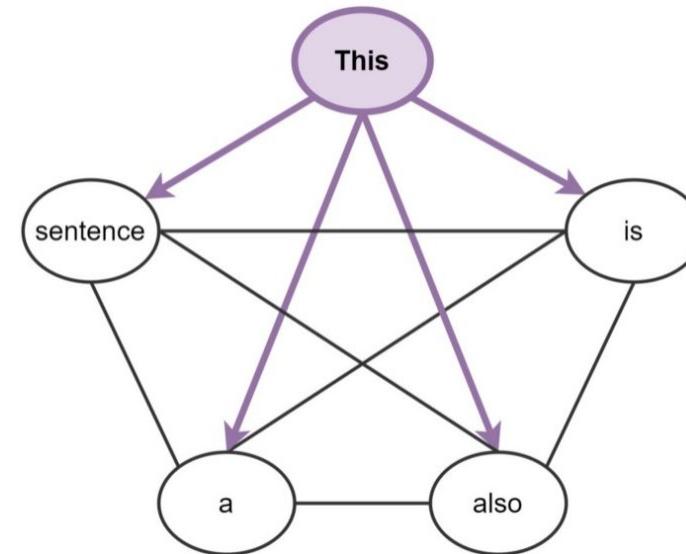
The attention mechanism



The multi-head attention. Different arrow styles and colours denote independent attention computations. The features from each head are concatenated or averaged.

Transformers are Graph Neural Networks

This is also a sentence



GA

- The attention weight can be computed as follows:

$$a_{ij} = a(h_i, h_j, e_{ij})$$

- a is a single-layer feedforward neural network.

$$\alpha_{ij} = \frac{e^{a_{ij}}}{\sum_{k \in N_i} e^{a_{ik}}}$$

Transformer

- The attention weight can be computed as follows:

$$a_{ij} = a(q_i, k_j)$$

$$a = \frac{1}{\sqrt{p}} (q_i^T k_j)$$

$$\alpha_{ij} = \frac{e^{a_{ij}}}{\sum_k e^{a_{ik}}}$$

$$v_i = \sum_j \alpha_{ij} v_j$$

We then describe the concrete operations of the adaptations of the normalization methods. Consider a batch of graphs $\{G_1, \dots, G_b\}$ where b is the batch size. Let n_g be the number of nodes in graph G_g . We generally denote $\hat{h}_{i,j,g}$ as the inputs to the normalization module, e.g., the j -th feature value of node v_i of graph G_g , $i = 1, \dots, n_g, j = 1, \dots, d, g = 1, \dots, b$. The adaptations take the general

$$\text{Norm}(\hat{h}_{i,j,g}) = \gamma \cdot \frac{\hat{h}_{i,j,g} - \mu}{\sigma} + \beta,$$

For BatchNorm, normalization and the computation of μ and σ are applied to all values in the same feature dimension across the nodes of all graphs in the batch as in Xu et al. (2019), i.e., over dimensions g, i of $h^i j, g$

Graph G_1 (3 nodes):

$$\hat{H}^{(G_1)} = \begin{bmatrix} 1 & 3 \\ 2 & 4 \\ 3 & 5 \end{bmatrix} \quad \begin{aligned} (1, 3) &\mapsto ((1 - 6)/5.099, (3 - 4)/2) \approx (-0.980, -0.5) \\ (2, 4) &\mapsto (-0.784, 0) \\ (3, 5) &\mapsto (-0.588, 0.5) \end{aligned}$$

Graph G_2 (2 nodes):

$$\hat{H}^{(G_2)} = \begin{bmatrix} 10 & 1 \\ 14 & 7 \end{bmatrix} \quad \begin{aligned} (10, 1) &\mapsto ((10 - 6)/5.099, (1 - 4)/2) \approx (0.784, -1.5) \\ (14, 7) &\mapsto (1.569, 1.5) \end{aligned}$$

All values in **channel 1** across the batch: [1, 2, 3, 10, 14]

All values in **channel 2** across the batch: [3, 4, 5, 1, 7]

Batch statistics (per channel):

- $\mu_1 = (1 + 2 + 3 + 10 + 14)/5 = 6, \quad \sigma_1 = \sqrt{\frac{(1-6)^2 + (2-6)^2 + (3-6)^2 + (10-6)^2 + (14-6)^2}{5}} = \sqrt{26} \approx 5.099.$
- $\mu_2 = (3 + 4 + 5 + 1 + 7)/5 = 4, \quad \sigma_2 = \sqrt{\frac{(3-4)^2 + (4-4)^2 + (5-4)^2 + (1-4)^2 + (7-4)^2}{5}} = \sqrt{4} = 2.$

For BatchNorm, normalization and the computation of μ and σ are applied to all values in the same feature dimension across the nodes of all graphs in the batch as in Xu et al. (2019), i.e., over dimensions g, i of $h^{\hat{}}_{\cdot i,j,g}$

To adapt LayerNorm to GNNs, we view each node as a basic component, resembling words in a sentence, and apply normalization to all feature values across different dimensions of each node, i.e., over dimension j of $h^{\hat{}}_{\cdot i,j,g}$.

For InstanceNorm, we regard each graph as an instance. The normalization is then applied to the feature values across all nodes for each individual graph, i.e., over dimension i of $h^{\hat{}}_{\cdot i,j,g}$.

$$\text{GraphNorm} \left(\hat{h}_{i,j} \right) = \gamma_j \cdot \frac{\hat{h}_{i,j} - \alpha_j \cdot \mu_j}{\hat{\sigma}_j} + \beta_j, \quad \mu_j = \frac{\sum_{i=1}^n \hat{h}_{i,j}}{n}, \quad \hat{\sigma}_j^2 = \frac{\sum_{i=1}^n (\hat{h}_{i,j} - \alpha_j \cdot \mu_j)^2}{n}.$$

InstanceNorm

$$\hat{H}^{(G_1)} = \begin{bmatrix} 1 & 3 \\ 2 & 4 \\ 3 & 5 \end{bmatrix}, \quad \mu^{(G_1)} = [2, 4]. \quad \hat{H} - \mu = \begin{bmatrix} -1 & -1 \\ 0 & 0 \\ 1 & 1 \end{bmatrix}.$$

Var.
2/3

$$(\quad) \quad \frac{\hat{H} - \mu}{\sigma} \approx \begin{bmatrix} -1.2247 & -1.2247 \\ 0 & 0 \\ 1.2247 & 1.2247 \end{bmatrix}.$$

GraphNorm

$$\alpha=[1, 0.5] \quad \hat{H} - \alpha \odot \mu = \begin{bmatrix} 1-2 & 3-2 \\ 2-2 & 4-2 \\ 3-2 & 5-2 \end{bmatrix} = \begin{bmatrix} -1 & 1 \\ 0 & 2 \\ 1 & 3 \end{bmatrix} \text{ Var. } 2/3, 14/3$$

$$\frac{\hat{H} - \alpha\mu}{\sigma} \approx \begin{bmatrix} -1.2247 & 0.4629 \\ 0 & 0.9258 \\ 1.2247 & 1.3887 \end{bmatrix}$$