

# Attention Layer

## Inputs:

**Query vectors:**  $\mathbf{Q}$  (Shape:  $N_Q \times D_Q$ )

**Input vectors:**  $\mathbf{X}$  (Shape:  $N_X \times D_X$ )

**Key matrix:**  $\mathbf{W}_K$  (Shape:  $D_X \times D_Q$ )

**Value matrix:**  $\mathbf{W}_V$  (Shape:  $D_X \times D_V$ )

## Computation:

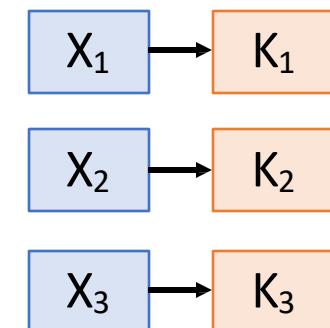
**Key vectors:**  $\mathbf{K} = \mathbf{X}\mathbf{W}_K$  (Shape:  $N_X \times D_Q$ )

**Value Vectors:**  $\mathbf{V} = \mathbf{X}\mathbf{W}_V$  (Shape:  $N_X \times D_V$ )

**Similarities:**  $E = \mathbf{Q}\mathbf{K}^T$  (Shape:  $N_Q \times N_X$ )  $E_{i,j} = \mathbf{Q}_i \cdot \mathbf{K}_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_Q \times N_X$ )

**Output vectors:**  $\mathbf{Y} = A\mathbf{V}$  (Shape:  $N_Q \times D_V$ )  $Y_i = \sum_j A_{i,j} \mathbf{V}_j$



# Attention Layer

## Inputs:

**Query vectors:**  $Q$  (Shape:  $N_Q \times D_Q$ )

**Input vectors:**  $X$  (Shape:  $N_X \times D_X$ )

**Key matrix:**  $W_K$  (Shape:  $D_X \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_X \times D_V$ )

## Computation:

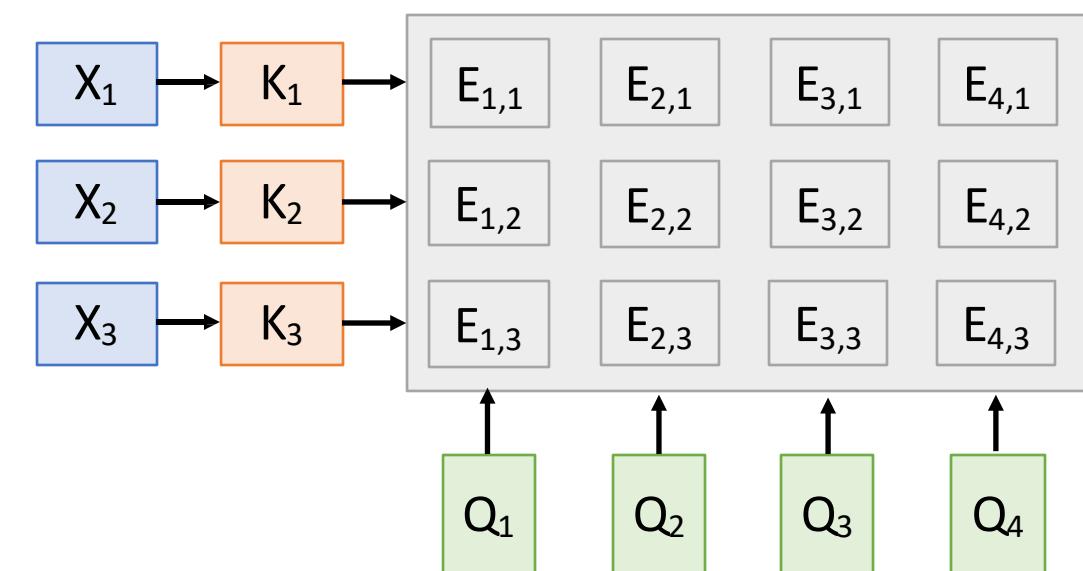
**Key vectors:**  $K = XW_K$  (Shape:  $N_X \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_X \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_Q \times N_X$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_Q \times N_X$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_Q \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Attention Layer

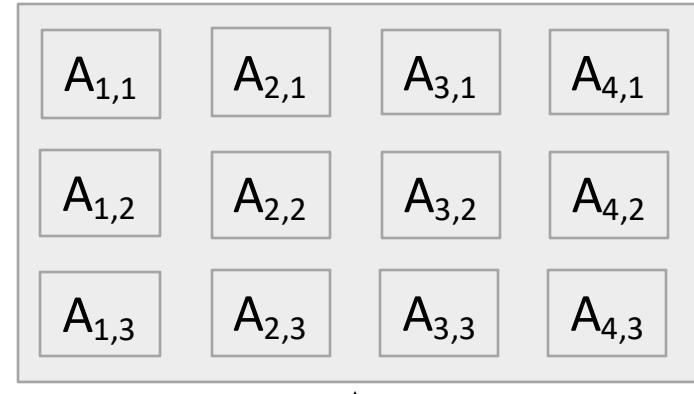
## Inputs:

**Query vectors:**  $Q$  (Shape:  $N_Q \times D_Q$ )

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )



Softmax(↑)

## Computation:

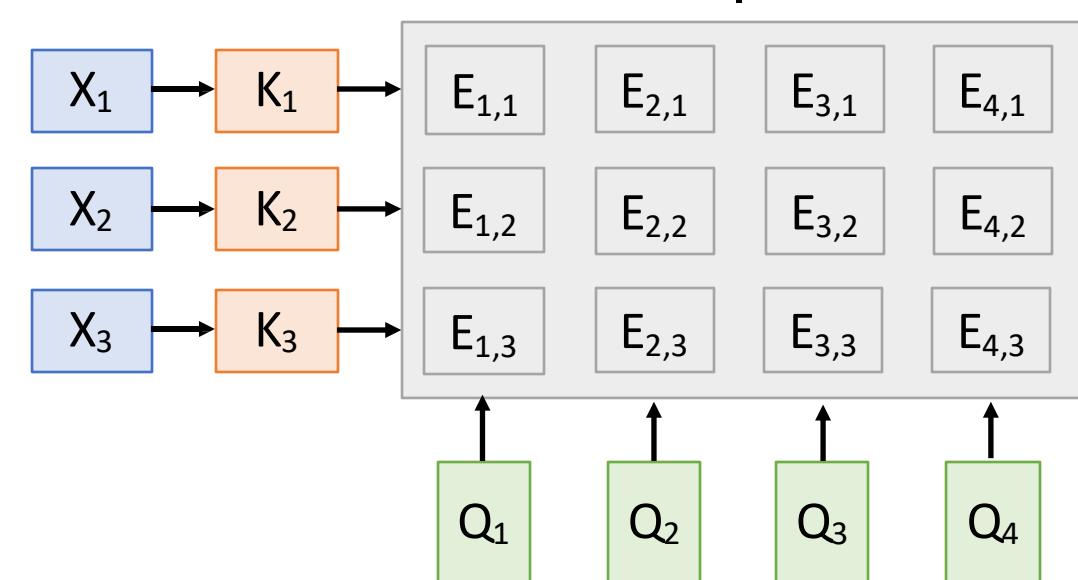
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_Q \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E)$  (Shape:  $N_Q \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_Q \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Attention Layer

## Inputs:

**Query vectors:**  $\mathbf{Q}$  (Shape:  $N_Q \times D_Q$ )

**Input vectors:**  $\mathbf{X}$  (Shape:  $N_X \times D_X$ )

**Key matrix:**  $\mathbf{W}_K$  (Shape:  $D_X \times D_Q$ )

**Value matrix:**  $\mathbf{W}_V$  (Shape:  $D_X \times D_V$ )

## Computation:

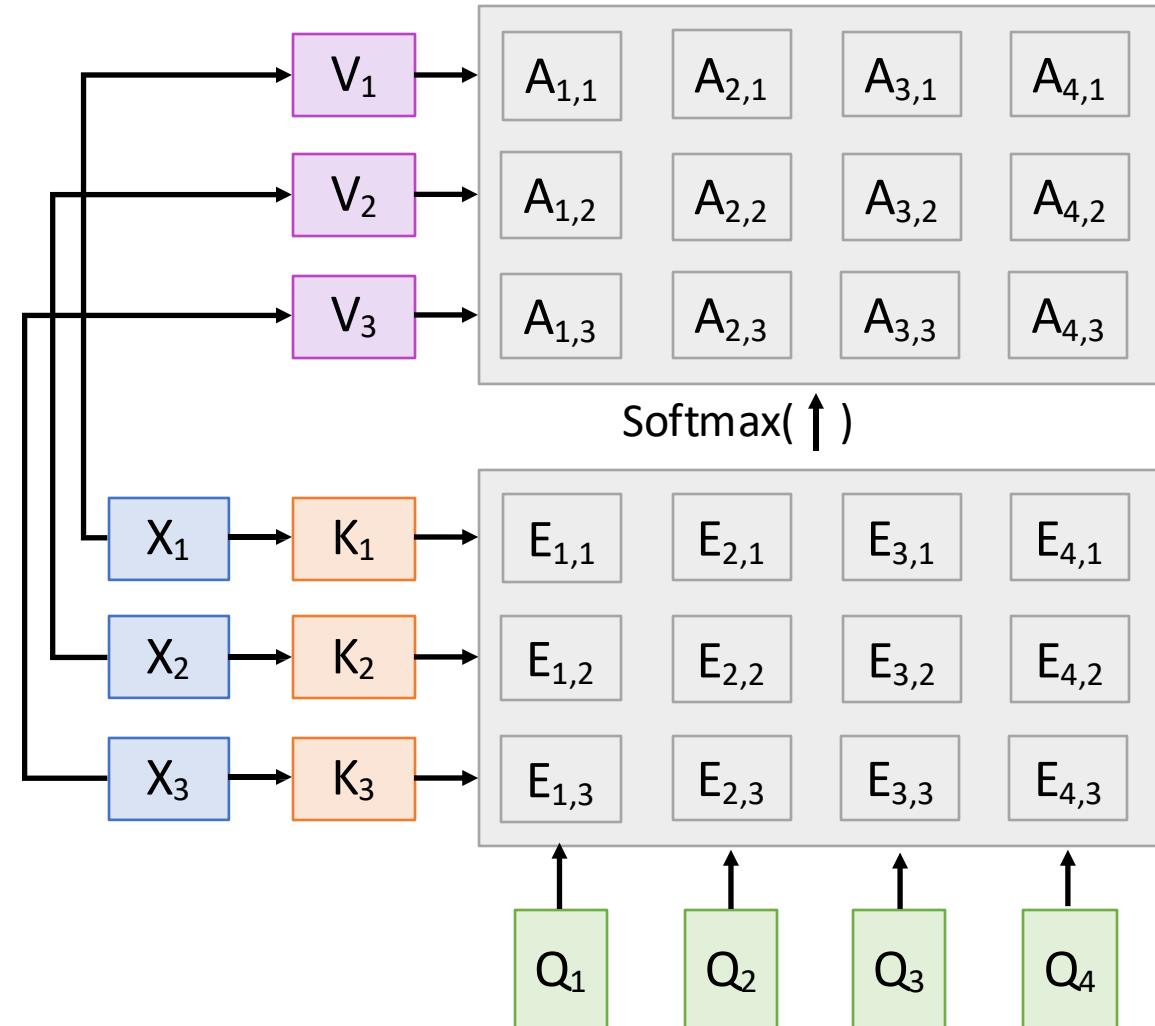
**Key vectors:**  $\mathbf{K} = \mathbf{X}\mathbf{W}_K$  (Shape:  $N_X \times D_Q$ )

**Value Vectors:**  $\mathbf{V} = \mathbf{X}\mathbf{W}_V$  (Shape:  $N_X \times D_V$ )

**Similarities:**  $\mathbf{E} = \mathbf{Q}\mathbf{K}^T$  (Shape:  $N_Q \times N_X$ )  $E_{i,j} = \mathbf{Q}_i \cdot \mathbf{K}_j / \text{sqrt}(D_Q)$

**Attention weights:**  $\mathbf{A} = \text{softmax}(\mathbf{E}, \text{dim}=1)$  (Shape:  $N_Q \times N_X$ )

**Output vectors:**  $\mathbf{Y} = \mathbf{A}\mathbf{V}$  (Shape:  $N_Q \times D_V$ )  $Y_i = \sum_j A_{i,j} \mathbf{V}_j$



# Attention Layer

## Inputs:

**Query vectors:**  $\mathbf{Q}$  (Shape:  $N_Q \times D_Q$ )

**Input vectors:**  $\mathbf{X}$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $\mathbf{W}_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $\mathbf{W}_V$  (Shape:  $D_x \times D_V$ )

## Computation:

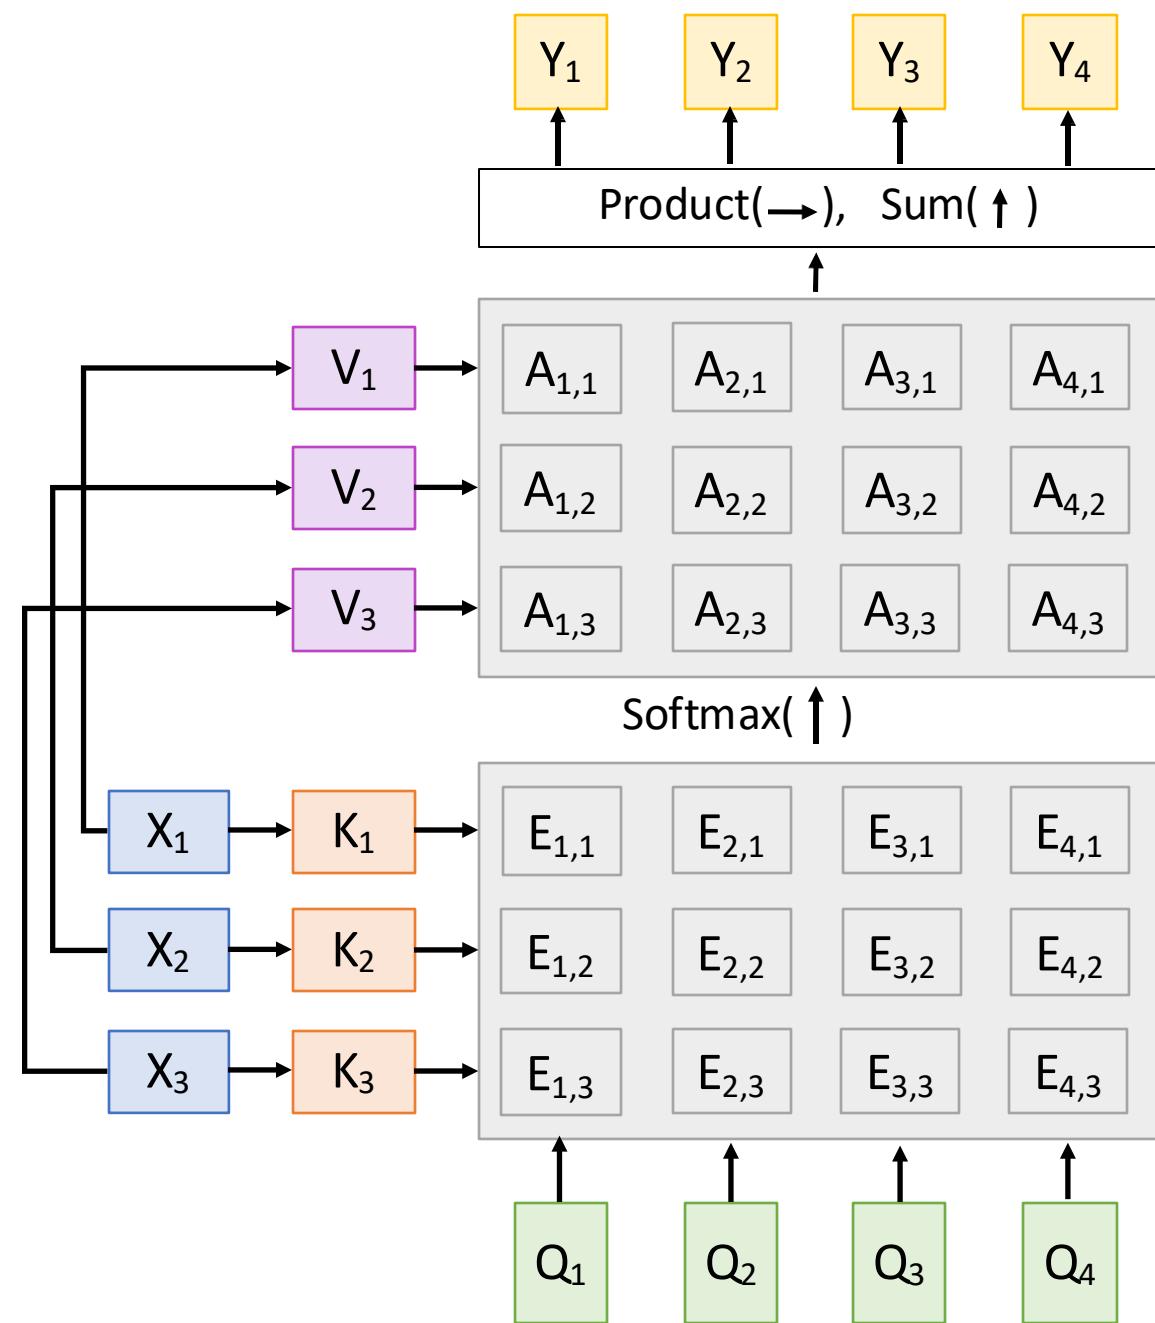
**Key vectors:**  $\mathbf{K} = \mathbf{X}\mathbf{W}_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $\mathbf{V} = \mathbf{X}\mathbf{W}_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $\mathbf{E} = \mathbf{Q}\mathbf{K}^T$  (Shape:  $N_Q \times N_x$ )  $E_{i,j} = \mathbf{Q}_i \cdot \mathbf{K}_j / \text{sqrt}(D_Q)$

**Attention weights:**  $\mathbf{A} = \text{softmax}(\mathbf{E}, \text{dim}=1)$  (Shape:  $N_Q \times N_x$ )

**Output vectors:**  $\mathbf{Y} = \mathbf{A}\mathbf{V}$  (Shape:  $N_Q \times D_V$ )  $Y_i = \sum_j A_{i,j} \mathbf{V}_j$



# Self-Attention Layer

One **query** per **input vector**

- Inputs:
- Input vectors:  $X$  (Shape:  $N_x \times D_x$ ) Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )
  - Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )  
Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )
- Computation:
- Query vectors:  $Q = XW_Q$
- Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )
- Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )
- Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,i} = Q_i \cdot K_i / \sqrt{D_Q}$

$X_1$

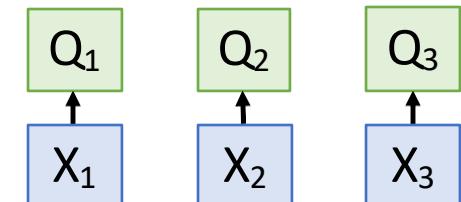
$X_2$

$X_3$

# Self-Attention Layer

One **query** per **input vector**

- Inputs:
  - Input vectors:  $X$   
(Shape:  $N_x \times D_x$ ) Key matrix:  
 $W_K$  (Shape:  $D_x \times D_Q$ ) Value matrix:  
 $W_V$  (Shape:  $D_x \times D_V$ ) Query matrix:  
 $W_Q$  (Shape:  $D_x \times D_Q$ )
- Computation:
  - Query vectors:  $Q = XW_Q$
  - Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )
  - Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )
  - Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j$



# Self-Attention Layer

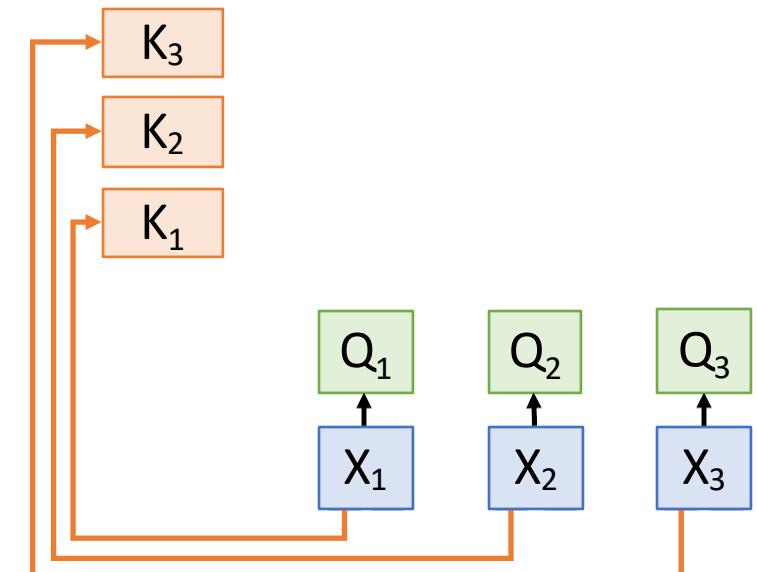
One **query** per **input vector**

- Inputs:
- Input vectors:  $X$  (Shape:  $N_x \times D_x$ ) Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

- Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )
- Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

- Computation:

- Query vectors:  $Q = XW_Q$
- Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )
- Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )
- Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{scale}(D_Q)$  Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )



# Self-Attention Layer

One **query** per **input vector**

- Inputs:
- Input vectors:  $X$

(Shape:  $N_x \times D_x$ ) Key matrix:

$W_K$  (Shape:  $D_x \times D_Q$ ) Value matrix:

$W_V$  (Shape:  $D_x \times D_V$ ) Query matrix:

$W_Q$  (Shape:  $D_x \times D_Q$ )

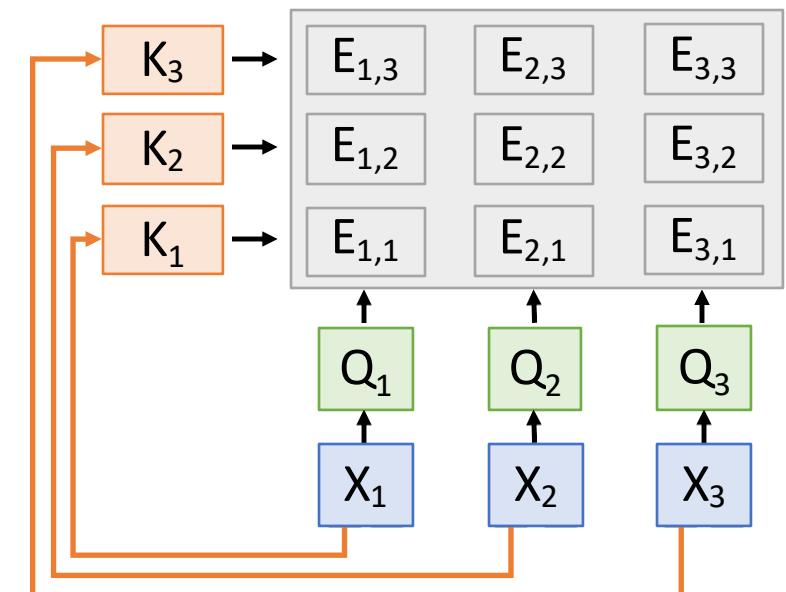
- Computation:

- Query vectors:  $Q = XW_Q$

- Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

- Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

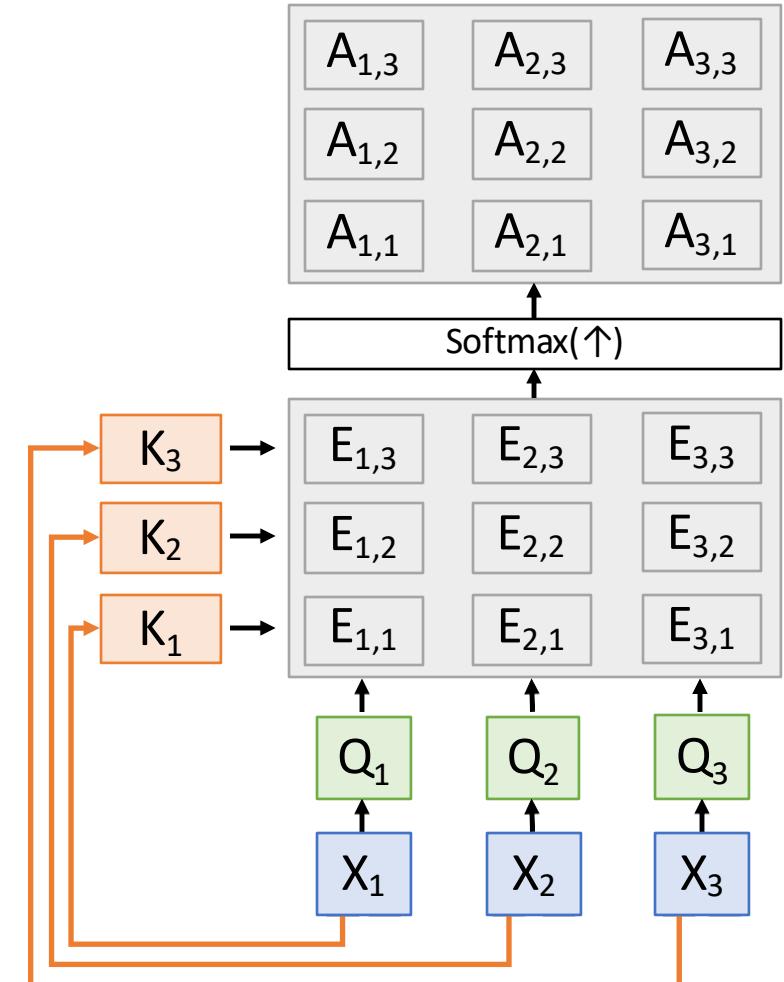
- Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{ij} = Q_i \cdot K_j$



# Self-Attention Layer

One **query** per **input vector**

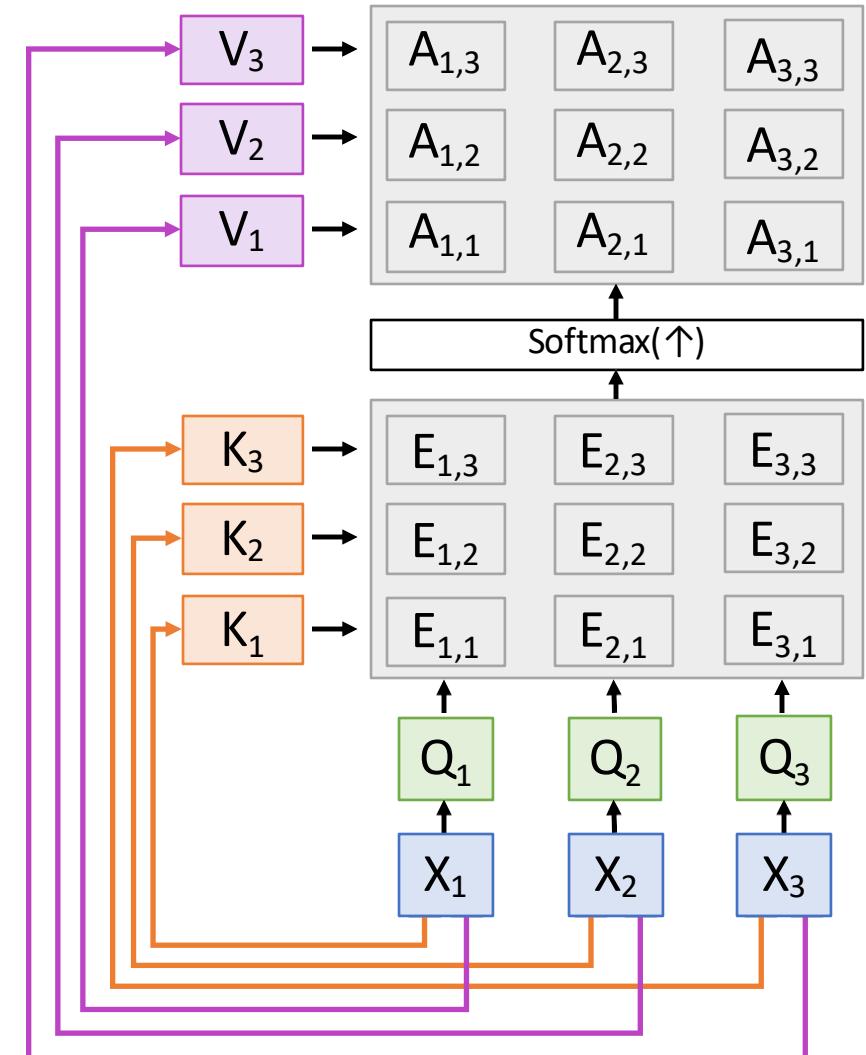
- Inputs:
- Input vectors:  $X$   
(Shape:  $N_x \times D_x$ ) Key matrix:  
 $W_K$  (Shape:  $D_x \times D_Q$ ) Value matrix:  
 $W_V$  (Shape:  $D_x \times D_V$ ) Query matrix:  
 $W_Q$  (Shape:  $D_x \times D_Q$ )
- Computation:
- Query vectors:  $Q = XW_Q$
- Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )
- Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )
- Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j$



# Self-Attention Layer

One **query** per **input vector**

- Inputs:
- Input vectors:  $X$   
(Shape:  $N_x \times D_x$ ) Key matrix:  
 $W_K$  (Shape:  $D_x \times D_Q$ ) Value matrix:  
 $W_V$  (Shape:  $D_x \times D_V$ ) Query matrix:  
 $W_Q$  (Shape:  $D_x \times D_Q$ )
- Computation:
- Query vectors:  $Q = XW_Q$
- Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )
- Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )
- Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j$



# Self-Attention Layer

One **query** per **input vector**

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

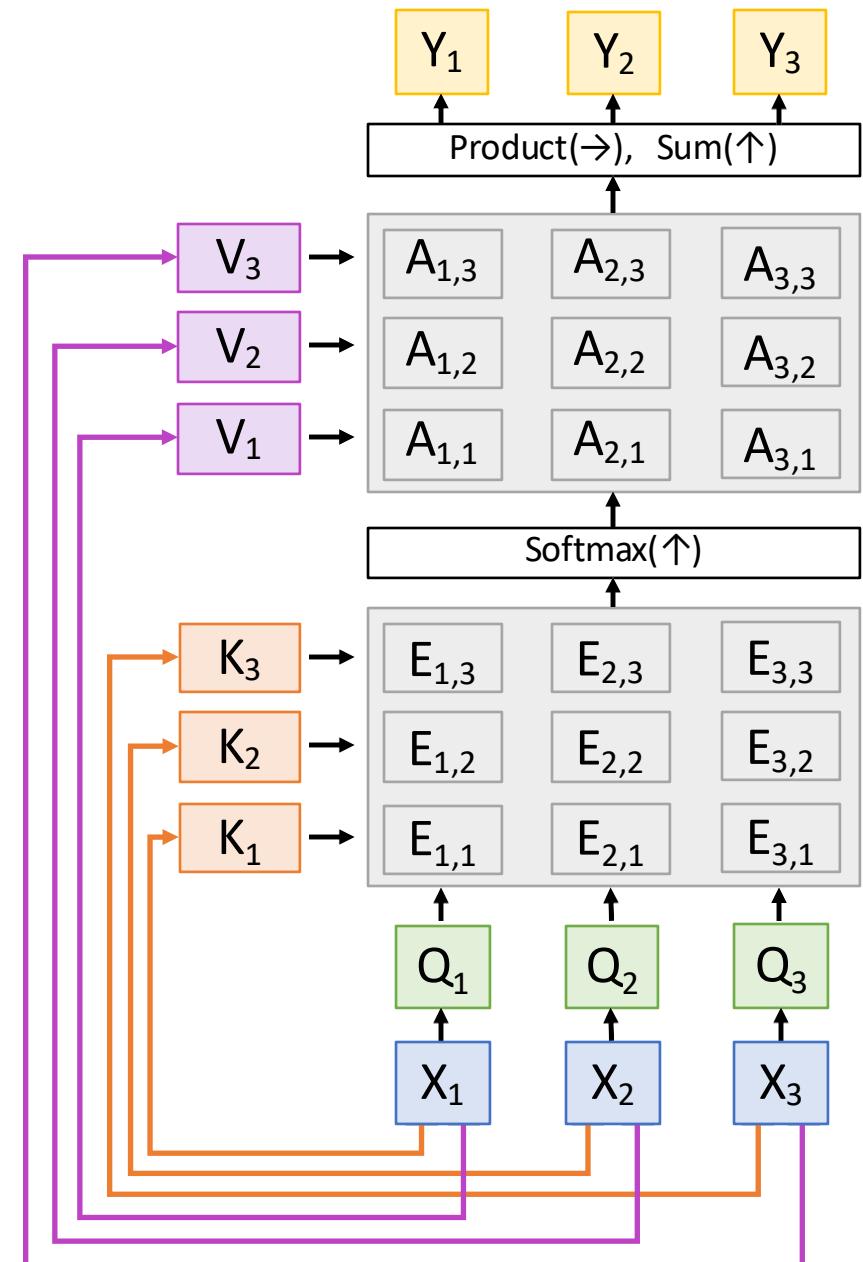
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Self-Attention Layer

## Inputs:

Input vectors:  $X$  (Shape:  $N_x \times D_x$ )

Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )

Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

Query vectors:  $Q = XW_Q$

Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

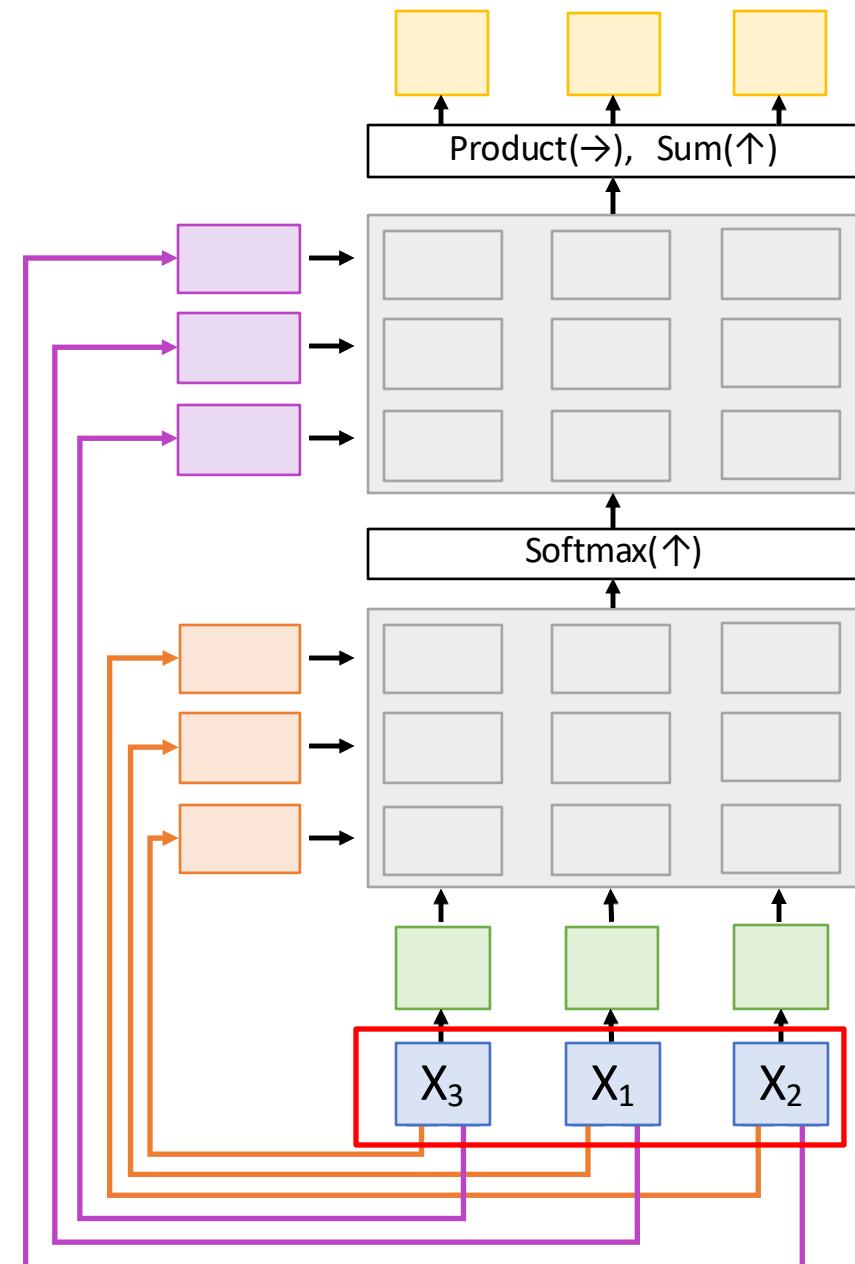
Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

Output vectors:  $Y = A V$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting**  
the input vectors:



# Self-Attention Layer

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

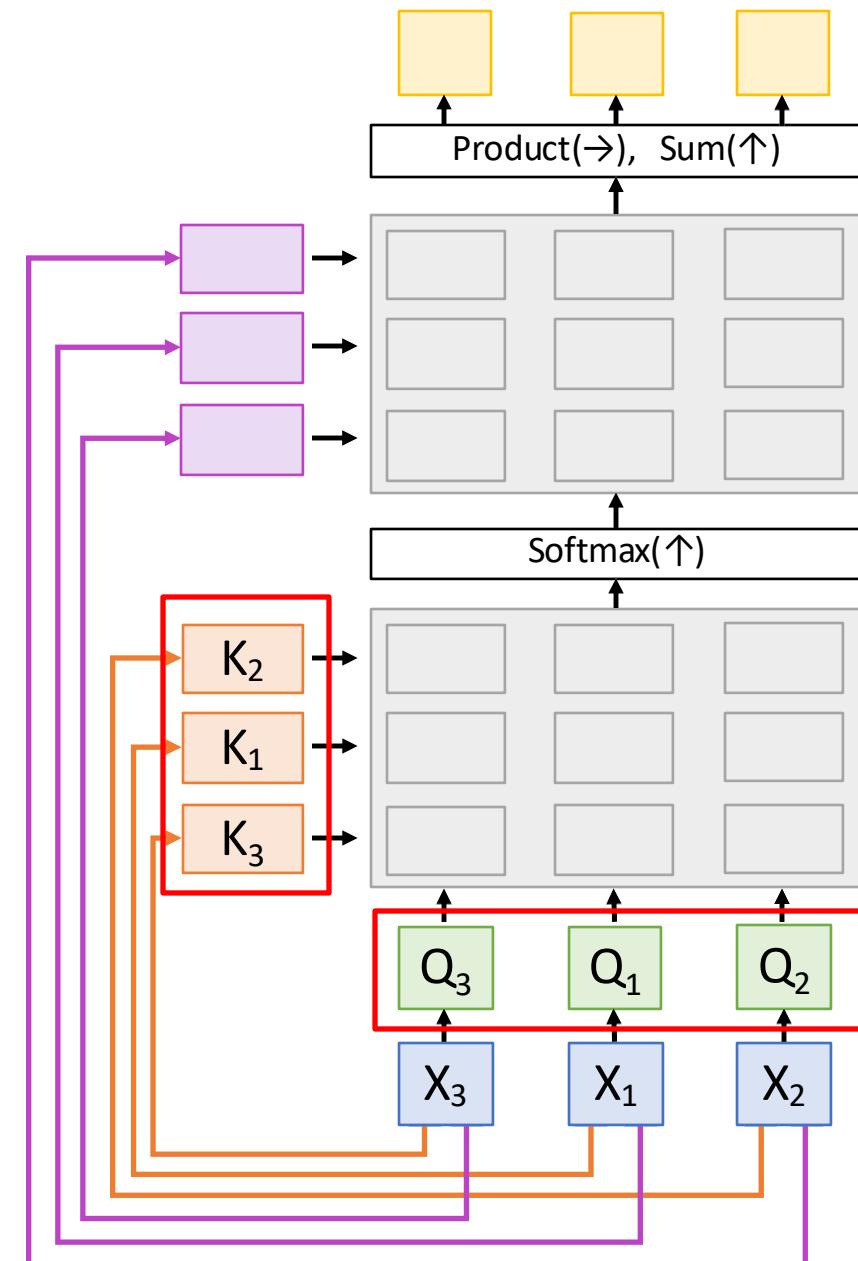
**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Queries and Keys will be the same, but permuted



# Self-Attention Layer

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

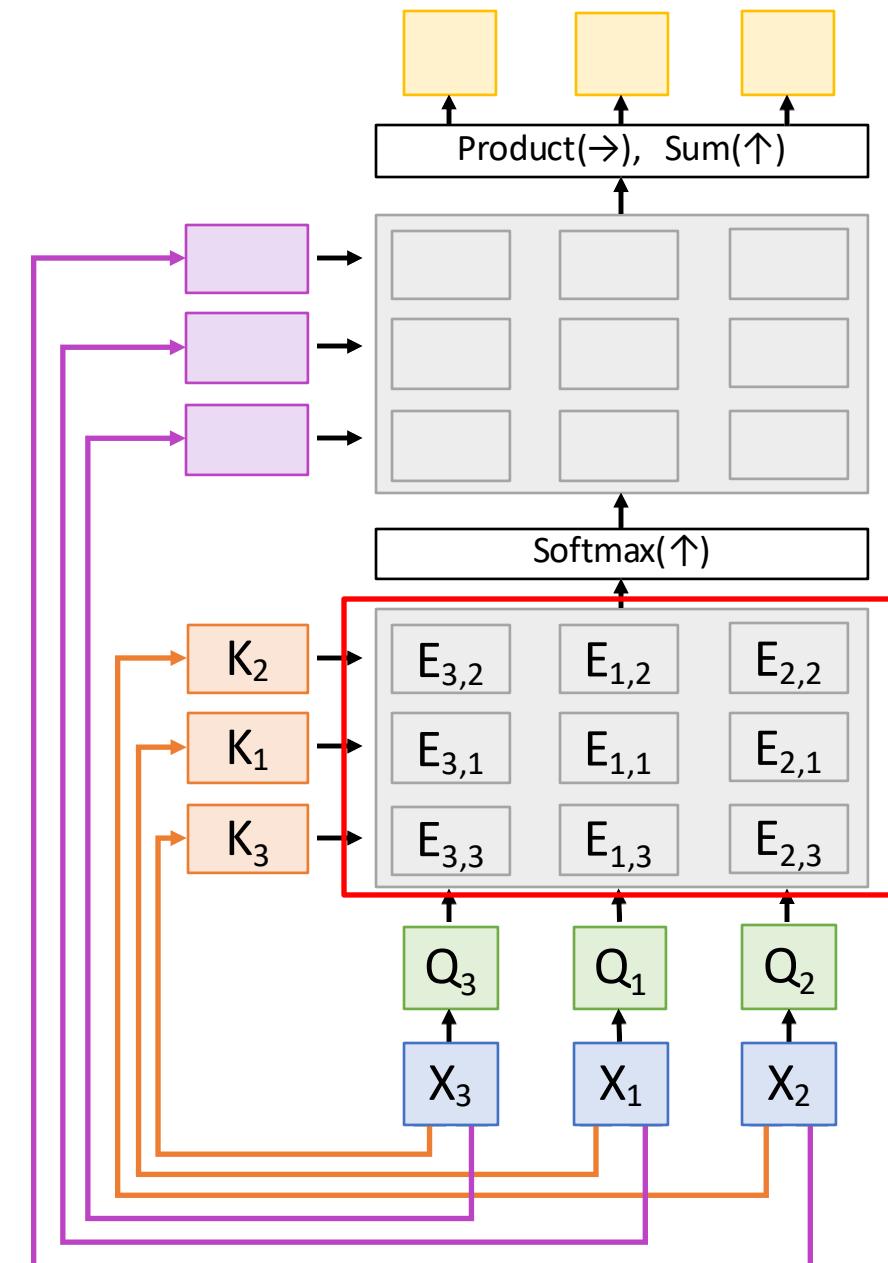
**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Similarities will be the same, but permuted



# Self-Attention Layer

## Inputs:

Input vectors:  $X$  (Shape:  $N_x \times D_x$ )

Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )

Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

Query vectors:  $Q = XW_Q$

Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

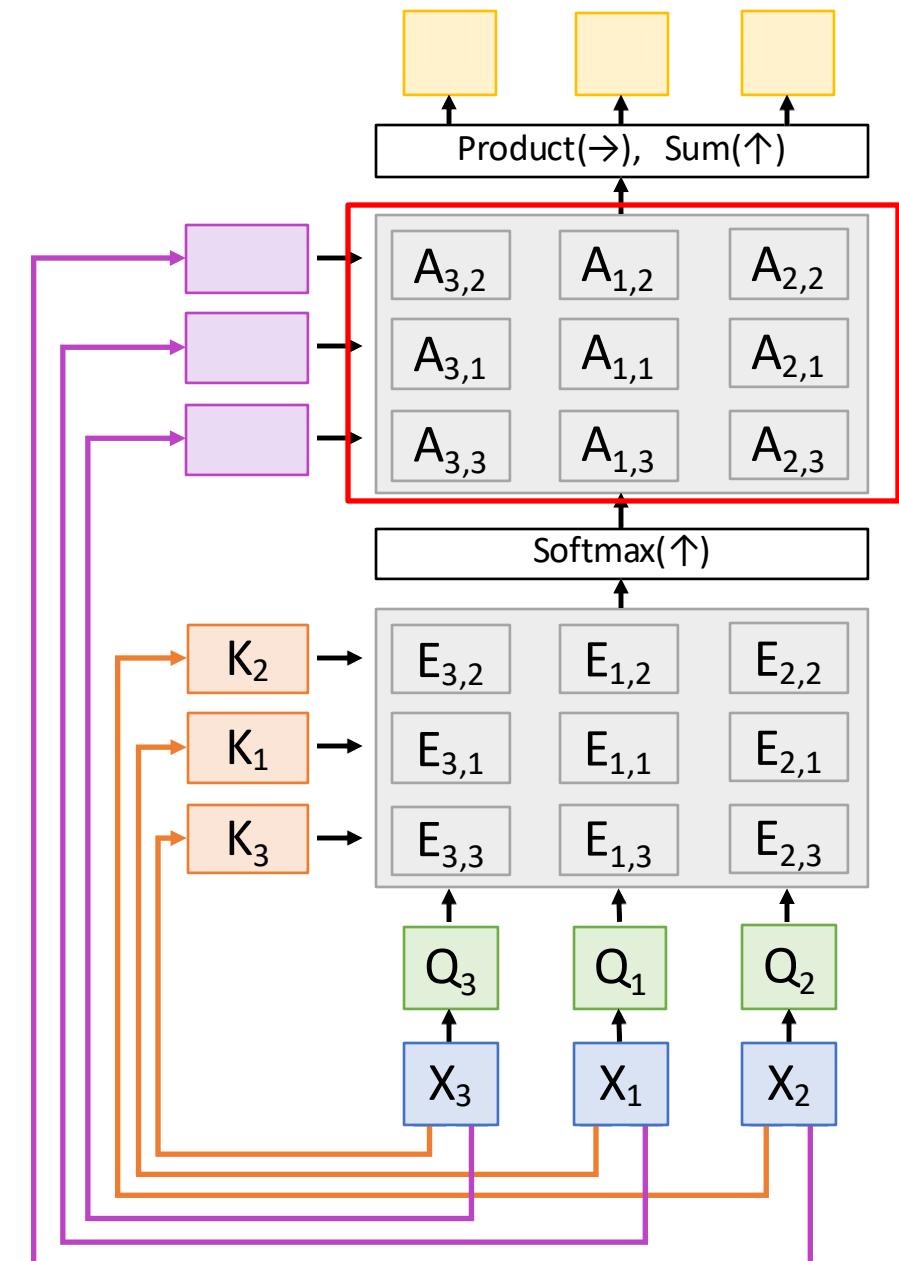
Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

Output vectors:  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Attention weights will be the same, but permuted



# Self-Attention Layer

## Inputs:

Input vectors:  $X$  (Shape:  $N_x \times D_x$ )

Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )

Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

Query vectors:  $Q = XW_Q$

Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

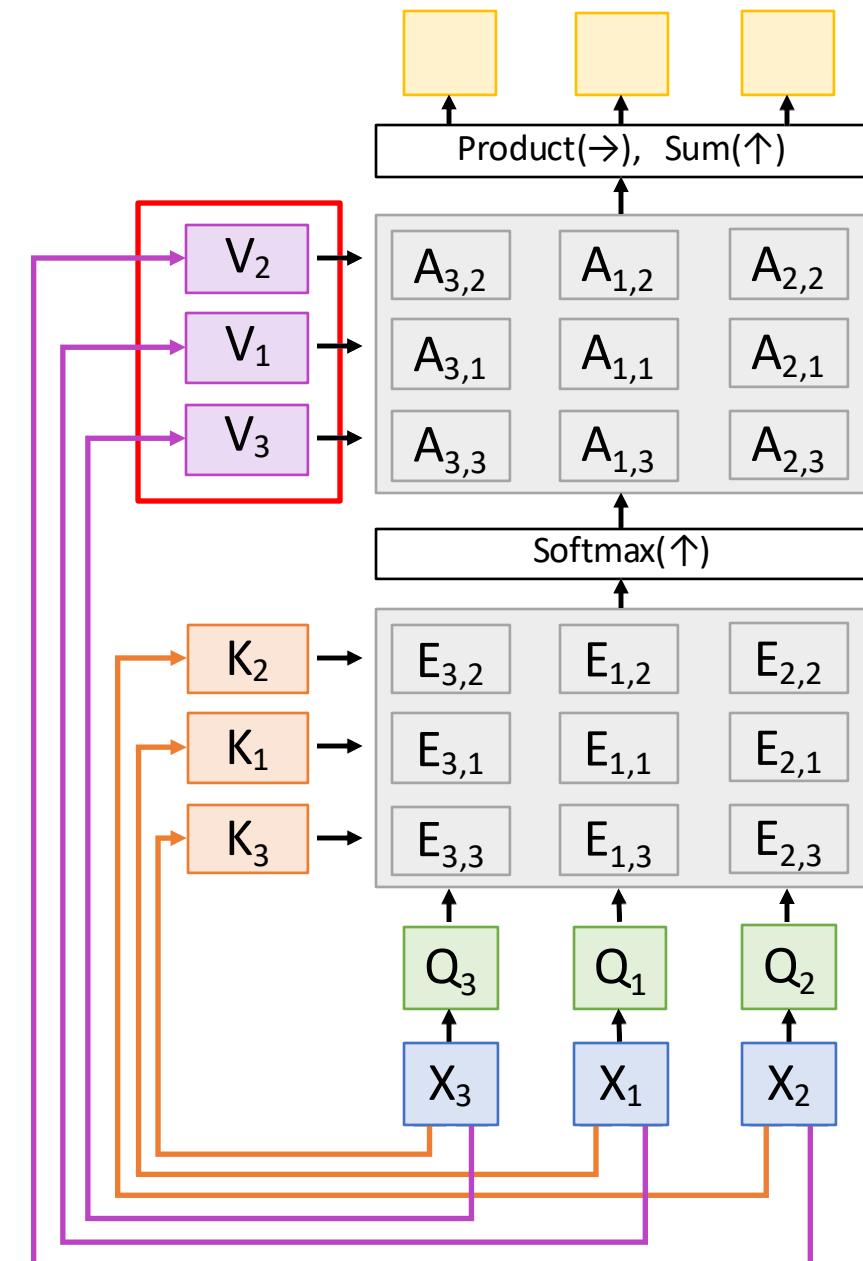
Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

Output vectors:  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Values will be the same, but permuted



# Self-Attention Layer

## Inputs:

Input vectors:  $X$  (Shape:  $N_x \times D_x$ )

Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )

Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

Query vectors:  $Q = XW_Q$

Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

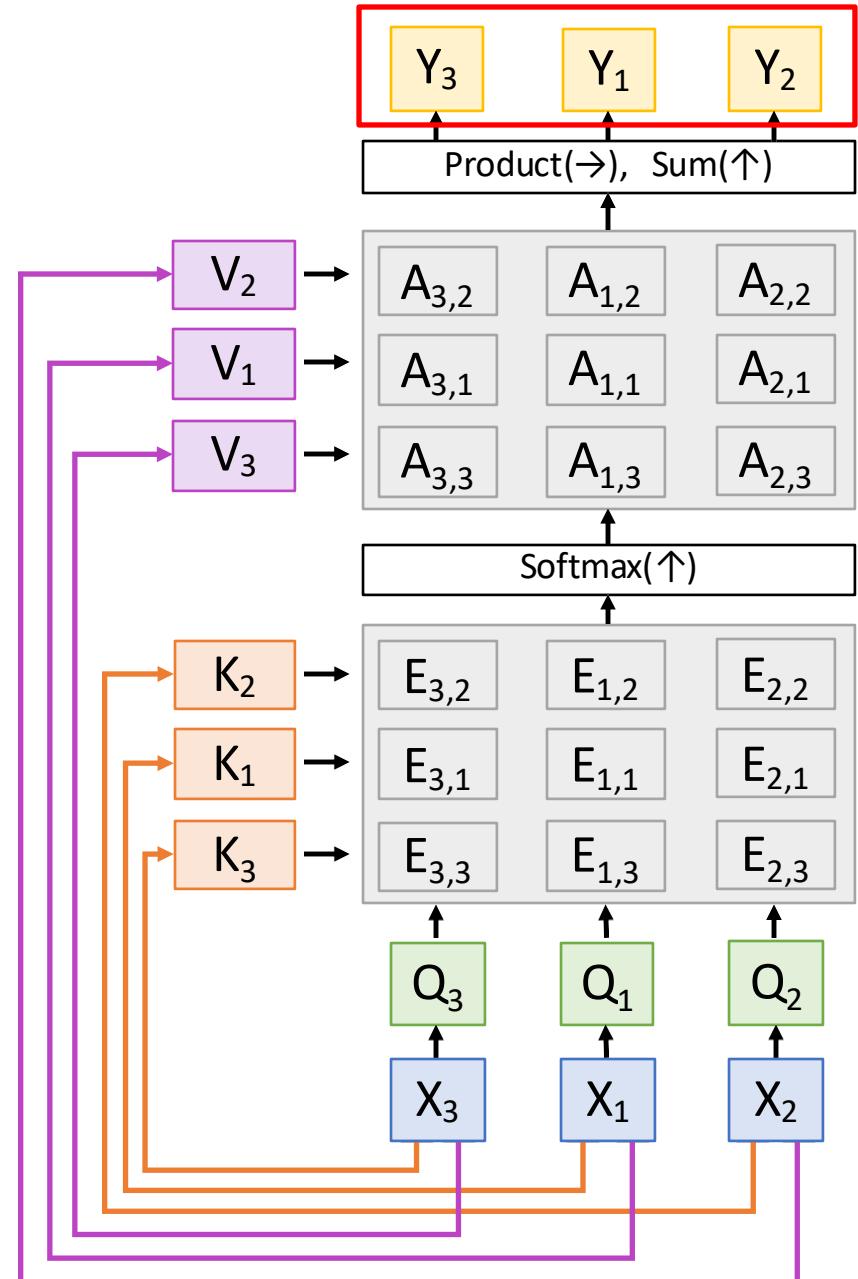
Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

Output vectors:  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Outputs will be the same, but permuted



# Self-Attention Layer

## Inputs:

Input vectors:  $X$  (Shape:  $N_x \times D_x$ )

Key matrix:  $W_K$  (Shape:  $D_x \times D_Q$ )

Value matrix:  $W_V$  (Shape:  $D_x \times D_V$ )

Query matrix:  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

Query vectors:  $Q = XW_Q$

Key vectors:  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

Value Vectors:  $V = XW_V$  (Shape:  $N_x \times D_V$ )

Similarities:  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

Attention weights:  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

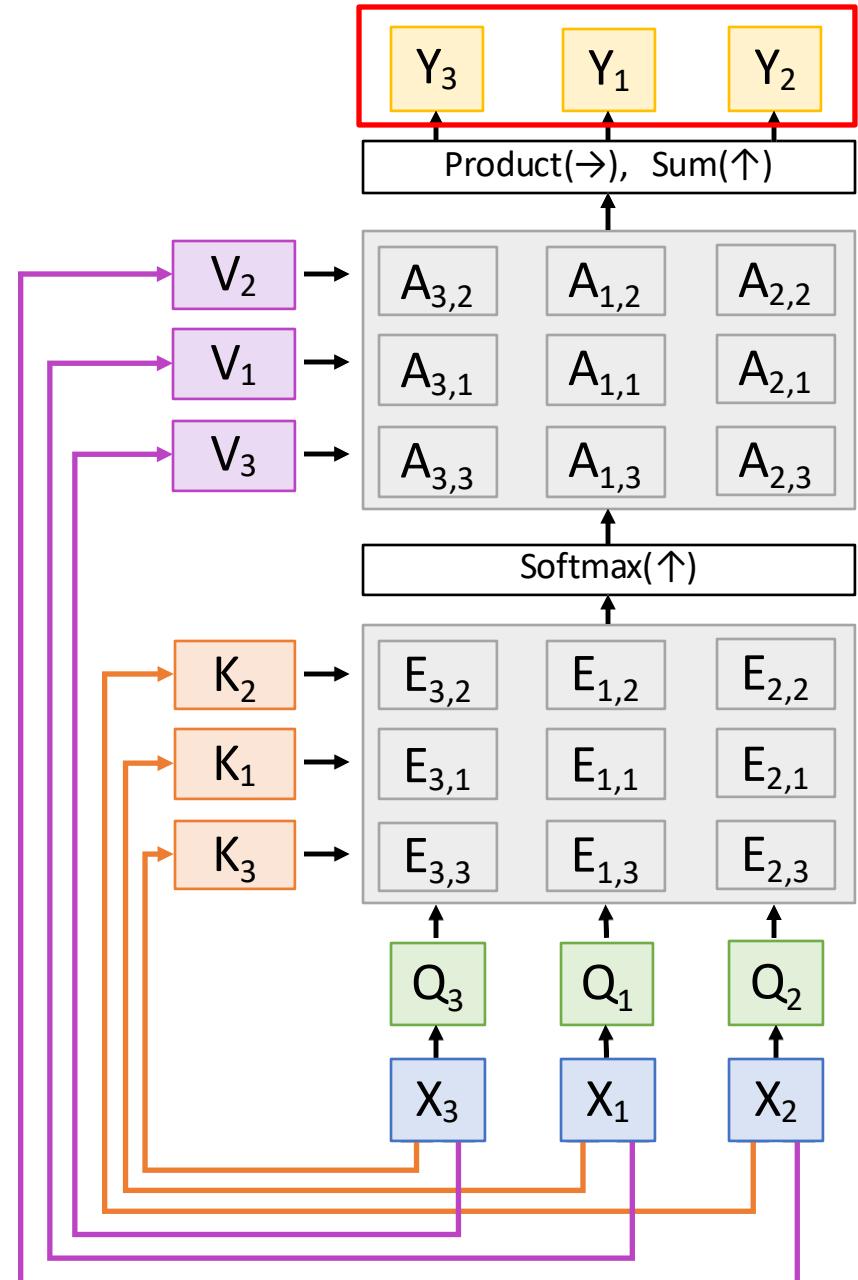
Output vectors:  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

Consider **permuting** the input vectors:

Outputs will be the same, but permuted

Self-attention layer is **Permutation Equivariant**  
 $f(s(x)) = s(f(x))$

Self-Attention layer works on **sets** of vectors



# Self-Attention Layer

Self attention doesn't  
“know” the order of the  
vectors it is processing!

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

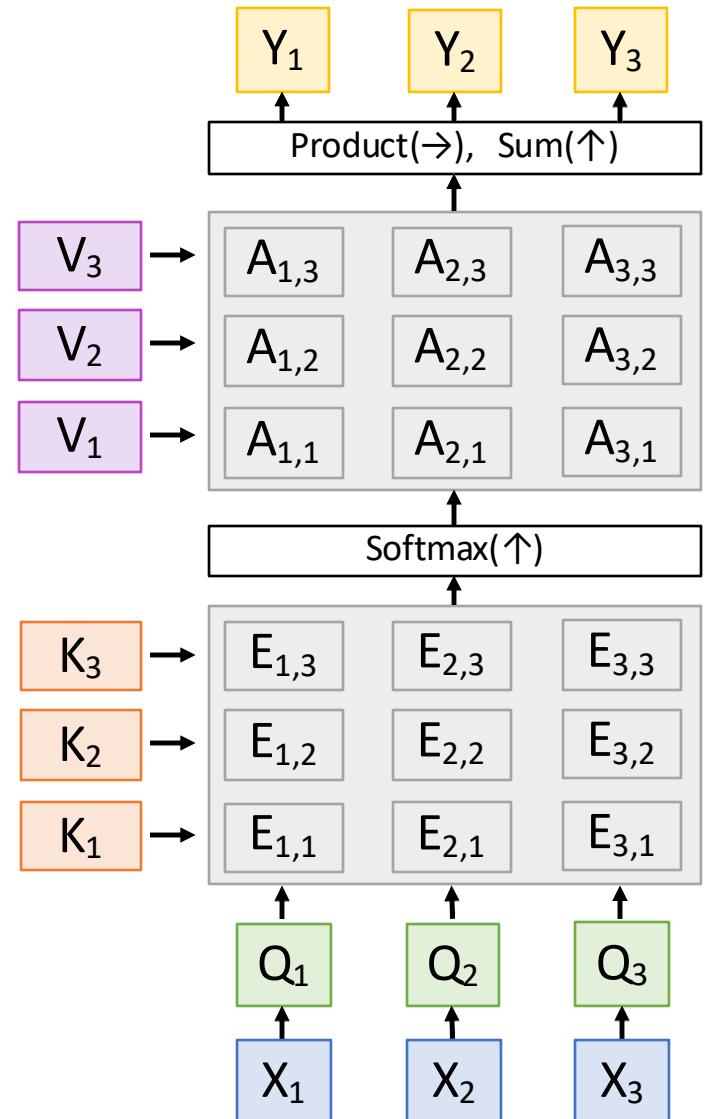
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Self-Attention Layer

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

Self attention doesn't  
“know” the order of the  
vectors it is processing!

In order to make  
processing position-  
aware, concatenate input  
with **positional encoding**

$E$  can be learned lookup  
table, or fixed function

## Computation:

**Query vectors:**  $Q = XW_Q$

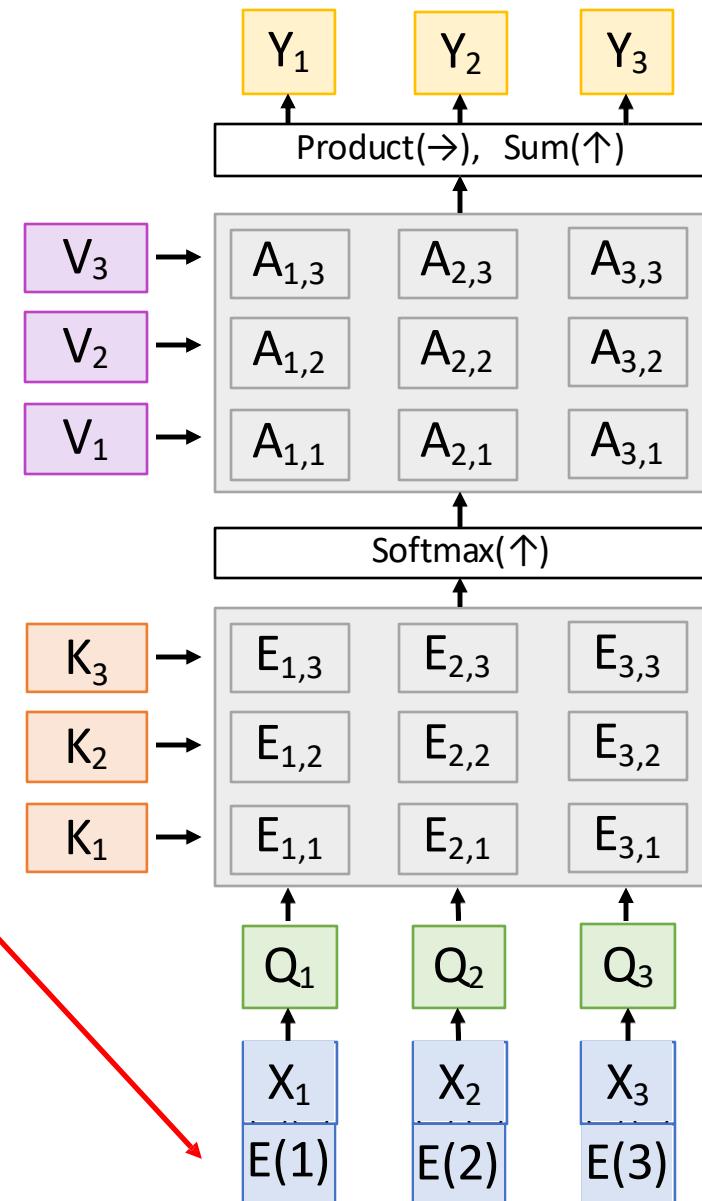
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Masked Self-Attention Layer

Don't let vectors "look ahead" in the sequence

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

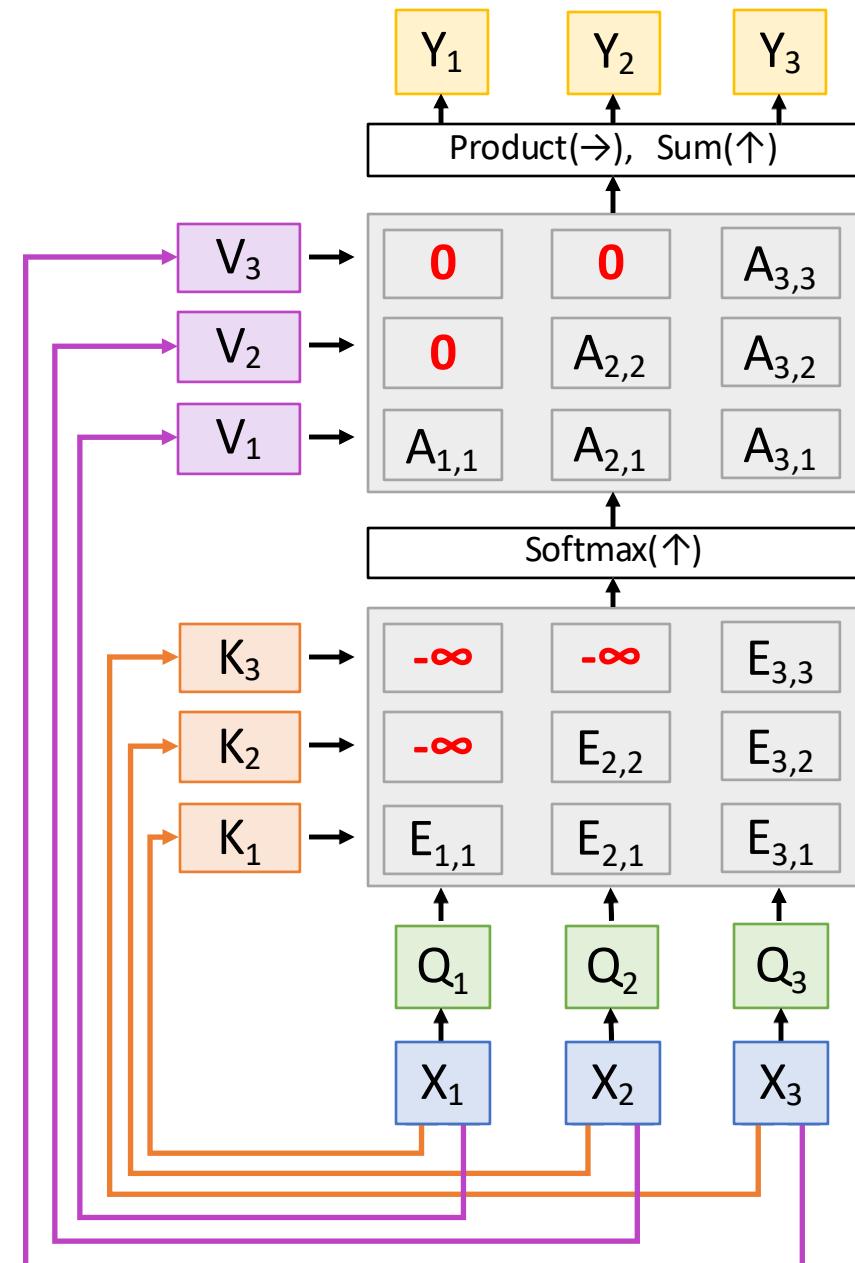
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Masked Self-Attention Layer

Don't let vectors "look ahead" in the sequence

Used for language modeling (predict next word)

## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

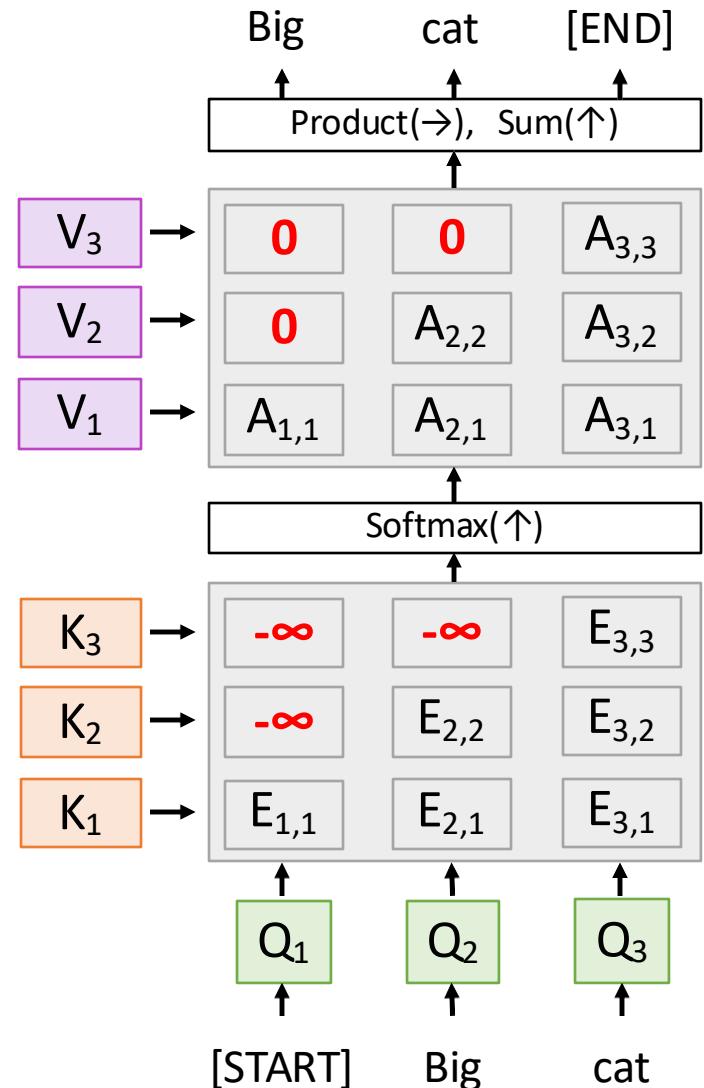
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \text{sqrt}(D_Q)$

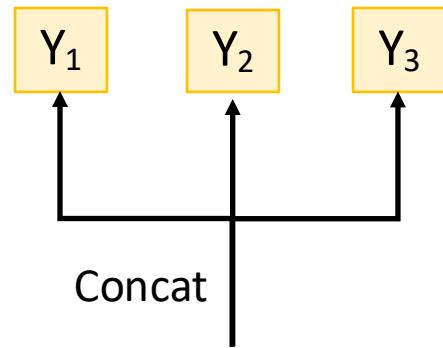
**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$



# Multihead Self-Attention Layer

Use H independent  
“Attention Heads” in parallel



## Inputs:

**Input vectors:**  $X$  (Shape:  $N_x \times D_x$ )

**Key matrix:**  $W_K$  (Shape:  $D_x \times D_Q$ )

**Value matrix:**  $W_V$  (Shape:  $D_x \times D_V$ )

**Query matrix:**  $W_Q$  (Shape:  $D_x \times D_Q$ )

## Computation:

**Query vectors:**  $Q = XW_Q$

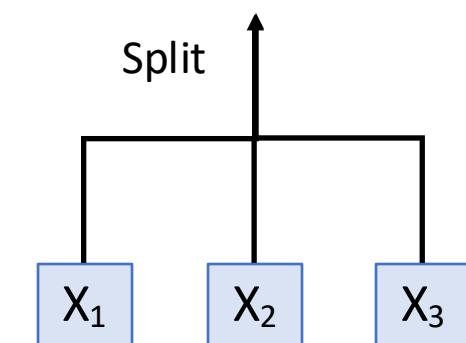
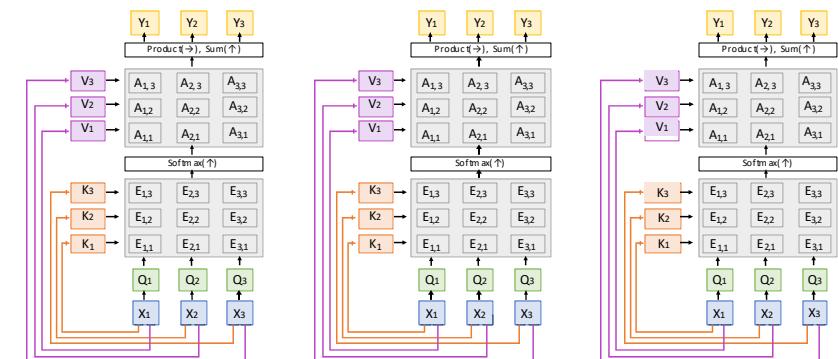
**Key vectors:**  $K = XW_K$  (Shape:  $N_x \times D_Q$ )

**Value Vectors:**  $V = XW_V$  (Shape:  $N_x \times D_V$ )

**Similarities:**  $E = QK^T$  (Shape:  $N_x \times N_x$ )  $E_{i,j} = Q_i \cdot K_j / \sqrt{D_Q}$

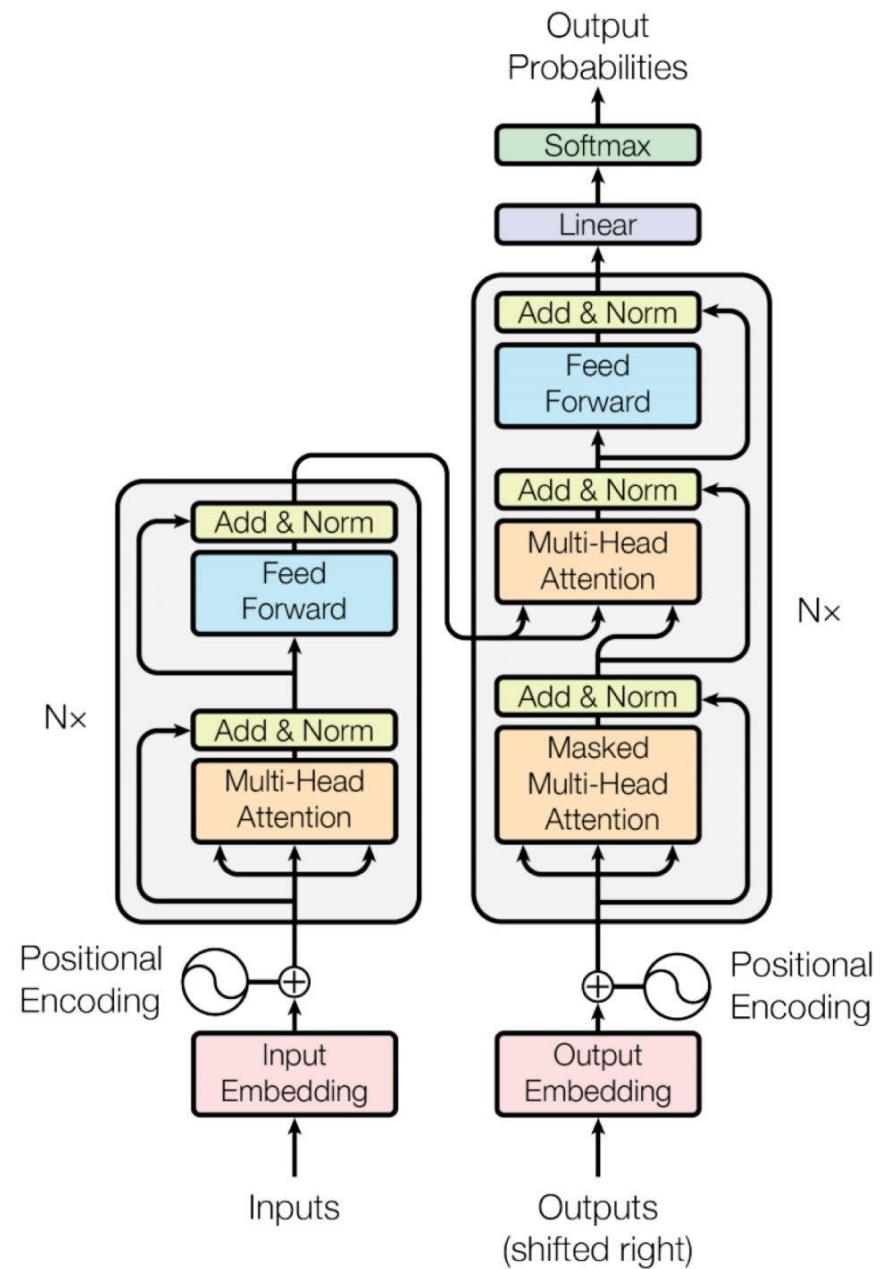
**Attention weights:**  $A = \text{softmax}(E, \text{dim}=1)$  (Shape:  $N_x \times N_x$ )

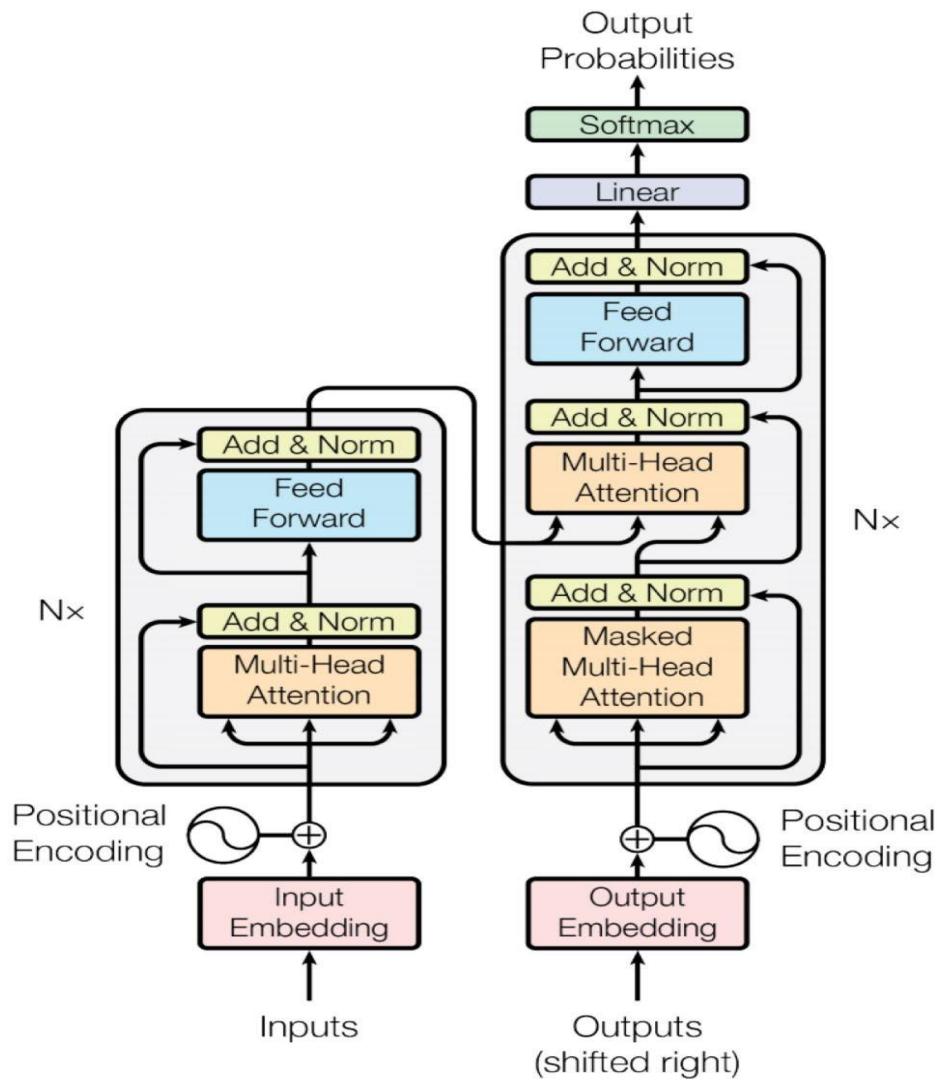
**Output vectors:**  $Y = AV$  (Shape:  $N_x \times D_V$ )  $Y_i = \sum_j A_{i,j} V_j$

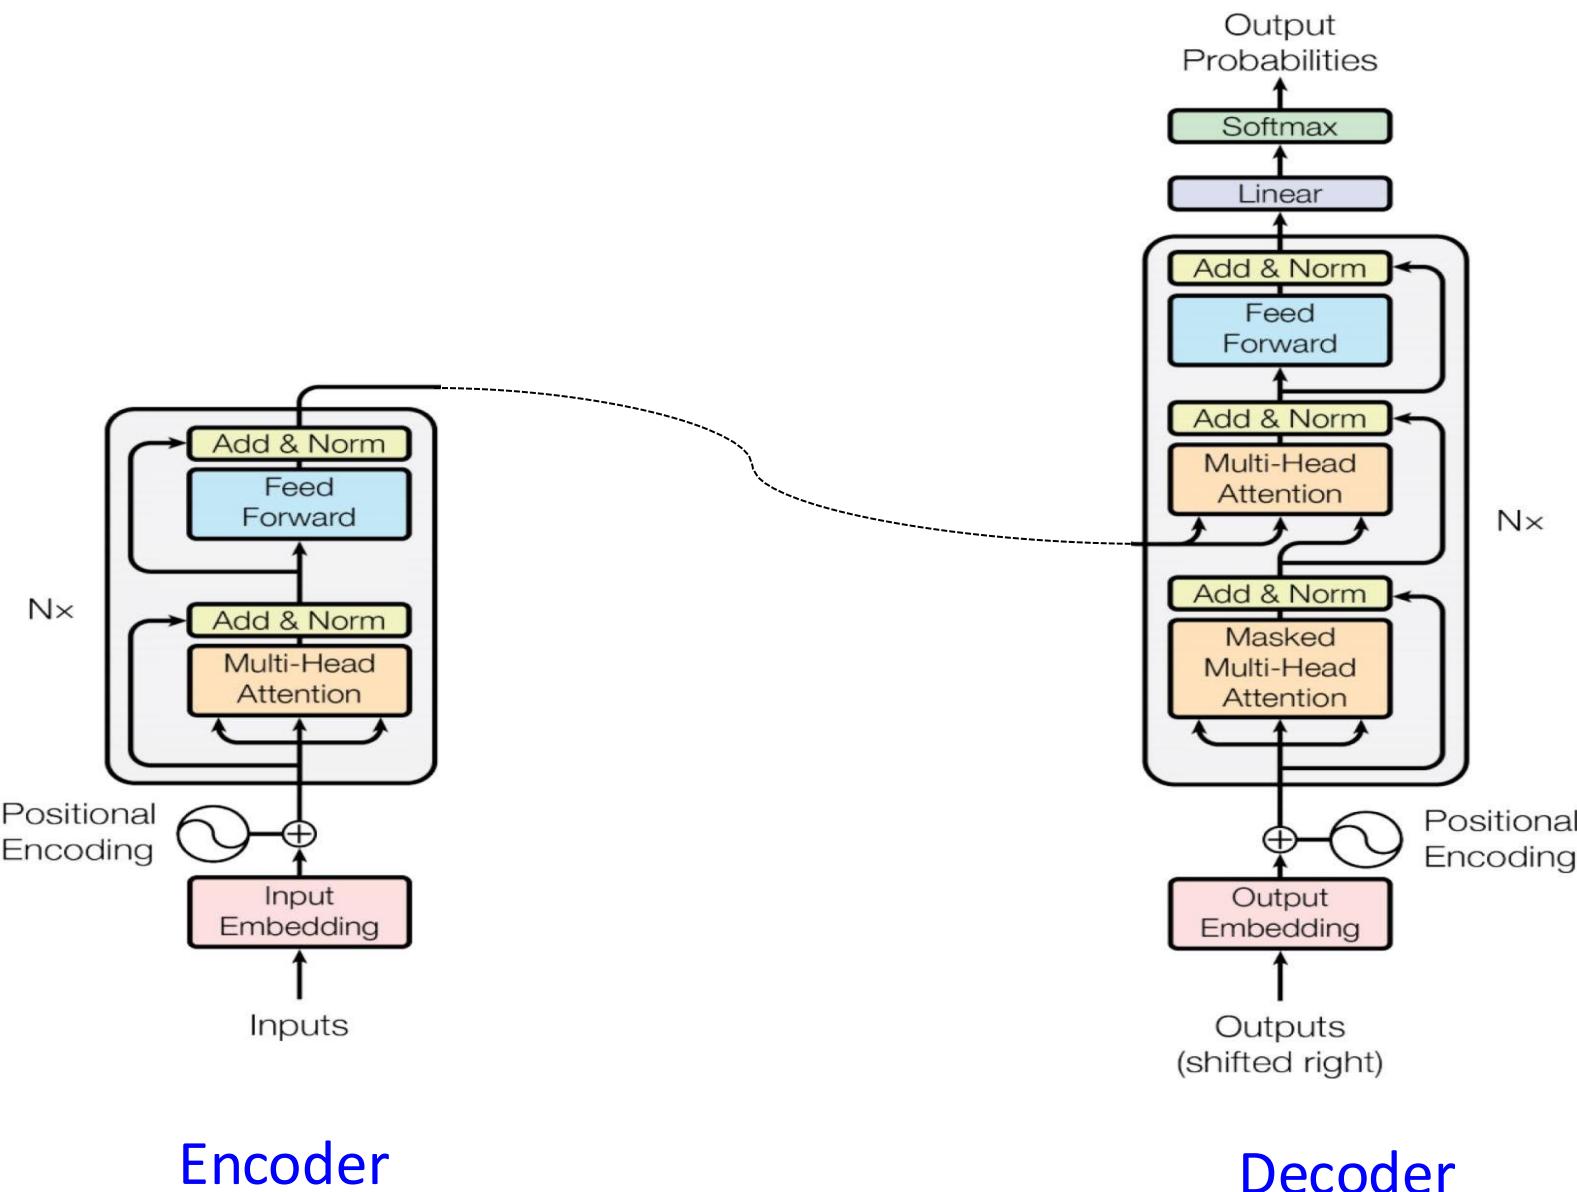


# Vision Transformers

## Attention Is All You Need (2017)





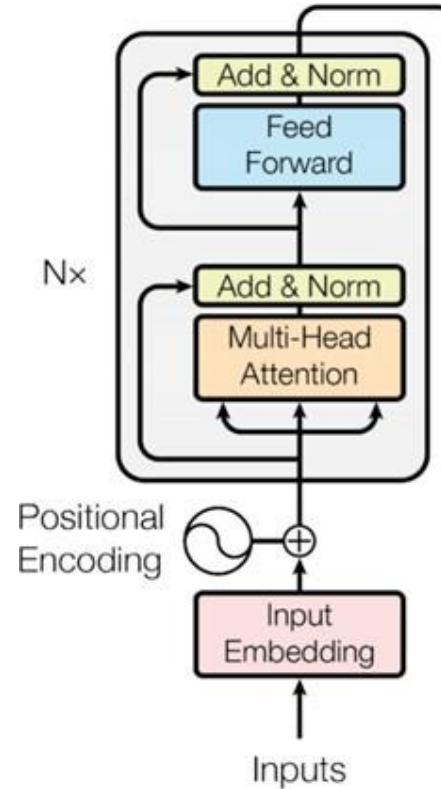


Encoder

Decoder

# Encoder

The encoder part of the transformer embeds the input sequence of  $n$  words  $X \in \mathbb{R}^{d \times n}$  into context vectors with the attention mechanism.



# Encoder

- ▶ The encoder consists of two main components: Self-Attention and Feedforward Neural Network (FFN).

- ▶ **Self-Attention:**

- ▶ Input: Matrix  $X$
- ▶ Linear Transformations to generate Query ( $Q$ ), Key ( $K$ ), and Value ( $V$ ) matrices:

$$Q = W_Q^T X, \quad K = W_K^T X, \quad V = W_V^T X$$

- ▶ Compute attention output  $Z$  using the formula:

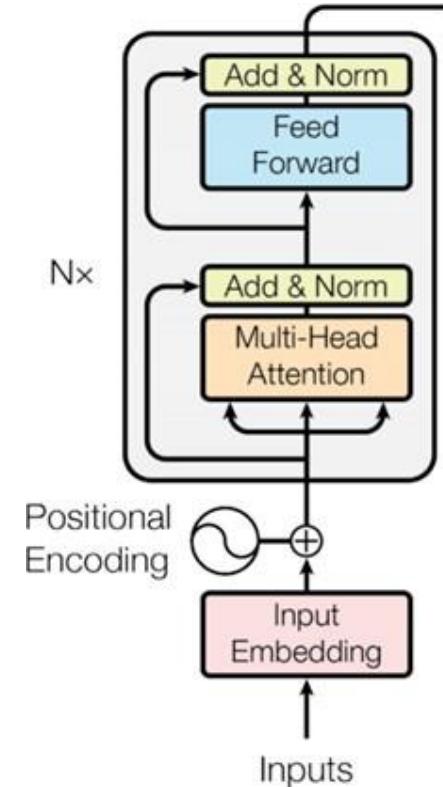
$$Z = V \text{softmax} \left( \frac{Q^T K}{\sqrt{p}} \right)$$

- ▶ Residual Connection:

$$X + Z$$

- ▶ Normalization:

$$(X + Z)$$



# Multi-Headed Attention

$$Q_1 = W_Q^1 X$$

$$K_1 = W_K^1 X$$

$$V_1 = W_V^1 X$$

$$Z_1 = V \text{softmax} \left( \frac{1}{\sqrt{p}} Q_1^T K_1 \right)$$

:

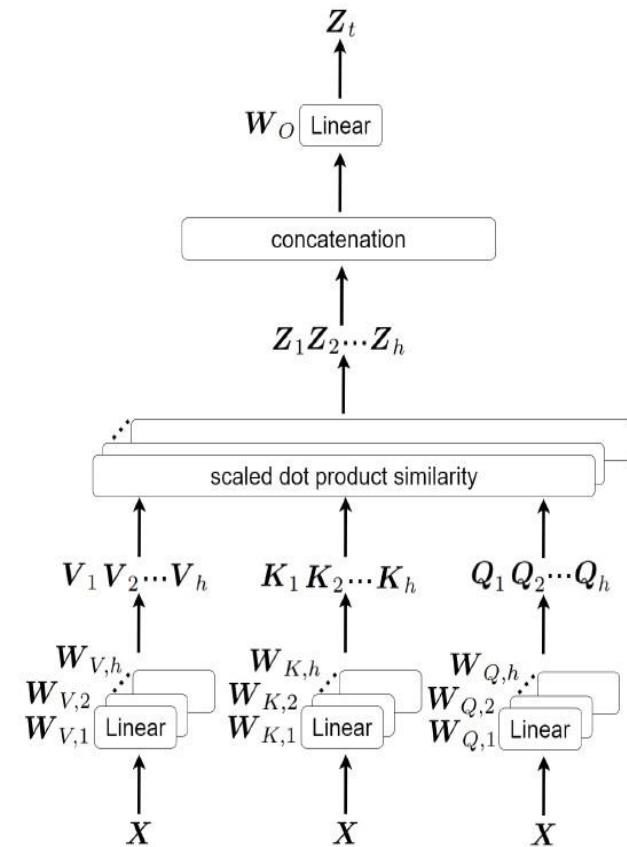
$$Q_h = W_Q^h X$$

$$K_h = W_K^h X$$

$$V_h = W_V^h X$$

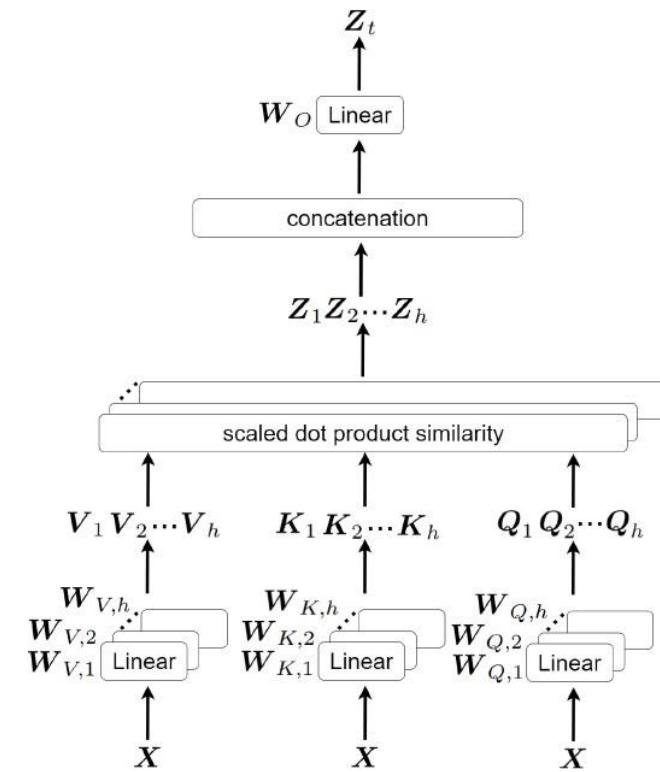
$$Z_h = V \text{softmax} \left( \frac{1}{\sqrt{p}} Q_h^T K_h \right)$$

Multi-Head Attention



# Multi-Headed Attention

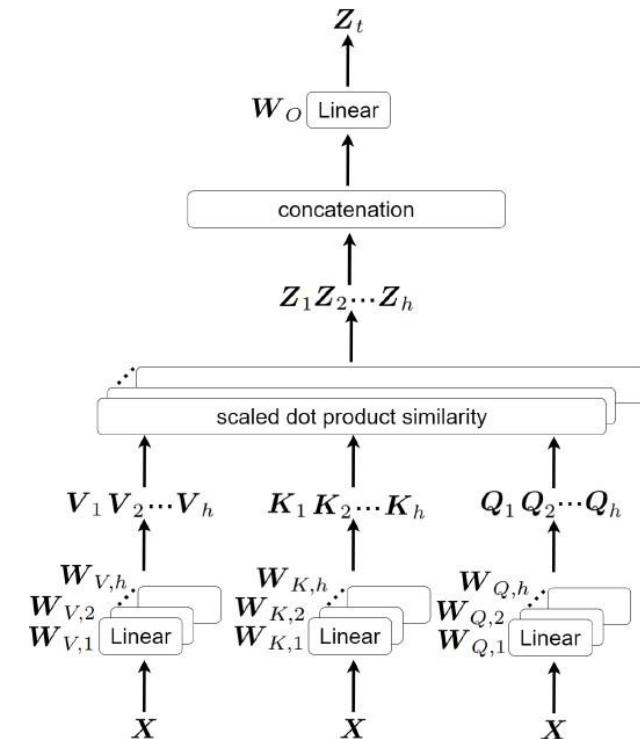
$$\begin{pmatrix} Z_1 \\ Z_2 \\ \vdots \\ Z_h \end{pmatrix} = \text{Concat}(\text{head}_1, \dots, \text{head}_h)$$



# Multi-Headed Attention

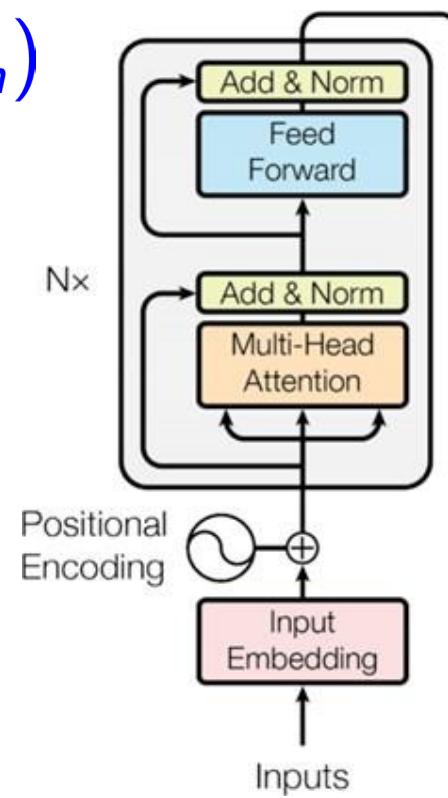
$$\begin{pmatrix} Z_1 \\ Z_2 \\ \vdots \\ Z_h \end{pmatrix} = \text{Concat}(\text{head}_1, \dots, \text{head}_h)$$

$$Z = \text{MultiHead}(Q, K, V) = W_0^T \begin{pmatrix} Z_1 \\ Z_2 \\ \vdots \\ Z_h \end{pmatrix}$$



# Multi-Headed Attention

$$\text{MultiHead}(Q, K, V) = W_O^T \text{Concat}(\text{head}_1, \dots, \text{head}_h)$$



# Structure of the Feed Forward Network

- ▶ Linear Layer 1
- ▶ ReLU Activation
- ▶ Linear Layer 2

$$FFN(x) = W_2^T \max(0, W_1^T X + b_1) + b_2$$

Two linear transformations with ReLU activation in between.

# Application of FFN to Each Position

- ▶ The Feed Forward Network (FFN) is applied independently to each position in the input sequence.
- ▶ Despite individual processing, all positions share the same set of weights and biases in the FFN.
- ▶ Key Points:
  - ▶ Shared parameters ensure consistency in processing across all positions.
  - ▶ Enables the model to generalize learnings from one position to all positions.
  - ▶ Facilitates parallel processing of the sequence, enhancing computational efficiency.

# Global vs Local

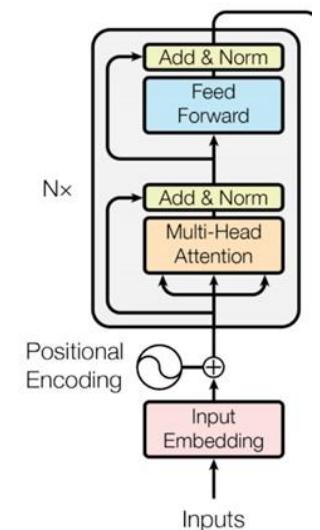
- Attention Mechanism:
- **Global Understanding:** Captures relationships among different positions in the sequence.
- **Context Aggregation:** Spreads relevant information across the sequence, enabling each position to see a broader context.

# Global vs Local

- **Attention Mechanism:**
- **Global Understanding:** Captures relationships among different positions in the sequence.
- **Context Aggregation:** Spreads relevant information across the sequence, enabling each position to see a broader context.
- **Feed-Forward Networks (FFN):**
- **Local Processing:** While attention looks across the entire sequence, FFN zooms back in to process each position independently.
- **Individual Refinement:** Enhances the representation of each position based on its own value, refining the information gathered so far.

# Encoder

If the output of the FFN is denoted by  $R$ , then a residual connection is established from the output of the previous layer  $(X + Z)$  to the output of the FFN, resulting in  $(X + Z) + R$ . This will be normalized  $((X + Z) + R)$  to form the output of the encoder.



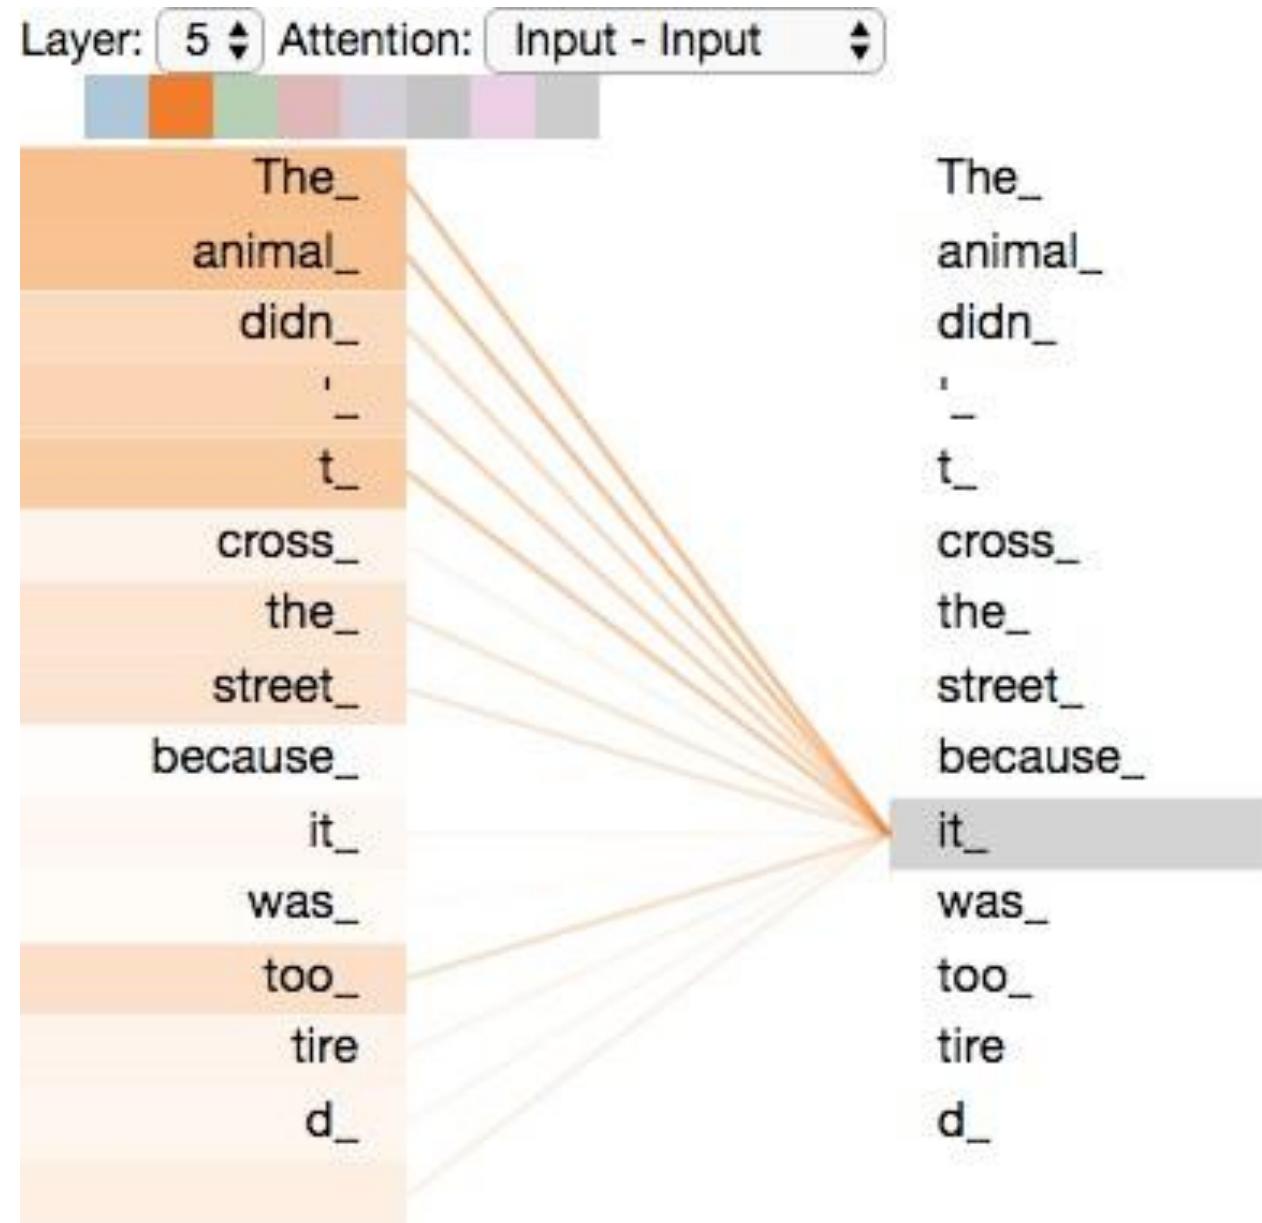
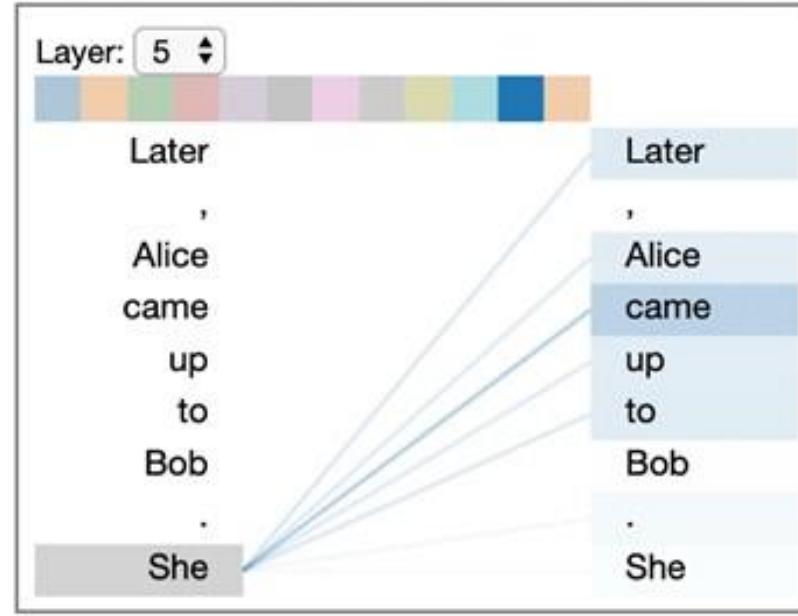
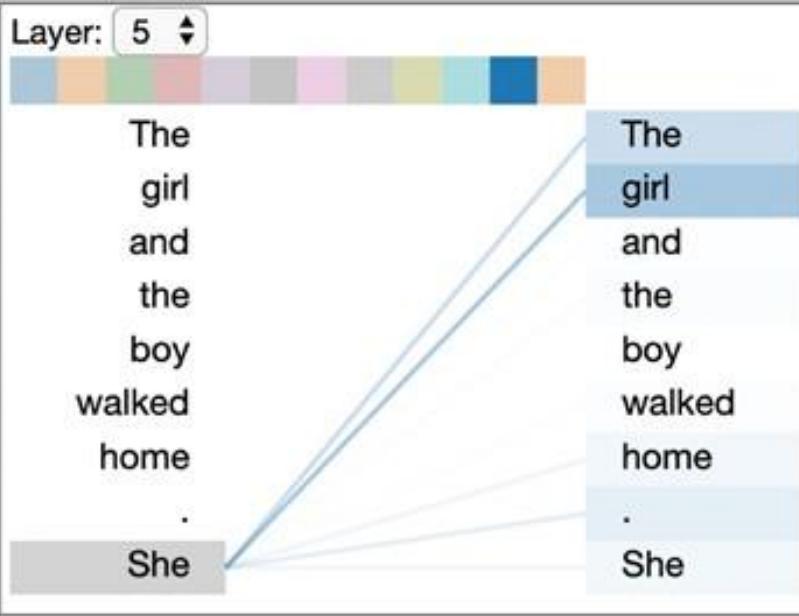


Figure: Jay Alammarz

**She**



**He**

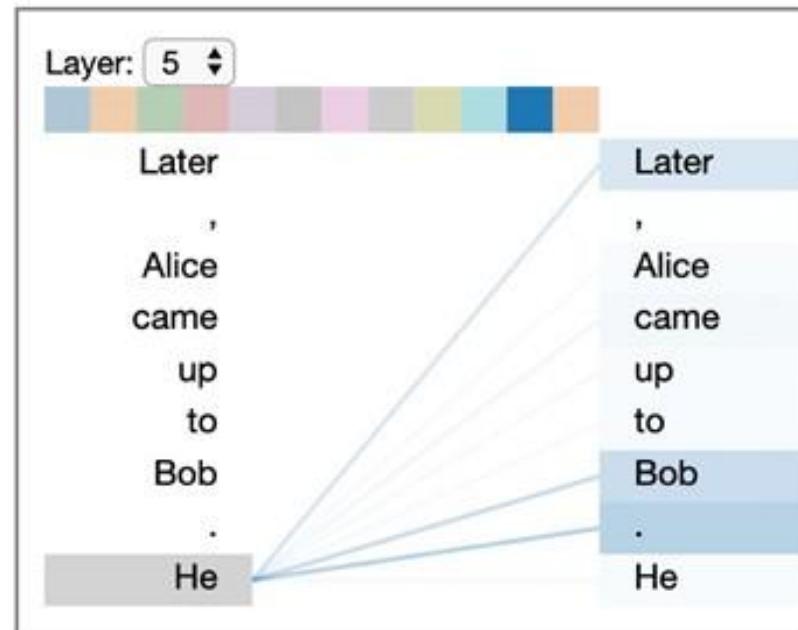
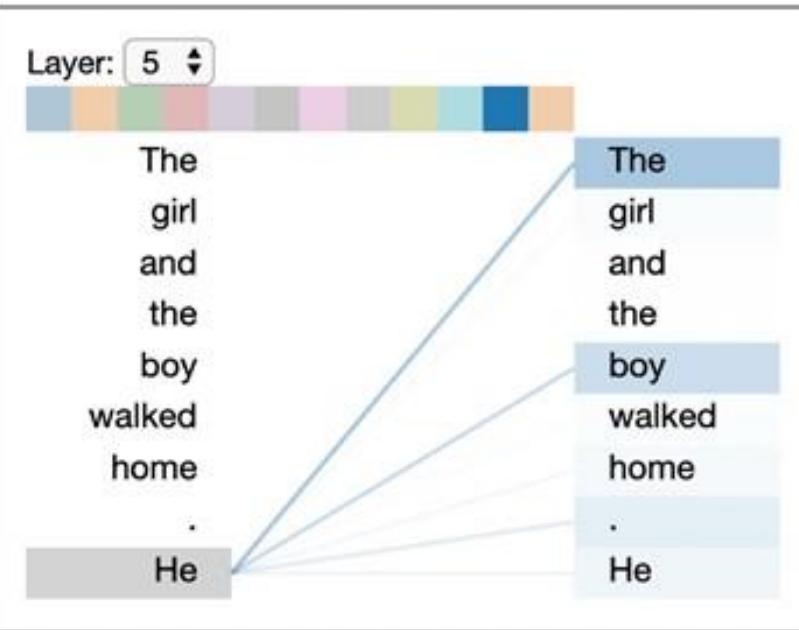
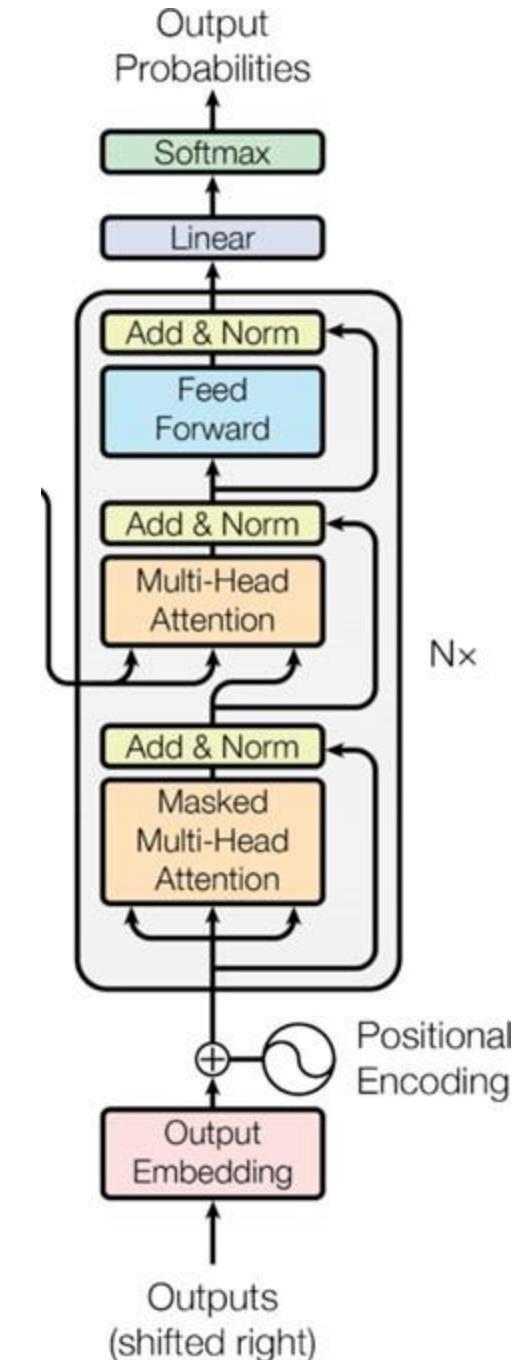


Figure: Jesse Vig

# Decoder

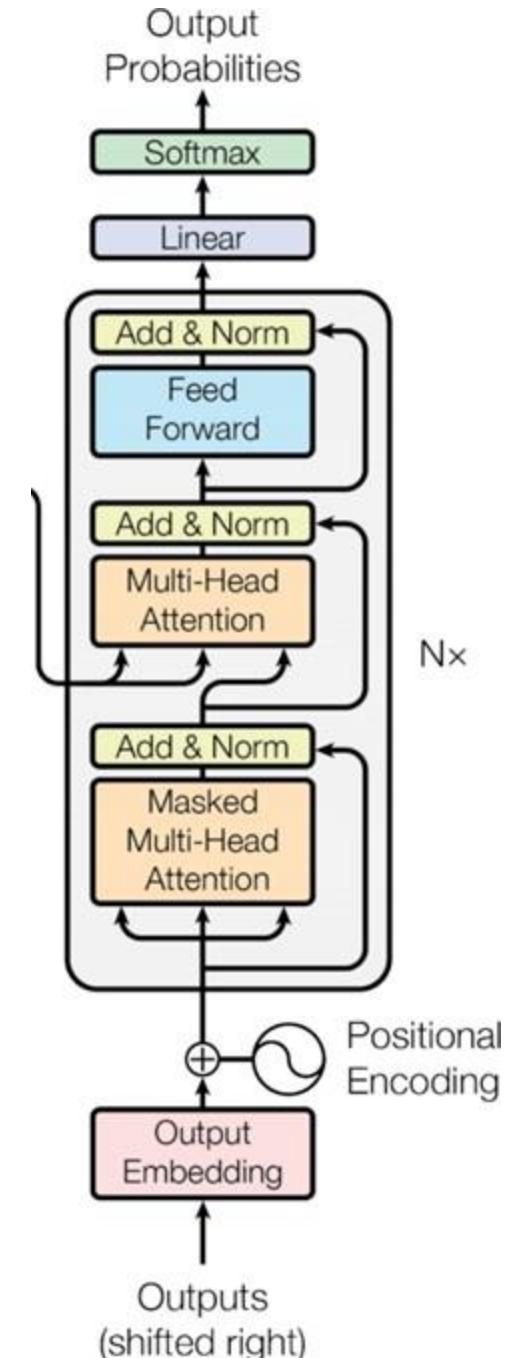
- **Masked self attention.**



# Decoder

- **Masked self attention.**

$$Z := \text{Attention}(Q, K, V) = V \text{softmax} \left( \frac{1}{\sqrt{P}} (Q^\top K) \right),$$



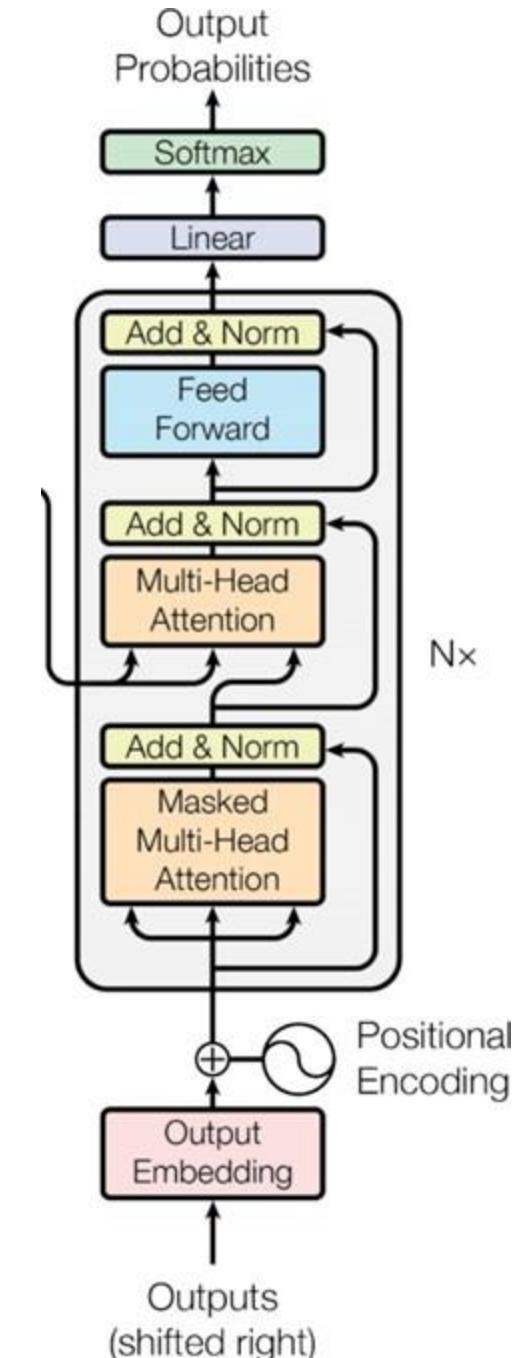
# Decoder

- **Masked self attention.**

$$Z := \text{maskedAttention}(Q, K, V) = V \text{softmax} \left( \frac{1}{\sqrt{p}} (Q^T K + M) \right),$$

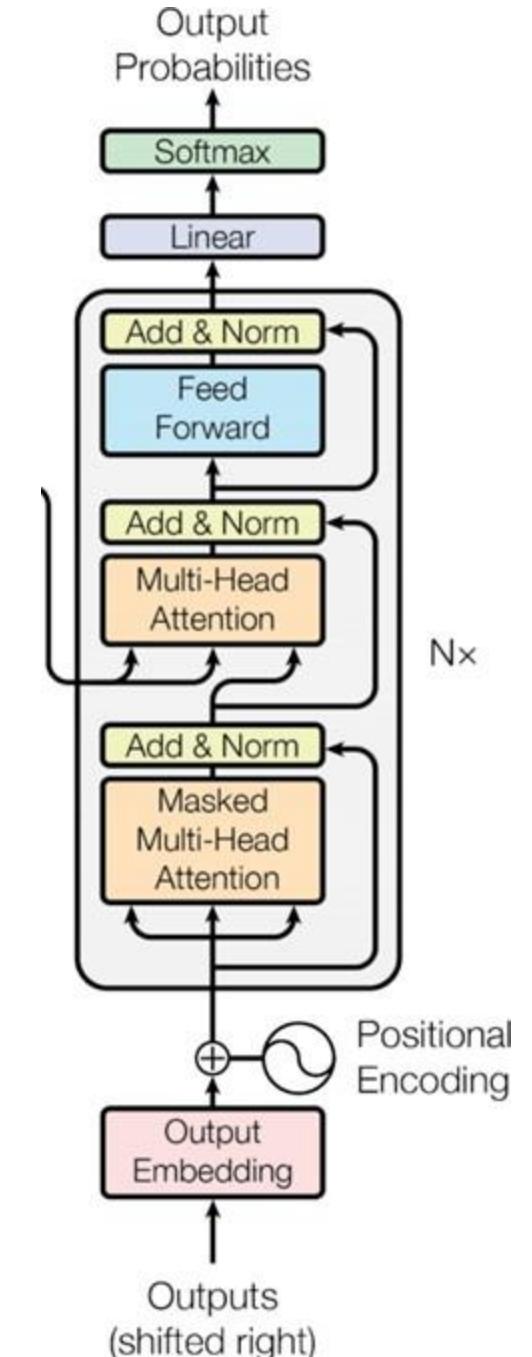
where the mask matrix  $M \in \mathbb{R}^{n \times n}$  is:

$$M(i, j) := \begin{cases} 0 & \text{if } j \leq i, \\ -\infty & \text{if } j > i. \end{cases}$$



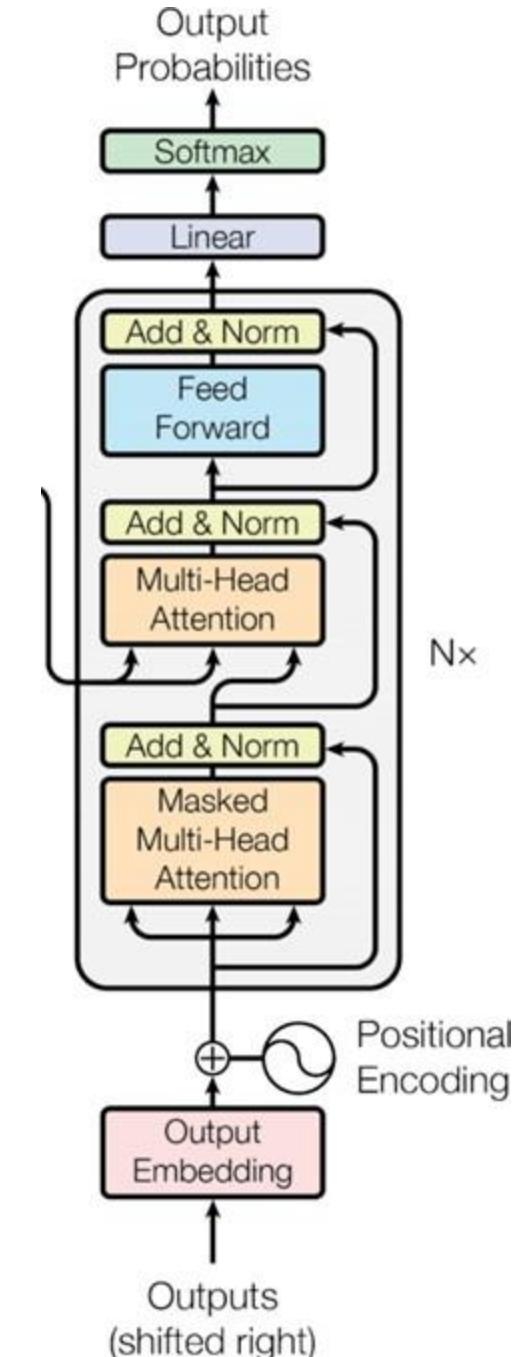
# Decoder

- **Cross Attention**
- Cross attention allows each position in one sequence to attend over all positions in another sequence.
- **Query (Q)**: Originates from a position in the first sequence, i.e. the output of a previous layer in the decoder.
- **Memory Keys (K) and Values (V)**: Both come from all positions in the second sequence, i.e. the output of the encoder.



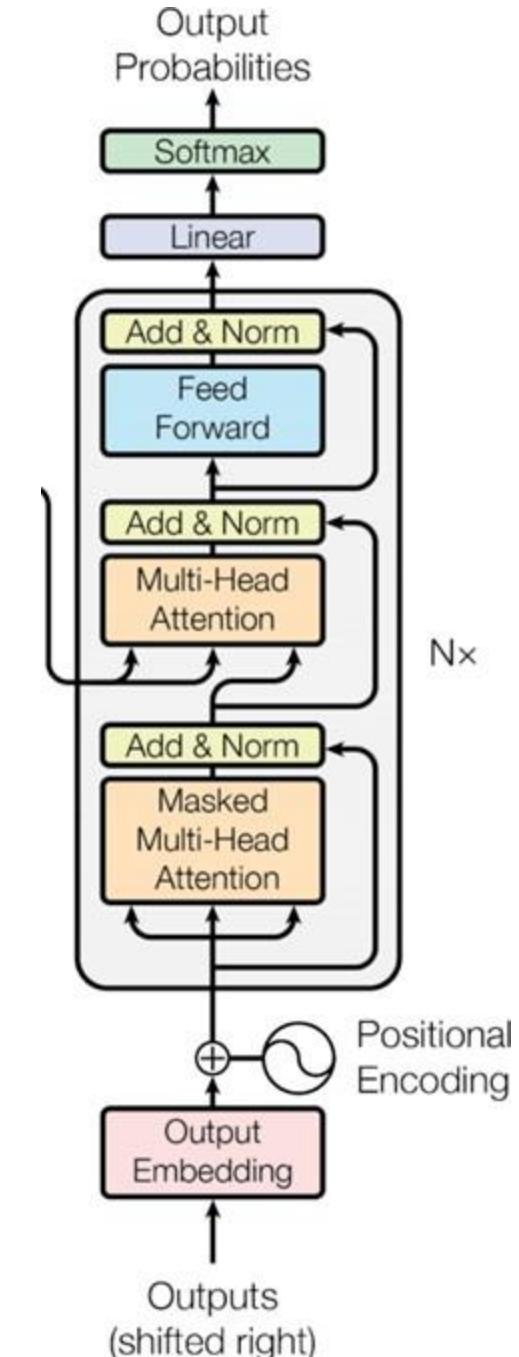
# Decoder

- Masked self attention.
- Cross attention layer is like what attention does in sequence-to-sequence models.



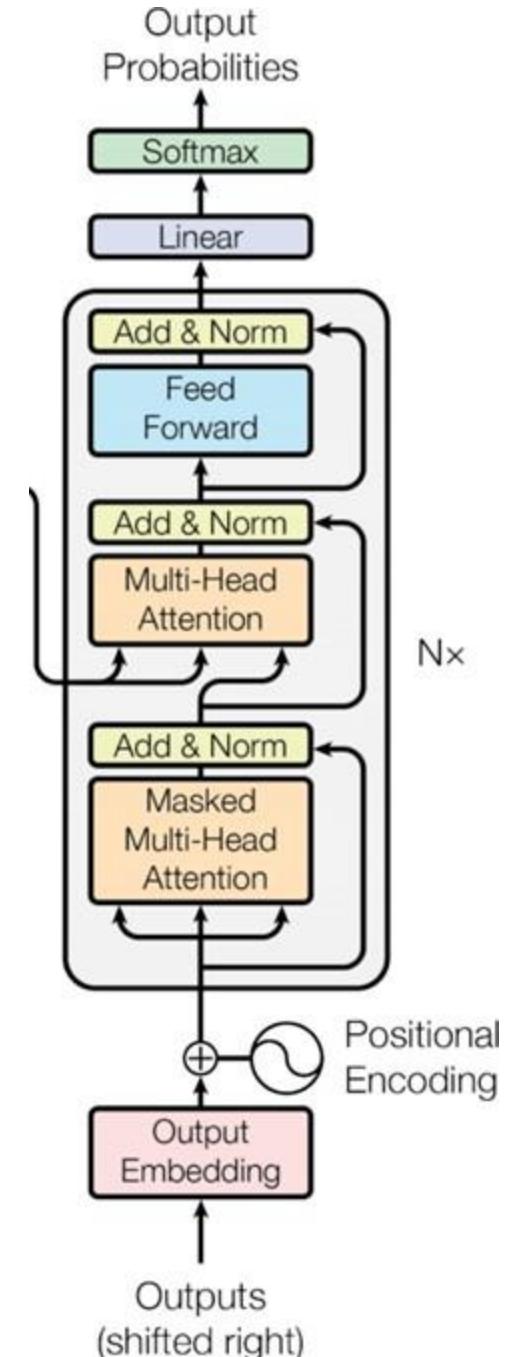
# Decoder

- Masked self attention.
- Cross attention attention layer is like what attention does in sequence-to-sequence models.
- It helps the decoder emphasize on relevant parts of the input.



# Decoder

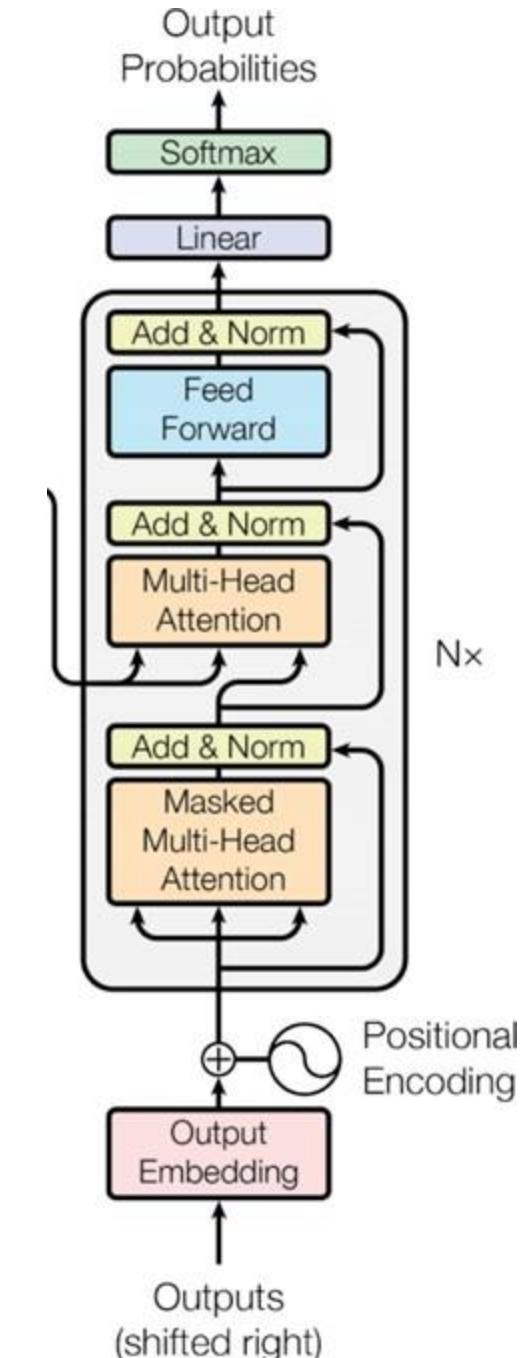
- Masked self attention.
- Cross attention attention layer is like what attention does in sequence-to-sequence models.
- It helps the decoder emphasize on relevant parts of the input.
- The same feed-forward network is applied to each position.



## From Feedforward Network to Word Prediction

### Linear Projection:

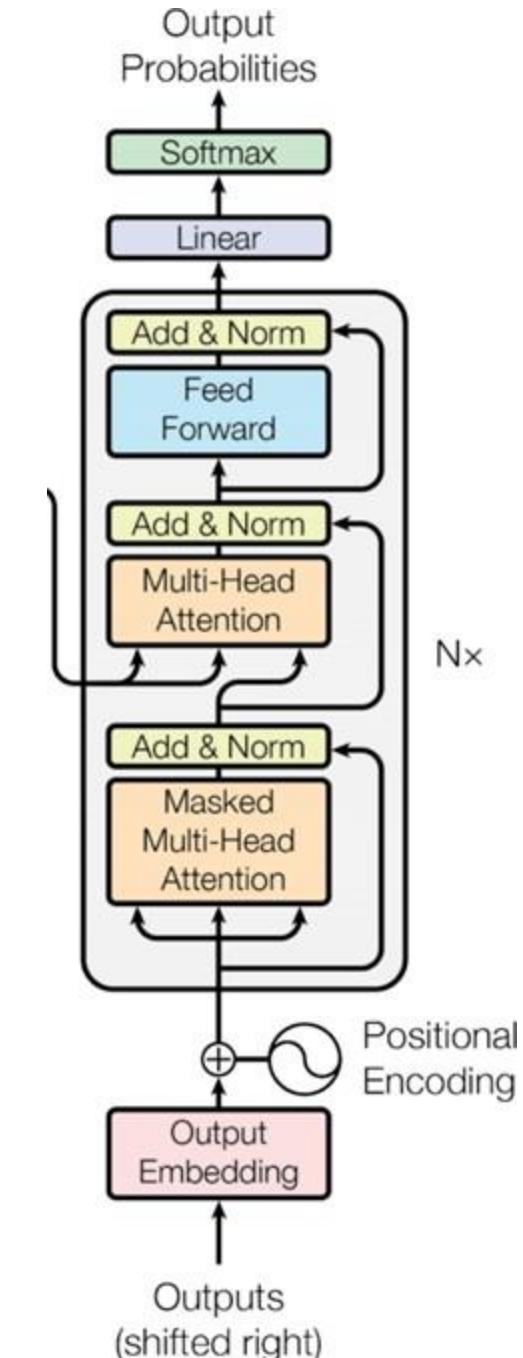
- Primary Role: Adjusting dimensionality.
- The linear layer serves to change the dimensionality of the feedforward network's output to match the size of the vocabulary.
- This ensures that the output has a dimension corresponding to every word in the dictionary.



## From Feedforward Network to Word Prediction

### Softmax Activation:

- This function transforms the linear layer's output into probabilities.
- Representing the likelihood of a respective word being the next word in the sequence.



# Positional Encoding

- **Problem:** no recurrence and no convolution, the model has no sense of the sequence.

# Positional Encoding

- **Problem:** no recurrence and no convolution, the model has no sense of the sequence.
- We need a way to account for the order of the tokens in the sequence.

# Positional Encoding

- **Problem:** no recurrence and no convolution, the model has no sense of the sequence.
- We need a way to account for the order of the tokens in the sequence.
- **Solution:** Adds a vector accounting for the position to each input embedding.

# Positional Encoding

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

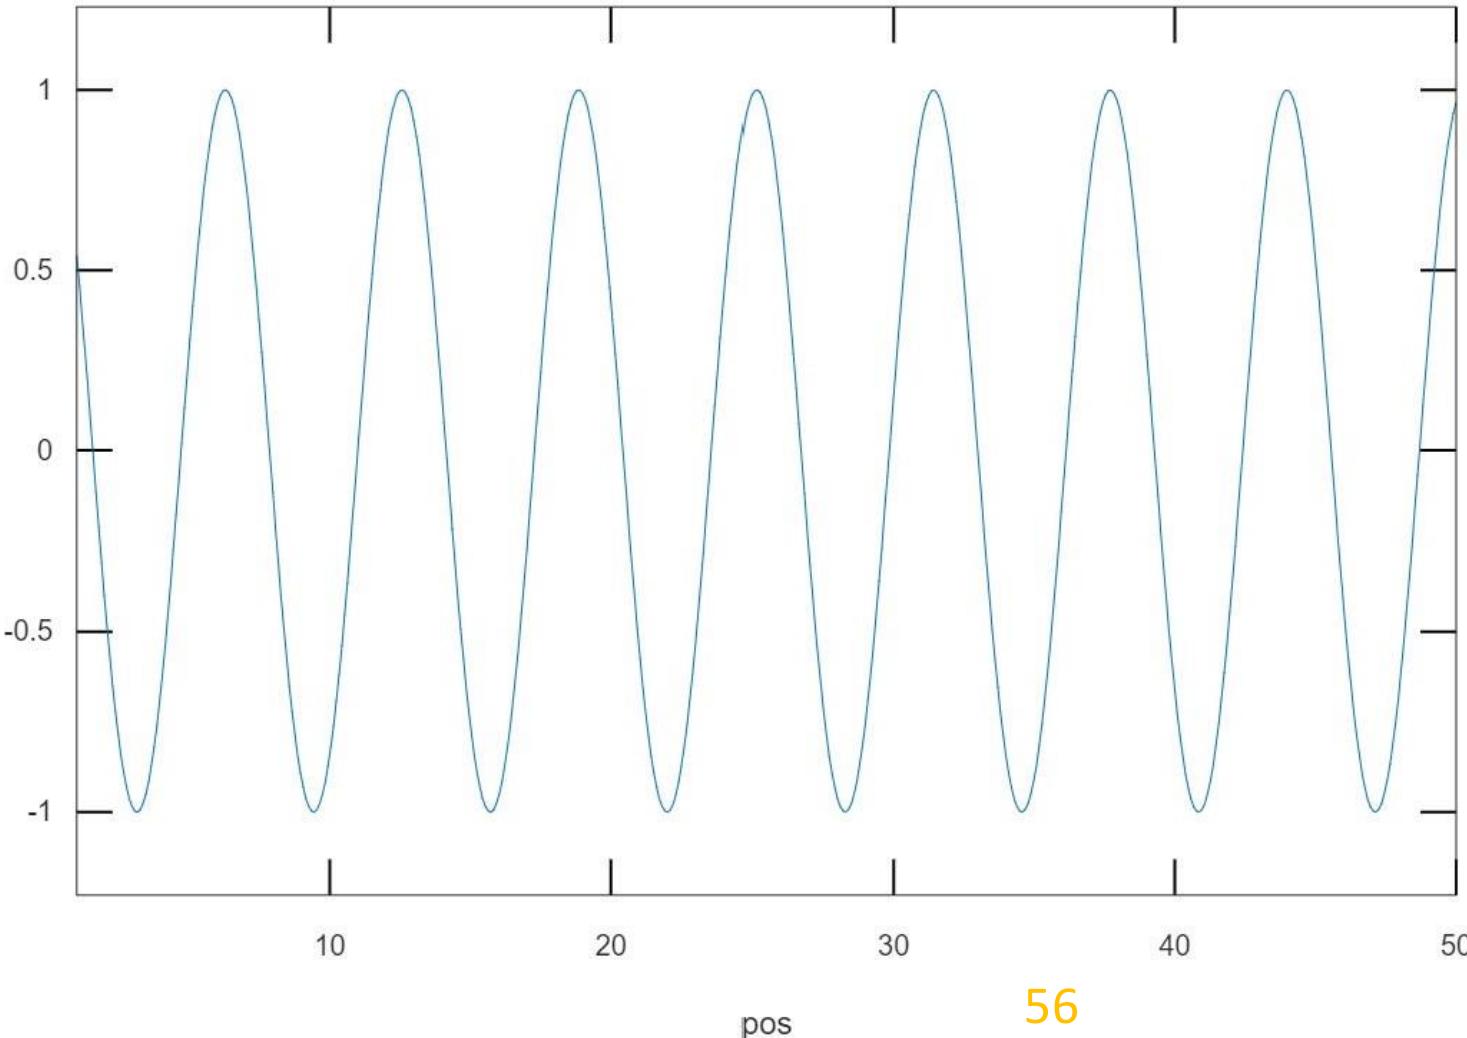
$p$  is the dimension to which we project

$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

# Positional Encoding

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

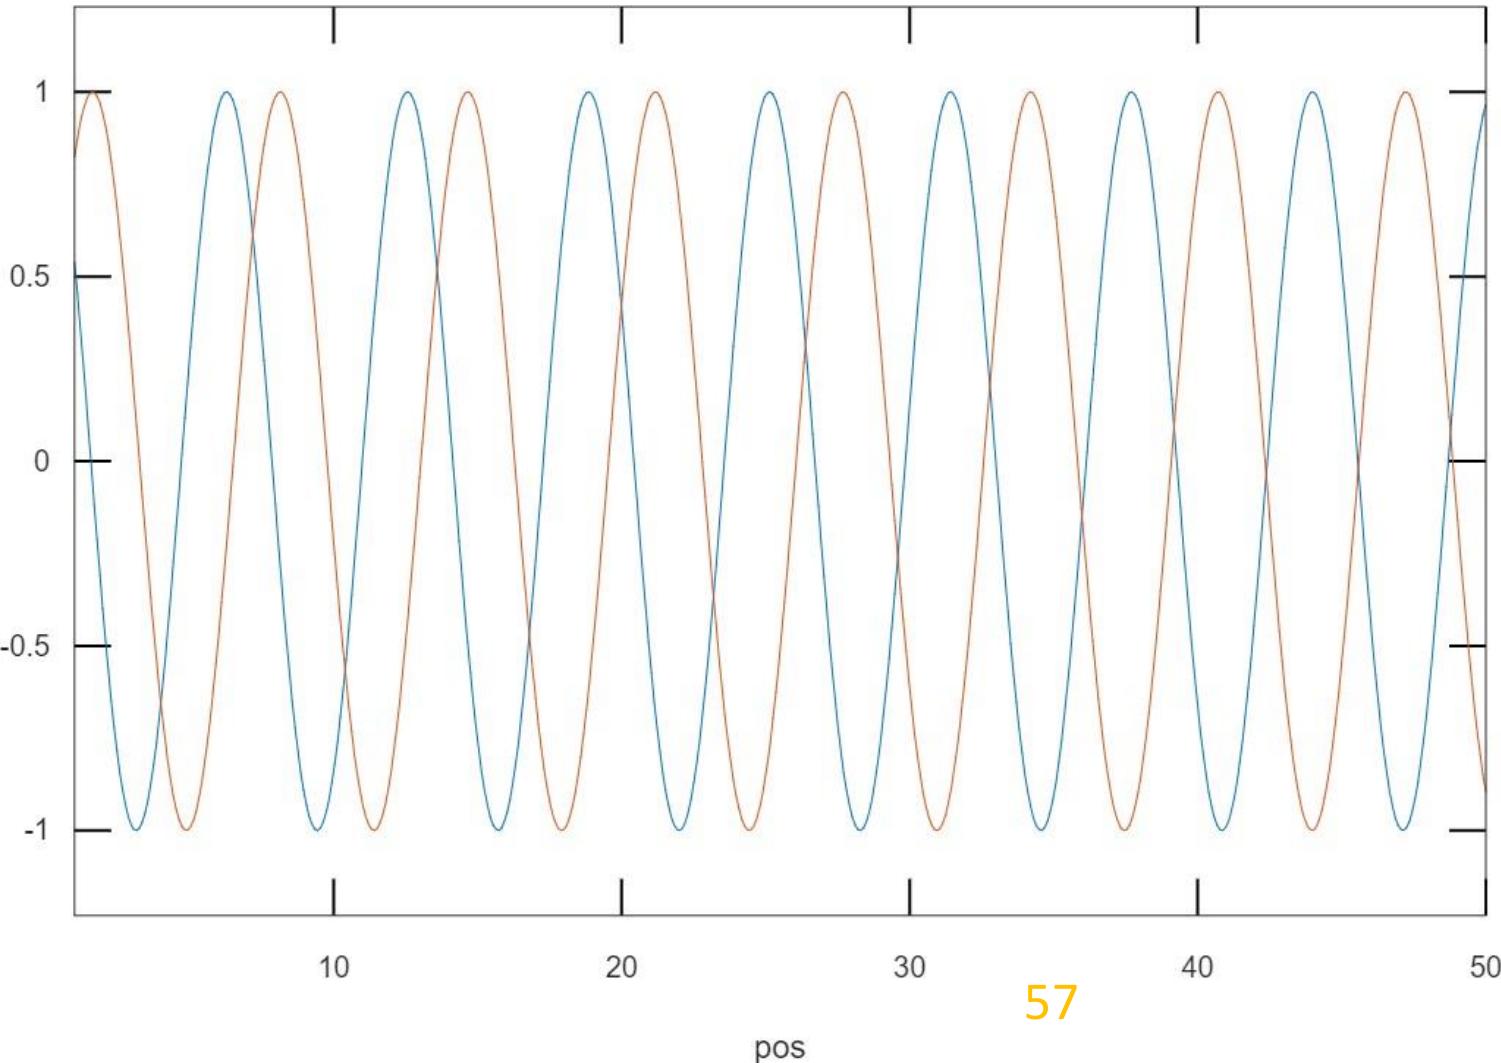
# Positional Encoding



$i = 0$

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

# Positional Encoding



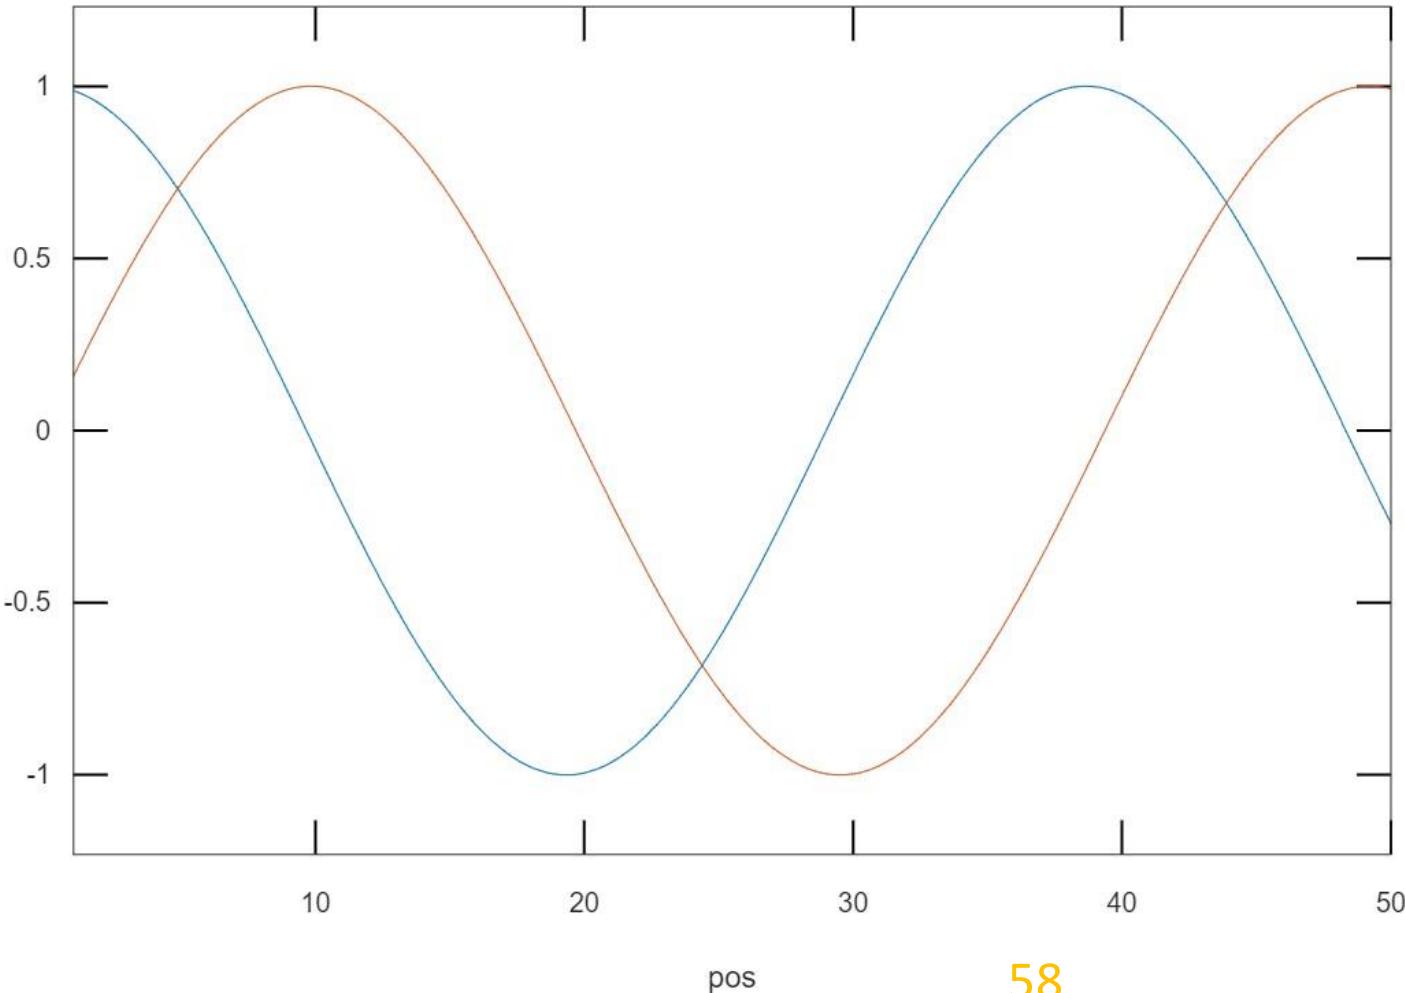
$i = 0$

$$PE(pos, 2i) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

$i = 1$

$$PE(pos, 2i+1) = \sin\left(\frac{pos}{10000^{\frac{2i+1}{p}}}\right)$$

# Positional Encoding

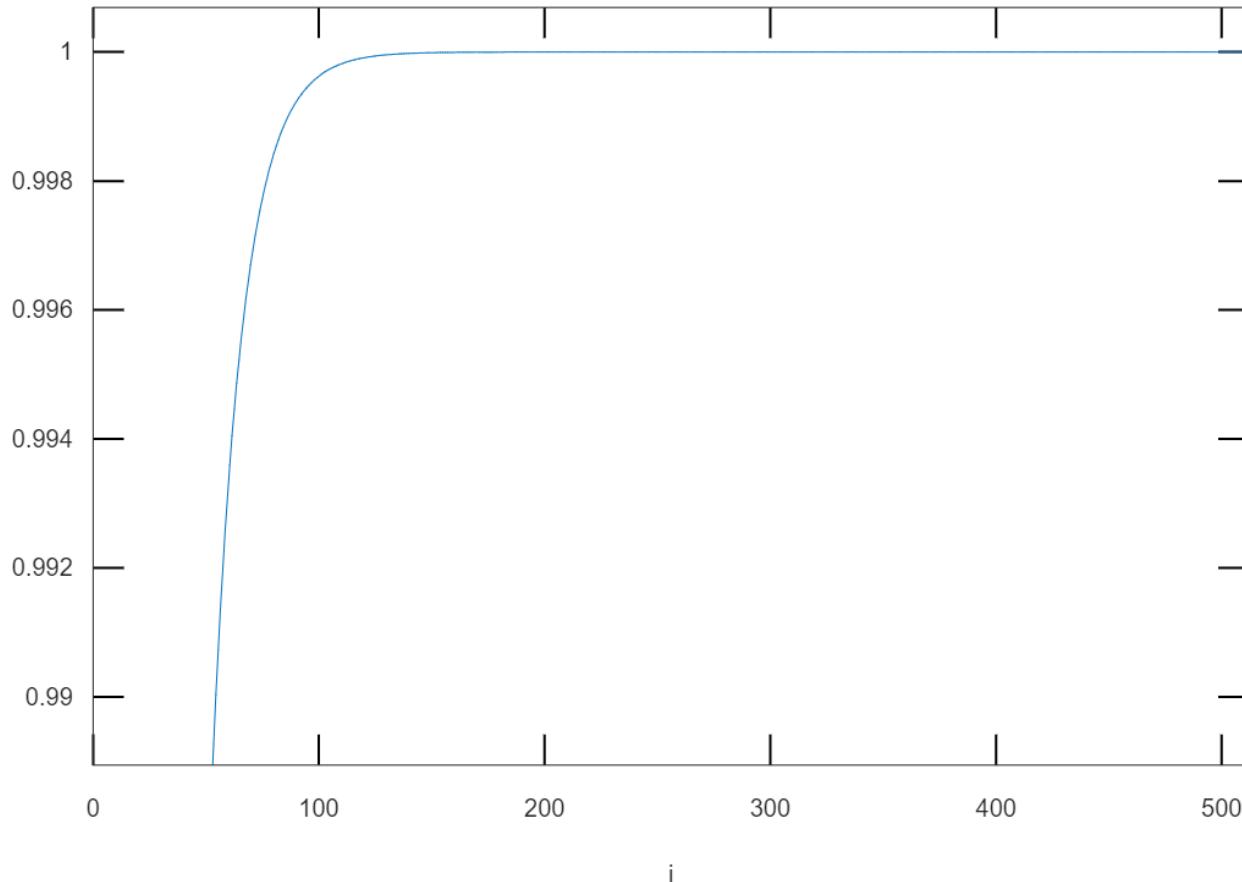


$i = 50$

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

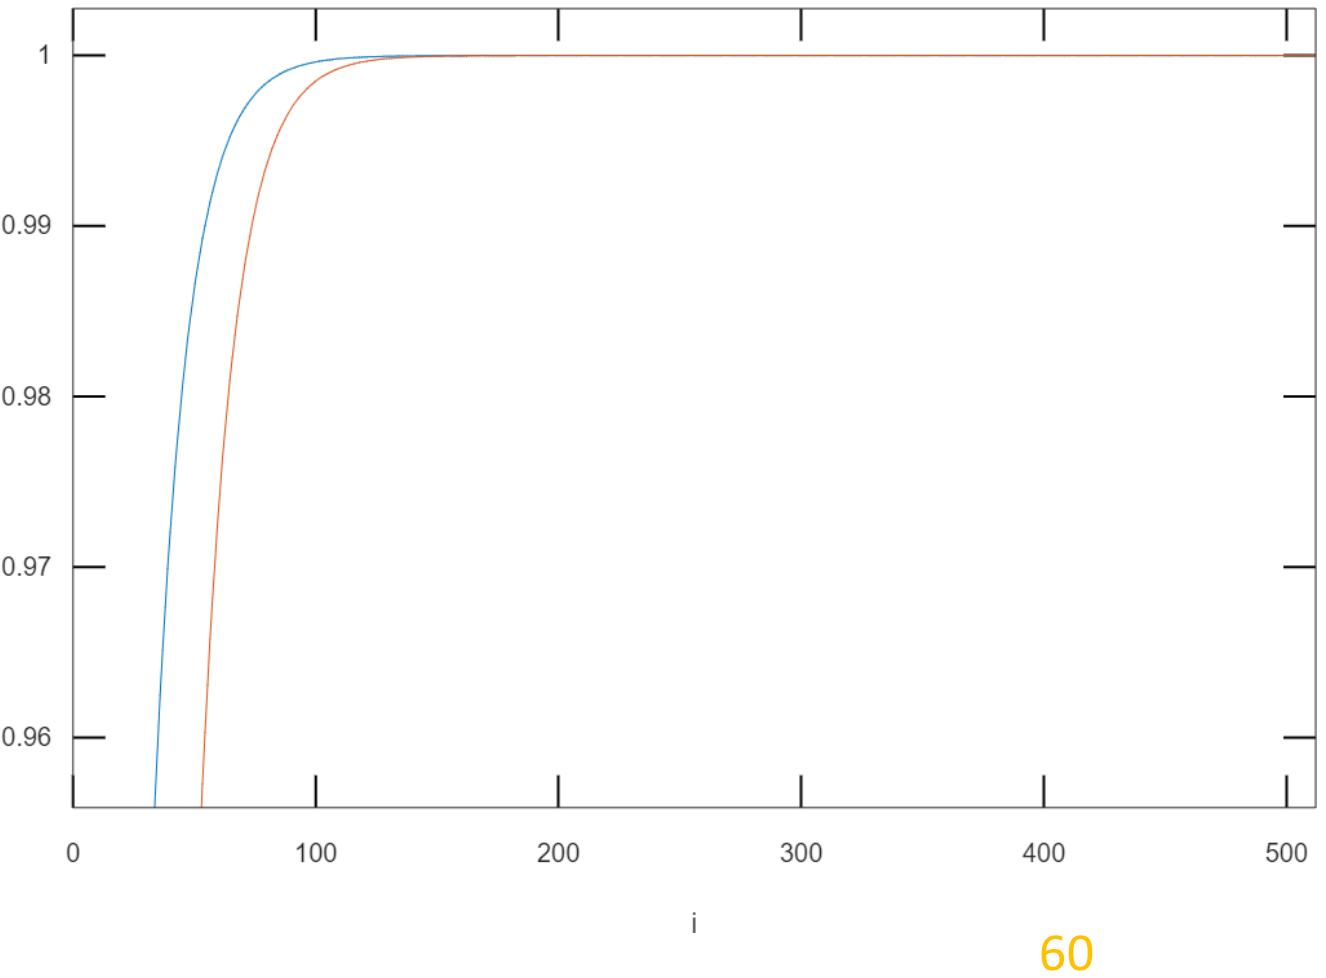
$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$

# Positional Encoding



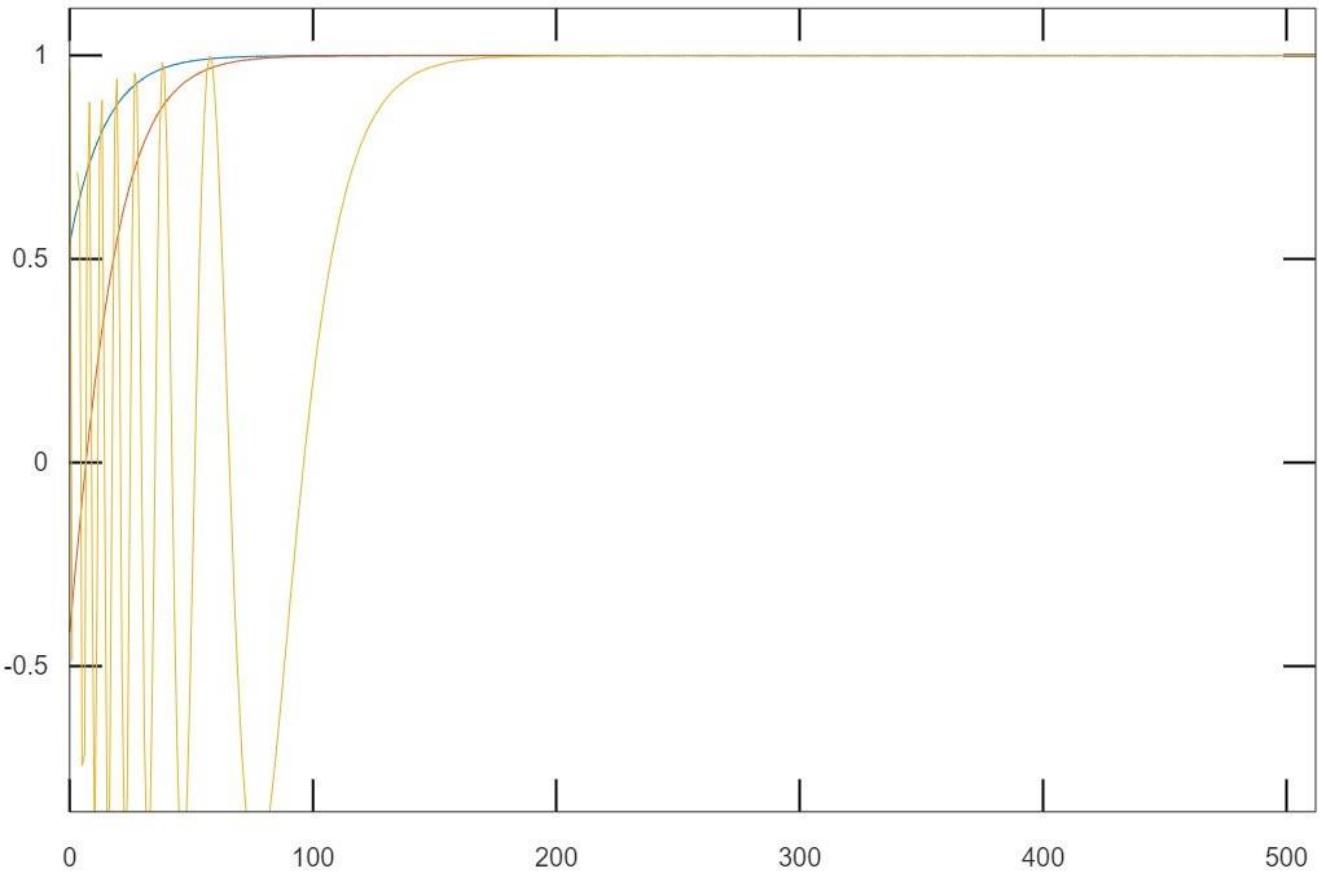
$pos = 1$

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000} \frac{2i}{p}\right)$$



$pos = 2$

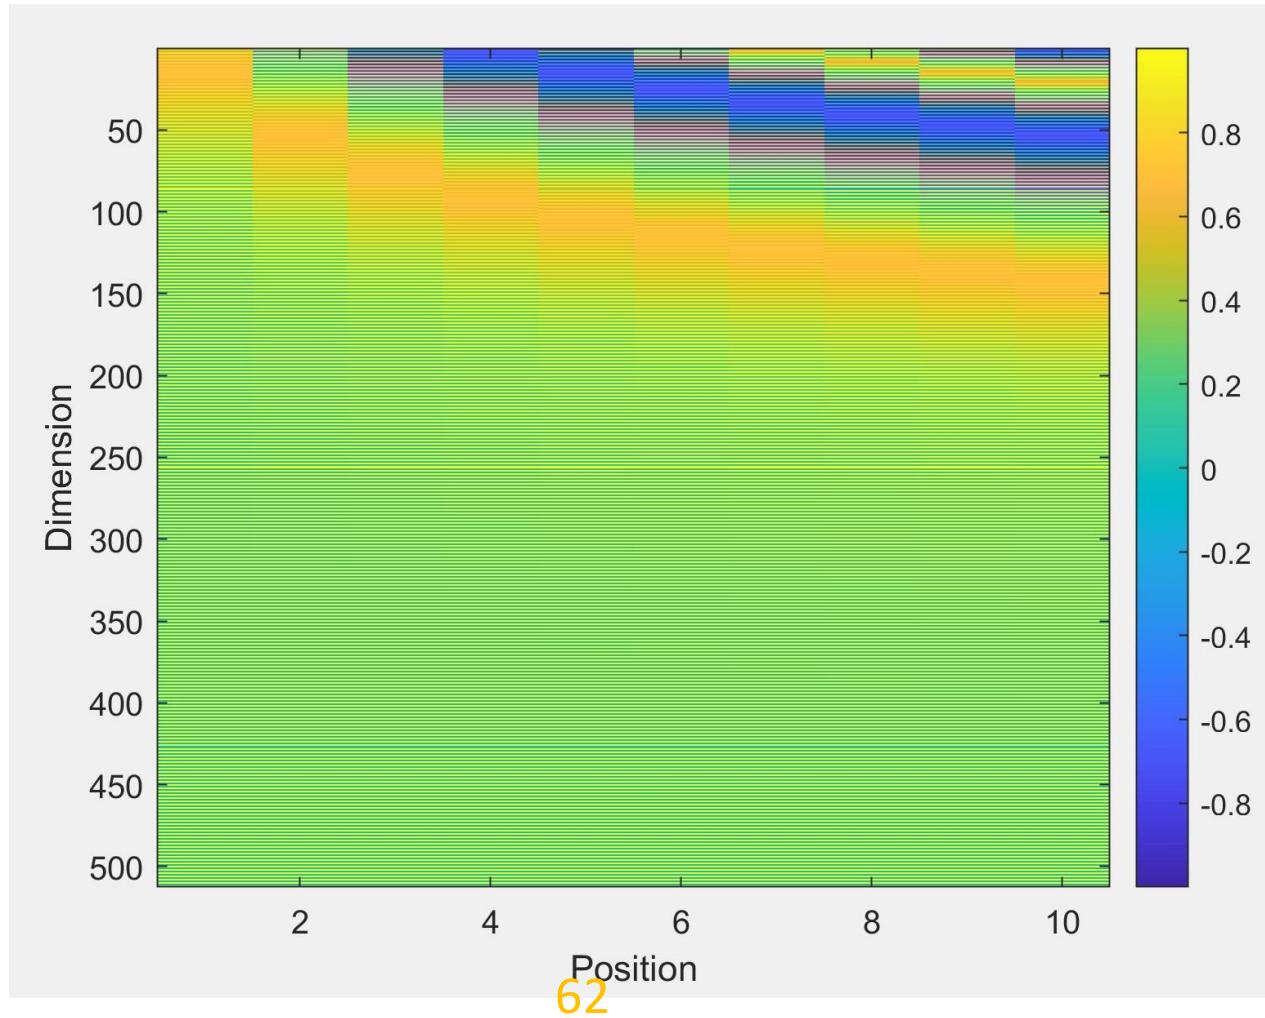
$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000^{\frac{2i}{p}}}\right)$$



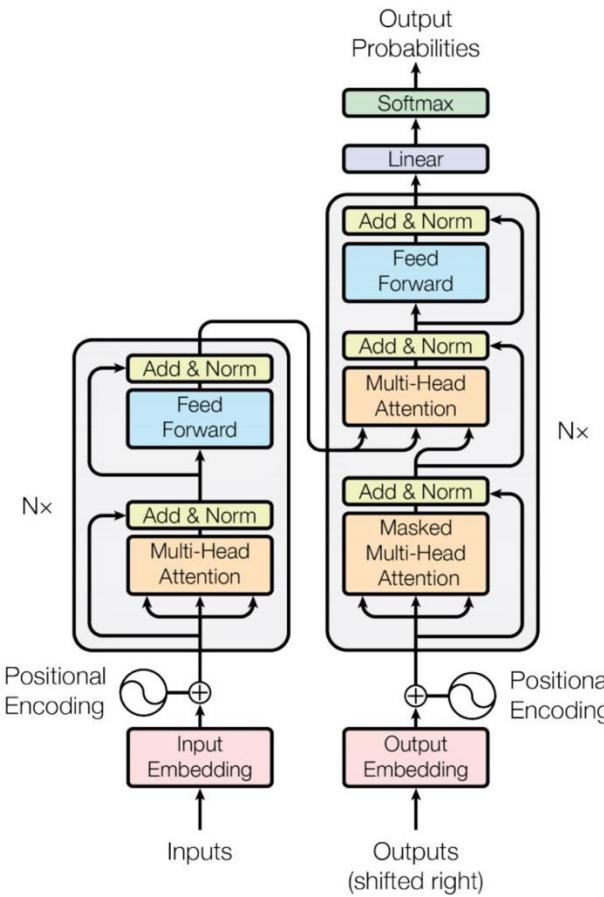
$pos = 50$

$$PE(pos, 2i+1) = \cos\left(\frac{pos}{10000} \cdot \frac{2i}{p}\right)$$

# Positional Encoding Visualization



BERT,  
GPT



# Generative Pre-trained Transformer (GPT)

- Improving Language Understanding by Generative Pre-Training (2018)

Alec Radford

Karthik Narasimhan

Tim Salimans

Ilya Sutskever

# Bidirectional Encoder Representations from Transformers (BERT)

- BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding (2018)

Jacob Devlin

Ming-Wei Chang

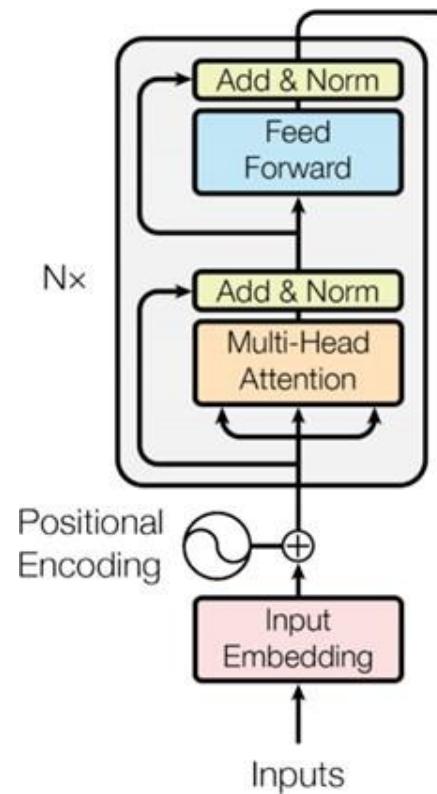
Kenton Lee

Kristina Toutanova

# BERT and GPT

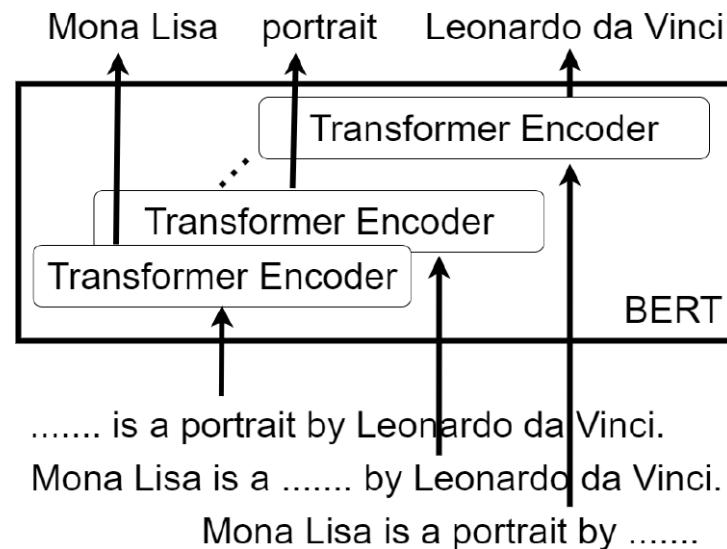
- The GPT is built using transformer decoder blocks.
- BERT is built using transformer encoder blocks.

# BERT



# Masked language model

- Masks words in the input and asks the model to predict the missing word.



# BERT: Bidirectional language model

- Masks words in the input and asks the model to predict the missing word.
- Additional task: Given two sentences (A and B), is B likely to be the sentence that follows A, or not?

# BERT

- BERT is designed to pretrain bidirectional representations from unlabeled text.

# BERT

- BERT is designed to pretrain bidirectional representations from unlabeled text.
- Jointly conditioning on both left and right context.

# BERT

- BERT is designed to pretrain bidirectional representations from unlabeled text.
- Jointly conditioning on both left and right context.
- The pre-trained BERT model can be finetuned with just one additional output layer.

# BERT

- BERT is designed to pretrain bidirectional representations from unlabeled text.
- Jointly conditioning on both left and right context.
- The pre-trained BERT model can be finetuned with just one additional output layer.
- It creates state-of-the-art models for a wide range of tasks, such as question answering and language inference.

# [CLS] Token in BERT

- The [CLS] token is prepended to the input text and travels through the Transformer layers alongside other tokens.
- All tokens, including [CLS], gather contextual information from the entire sequence due to the self-attention mechanism.
- For sentence-level tasks, the final hidden state of the [CLS] token is used as the sentence representation.
- During fine-tuning on a specific task, the model learns to imbue the [CLS] token with a meaningful representation of the entire sentence, optimized for that task.
- Example Usage: In classification tasks, the [CLS] token representation is fed into a classifier to determine the sentence's class.

# BERT

BERT is basically a trained Transformer Encoder stack

## **1. Transformer:**

1. Encoder Layers: 6
2. FFNN Hidden Layer Units: 512
3. Attention Heads: 8

## **2. BERT Base:**

1. Encoder Layers: 12
2. FFNN Hidden Layer Units: 768
3. Attention Heads: 12
4. Total Parameters: 110 million

## **3. BERT Large:**

1. Encoder Layers: 24
2. FFNN Hidden Layer Units: 1024
3. Attention Heads: 16
4. Total Parameters: 340 million

# BERT

- **RoBERTa:**
  - Optimizes BERT's training process by using more data, larger batch sizes, and longer training times, resulting in improved performance on NLP tasks.
- **TinyBERT:**
  - A smaller and faster version of BERT designed for resource-constrained environments, retaining competitive performance with significantly fewer parameters.

# BERT

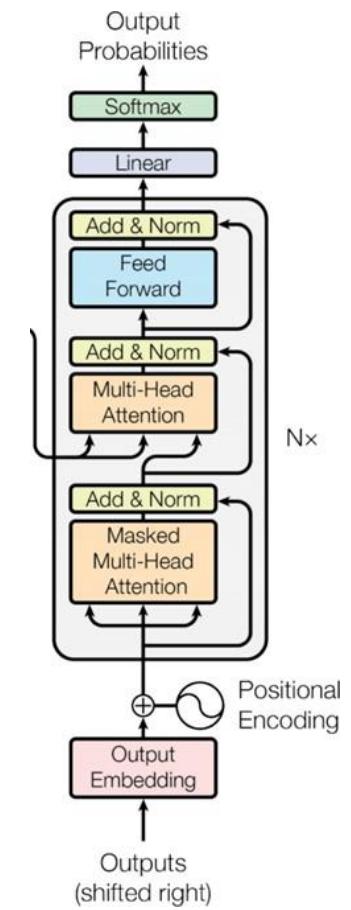
## 1. Multilingual BERT:

1. Trained on 104 different languages, capable of "zero-shot" adaptation to a new language domain.

## 2. Domain Specific BERT Variants:

1. BioBERT: Retrained on a biomedical corpus.
2. SciBERT: Trained on over one million published articles.
3. BERTweet: A RoBERTa model trained on 850 million tweets.
4. FinBERT: Adapted to the financial domain.

# GPT



# GP T

Predict the next word, given all of the previous words

- **Architecture:**

- Stack of Transformer decoder blocks.
- No encoder, hence no cross-attention module.
- Components: Positional encoding, masked multihead self-attention, and feedforward network.

- **Directionality:**

- Only considers previous (left) words in attention, not bidirectional like BERT.
- Utilizes masked multihead self-attention for this purpose.

# GPT 1

- Released: 2018
- Parameters: 117 Million
- Layers: 12
- Training Data: Books1 Corpus (7,000 unpublished books)
- Focus: Unsupervised pre-training, Transformer architecture, large-scale language modeling

# GPT 2

- Released: 2019
- Parameters: ~1.5 Billion
- Layers: 48
- Training Data: 40GB (English)
- Focus: Transformer architecture, self-attention mechanism

# GPT 3

- Released: 2020
- Parameters: 175 Billion
- Layers: 175
- Training Data: 570GB (Multilingual)
- Focus: Few-shot learning, prompt engineering, Python support

# GPT 4

- Release: Not Yet
- Parameters: ~100 Trillion (speculative)
- Layers: Unknown
- Training Data: Larger, more diverse (speculative)
- Focus: GPT-4 is known to be a multimodal model, capable of processing both text and image inputs to generate text outputs. Advanced few-shot learning, improved NLU and NLG, reasoning and inference