

# Video Processing

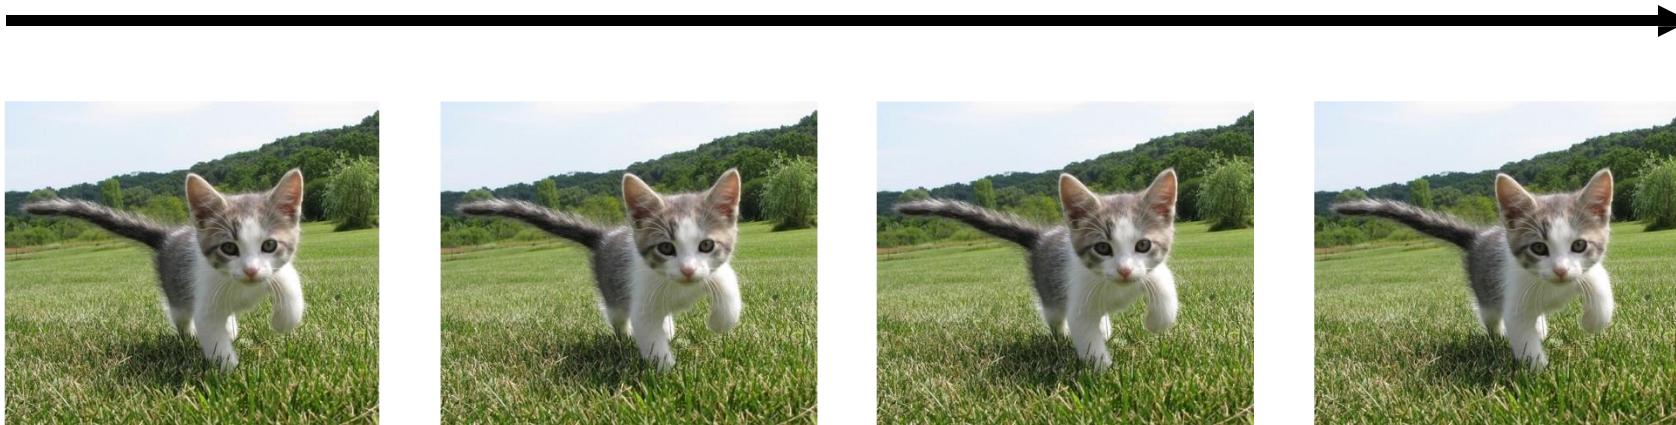
ViViT: A Video Vision Transformer, ICCV21

VideoMAE: Masked Autoencoders are Data-Efficient Learners for Self-Supervised Video Pre-Training, NeurIPS22

# Today: Video = 2D + Time

A video is a **sequence** of images

4D tensor:  $T \times 3 \times H \times W$   
(or  $3 \times T \times H \times W$ )



# Example task: Video Classification



Input video:

$T \times 3 \times H \times W$

Swimming  
**Running**  
Jumping  
Eating  
Standing

# Example task: Video Classification



Images: Recognize **objects**



Dog  
**Cat**  
Fish  
Truck



Videos: Recognize **actions**



Swimming  
**Running**  
Jumping  
Eating  
Standing

# Problem: Videos are big!

Videos are ~30 frames per second (fps)



Size of uncompressed video  
(3 bytes per pixel):

SD (640 x 480): **~1.5 GB per minute**

HD (1920 x 1080): **~10 GB per minute**

Input video:

$T \times 3 \times H \times W$

# Problem: Videos are big!

Videos are ~30 frames per second (fps)



Input video:

$T \times 3 \times H \times W$

Size of uncompressed video  
(3 bytes per pixel):

SD (640 x 480): **~1.5 GB per minute**

HD (1920 x 1080): **~10 GB per minute**

Solution: Train on short **clips**: low  
fps and low spatial resolution  
e.g.  $T = 16$ ,  $H=W=112$   
(3.2 seconds at 5 fps, 588 KB)

# Training on Clips

**Raw video:** Long, high FPS



# Training on Clips

**Raw video:** Long, high FPS



**Training:** Train model to classify short **clips** with low FPS

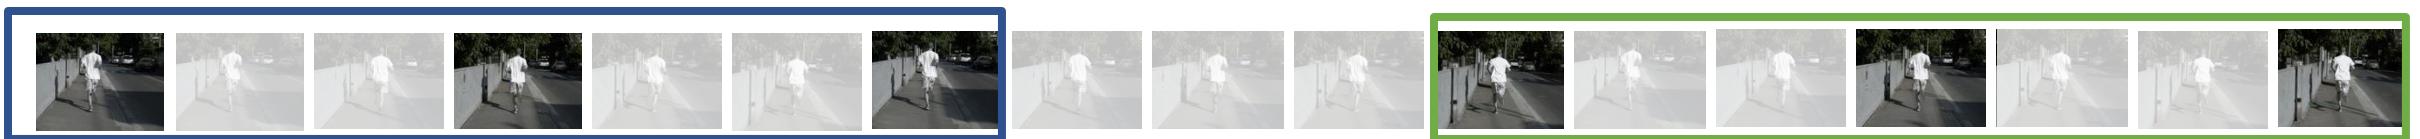


# Training on Clips

**Raw video:** Long, high FPS



**Training:** Train model to classify short **clips** with low FPS



**Testing:** Run model on different clips, average predictions



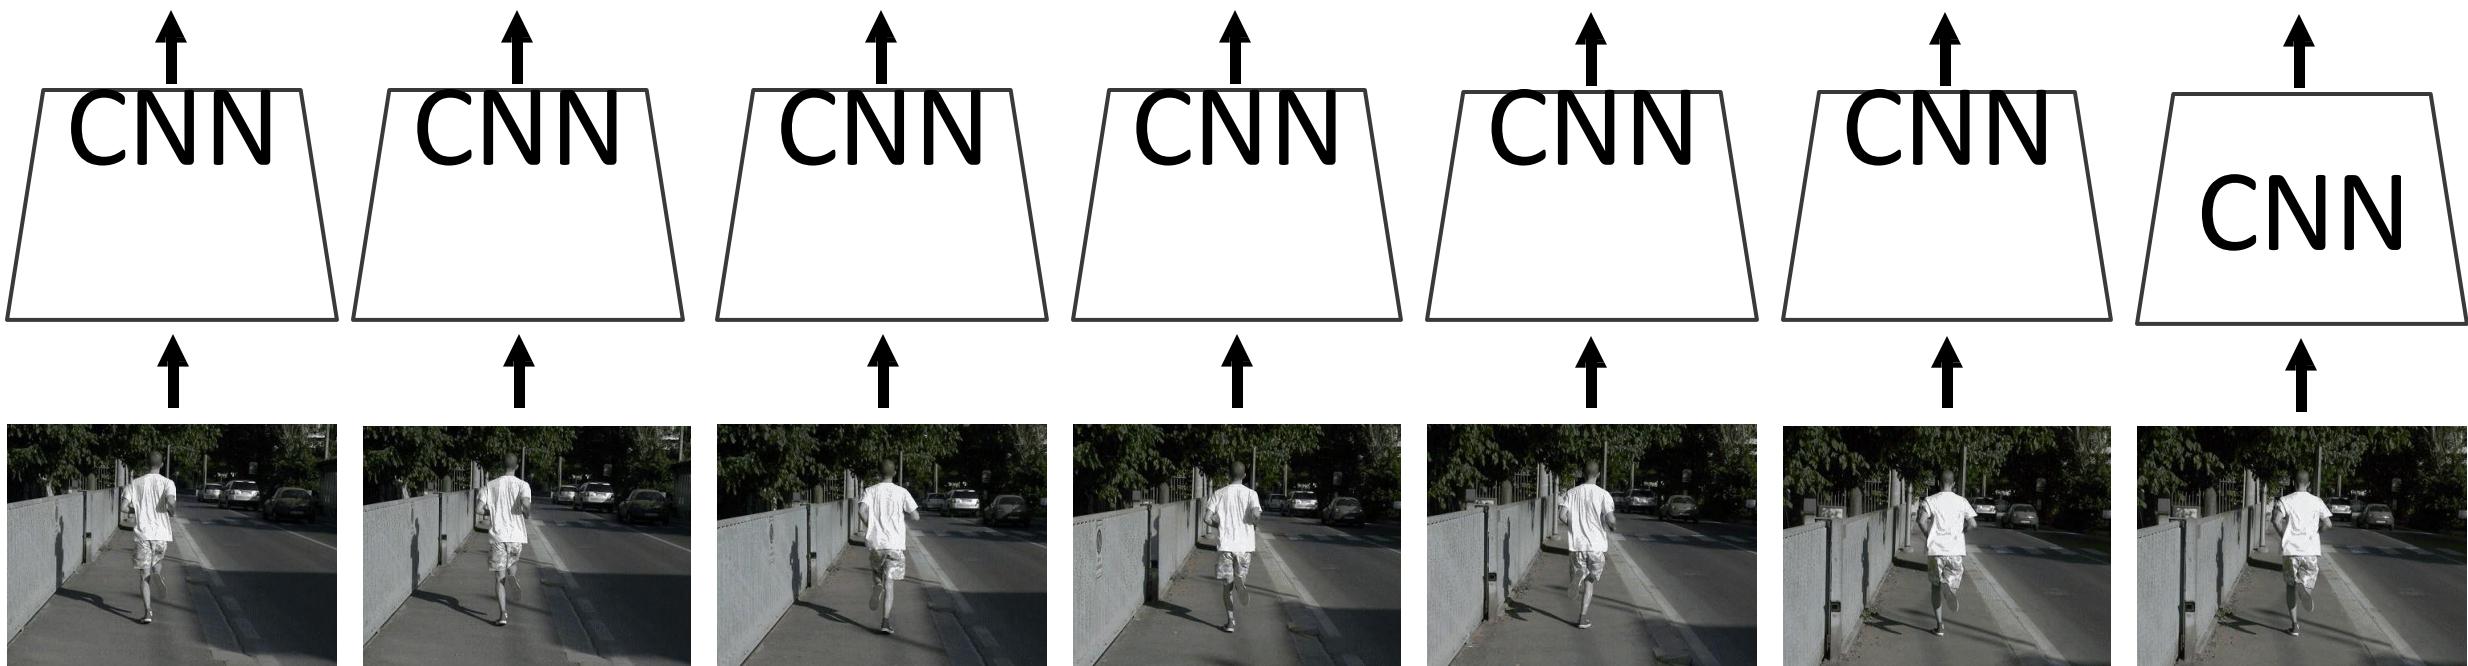
- What could be a very simple video classifier?

# Video Classification: Single-Frame CNN

Simple idea: train normal 2D CNN to classify video frames independently!

Often a **very** strong baseline for video classification

“Running” “Running” “Running” “Running” “Running” “Running” “Running”

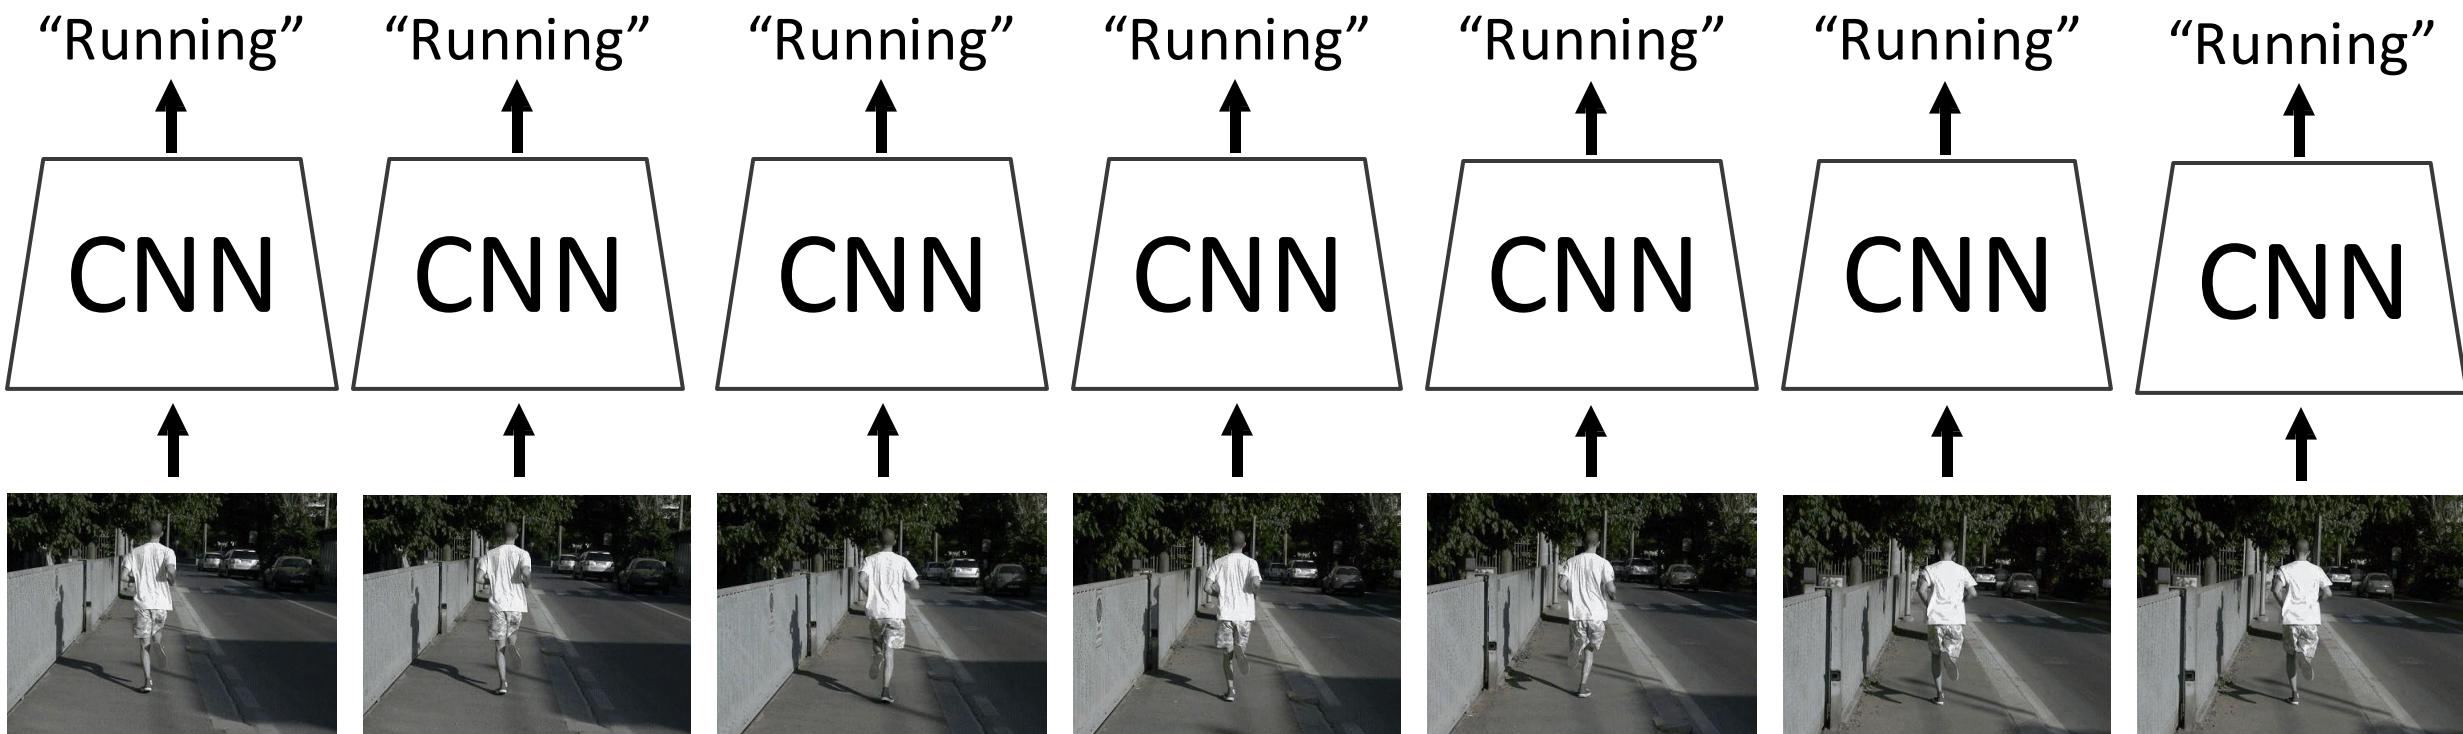


# Video Classification: Single-Frame CNN

Simple idea: train normal 2D CNN to classify video frames independently!

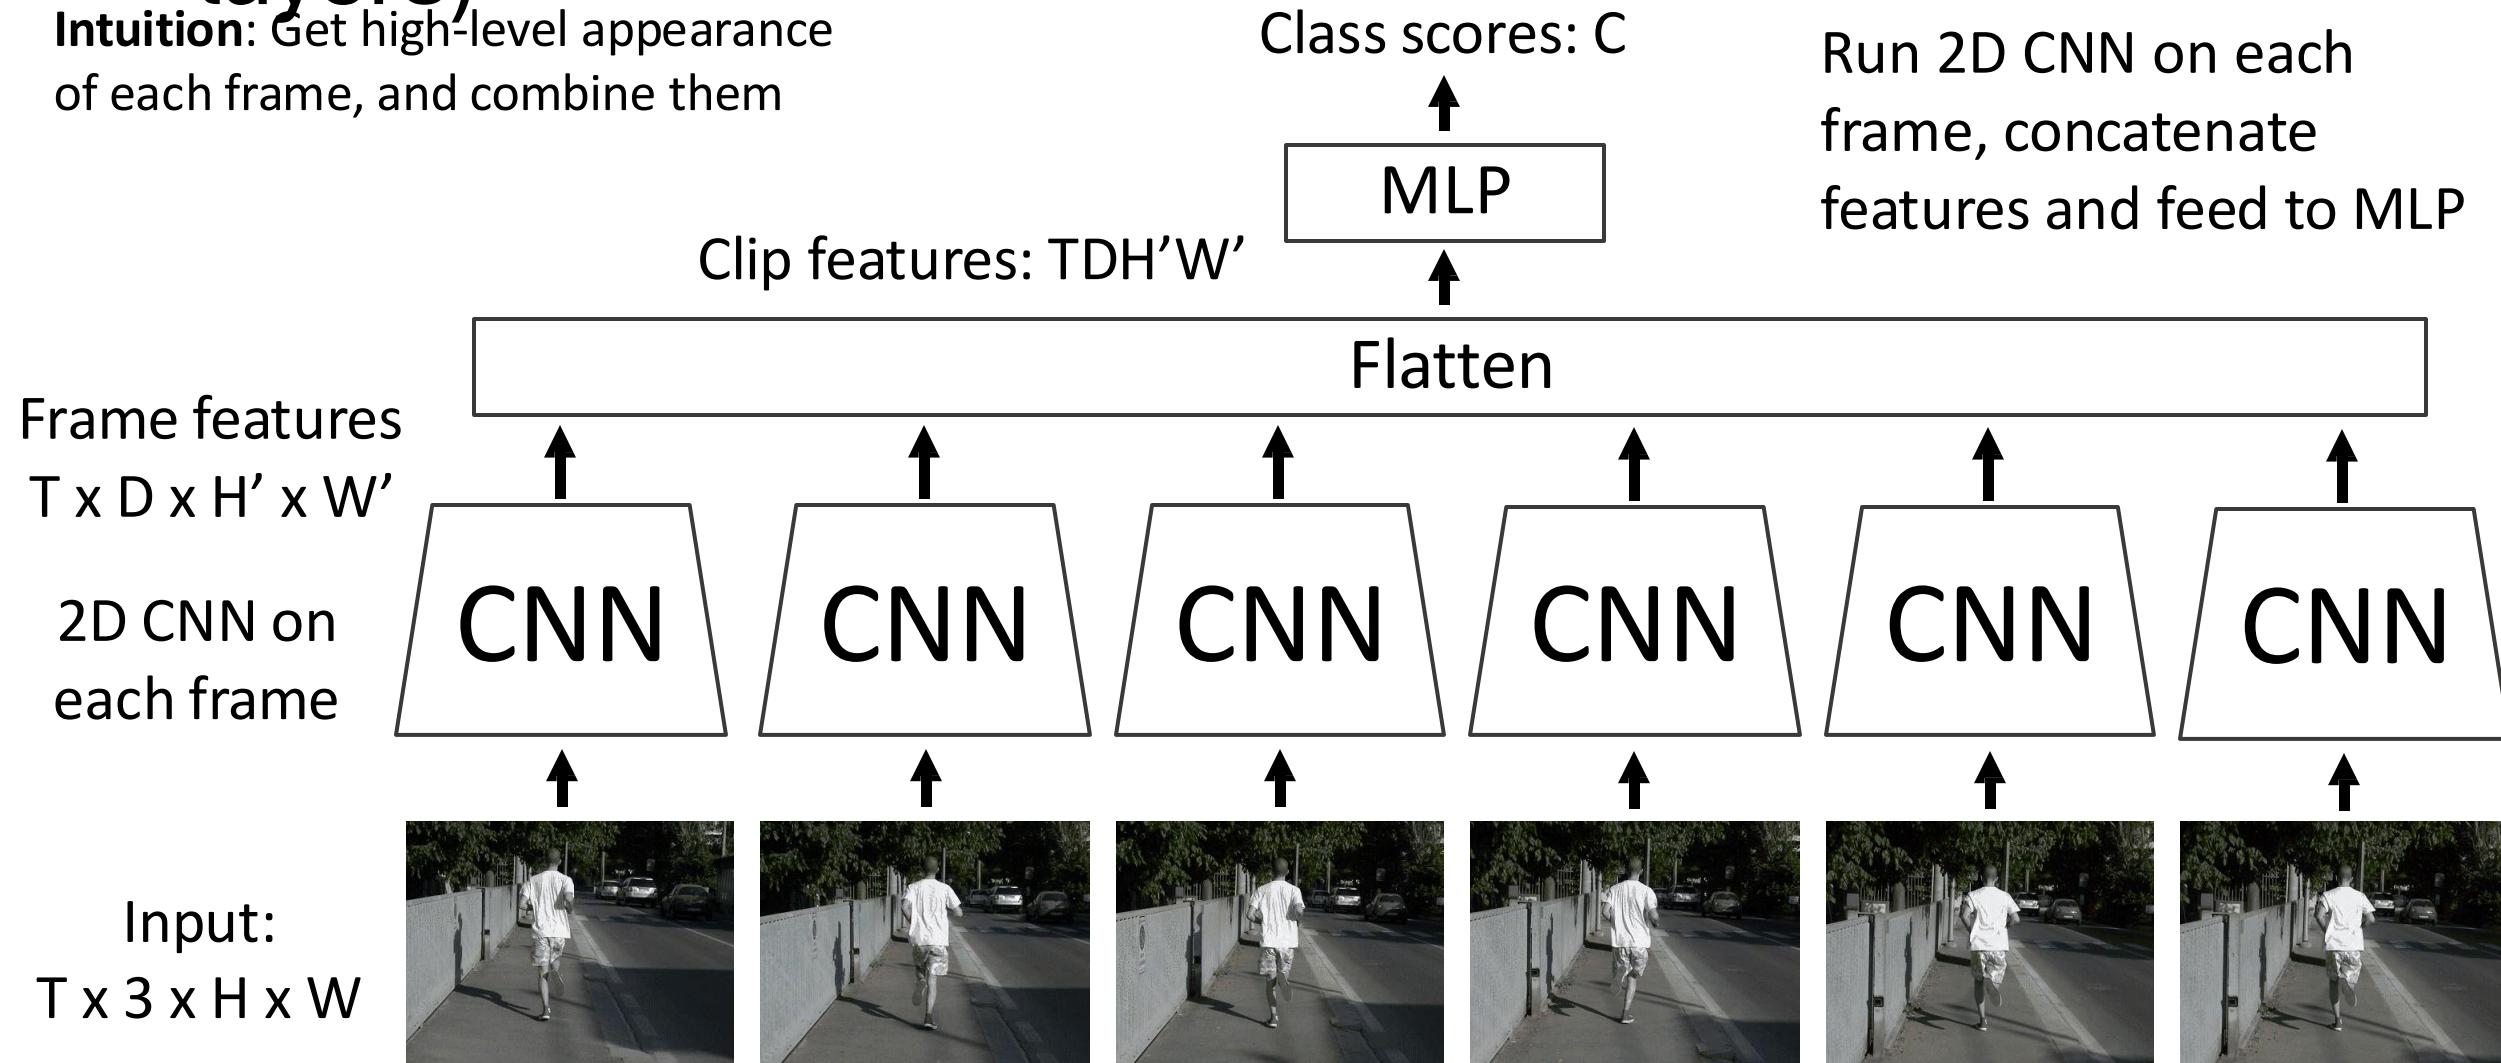
(Average predicted probs at test-time)

Often a **very** strong baseline for video classification



# Video Classification: Late Fusion (with FC layers)

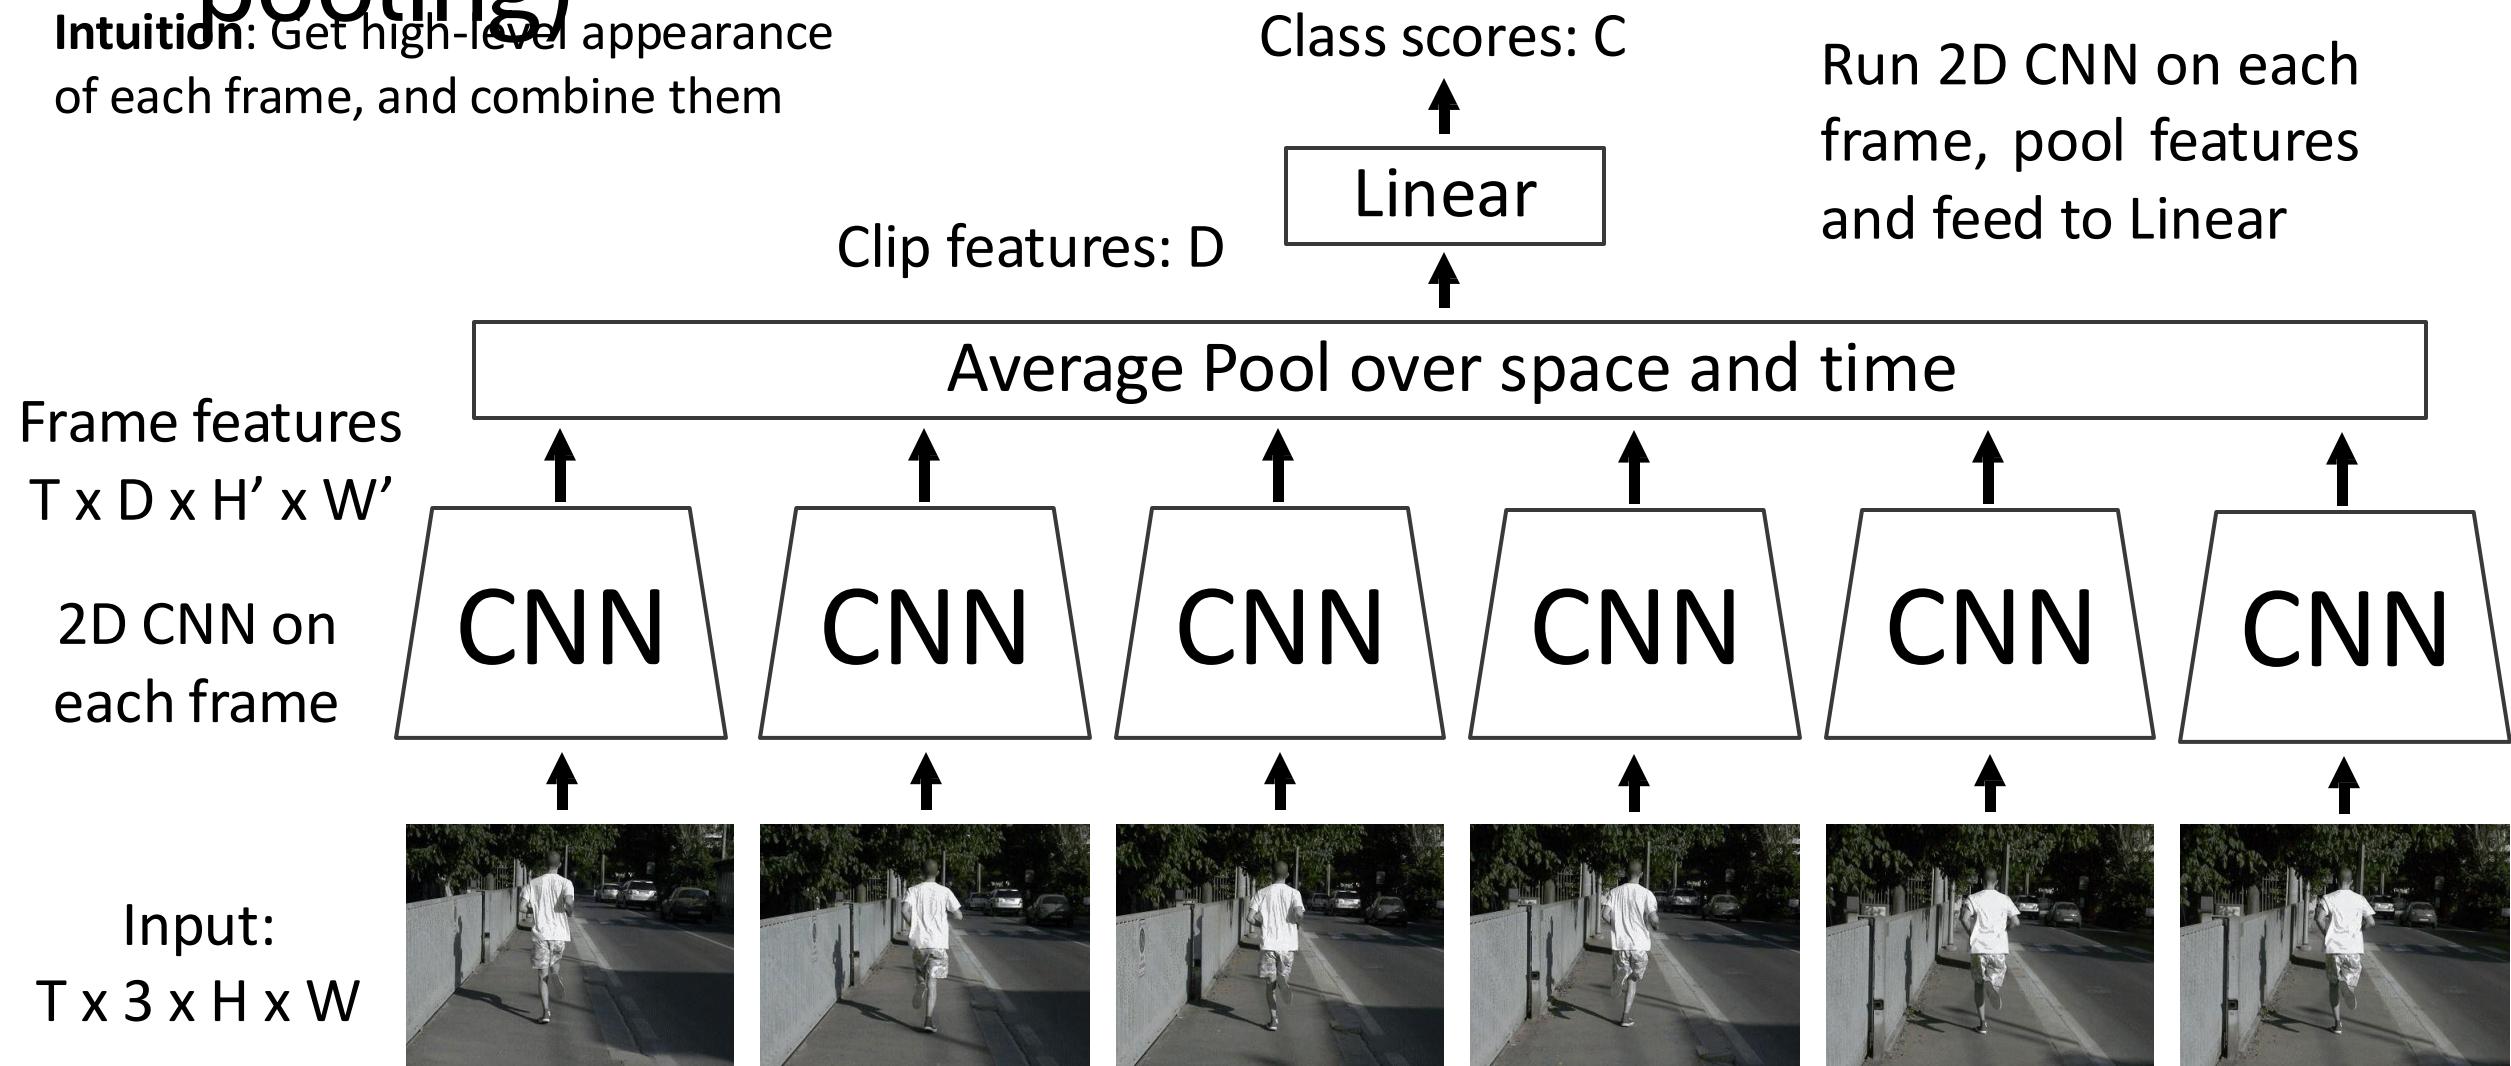
**Intuition:** Get high-level appearance of each frame, and combine them



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

# Video Classification: Late Fusion (with pooling)

Intuition: Get high-level appearance of each frame, and combine them



# Video Classification: Late Fusion (with pooling)

**Intuition:** Get high-level appearance of each frame, and combine them

**Problem:** Hard to compare low-level motion between frames

Class scores: C

Linear

Run 2D CNN on each frame, pool features and feed to Linear

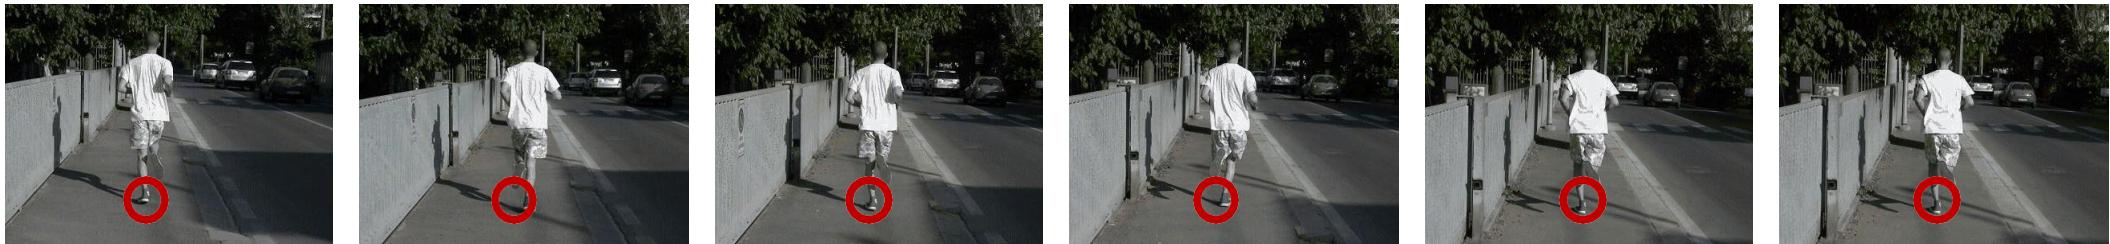
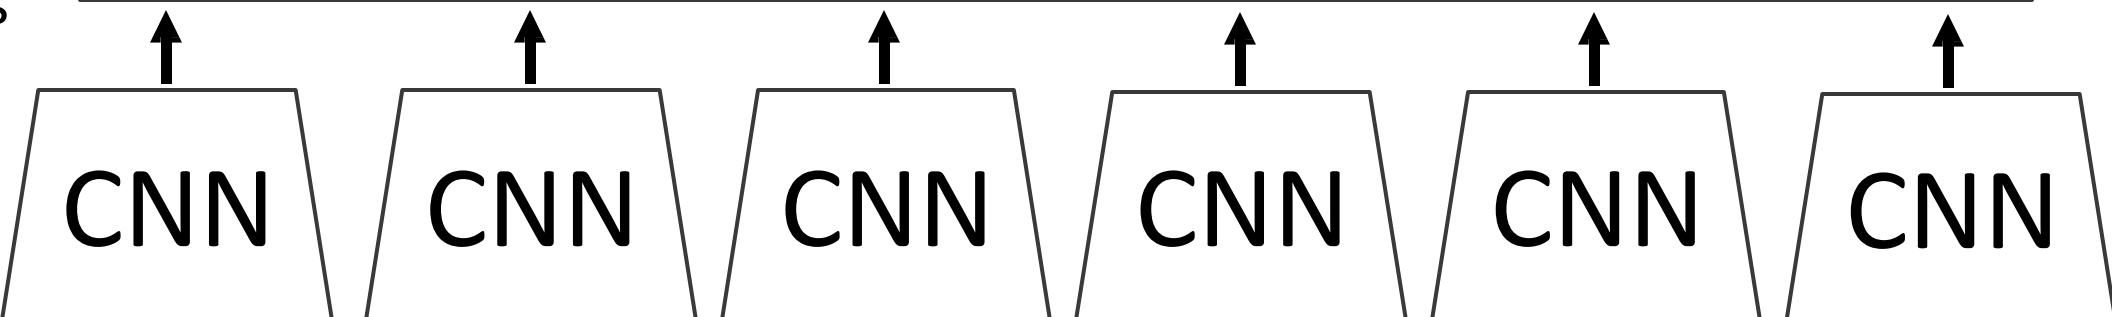
Clip features: D

Frame features

$T \times D \times H' \times W'$

2D CNN on each frame

Average Pool over space and time



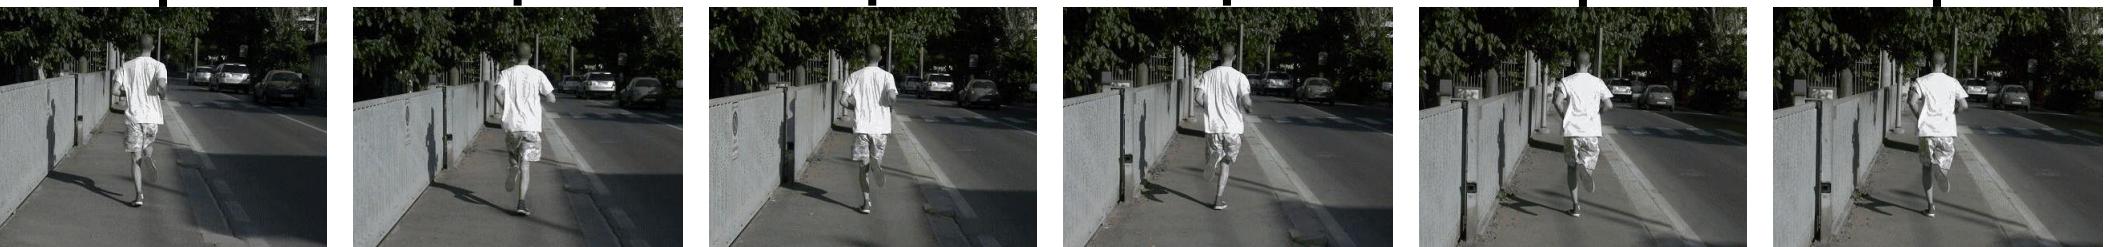
# Video Classification: Early Fusion

**Intuition:** Compare frames with very first conv layer, after that normal 2D CNN

First 2D convolution collapses all temporal information:  
**Input:**  $3T \times H \times W$   
**Output:**  $D \times H \times W$

Reshape:  
 $3T \times H \times W$

**Input:**  
 $T \times 3 \times H \times W$



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

# Video Classification: Early Fusion

**Intuition:** Compare frames with very first conv layer, after that normal 2D CNN

**Problem:** One layer of temporal processing may not be enough!

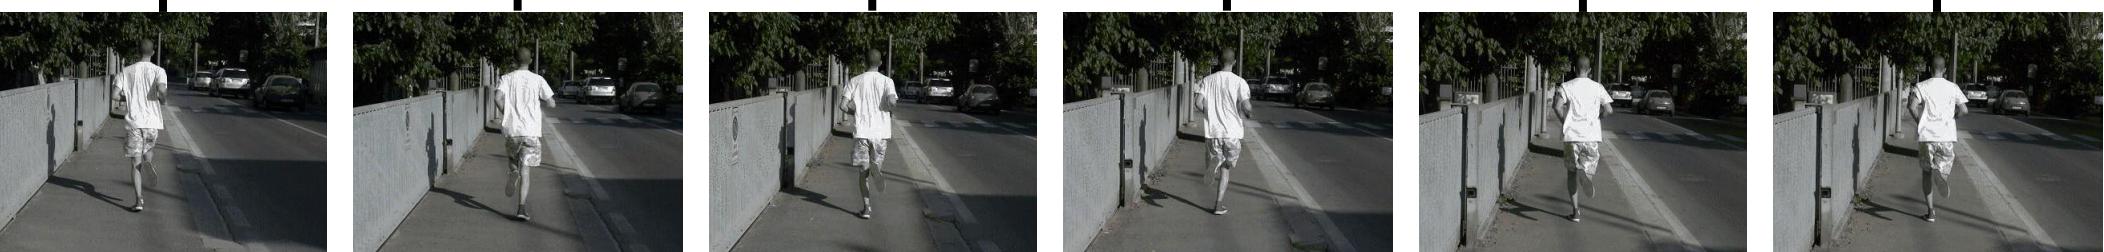
First 2D convolution collapses all temporal information:

**Input:**  $3T \times H \times W$

**Output:**  $D \times H \times W$

Reshape:  
 $3T \times H \times W$

**Input:**  
 $T \times 3 \times H \times W$

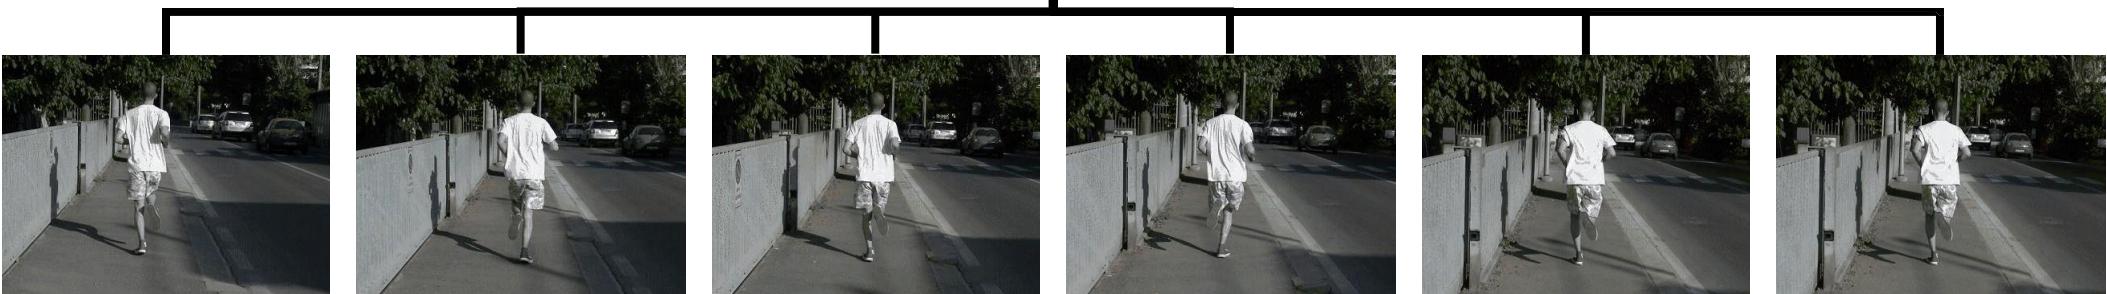


# Video Classification: 3D CNN

**Intuition:** Use 3D versions of convolution and pooling to slowly fuse temporal information over the course of the network

Each layer in the network is a 4D tensor:  $D \times T \times H \times W$   
Use 3D conv and 3D pooling operations

**Input:**  
 $3 \times T \times H \times W$



Ji et al, "3D Convolutional Neural Networks for Human Action Recognition", TPAMI 2010 ; Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014

# Early Fusion vs Late Fusion vs 3D

## CNN

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3

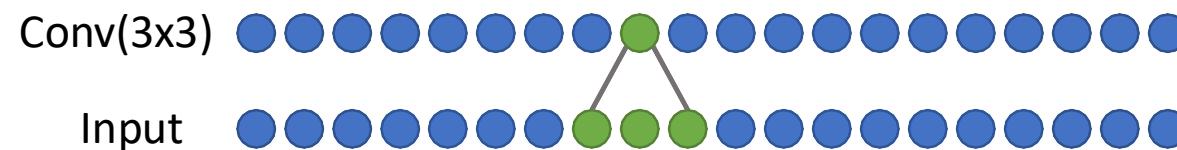
Late  
Fusion

# Early Fusion vs Late Fusion vs 3D

## CNN

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3

Late  
Fusion

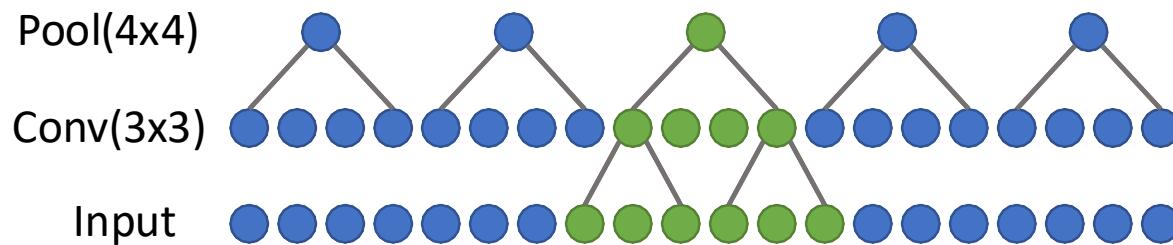


# Early Fusion vs Late Fusion vs 3D

## CNN

Late  
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	$3 \times 20 \times 64 \times 64$	
Conv2D(3x3, 3->12)	$12 \times 20 \times 64 \times 64$	$1 \times 3 \times 3$
Pool2D(4x4)	$12 \times 20 \times 16 \times 16$	$1 \times 6 \times 6$



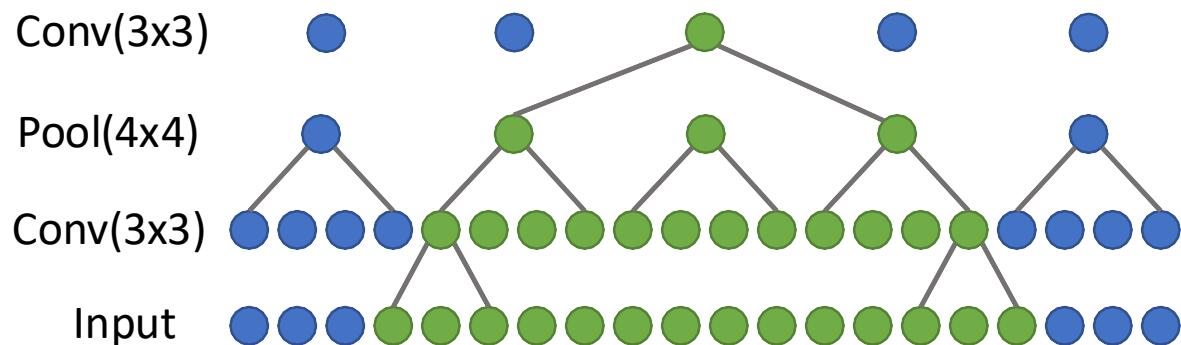
# Early Fusion vs Late Fusion vs 3D

## CNN

Late  
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14

Build slowly in space

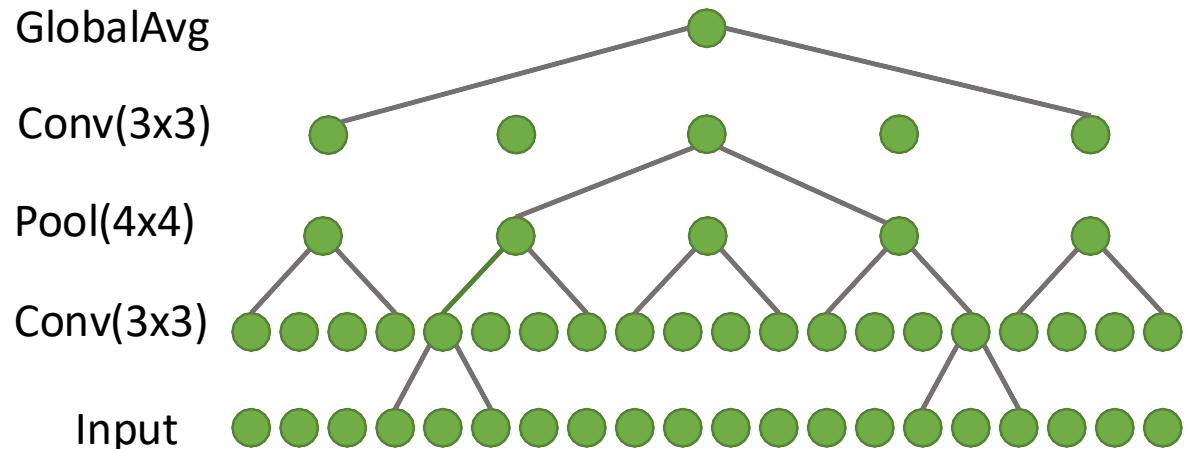


# Early Fusion vs Late Fusion vs 3D CNN

Late  
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64

Build slowly in space,  
All-at-once in time at end



(Small example)

# Early Fusion vs Late Fusion vs 3D CNN

Late Fusion  
Early Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3*20->12)	12 x 64 x 64	20 x 3 x 3
Pool2D(4x4)	12 x 16 x 16	20 x 6 x 6
Conv2D(3x3, 12->24)	24 x 16 x 16	20 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

Build slowly in space,  
All-at-once in time at end

Build slowly in space,  
All-at-once in time at start

# Early Fusion vs Late Fusion vs 3D

## CNN

Late  
Fusion

Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64
Input	3 x 20 x 64 x 64	
Conv2D(3x3, 3*10->12)	12 x 64 x 64	20 x 3 x 3
Pool2D(4x4)	12 x 16 x 16	20 x 6 x 6
Conv2D(3x3, 12->24)	24 x 16 x 16	20 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64
Input	3 x 20 x 64 x 64	
Conv3D(3x3x3, 3->12)	12 x 20 x 64 x 64	3 x 3 x 3
Pool3D(4x4x4)	12 x 5 x 16 x 16	6 x 6 x 6
Conv3D(3x3x3, 12->24)	24 x 5 x 16 x 16	14 x 14 x 14
GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

Build slowly in space,  
All-at-once in time at end

Early  
Fusion

Build slowly in space,  
All-at-once in time at start

3D CNN

Build slowly in space,  
Build slowly in time  
"Slow Fusion"

(Small example  
architectures,  
in practice  
much bigger)

# Early Fusion vs Late Fusion vs 3D

## CNN

	Layer	Size (C x T x H x W)	Receptive Field (T x H x W)
Late Fusion	Input	3 x 20 x 64 x 64	
	Conv2D(3x3, 3->12)	12 x 20 x 64 x 64	1 x 3 x 3
	Pool2D(4x4)	12 x 20 x 16 x 16	1 x 6 x 6
	Conv2D(3x3, 12->24)	24 x 20 x 16 x 16	1 x 14 x 14
	GlobalAvgPool	24 x 1 x 1 x 1	20 x 64 x 64
Early Fusion	Input	3 x 20 x 64 x 64	
	Conv2D(3x3, 3*10->12)	12 x 64 x 64	20 x 3 x 3
	Pool2D(4x4)	12 x 16 x 16	20 x 6 x 6
	Conv2D(3x3, 12->24)	24 x 16 x 16	20 x 14 x 14
	GlobalAvgPool	24 x 1 x 1	20 x 64 x 64
3D CNN	Input	3 x 20 x 64 x 64	
	Conv3D(3x3x3, 3->12)	12 x 20 x 64 x 64	3 x 3 x 3
	Pool3D(4x4x4)	12 x 5 x 16 x 16	6 x 6 x 6
	Conv3D(3x3x3, 12->24)	24 x 5 x 16 x 16	14 x 14 x 14
	GlobalAvgPool	24 x 1 x 1	20 x 64 x 64

What is the difference?

Build slowly in space,  
All-at-once in time at end

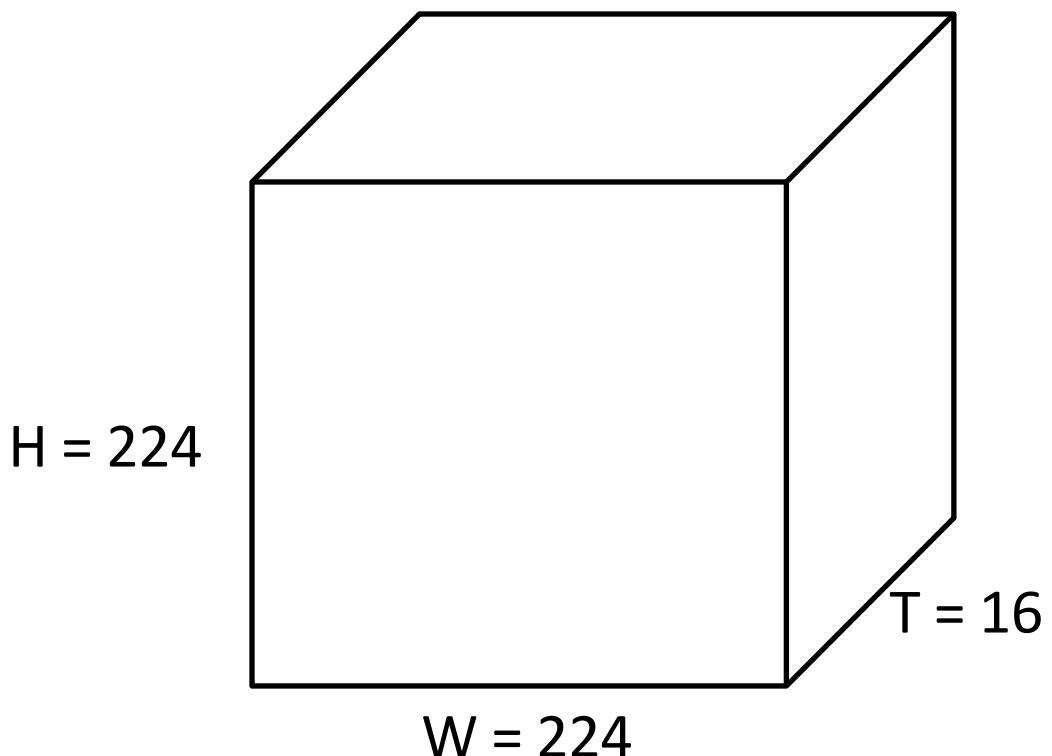
Build slowly in space,  
All-at-once in time at start

Build slowly in space,  
Build slowly in time  
"Slow Fusion"

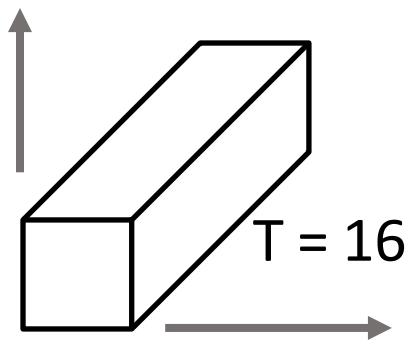
(Small example architectures,  
in practice much bigger)

# 2D Conv (Early Fusion) vs 3D Conv (3D CNN)

**Input:**  $C_{in} \times T \times H \times W$   
(3D grid with  $C_{in}$ -dim  
feat at each point)

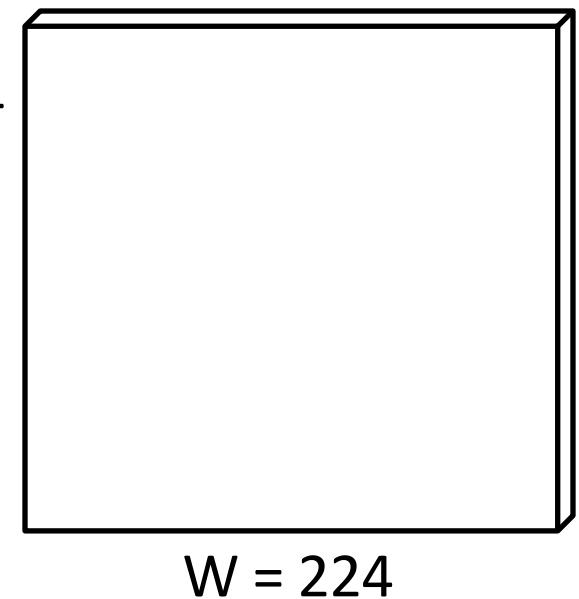


**Weight:**  
 $C_{out} \times C_{in} \times T \times 3 \times 3$   
Slide over x and y



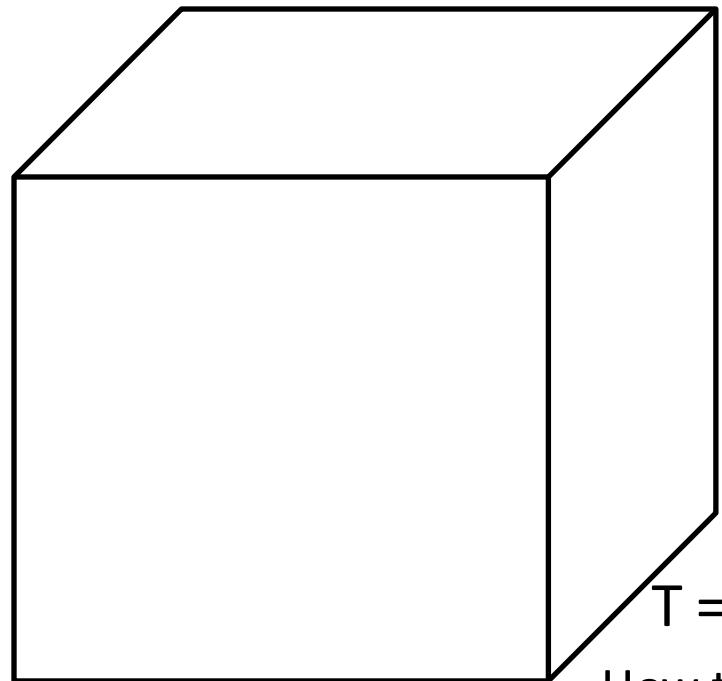
$C_{out}$  different filters

**Output:**  
 $C_{out} \times H \times W$   
2D grid with  $C_{out}$ -dim  
feat at each point

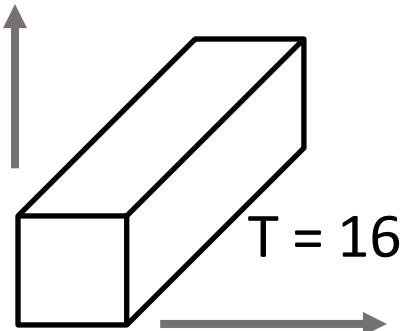


# 2D Conv (Early Fusion) vs 3D Conv (3D CNN)

**Input:**  $C_{in} \times T \times H \times W$   
(3D grid with  $C_{in}$ -dim  
feat at each point)



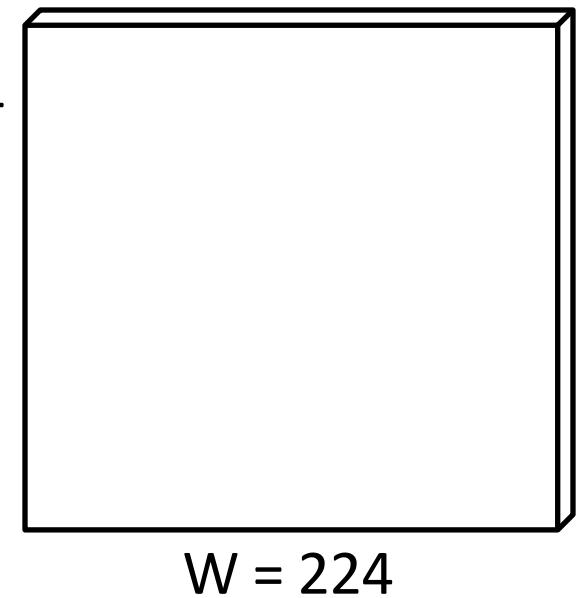
**Weight:**  
 $C_{out} \times C_{in} \times T \times 3 \times 3$   
Slide over x and y



$C_{out}$  different filters

How to recognize **blue** to **orange**  
transitions anywhere in space and time?

**Output:**  
 $C_{out} \times H \times W$   
2D grid with  $C_{out}$ -dim  
feat at each point

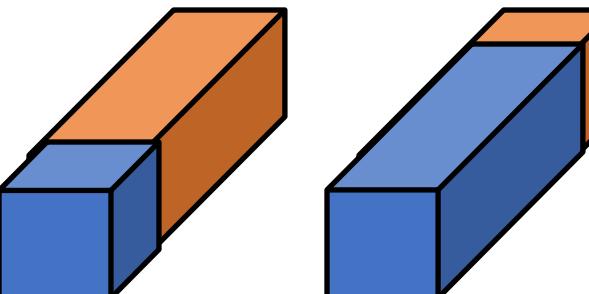
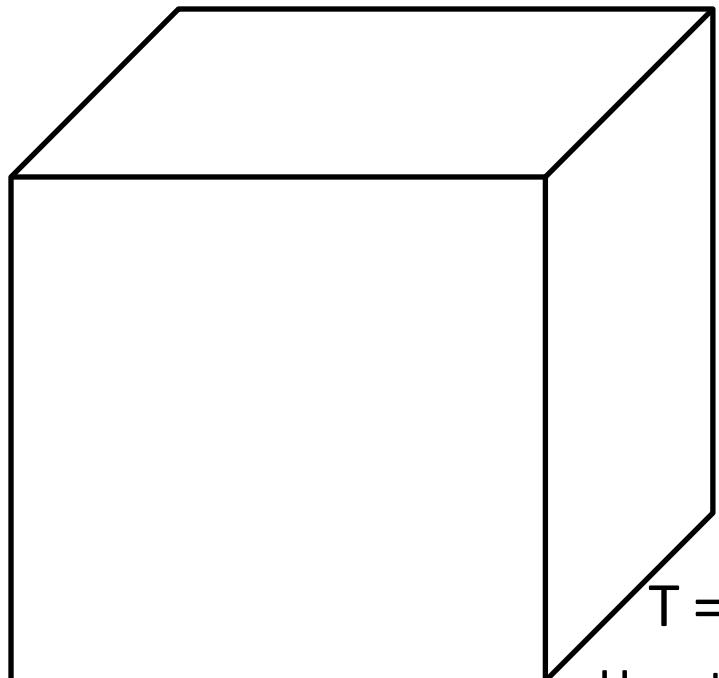


# 2D Conv (Early Fusion) vs 3D Conv (3D CNN)

**Input:**  $C_{in} \times T \times H \times W$   
(3D grid with  $C_{in}$ -dim  
feat at each point)

**Weight:**  
 $C_{out} \times C_{in} \times T \times 3 \times 3$   
Slide over x and y

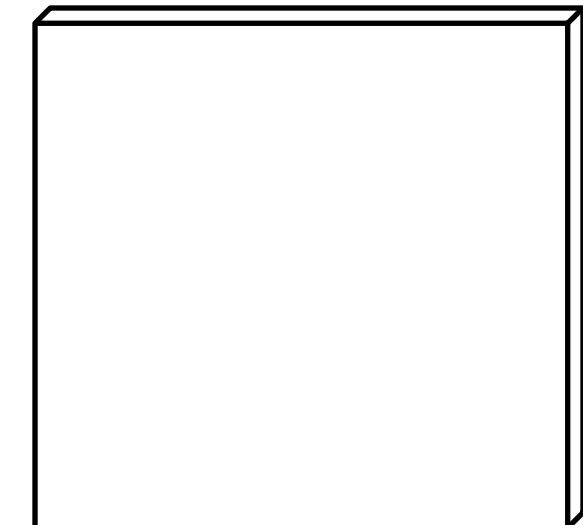
**Output:**  
 $C_{out} \times H \times W$   
2D grid with  $C_{out}$  -dim



$T = 16$

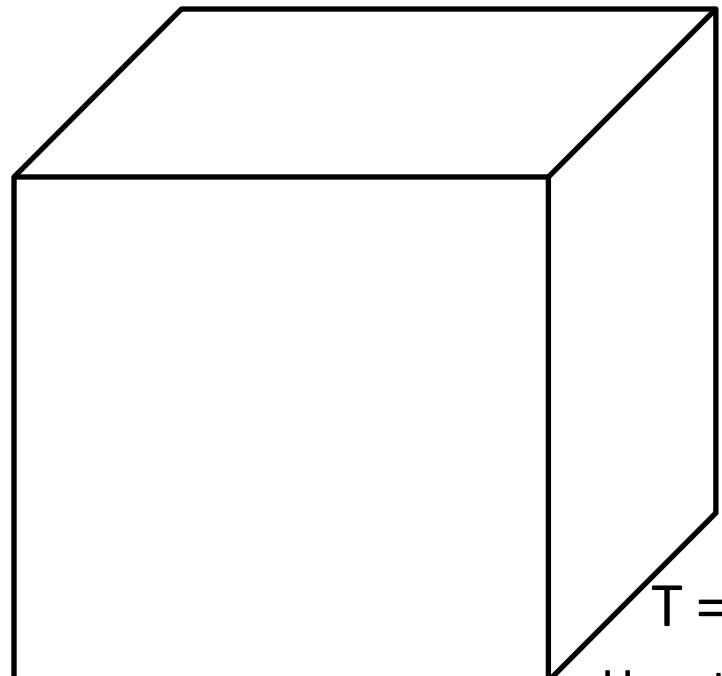
$C_{out}$  different filters

How to recognize **blue** to **orange**  
transitions anywhere in space and time?



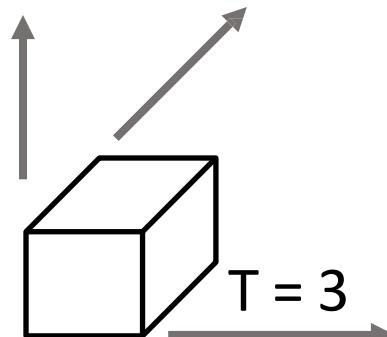
# 2D Conv (Early Fusion) vs 3D Conv (3D CNN)

**Input:**  $C_{in} \times T \times H \times W$   
(3D grid with  $C_{in}$ -dim  
feat at each point)



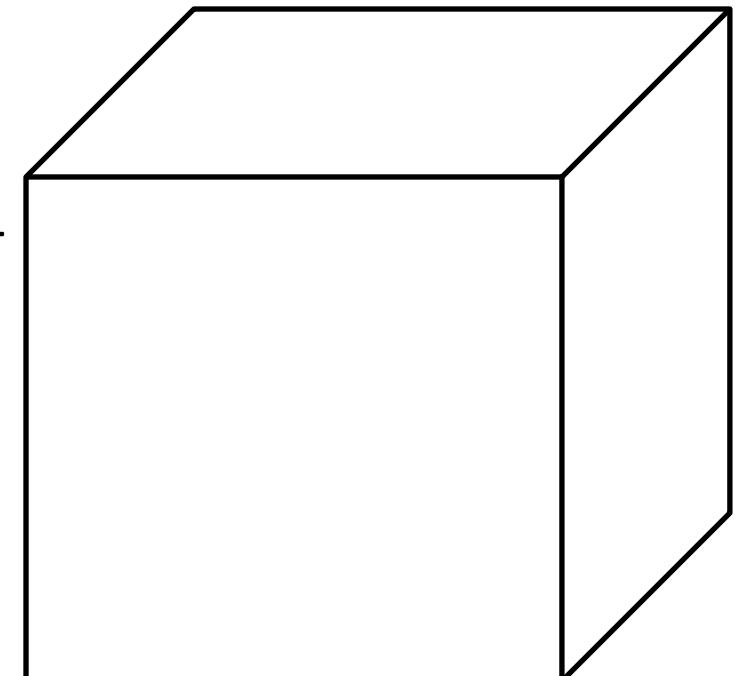
How to recognize **blue** to **orange**  
transitions anywhere in space and time?

**Weight:**  
 $C_{out} \times C_{in} \times 3 \times 3 \times 3$   
Slide over x and y



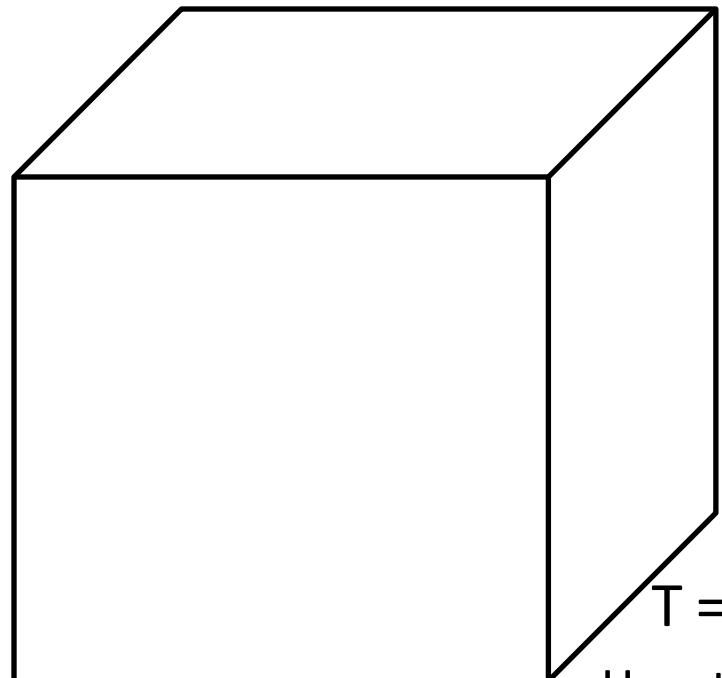
$C_{out}$  different filters

**Output:**  
 $C_{out} \times T \times H \times W$   
3D grid with  $C_{out}$ -dim  
feat at each point

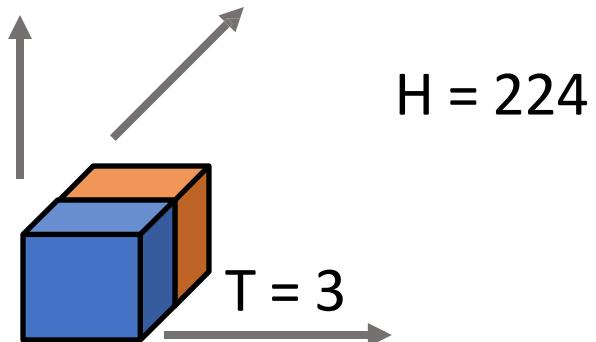


# 2D Conv (Early Fusion) vs 3D Conv (3D CNN)

**Input:**  $C_{in} \times T \times H \times W$   
(3D grid with  $C_{in}$ -dim  
feat at each point)

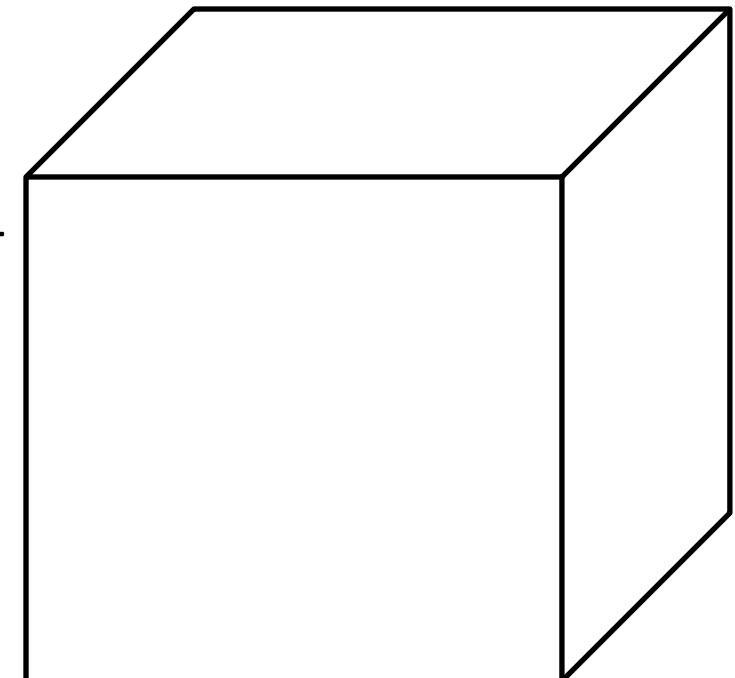


**Weight:**  
 $C_{out} \times C_{in} \times 3 \times 3 \times 3$   
Slide over x and y



$C_{out}$  different filters  
How to recognize **blue** to **orange**  
transitions anywhere in space and time?

**Output:**  
 $C_{out} \times T \times H \times W$   
3D grid with  $C_{out}$ -dim  
feat at each point



# Example Video Dataset: Sports-1M

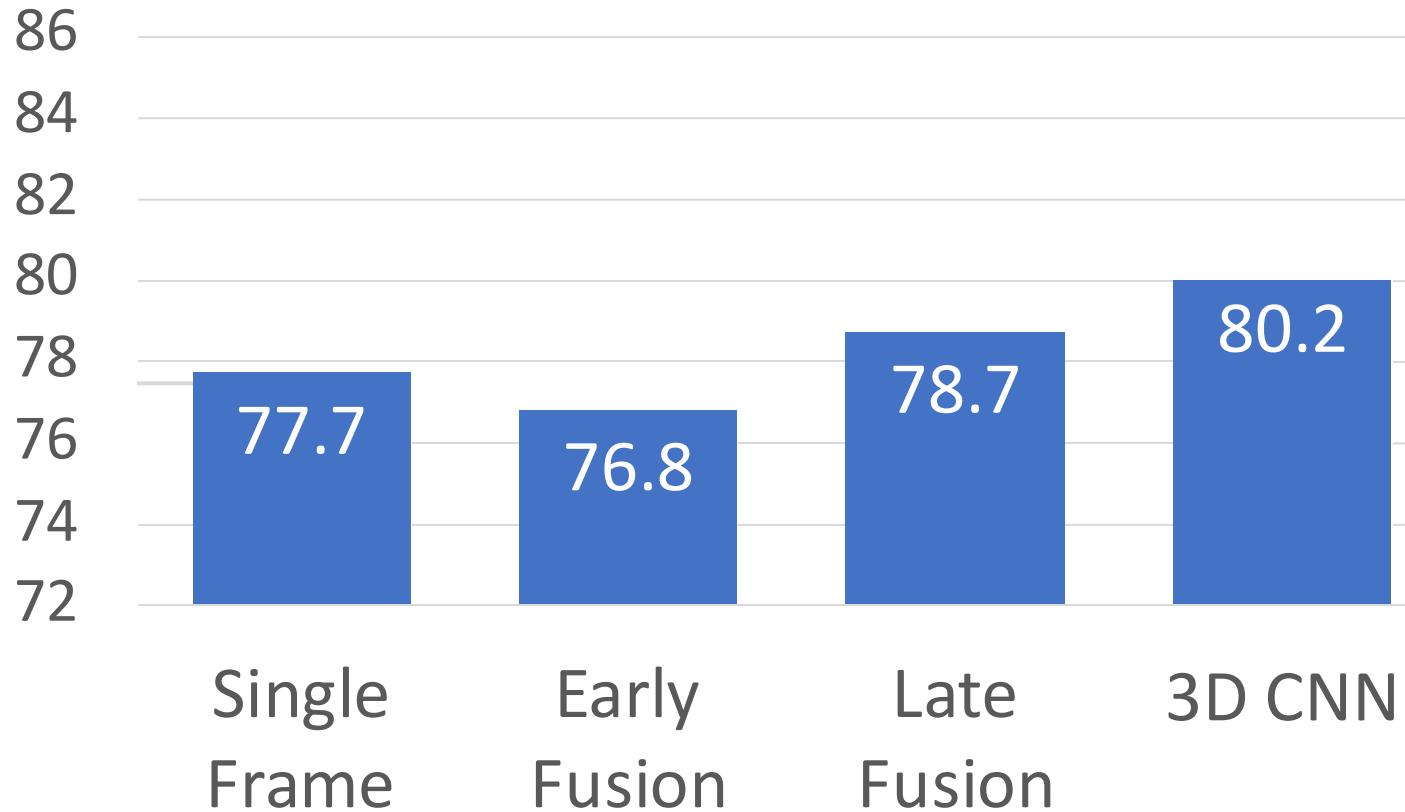


1 million YouTube videos  
annotated with labels for  
487 different types of sports

**Ground Truth**  
**Correct prediction**  
**Incorrect prediction**

# Early Fusion vs Late Fusion vs 3D CNN

Sports-1M Top-5 Accuracy



Single Frame model works well – always try this first!

3D CNNs have improved a lot since 2014!

# C3D: The VGG of 3D CNNs

3D CNN that uses all  $3 \times 3 \times 3$  conv and  $2 \times 2 \times 2$  pooling  
(except Pool1 which is  $1 \times 2 \times 2$ )

Released model pretrained on Sports-1M: Many people used this as a video feature extractor

Layer	Size
Input	$3 \times 16 \times 112 \times 112$
Conv1 ( $3 \times 3 \times 3$ )	$64 \times 16 \times 112 \times 112$
Pool1 ( $1 \times 2 \times 2$ )	$64 \times 16 \times 56 \times 56$
Conv2 ( $3 \times 3 \times 3$ )	$128 \times 16 \times 56 \times 56$
Pool2 ( $2 \times 2 \times 2$ )	$128 \times 8 \times 28 \times 28$
Conv3a ( $3 \times 3 \times 3$ )	$256 \times 8 \times 28 \times 28$
Conv3b ( $3 \times 3 \times 3$ )	$256 \times 8 \times 28 \times 28$
Pool3 ( $2 \times 2 \times 2$ )	$256 \times 4 \times 14 \times 14$
Conv4a ( $3 \times 3 \times 3$ )	$512 \times 4 \times 14 \times 14$
Conv4b ( $3 \times 3 \times 3$ )	$512 \times 4 \times 14 \times 14$
Pool4 ( $2 \times 2 \times 2$ )	$512 \times 2 \times 7 \times 7$
Conv5a ( $3 \times 3 \times 3$ )	$512 \times 2 \times 7 \times 7$
Conv5b ( $3 \times 3 \times 3$ )	$512 \times 2 \times 7 \times 7$
Pool5	$512 \times 1 \times 3 \times 3$
FC6	4096
FC7	4096
FC8	C

# C3D: The VGG of 3D CNNs

3D CNN that uses all  $3 \times 3 \times 3$  conv and  $2 \times 2 \times 2$  pooling  
(except Pool1 which is  $1 \times 2 \times 2$ )

Released model pretrained on Sports-1M: Many people used this as a video feature extractor

**Problem:**  $3 \times 3 \times 3$  conv is very expensive!

AlexNet: 0.7 GFLOP

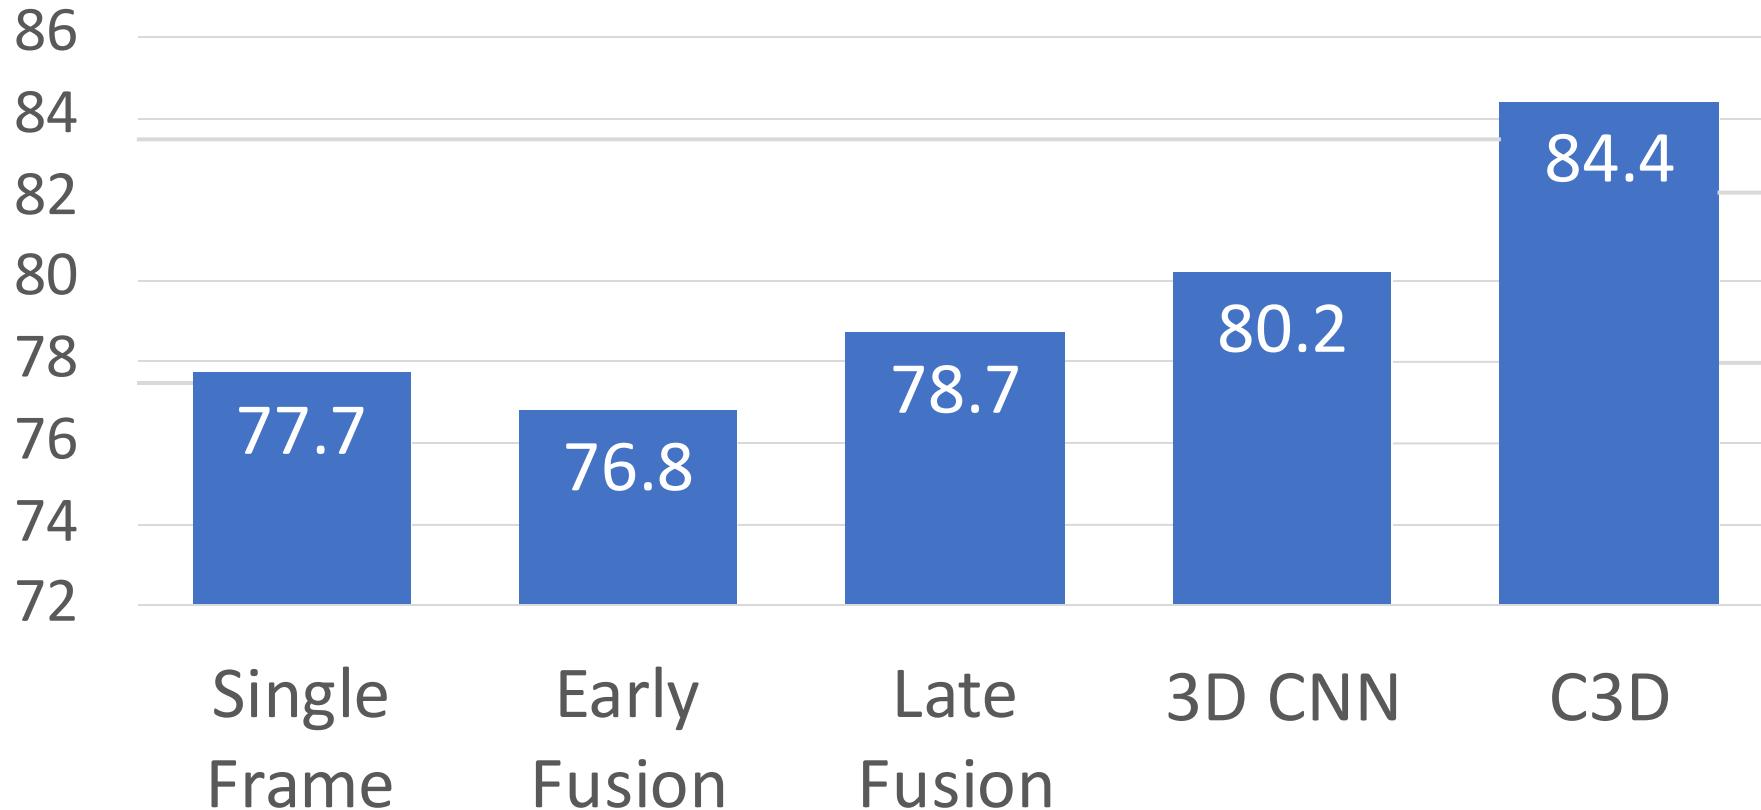
VGG-16: 13.6 GFLOP

**C3D: 39.5 GFLOP (2.9x VGG!)**

Layer	Size	MFLOPs
Input	$3 \times 16 \times 112 \times 112$	
Conv1 ( $3 \times 3 \times 3$ )	$64 \times 16 \times 112 \times 112$	1.04
Pool1 ( $1 \times 2 \times 2$ )	$64 \times 16 \times 56 \times 56$	
Conv2 ( $3 \times 3 \times 3$ )	$128 \times 16 \times 56 \times 56$	11.10
Pool2 ( $2 \times 2 \times 2$ )	$128 \times 8 \times 28 \times 28$	
Conv3a ( $3 \times 3 \times 3$ )	$256 \times 8 \times 28 \times 28$	5.55
Conv3b ( $3 \times 3 \times 3$ )	$256 \times 8 \times 28 \times 28$	11.10
Pool3 ( $2 \times 2 \times 2$ )	$256 \times 4 \times 14 \times 14$	
Conv4a ( $3 \times 3 \times 3$ )	$512 \times 4 \times 14 \times 14$	2.77
Conv4b ( $3 \times 3 \times 3$ )	$512 \times 4 \times 14 \times 14$	5.55
Pool4 ( $2 \times 2 \times 2$ )	$512 \times 2 \times 7 \times 7$	
Conv5a ( $3 \times 3 \times 3$ )	$512 \times 2 \times 7 \times 7$	0.69
Conv5b ( $3 \times 3 \times 3$ )	$512 \times 2 \times 7 \times 7$	0.69
Pool5	$512 \times 1 \times 3 \times 3$	
FC6	4096	0.51
FC7	4096	0.45
FC8	C	0.05

# Early Fusion vs Late Fusion vs 3D CNN

Sports-1M Top-5 Accuracy



Karpathy et al, "Large-scale Video Classification with Convolutional Neural Networks", CVPR 2014  
Tran et al, "Learning Spatiotemporal Features with 3D Convolutional Networks", ICCV 2015

# Measuring Motion: Optical Flow

Image at frame t

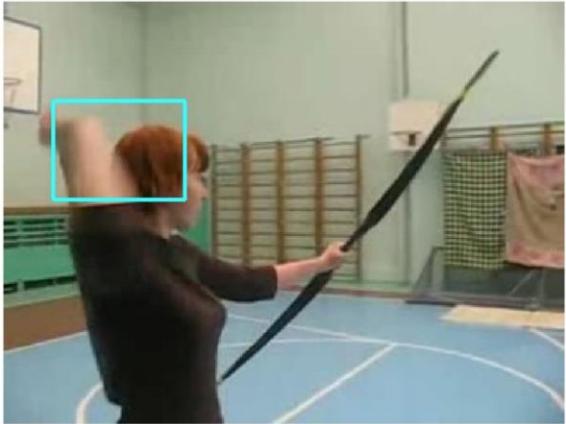


Image at frame t+1

# Measuring Motion: Optical Flow

Image at frame t

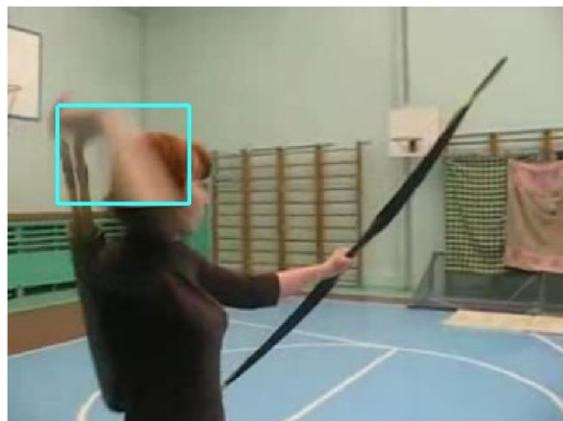
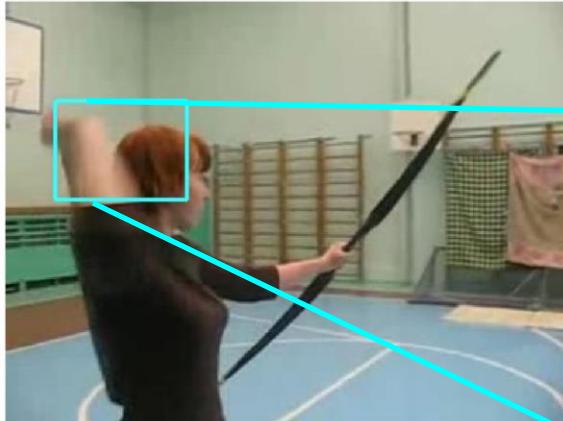
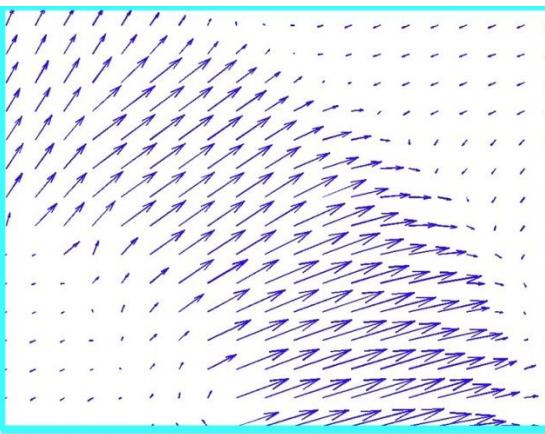


Image at frame t+1

Optical flow gives a displacement field  $F$  between images  $I_t$  and  $I_{t+1}$



Tells where each pixel will move in the next frame:  
 $F(x, y) = (dx, dy)$   
 $I_{t+1}(x+dx, y+dy) = I_t(x, y)$

# Measuring Motion: Optical Flow

Image at frame t

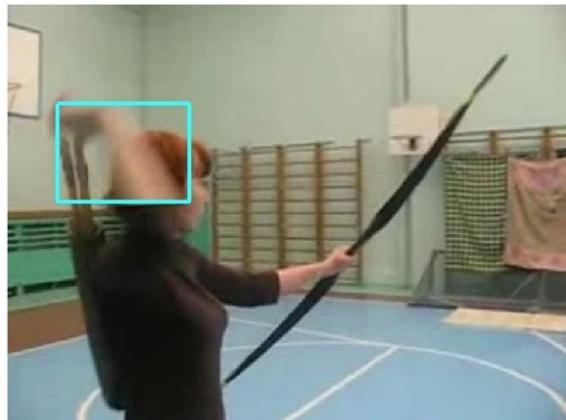
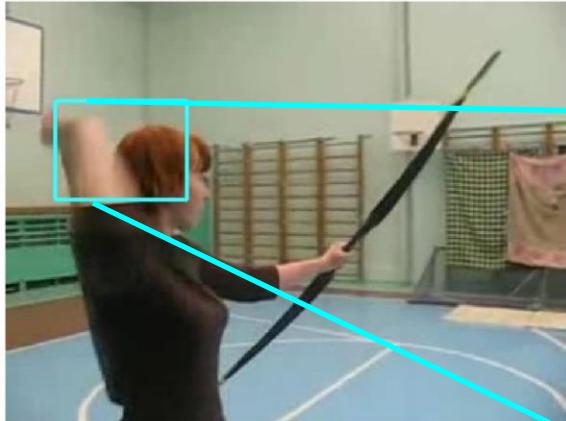
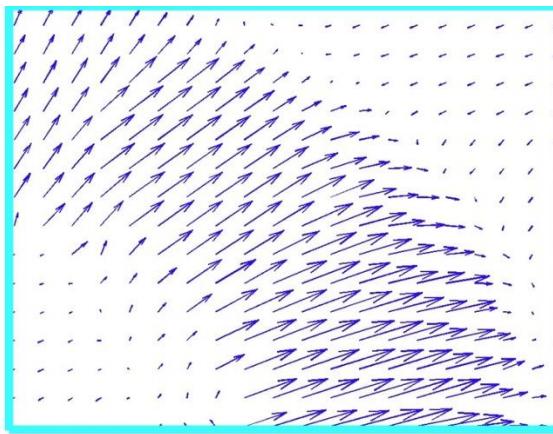


Image at frame  $t+1$

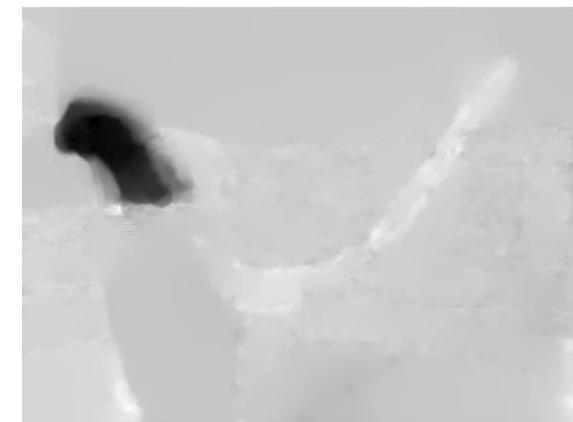
Optical flow gives a displacement field  $F$  between images  $I_t$  and  $I_{t+1}$



Tells where each pixel will move in the next frame:  
 $F(x, y) = (dx, dy)$   
 $I_{t+1}(x+dx, y+dy) = I_t(x, y)$

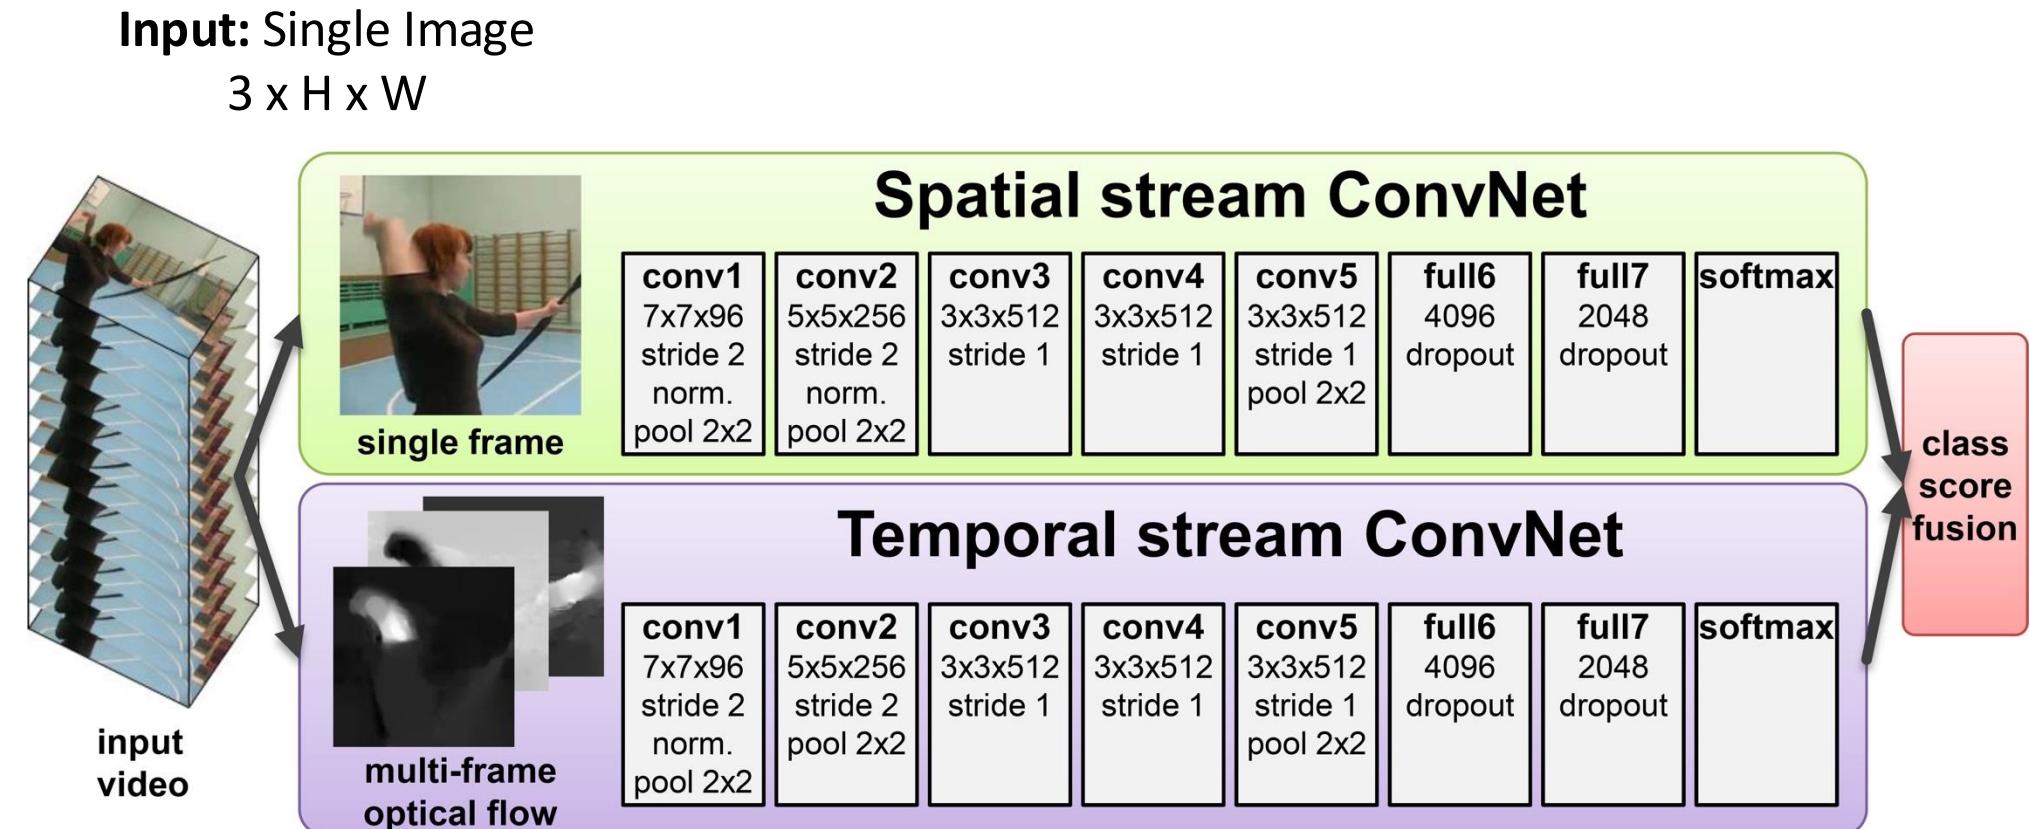
Optical Flow highlights local motion

Horizontal flow  $dx$



Vertical Flow  $dy$

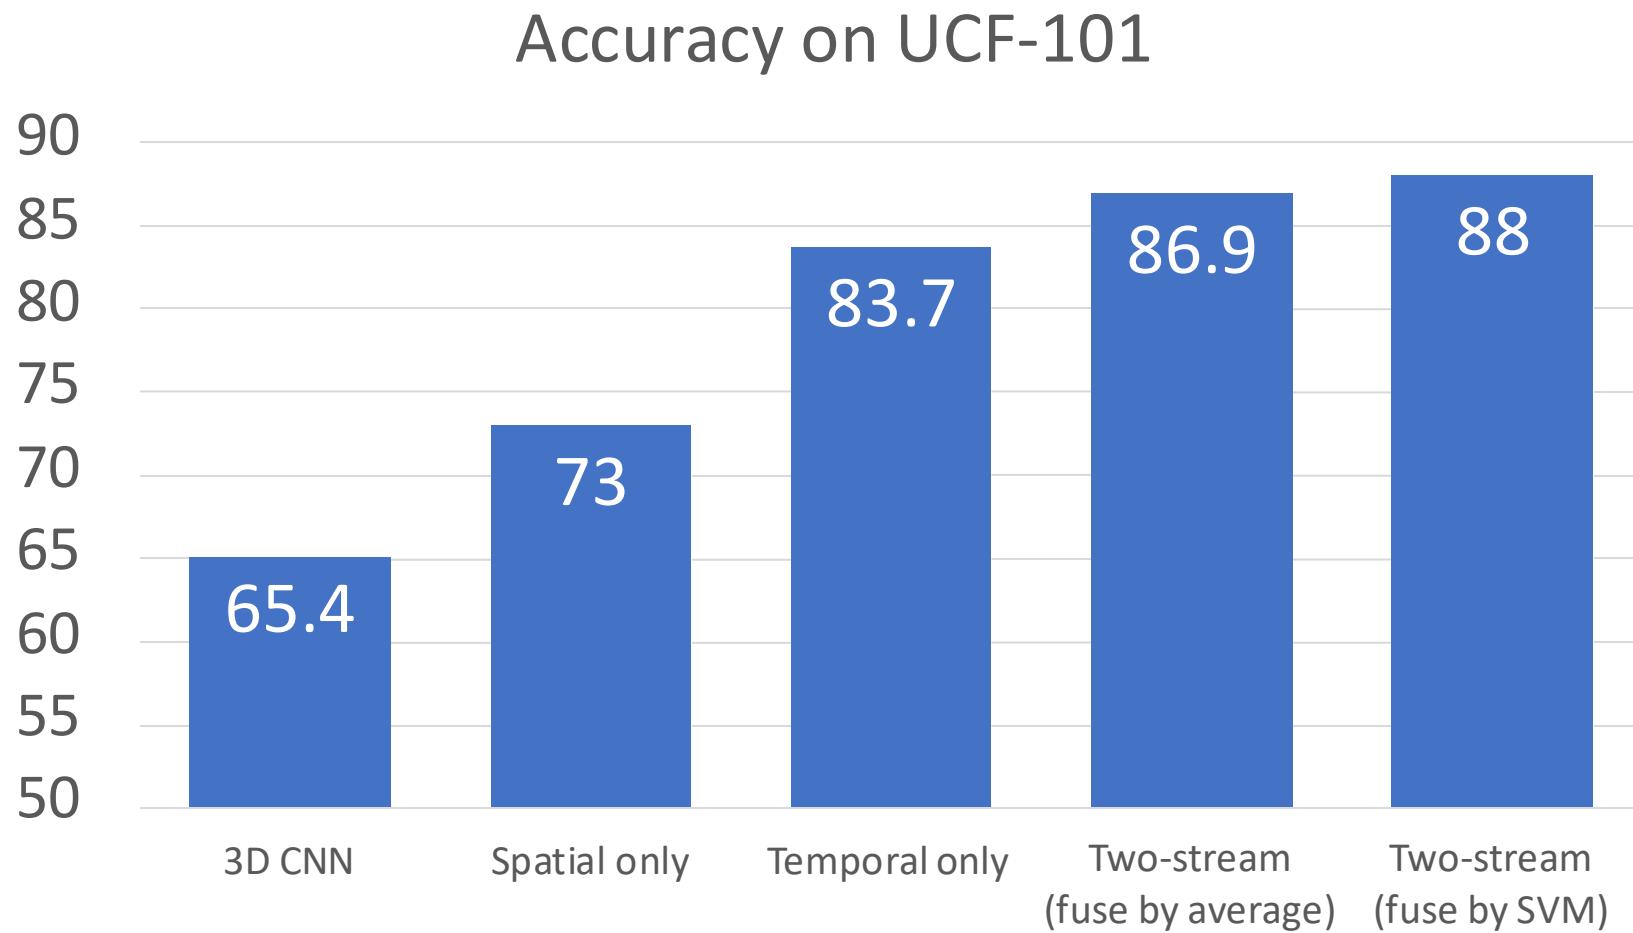
# Separating Motion and Appearance: Two-Stream Networks



**Input:** Stack of optical flow:  
 $[2*(T-1)] \times H \times W$

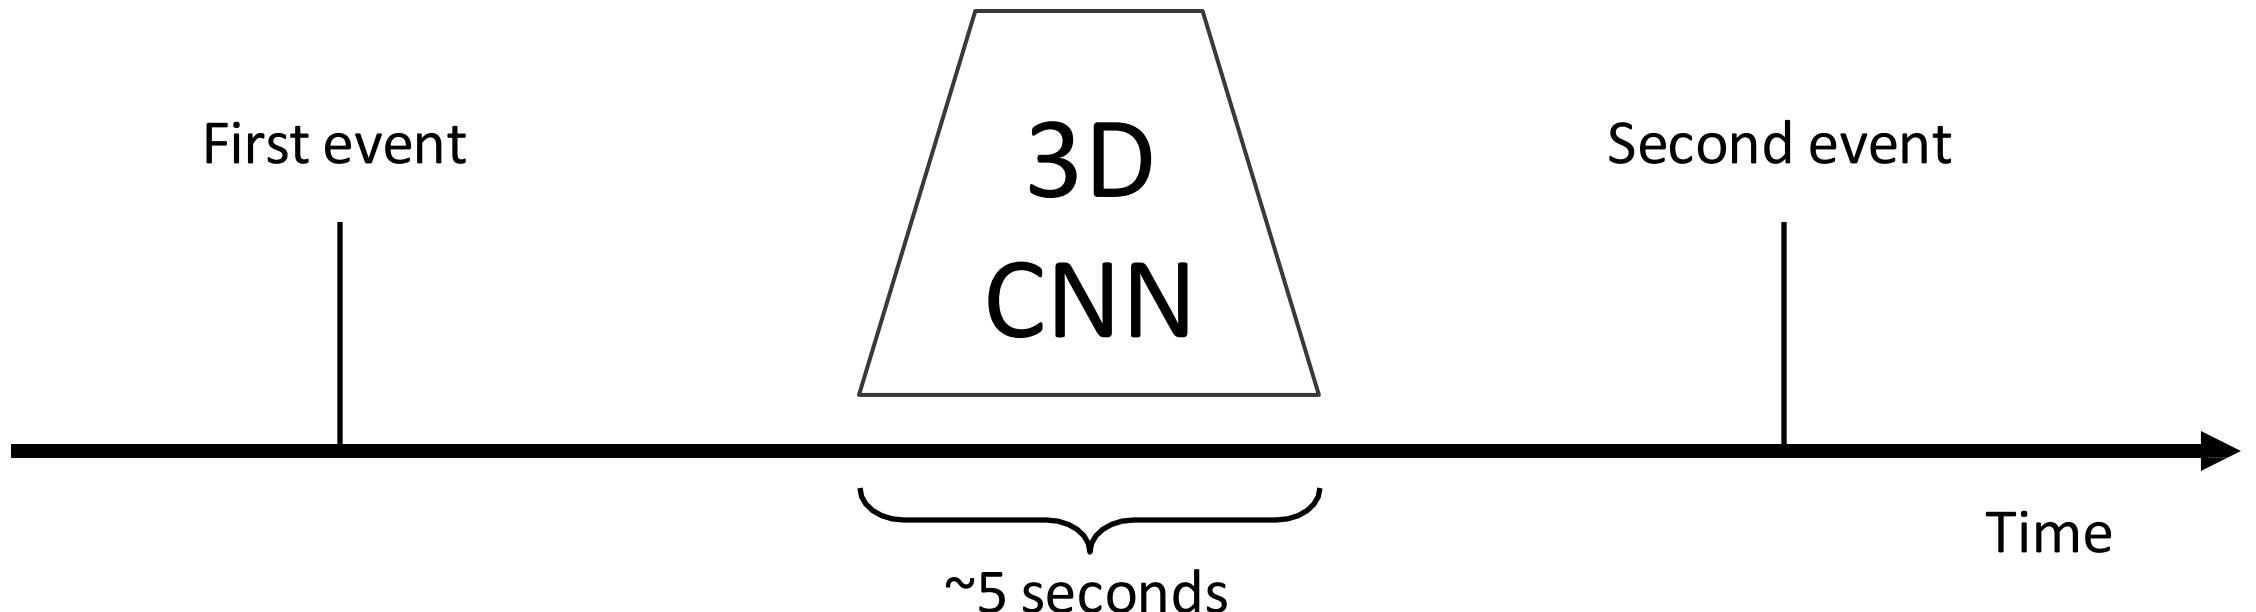
**Early fusion:** First 2D conv  
processes all flow images

# Separating Motion and Appearance: Two-Stream Networks



# Modeling long-term temporal structure

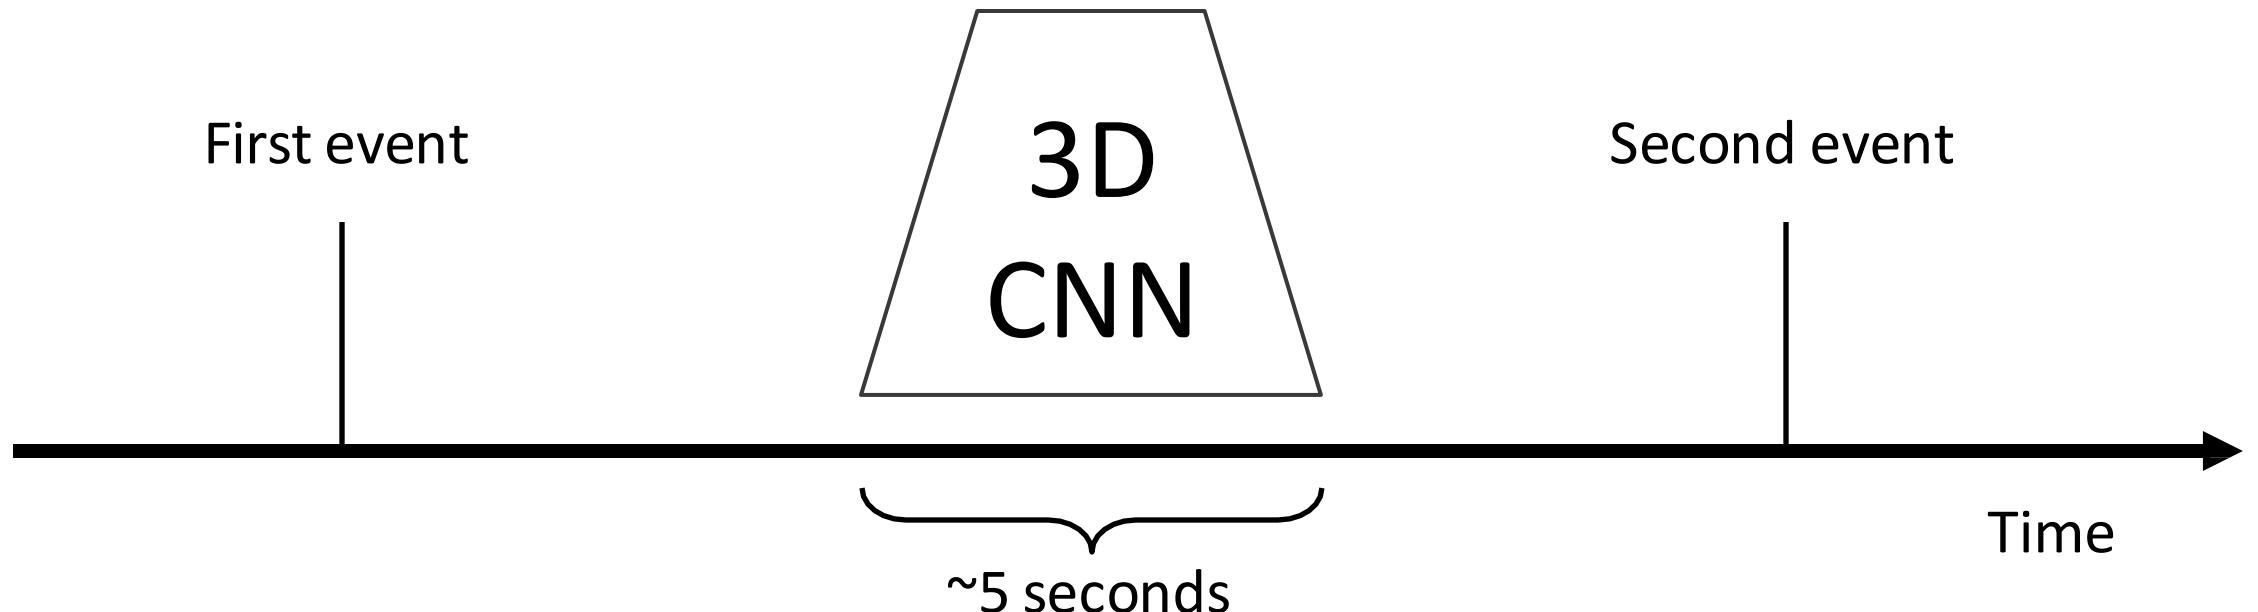
So far all our temporal CNNs only model local motion between frames in very short clips of ~2-5 seconds. What about long-term structure?



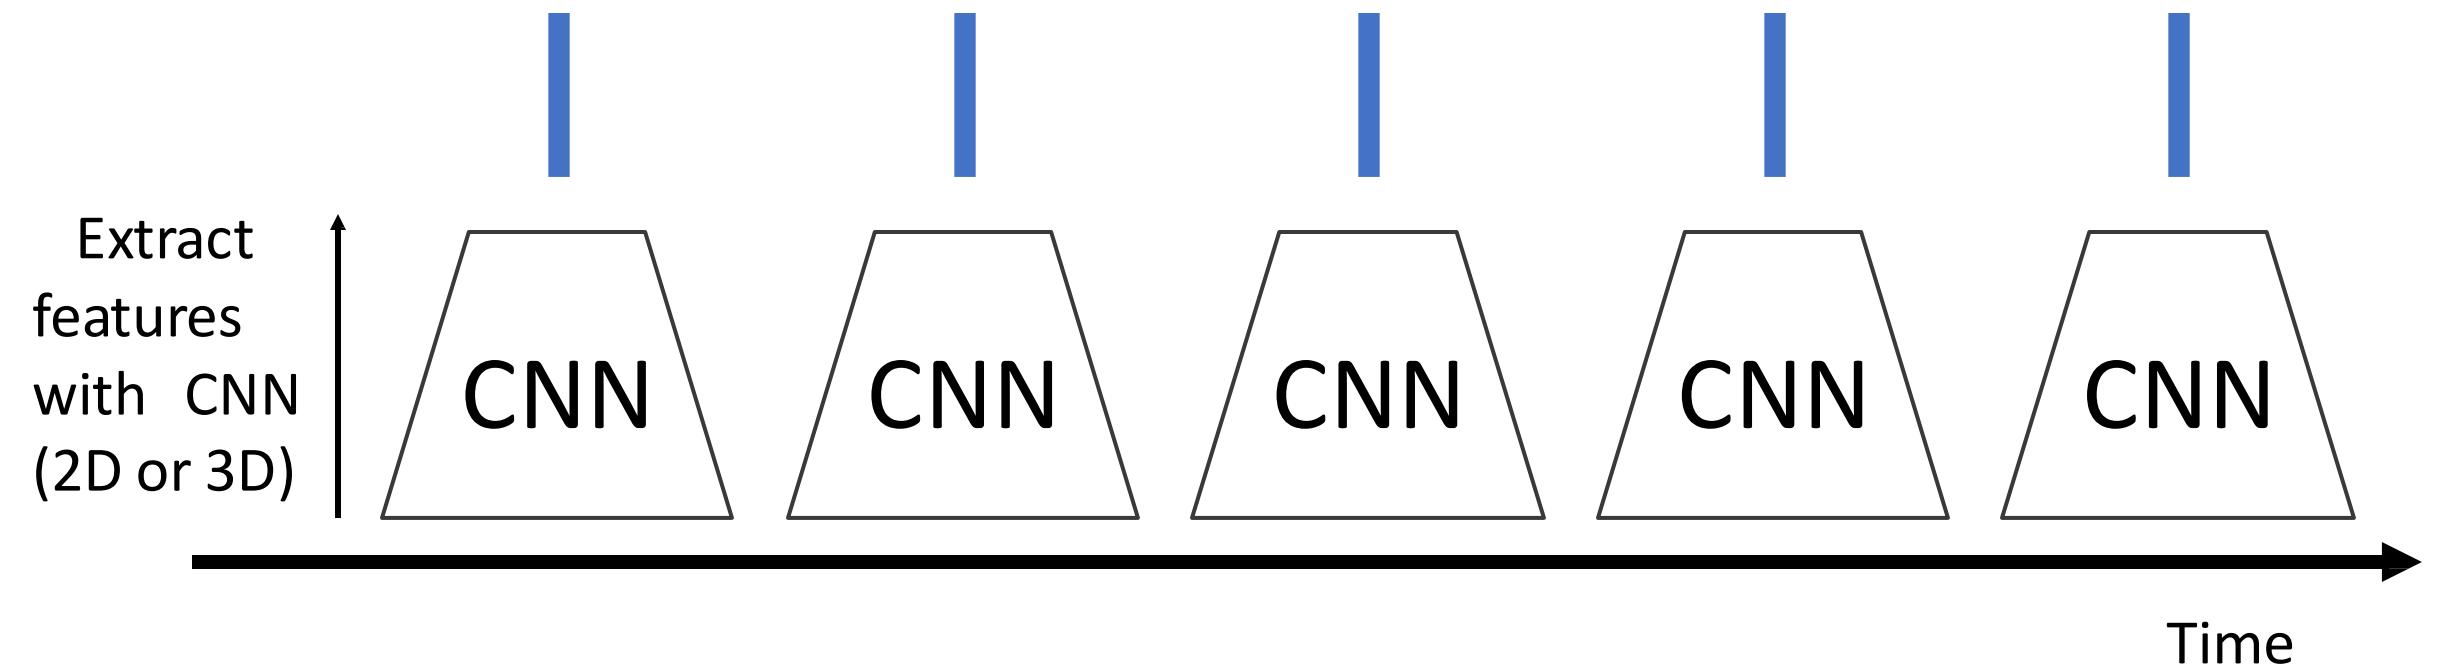
# Modeling long-term temporal structure

So far all our temporal CNNs only model local motion between frames in very short clips of ~2-5 seconds. What about long-term structure?

We know how to handle sequences!  
How about recurrent networks?

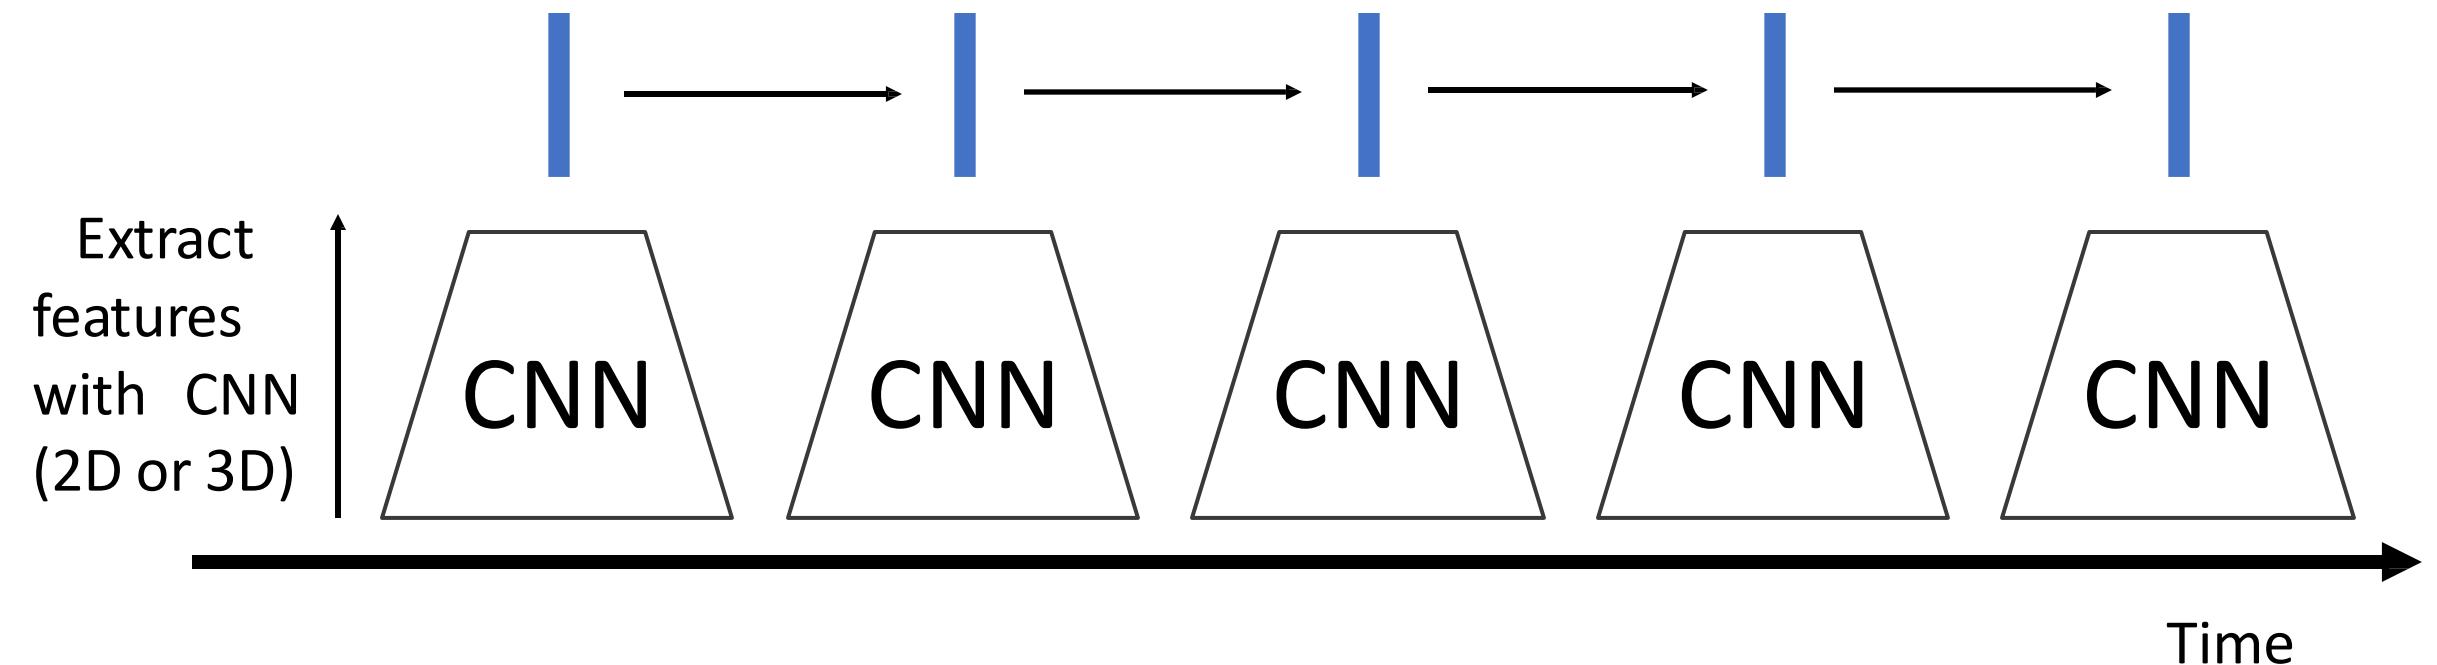


# Modeling long-term temporal structure



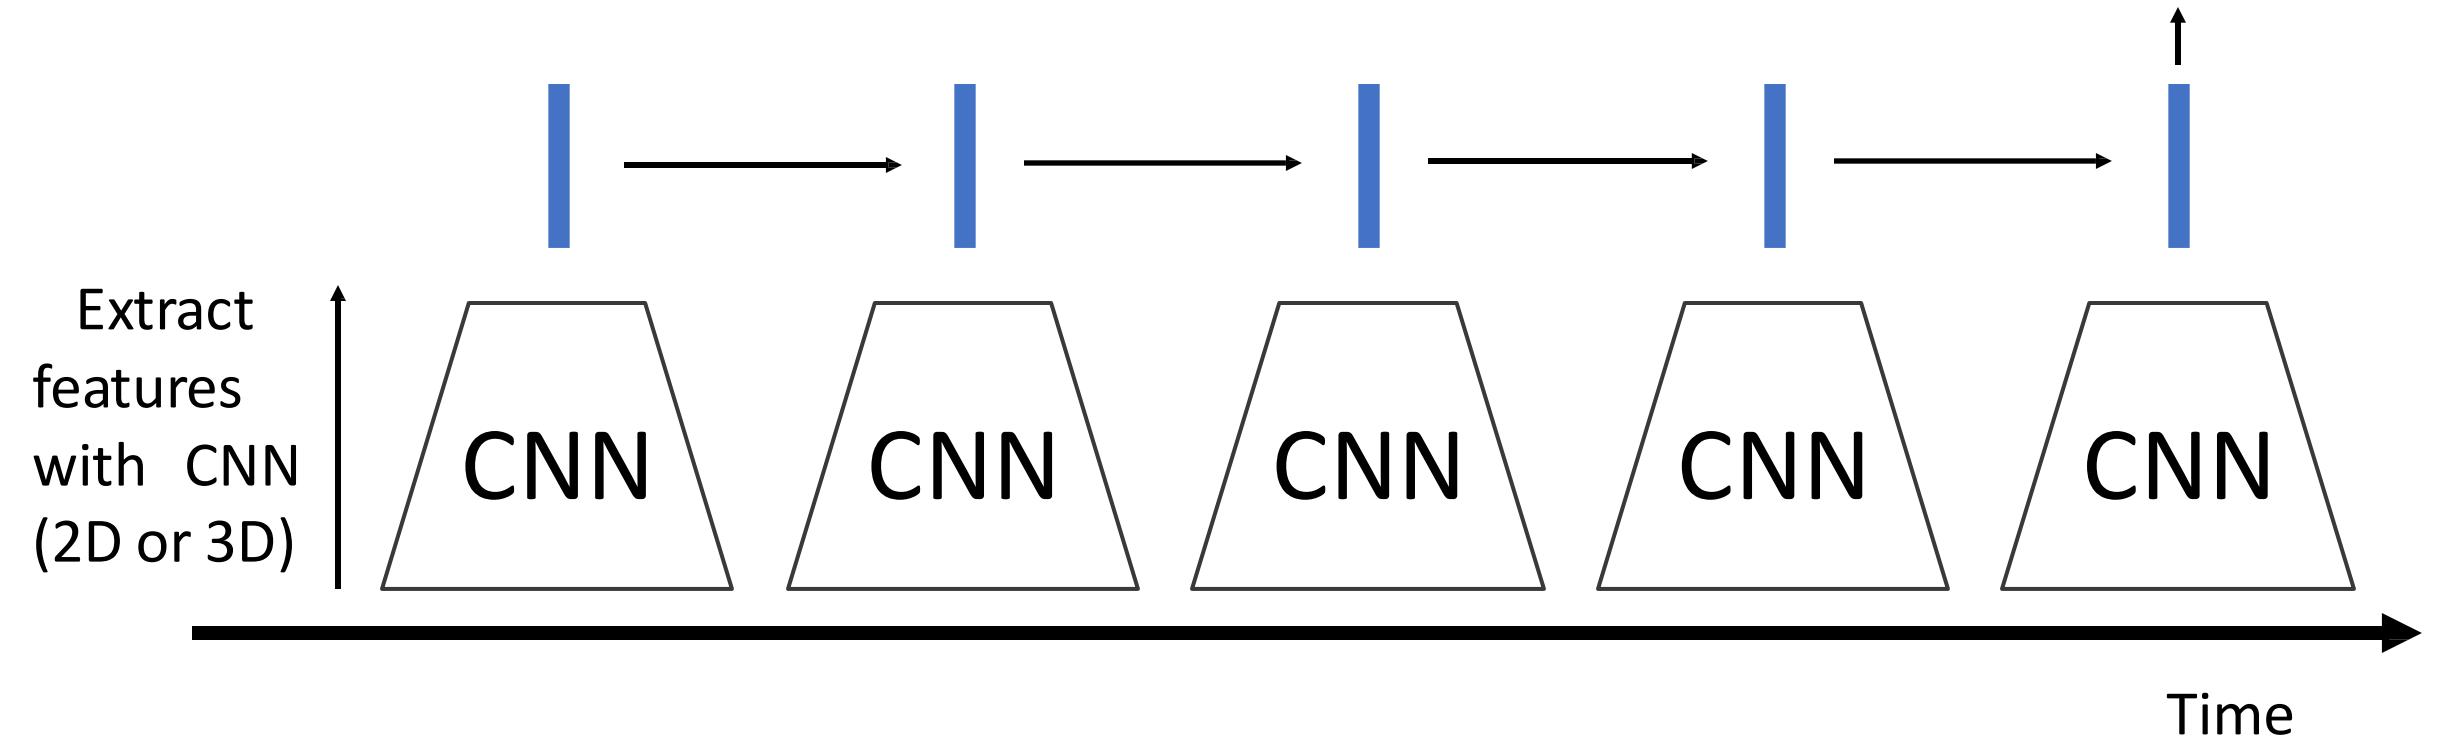
# Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)



# Modeling long-term temporal structure

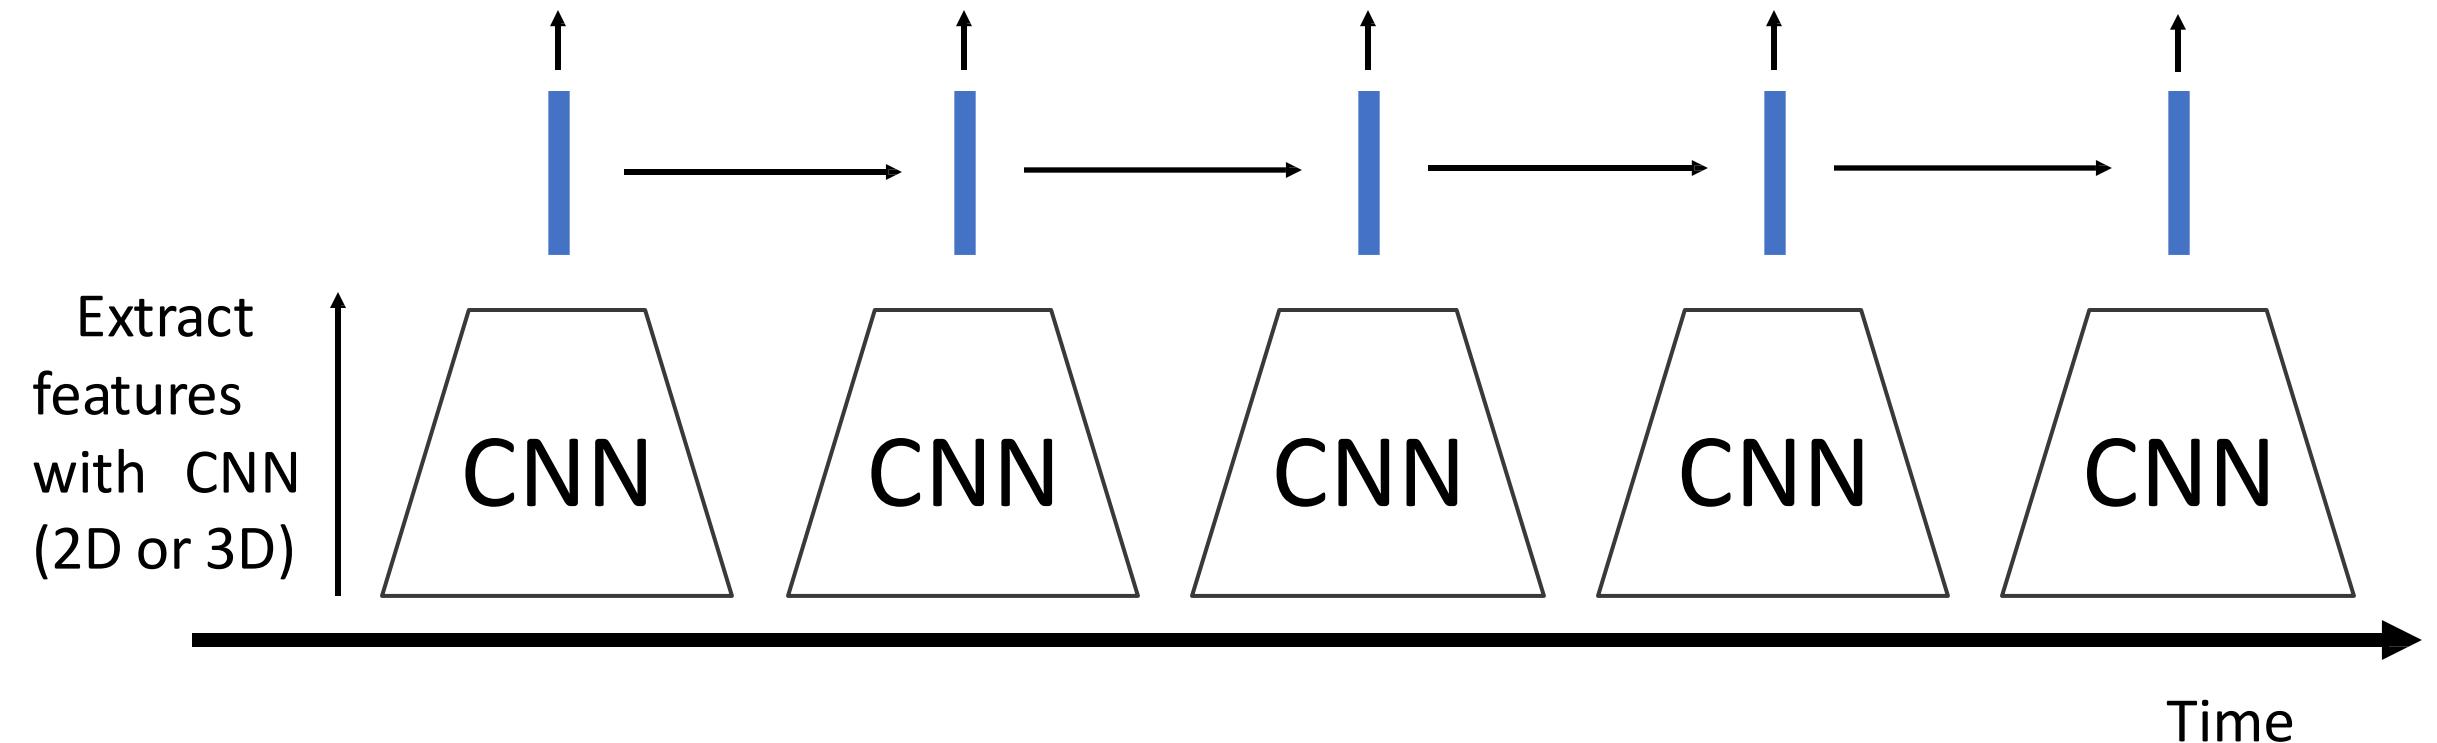
Process local features using recurrent network (e.g. LSTM)  
Many to one: One output at end of video



# Modeling long-term temporal structure

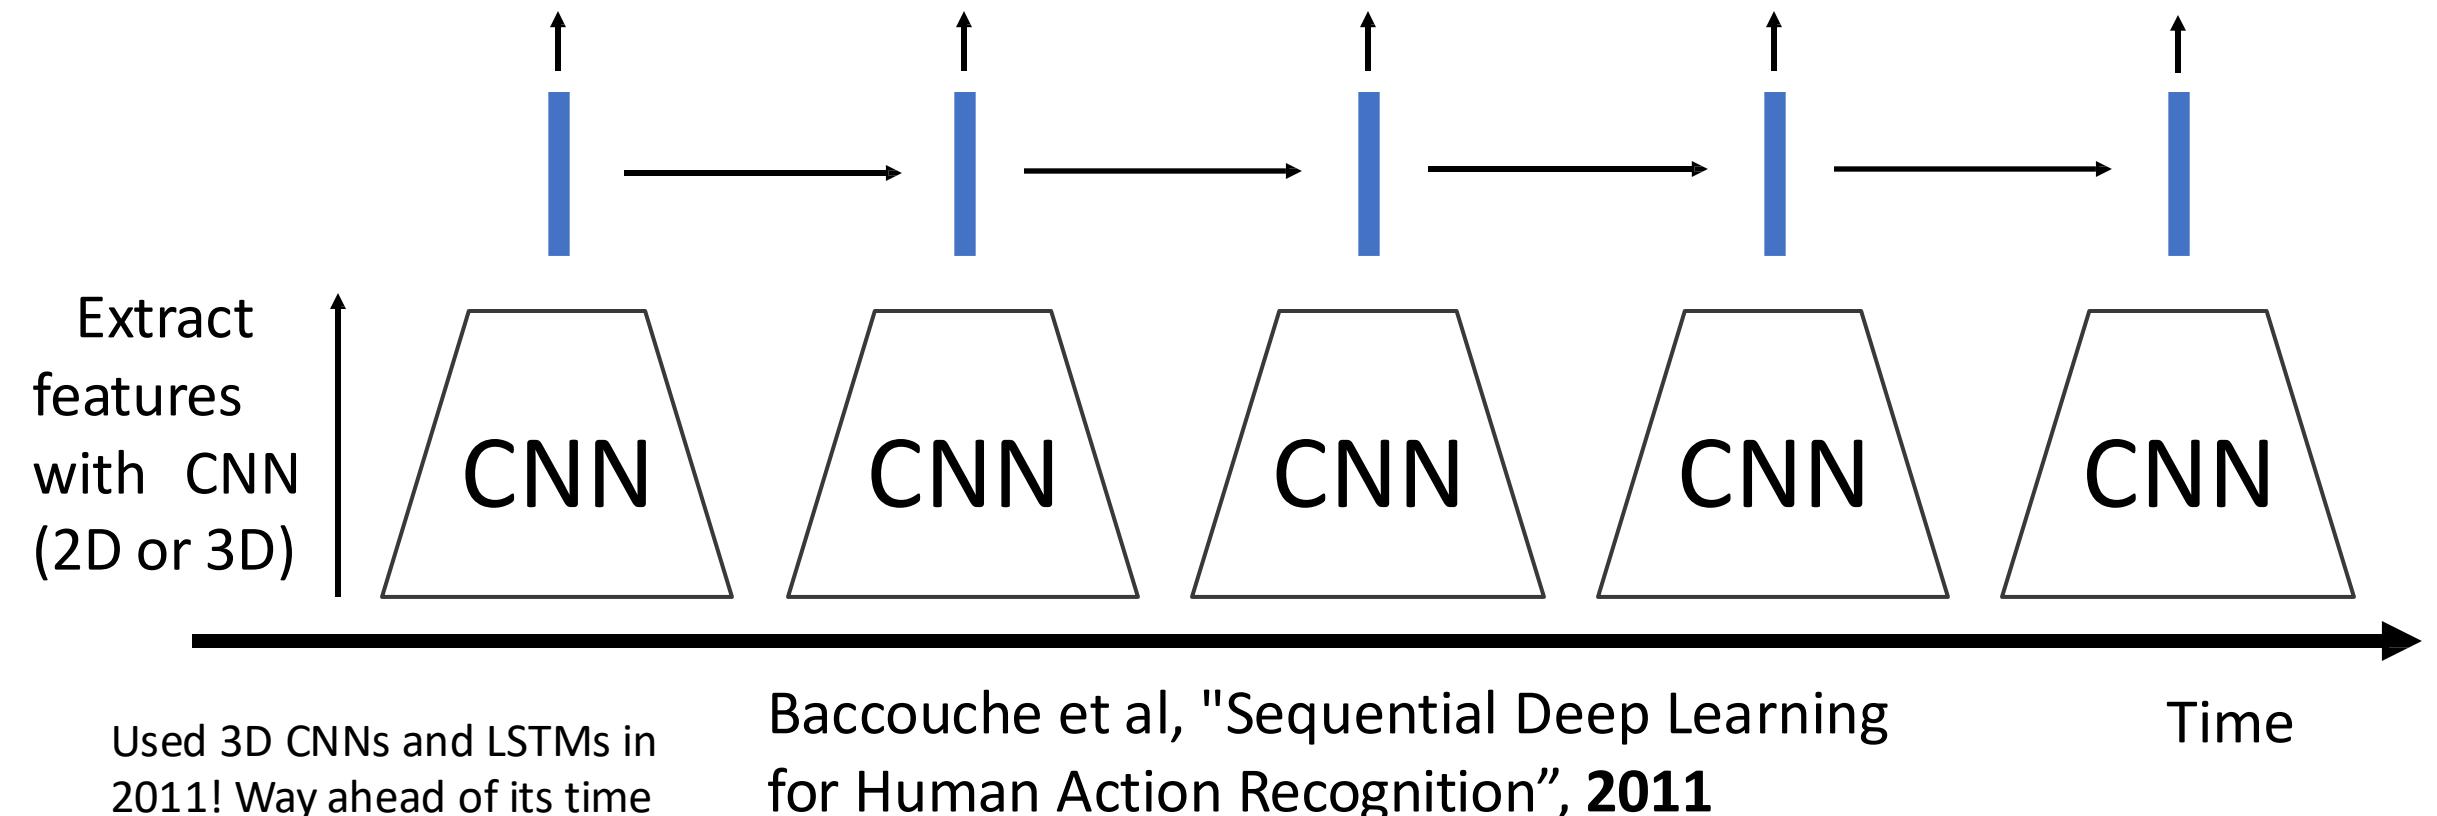
Process local features using recurrent network (e.g. LSTM)

Many to many: one output per video frame



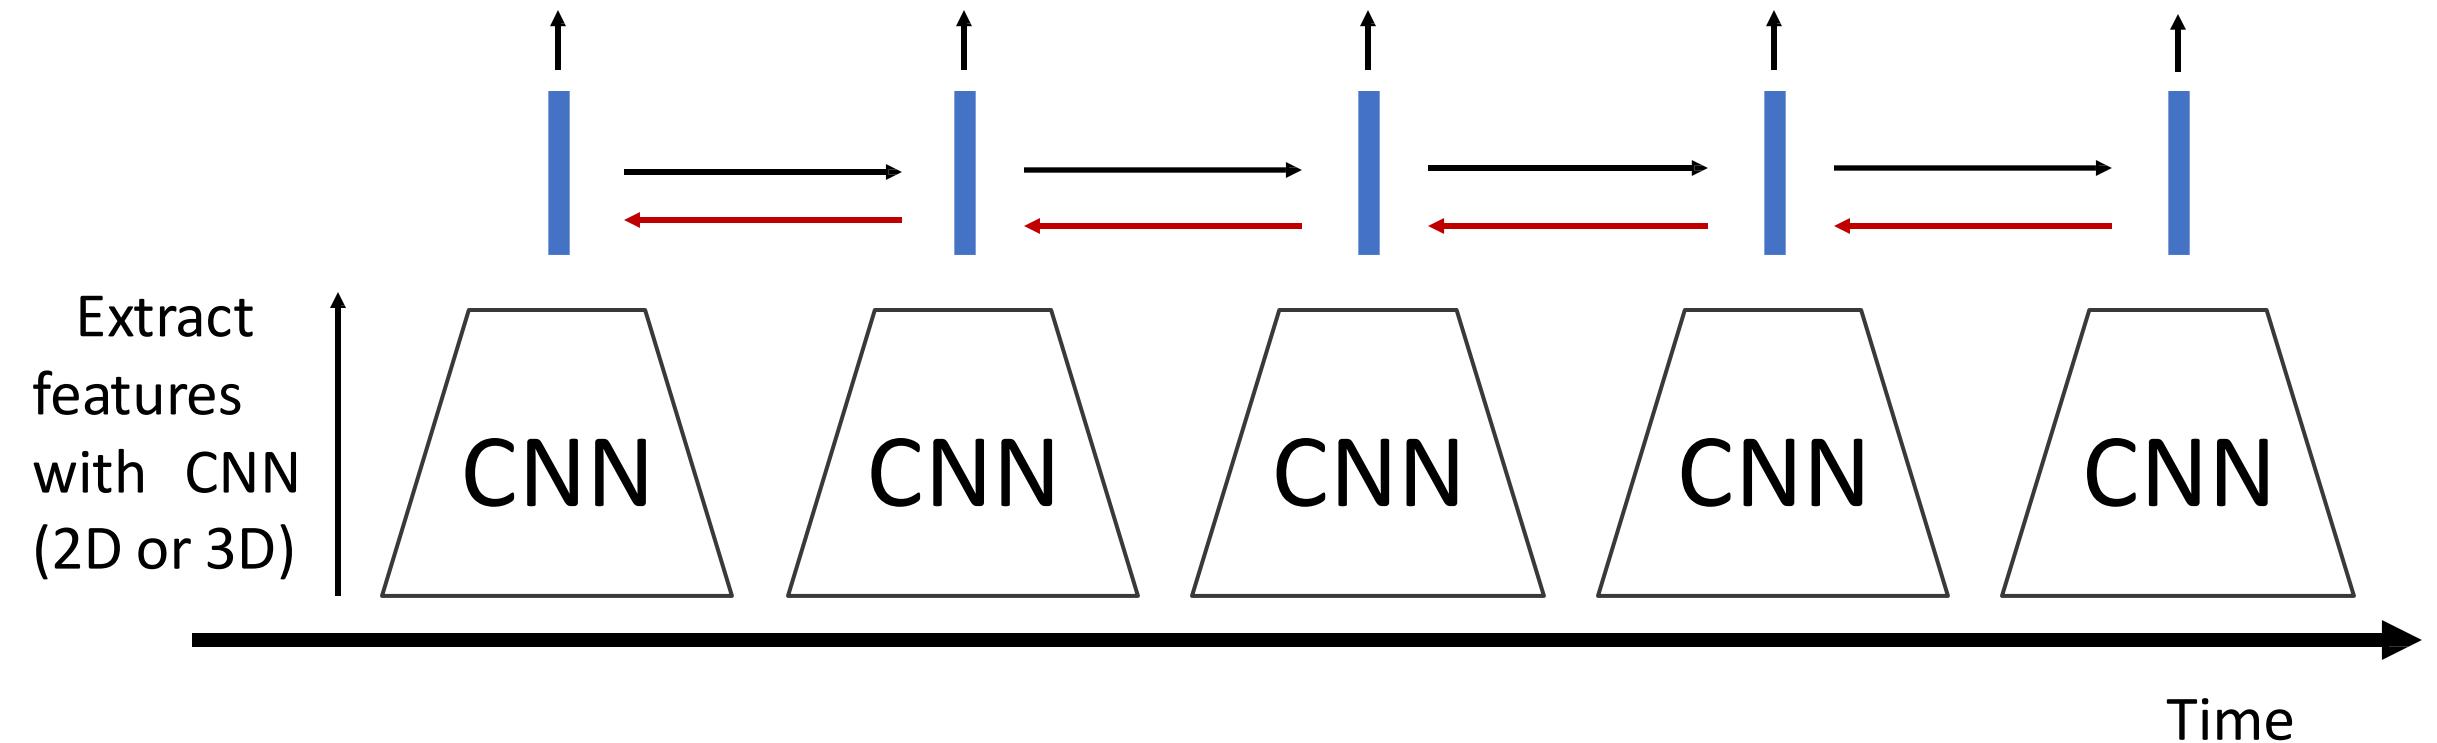
# Modeling long-term temporal structure

Process local features using recurrent network (e.g. LSTM)  
Many to many: one output per video frame



# Modeling long-term temporal structure

Sometimes don't backprop to CNN to save memory;  
pretrain and use it as a feature extractor



Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011

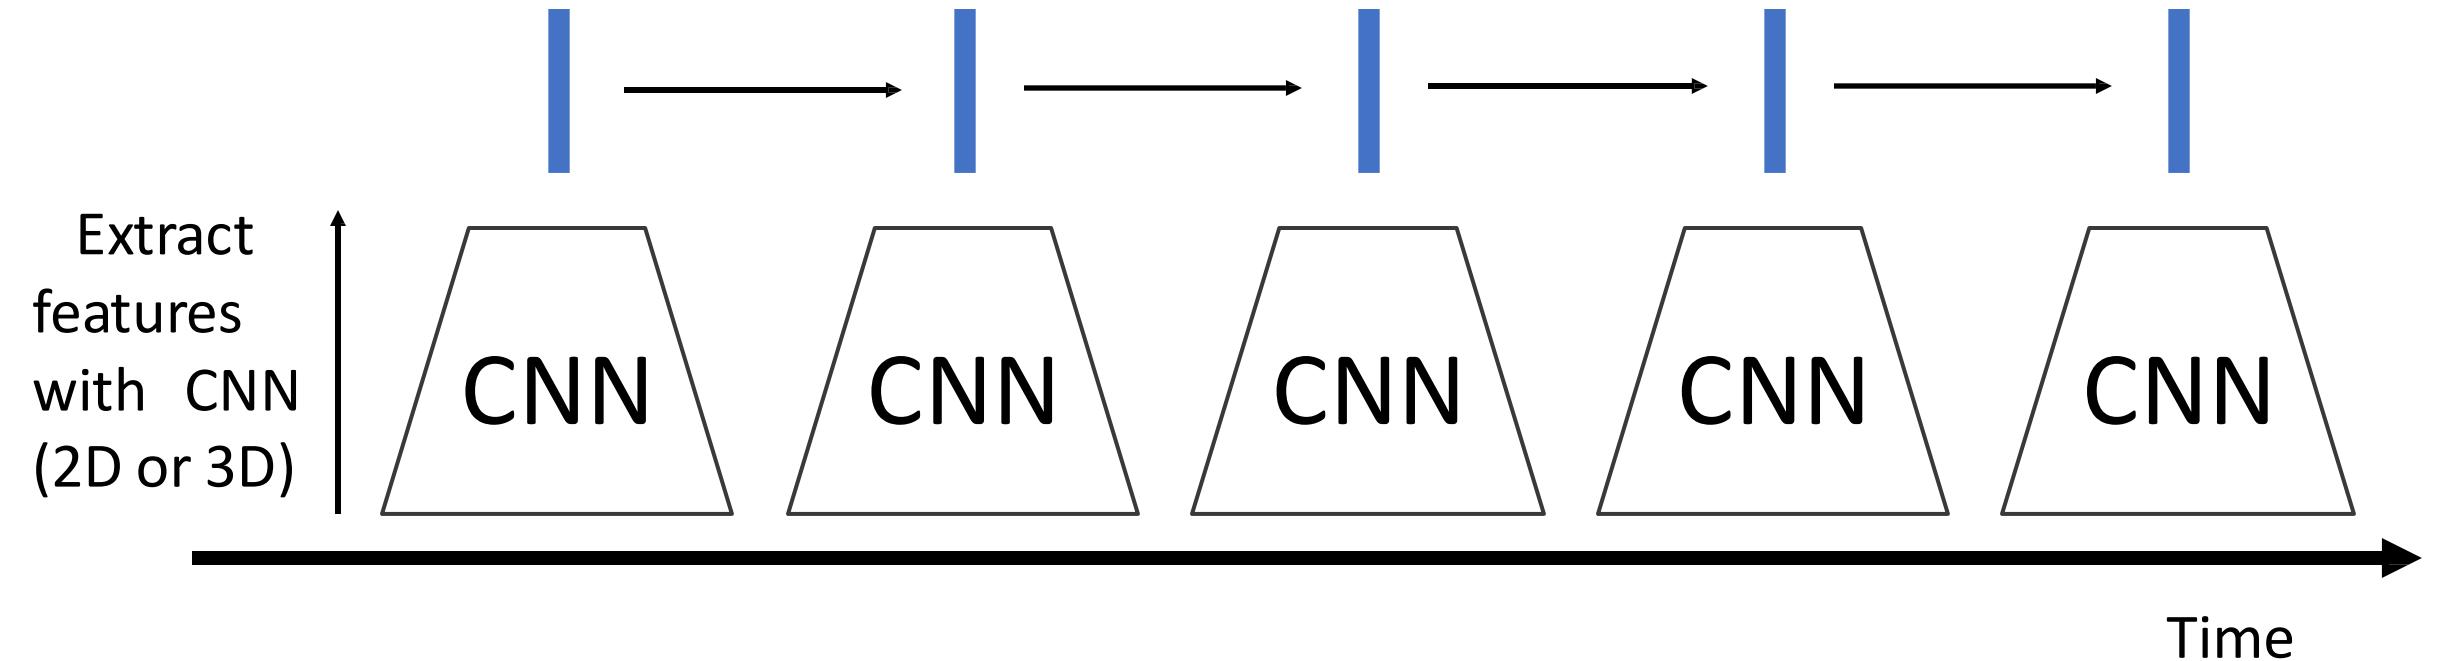
Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

# Modeling long-term temporal structure

Inside CNN: Each value a function of a fixed temporal window (local temporal structure)

Inside RNN: Each vector is a function of all previous vectors (global temporal structure)

Can we merge both approaches?



Baccouche et al, "Sequential Deep Learning for Human Action Recognition", 2011

Donahue et al, "Long-term recurrent convolutional networks for visual recognition and description", CVPR 2015

# Recall: Self-Attention

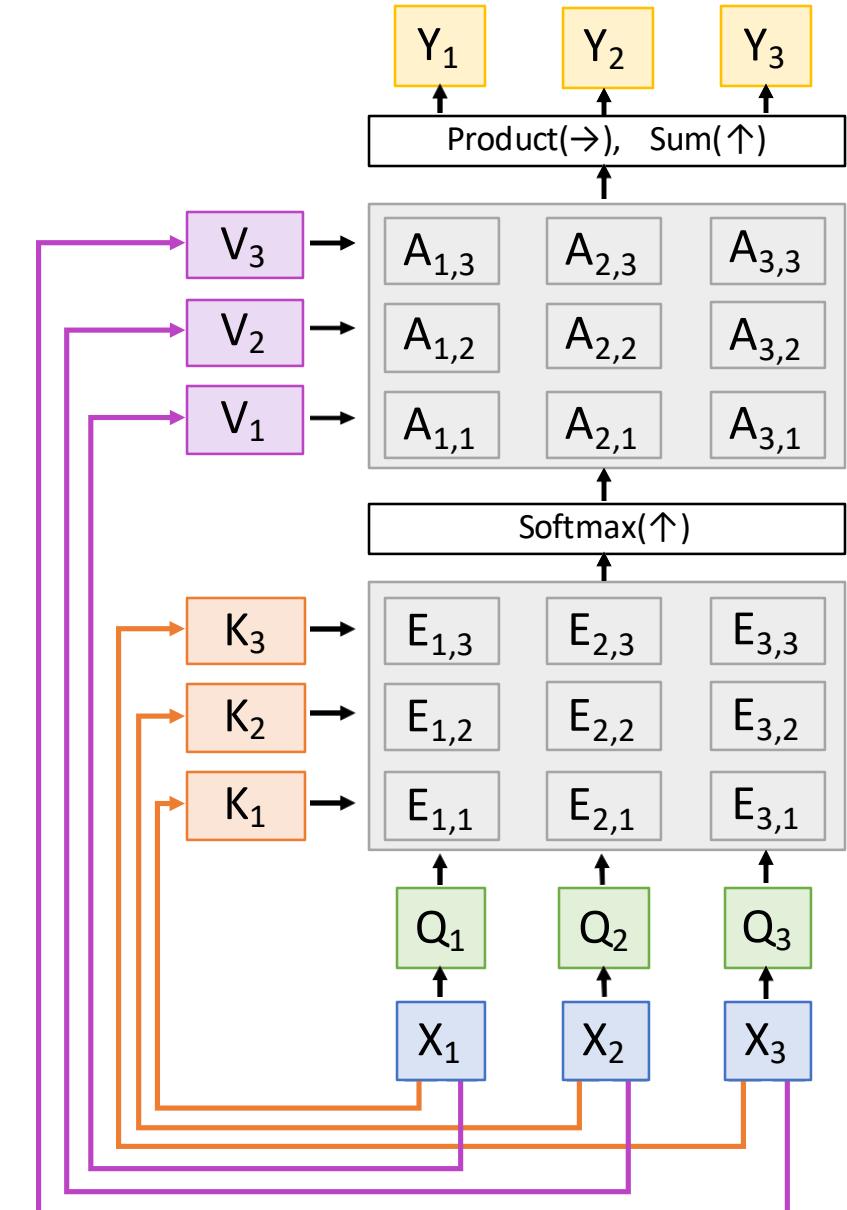
**Input:** Set of vectors  $x_1, \dots, x_N$

**Keys, Queries, Values:** Project each  $x$  to a key, query, and value using linear layer

**Affinity matrix:** Compare each pair of  $x$ , (using scaled dot-product between keys and values) and normalize using softmax

**Output:** Weighted sum of values, with weights given by affinity matrix

Features in 3D CNN:  $C \times T \times H \times W$   
Interpret as a set of THW vectors of dim C

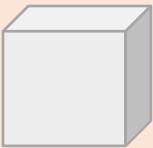


# Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



3D  
CNN

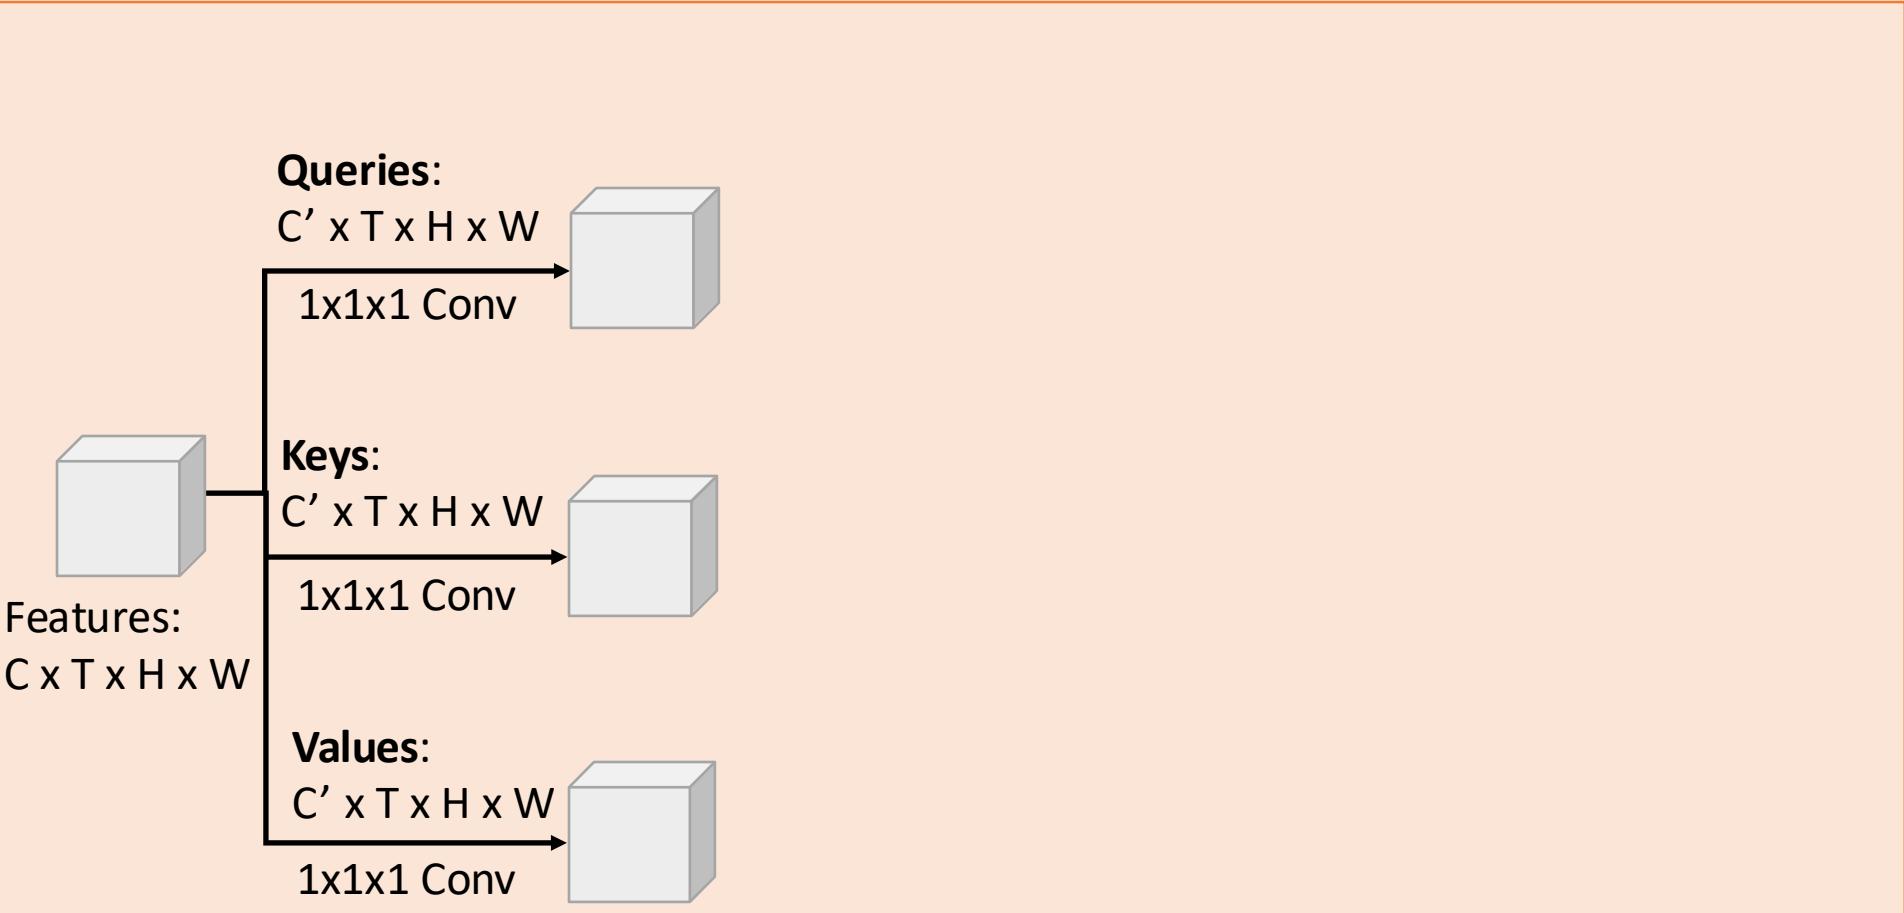


Features:  
 $C \times T \times H \times W$

Nonlocal Block

# Spatio-Temporal Self-Attention (Nonlocal Block)

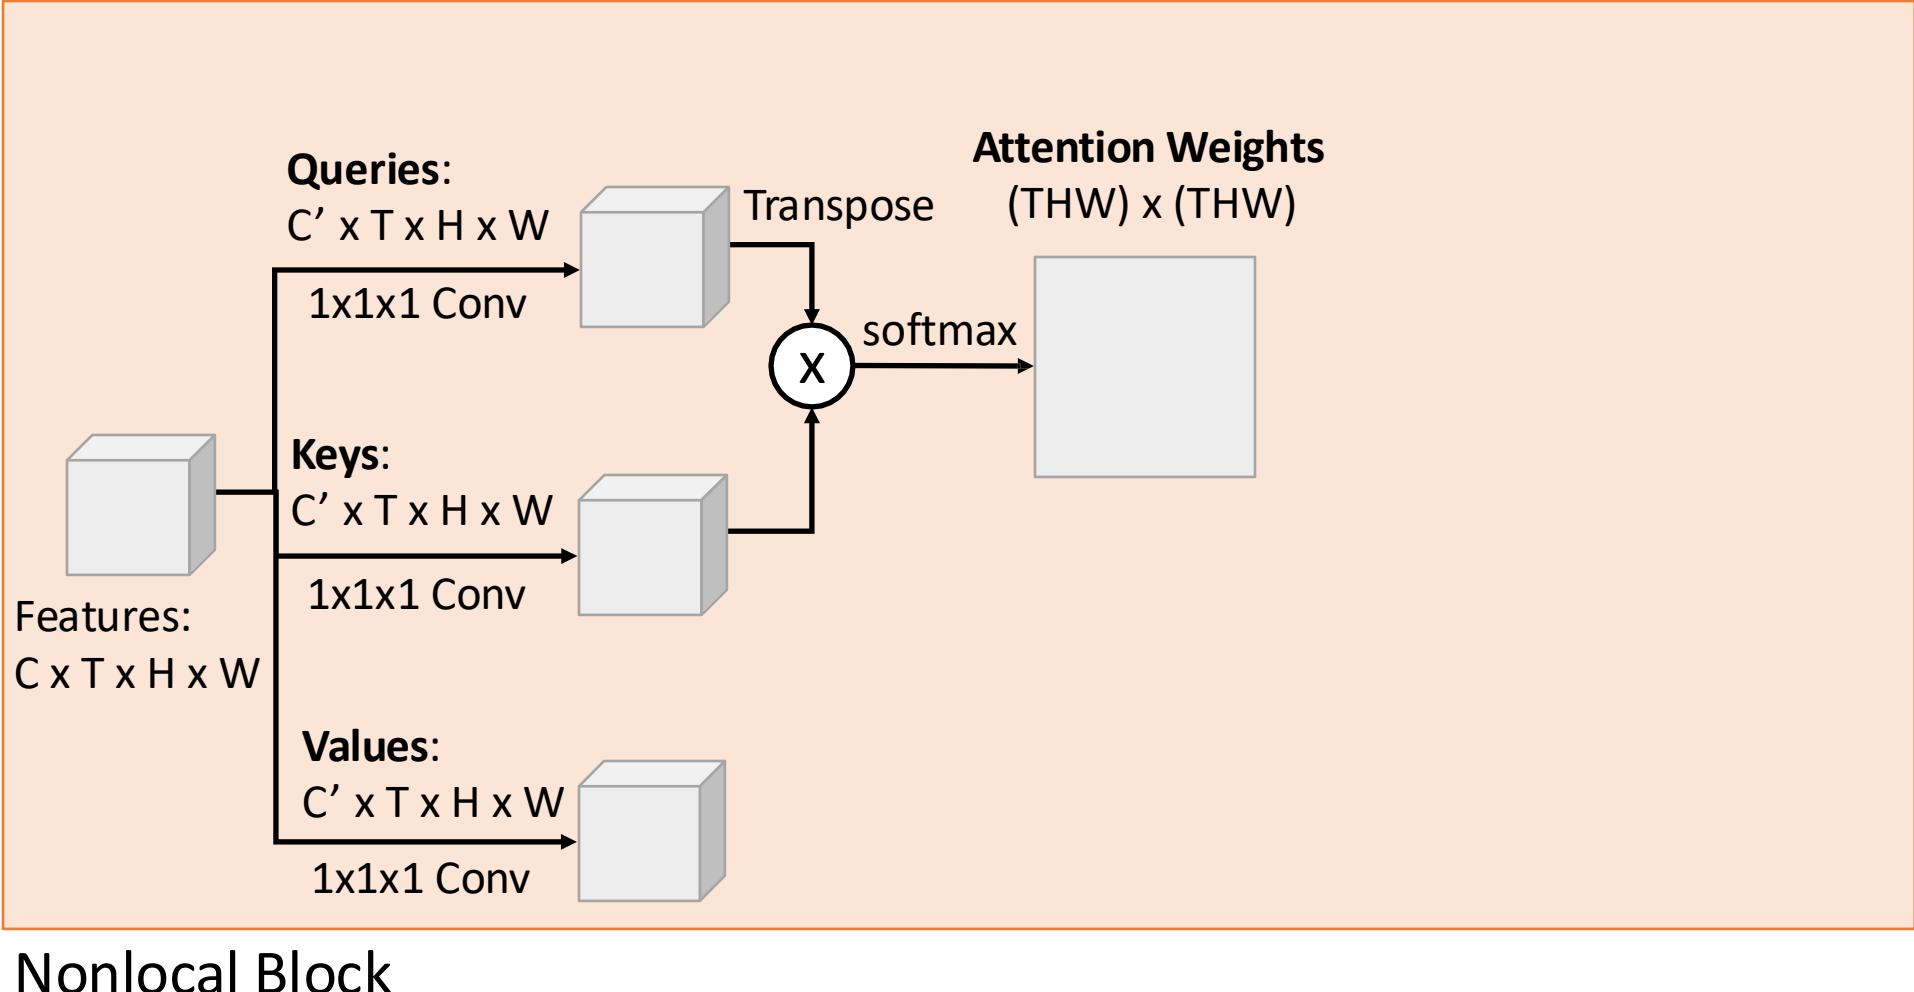
Input clip



Nonlocal Block

# Spatio-Temporal Self-Attention (Nonlocal Block)

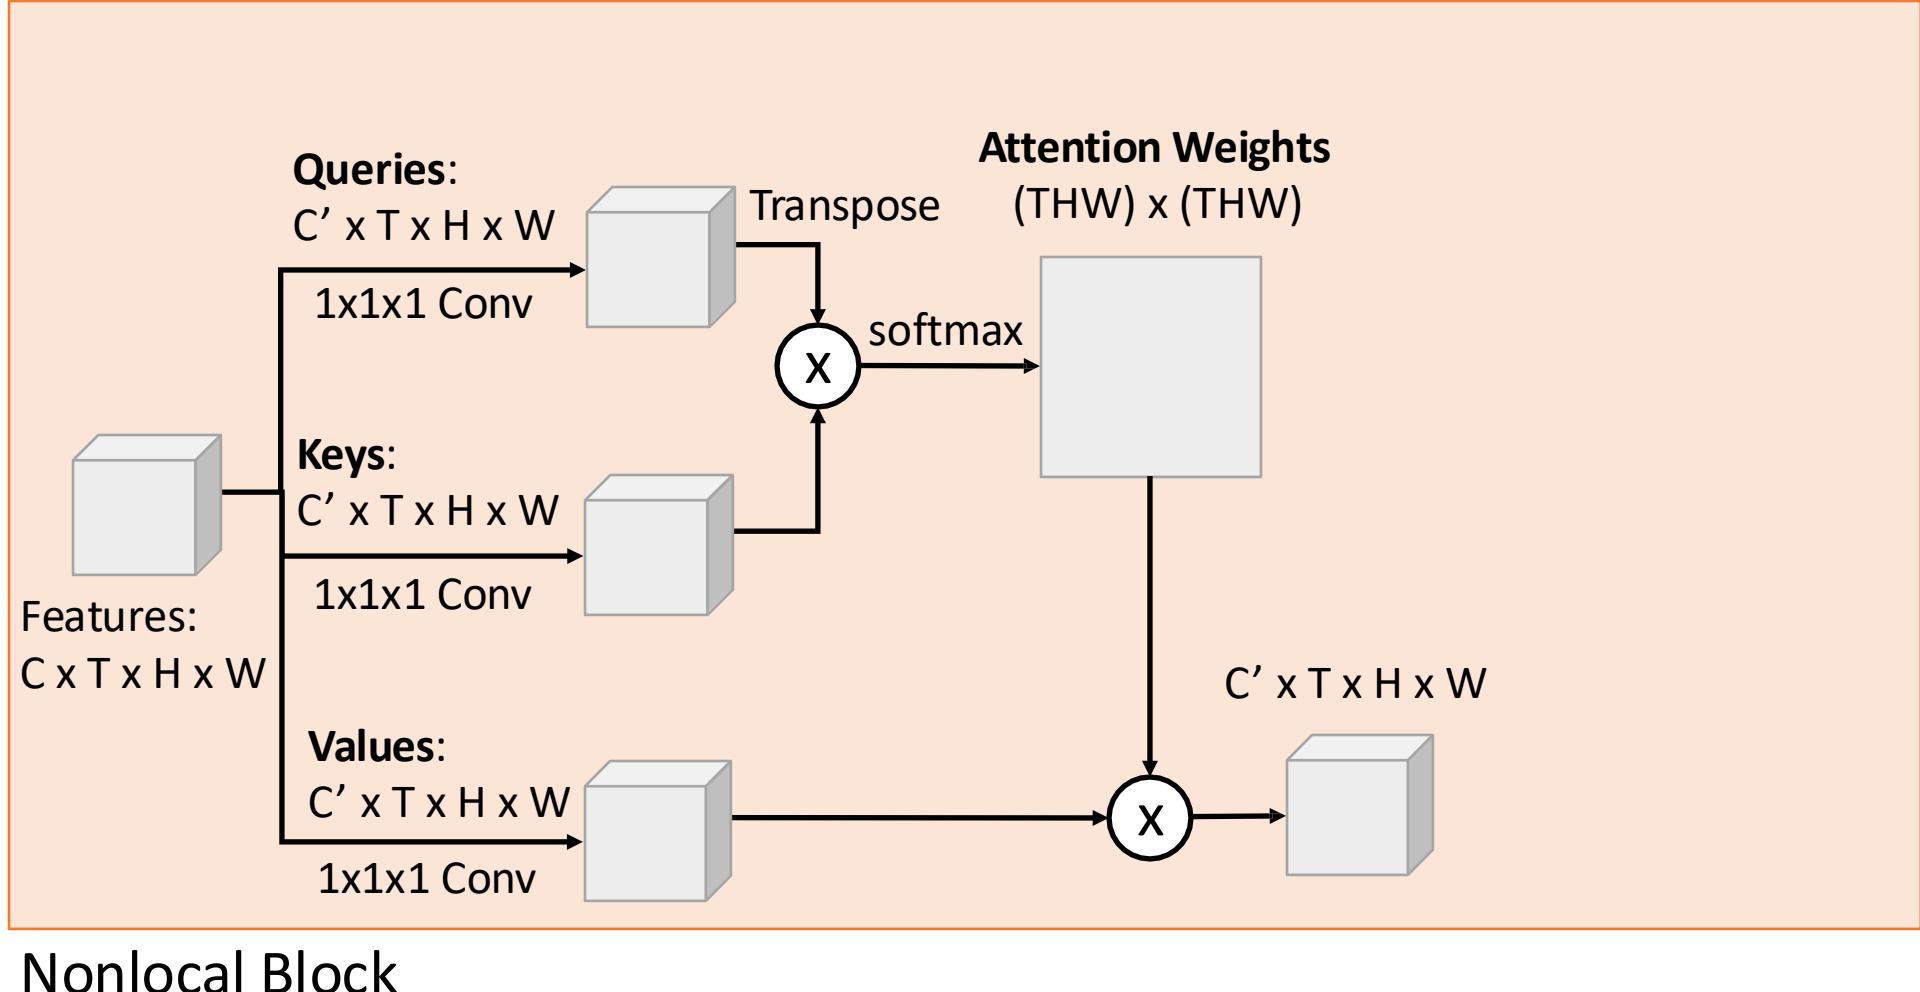
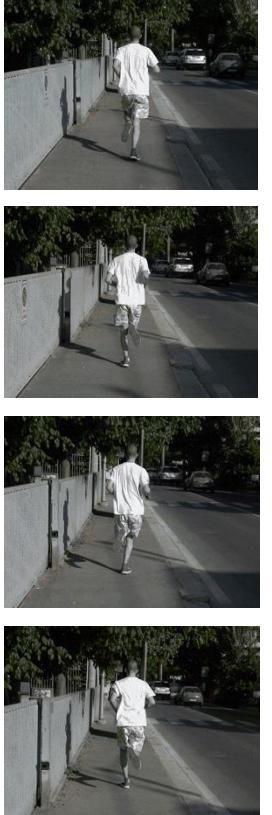
Input clip



Nonlocal Block

# Spatio-Temporal Self-Attention (Nonlocal Block)

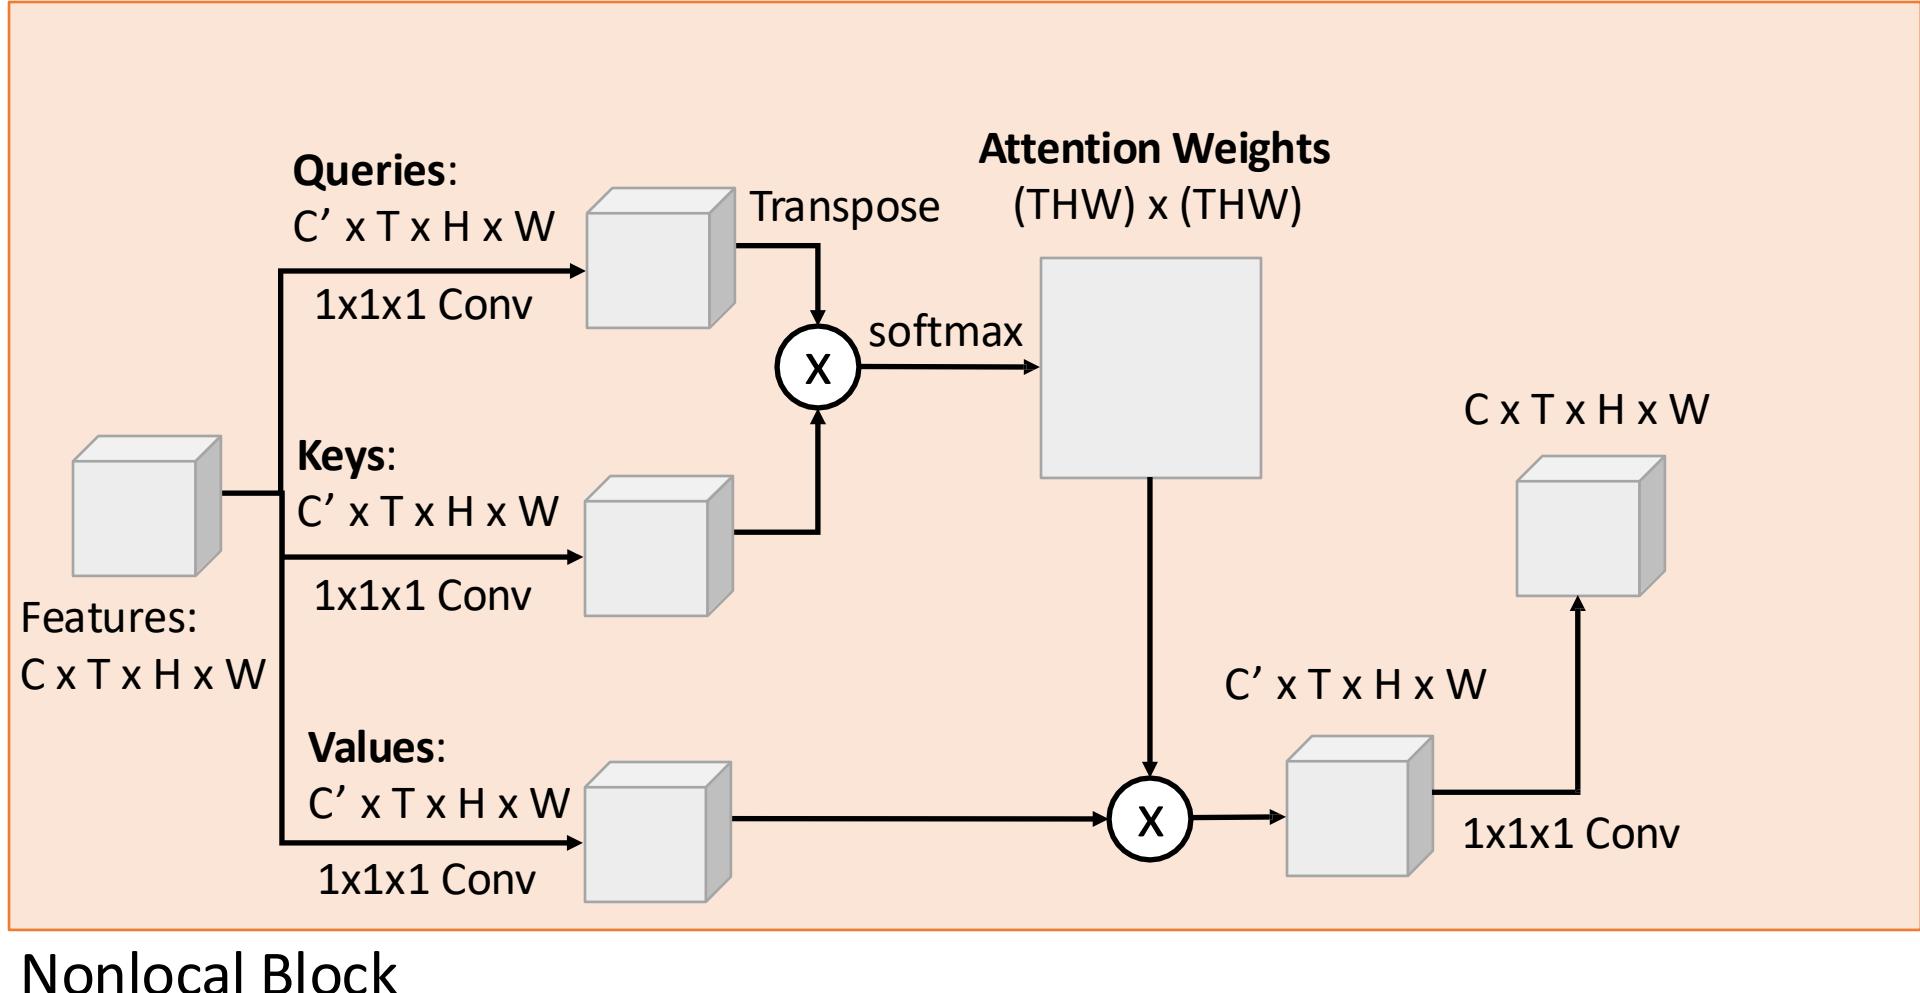
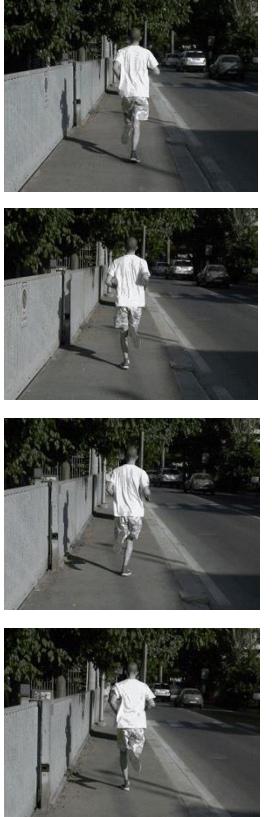
Input clip



Nonlocal Block

# Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



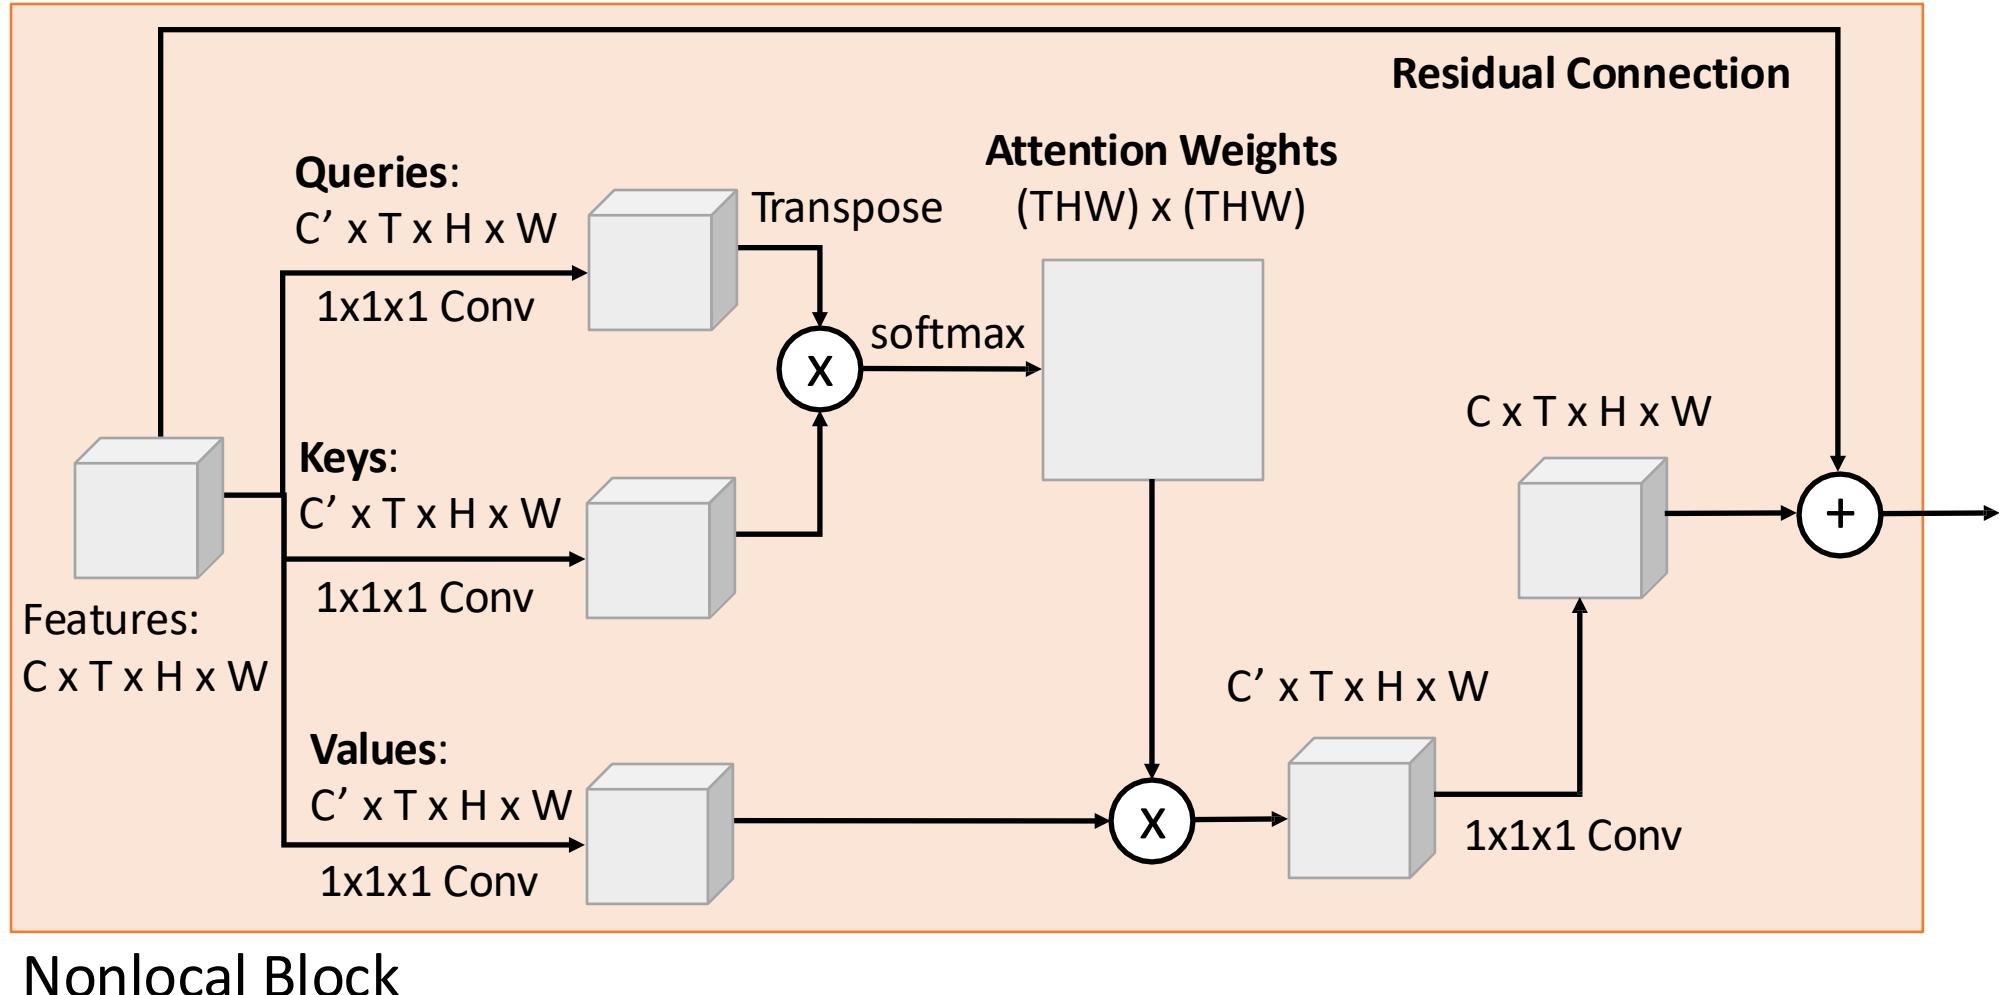
Nonlocal Block

# Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

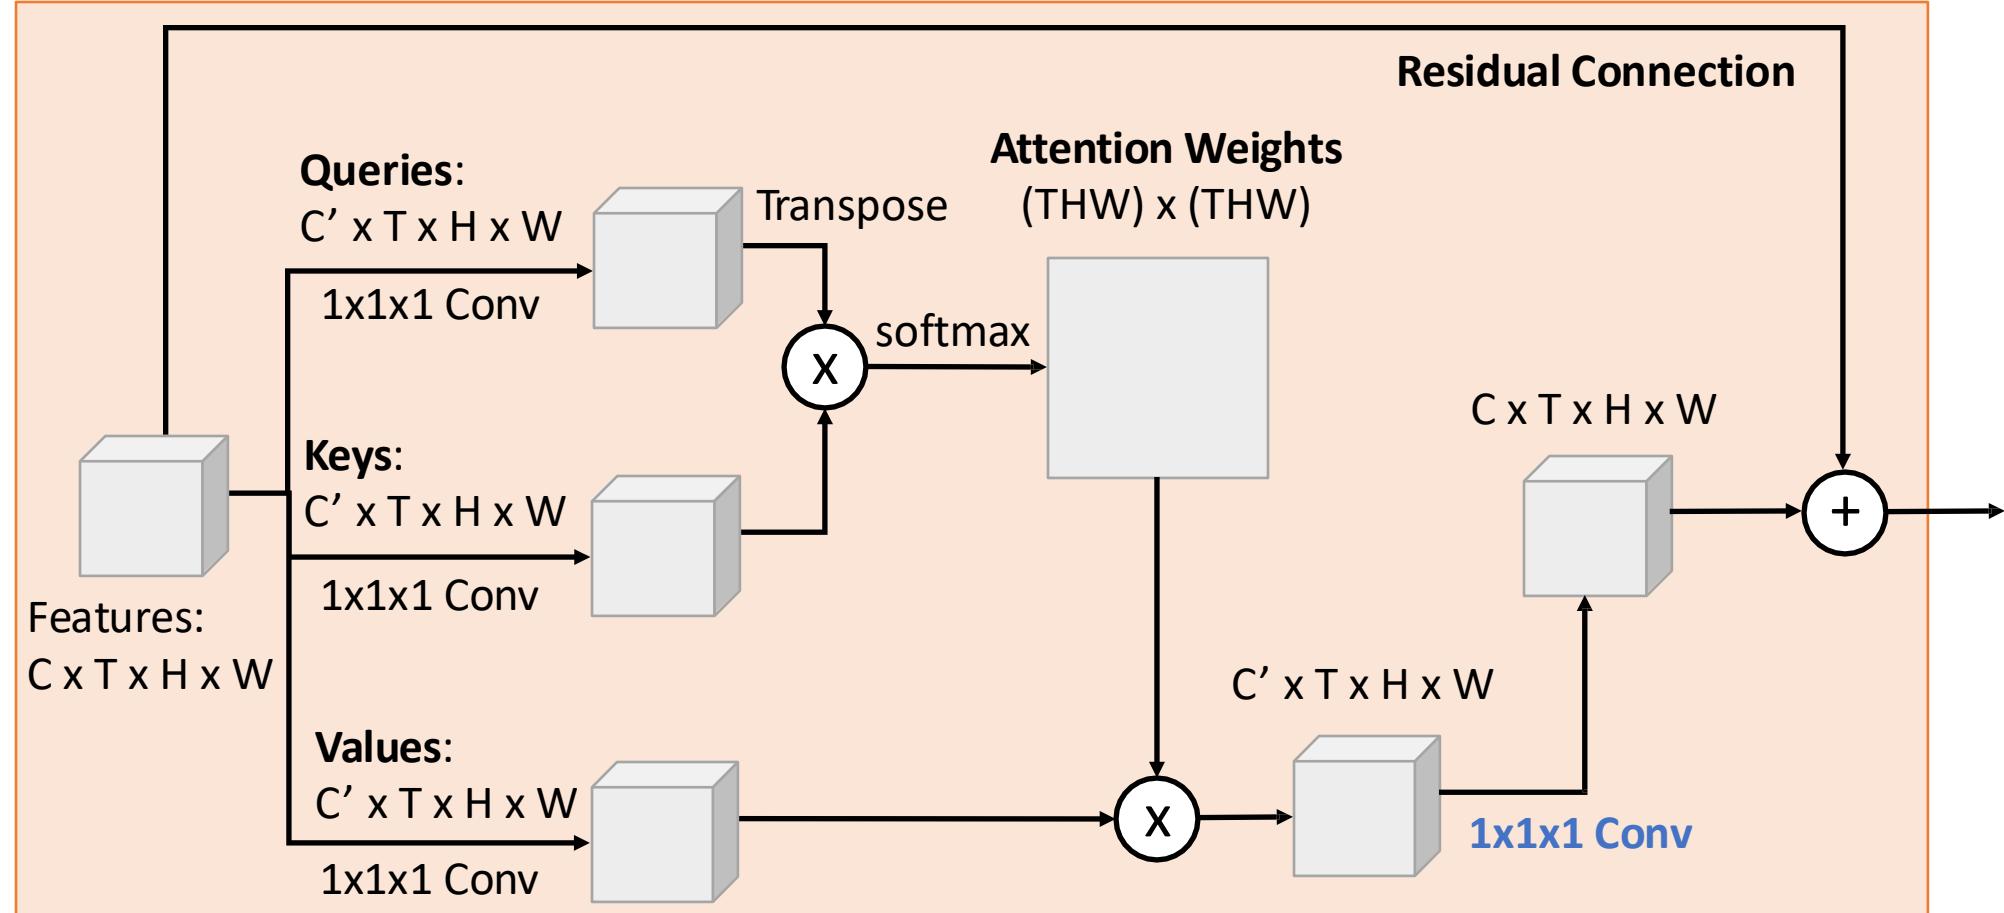
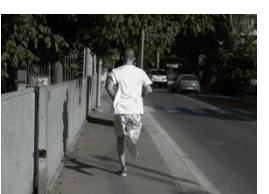


3D  
CNN



# Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip



Nonlocal Block

**Trick:** Initialize **last conv** to 0, then entire block computes identity. Can insert into existing 3D CNNs

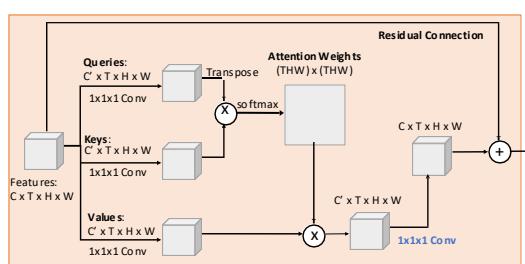
# Spatio-Temporal Self-Attention (Nonlocal Block)

Input clip

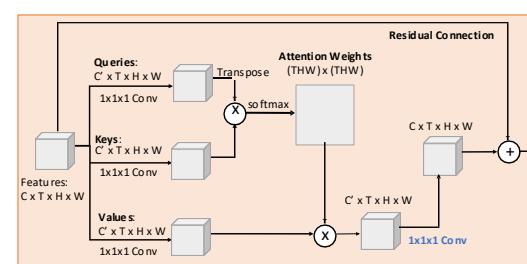
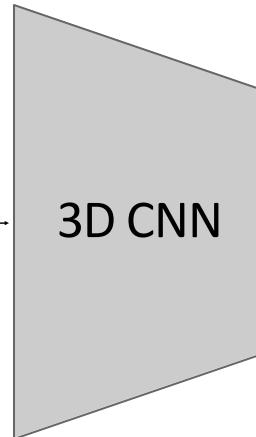


3D CNN

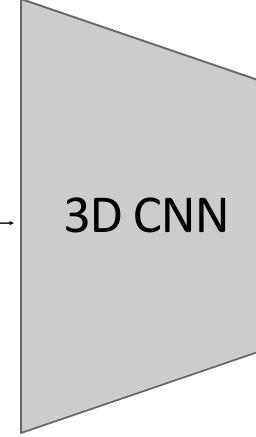
We can add nonlocal blocks into existing 3D CNN architectures.  
But what is the best 3D CNN architecture?



Nonlocal Block



Nonlocal Block



Running

# Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.  
Can we reuse image architectures for video?

**Idea:** take a 2D CNN architecture.

Replace each 2D  $K_h \times K_w$  conv/pool  
layer with a 3D  $K_t \times K_h \times K_w$  version

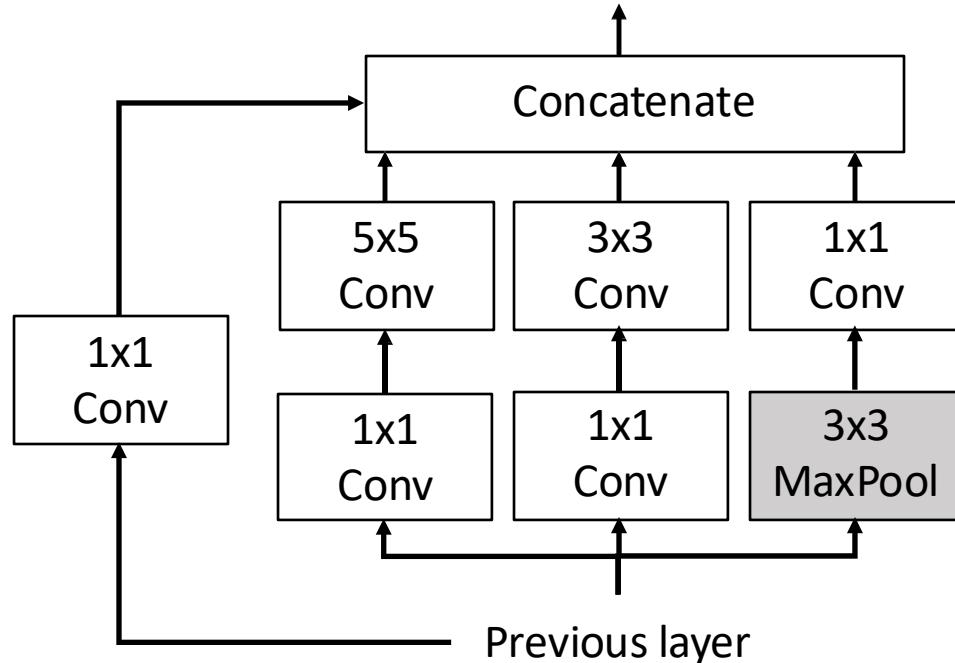
# Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.  
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each 2D  $K_h \times K_w$  conv/pool layer with a 3D  $K_t \times K_h \times K_w$  version

Inception Block: Original



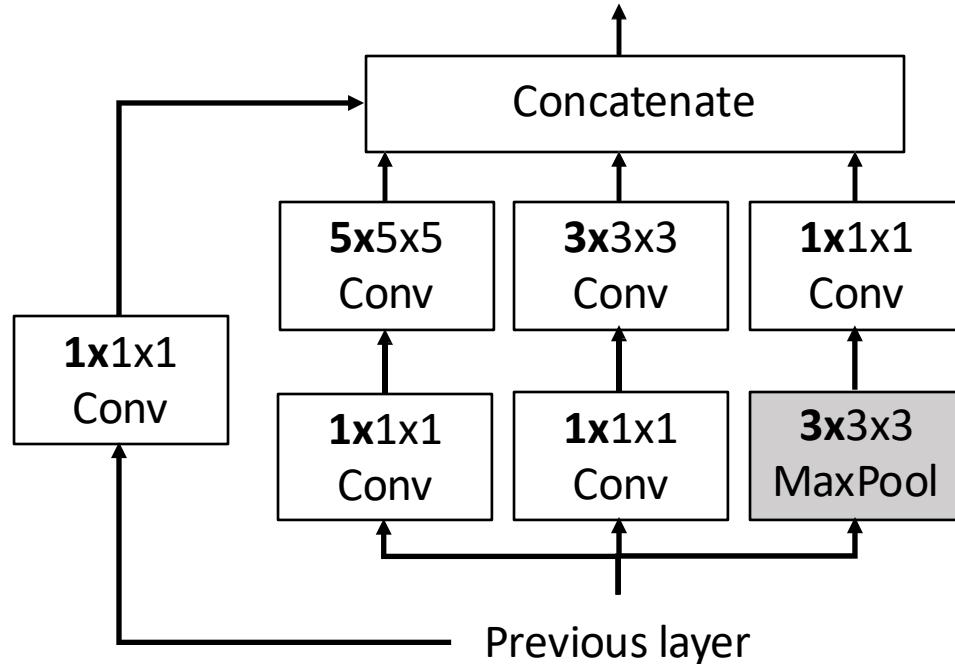
# Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.  
Can we reuse image architectures for video?

Idea: take a 2D CNN architecture.

Replace each 2D  $K_h \times K_w$  conv/pool layer with a 3D  $K_t \times K_h \times K_w$  version

Inception Block: Inflated



# Inflating 2D Networks to 3D (I3D)

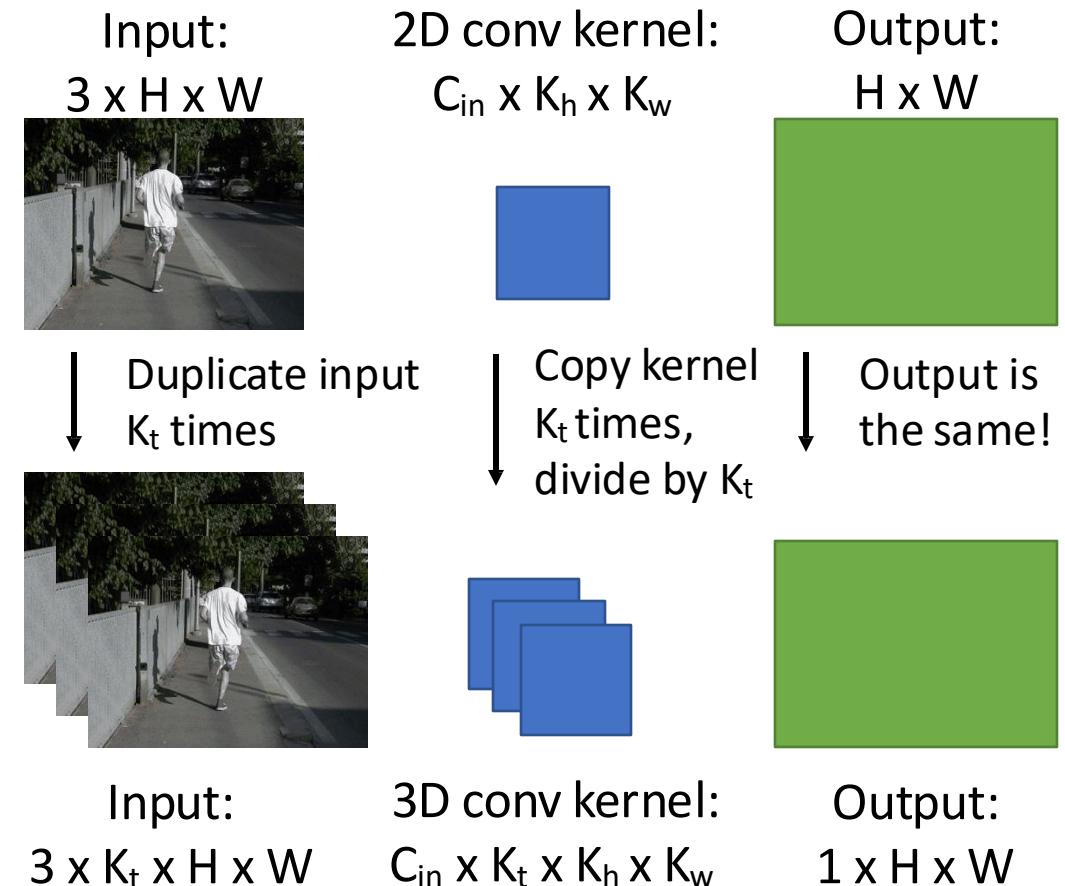
There has been a lot of work on architectures for images.  
Can we reuse image architectures for video?

**Idea:** take a 2D CNN architecture.

Replace each 2D  $K_h \times K_w$  conv/pool layer with a 3D  $K_t \times K_h \times K_w$  version

Can use weights of 2D conv to initialize 3D conv: copy  $K_t$  times in space and divide by  $K_t$

This gives the same result as 2D conv given “constant” video input



# Inflating 2D Networks to 3D (I3D)

There has been a lot of work on architectures for images.  
Can we reuse image architectures for video?

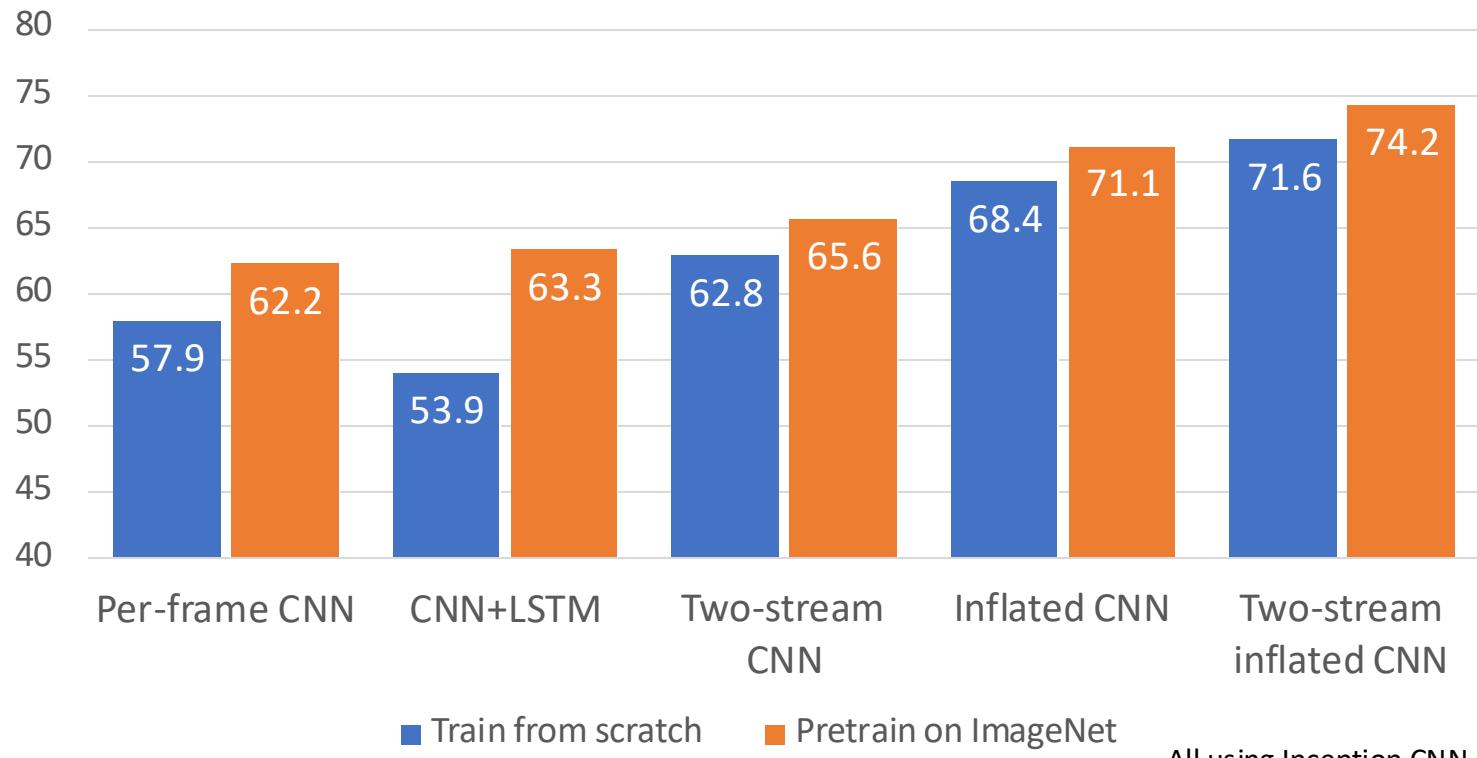
Idea: take a 2D CNN architecture.

Replace each 2D  $K_h \times K_w$  conv/pool layer with a 3D  $K_t \times K_h \times K_w$  version

Can use weights of 2D conv to initialize 3D conv: copy  $K_t$  times in space and divide by  $K_t$

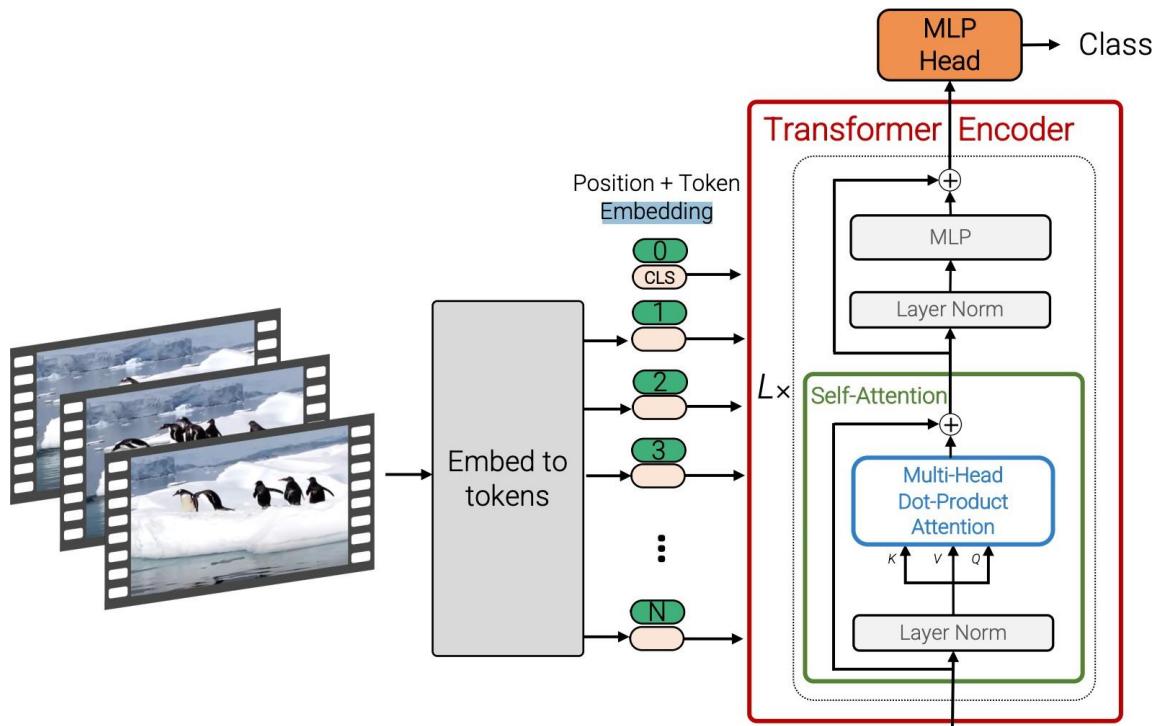
This gives the same result as 2D conv given “constant” video input

Top-1 Accuracy on Kinetics-400

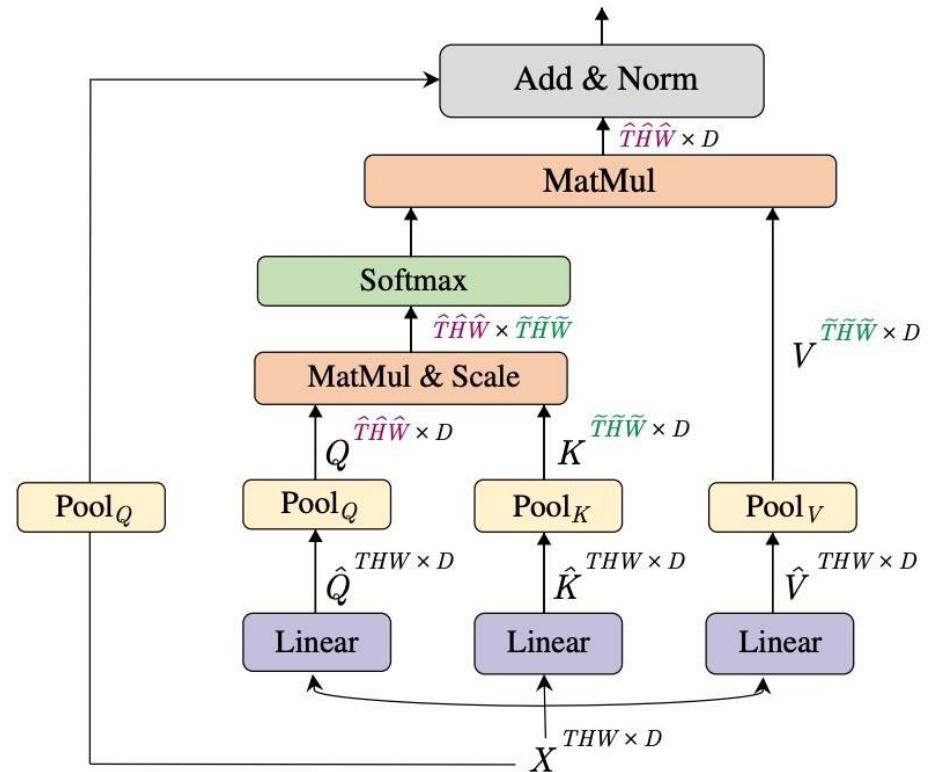


# Vision Transformers for Video

Factorized attention: Attend over space / time



Pooling module: Reduce number of tokens

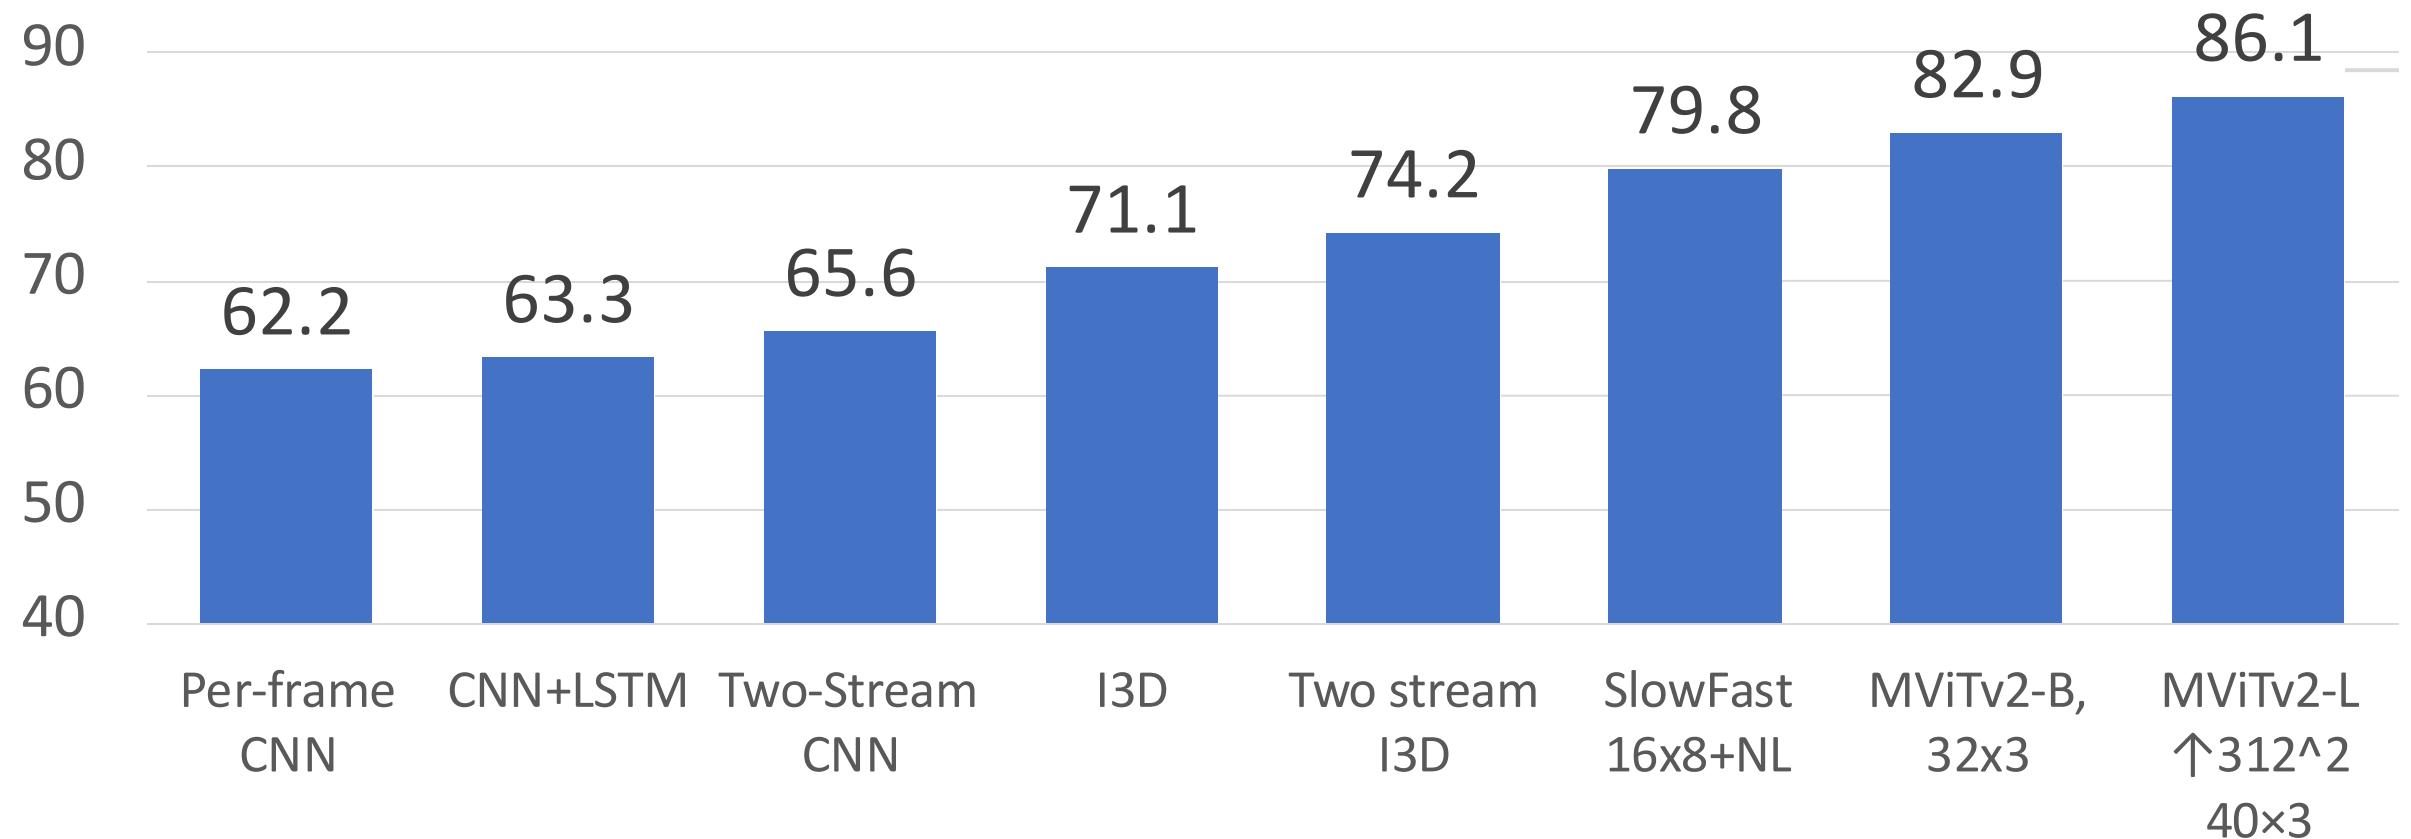


Bertasius et al, “Is Space-Time Attention All You Need for Video Understanding?”, ICML 2021  
Arnab et al, “ViViT: A Video Vision Transformer”, ICCV 2021

[Fan et al, “Multiscale Vision Transformers”, ICCV 2021](#)  
Li et al, “MViTv2: Improved Multiscale Vision Transformers”

# Vision Transformers for Video

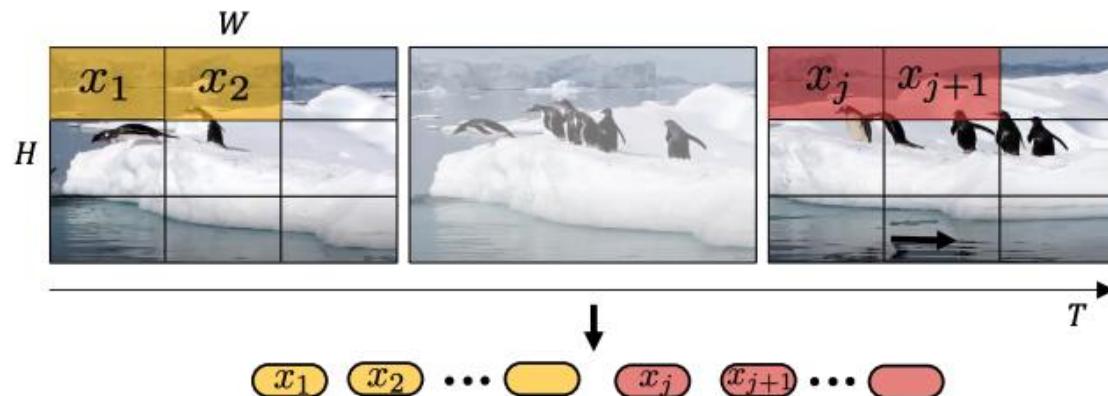
Top-1 Accuracy on Kinetics-400



Li et al, "MViTv2: Improved Multiscale Vision Transformers for Classification and Detection", CVPR 2022

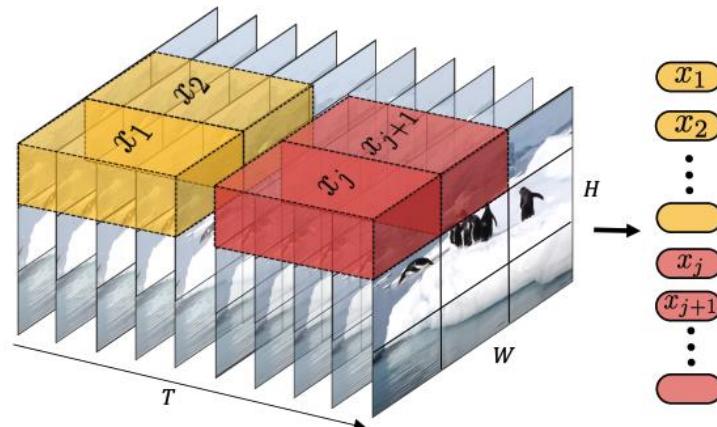
# ViViT: A Video Vision Transformer

- Tokenization: Uniform frame sampling
- uniformly sample  $nt$  frames from the input video clip, embed each 2D frame independently using the same method as ViT
- Concretely, if  $nh \times nw$  non-overlapping image patches are extracted from each frame, then a total of  $nt \times nh \times nw$  tokens will be forwarded through the transformer encoder.



# Tubelet embedding

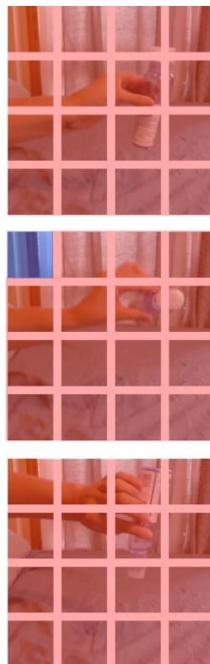
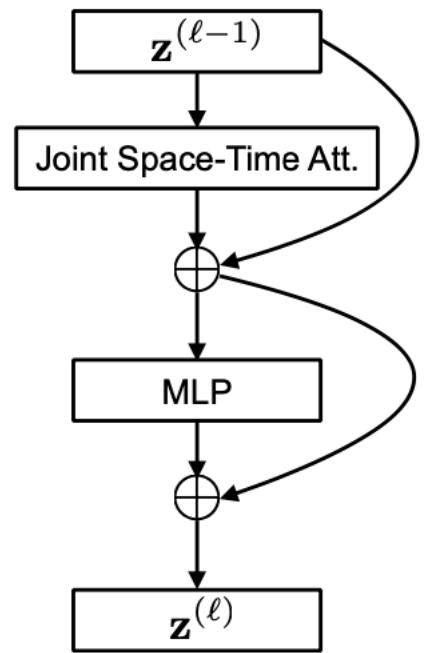
- Extract non-overlapping, spatio-temporal “tubes” from the input volume, and linearly project to  $d$  dimensions.
- For a tubelet of dimension  $t \times h \times w$ ,  $nt = \text{floor}(T/t)$ ,  $nh = \text{floor}(H/h)$
- Smaller tubelet dimensions thus result in more tokens which increases the computation.
- Intuitively, this method fuses spatio-temporal information during tokenisation



# Model 1: Spatio-temporal attention

- ST

NF + 1 keys

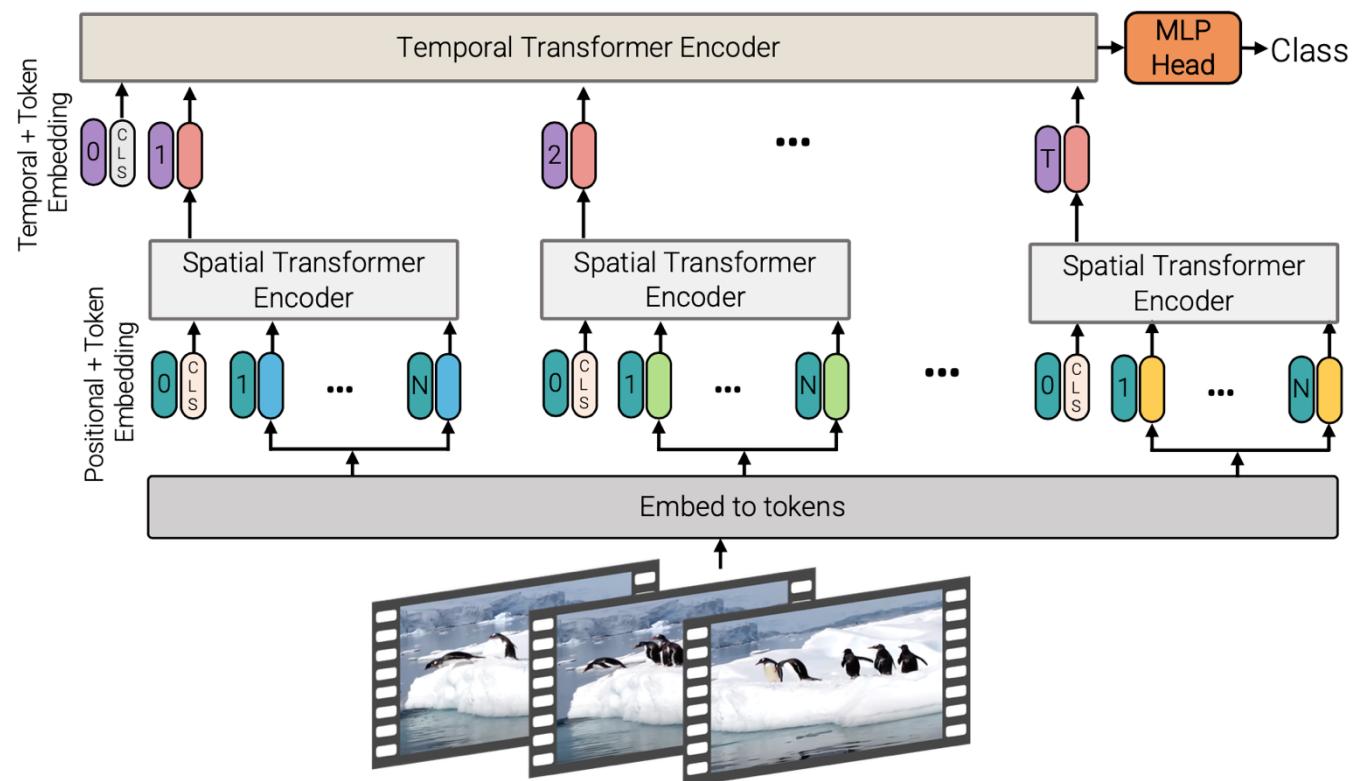


$$\alpha_{(p,t)}^{(\ell,a)} = \text{SM} \left( \frac{\mathbf{q}_{(p,t)}^{(\ell,a)}}{\sqrt{D_h}}^\top \cdot \left[ \mathbf{k}_{(0,0)}^{(\ell,a)} \left\{ \mathbf{k}_{(p',t')}^{(\ell,a)} \right\}_{\substack{p'=1,\dots,N \\ t'=1,\dots,F}} \right] \right)$$

$k_{0,0}$  is the cls token,  $\alpha \in R^{NF+1}$

## Model 2: Factorised encoder

- Two separate transformer encoders
- The first, spatial encoder, only models interactions between tokens extracted from the same temporal index (same frame)



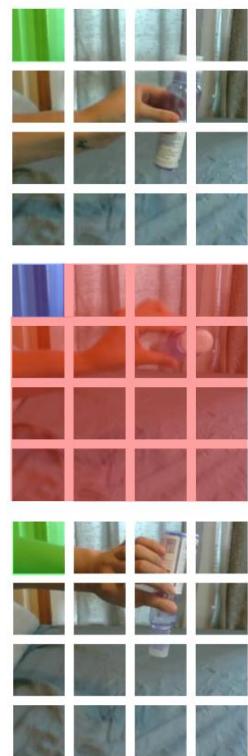
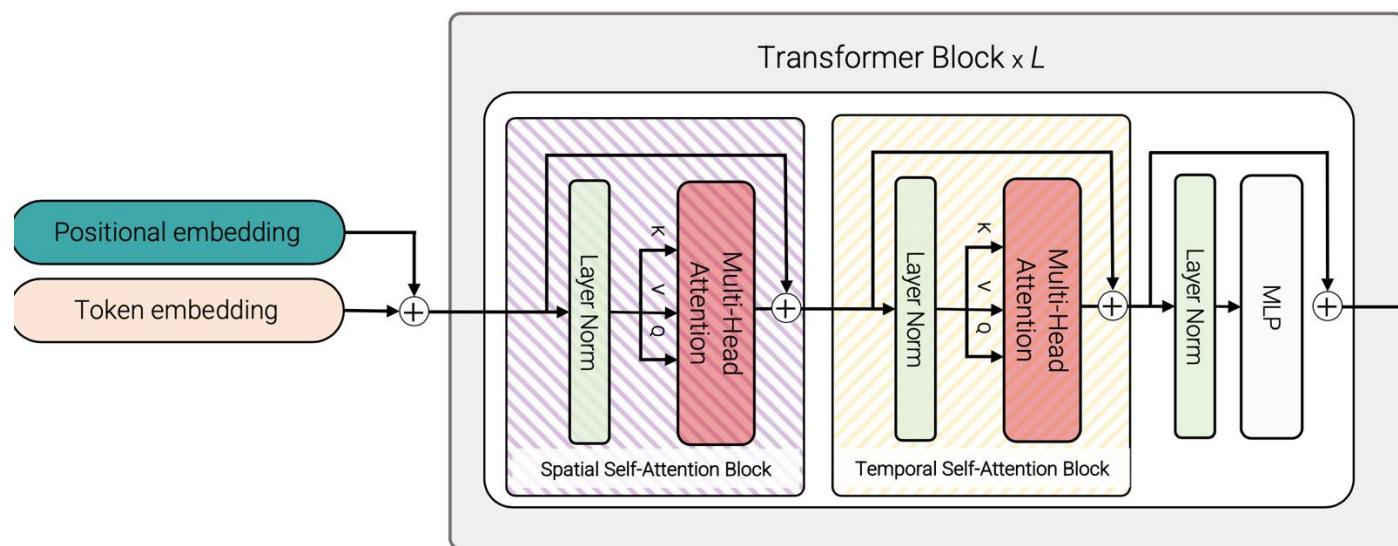
# Contd.

- The frame-level representations,  $h_i$ , are concatenated into  $nt \times d$ , and then forwarded through a temporal encoder consisting of  $L_t$  transformer layers to model interactions between tokens from different temporal indices. The output token of this encoder is then finally classified.
- Although this model has more transformer layers than Model 1 (and thus more parameters), it requires fewer floating point operations (FLOPs), as the two separate transformer blocks have a complexity of  $O((nh \cdot nw)^2 + nt^2)$  compared to  $O((nt \cdot nh \cdot nw)^2)$  of Model 1.

## Model 3: Factorised self-attention

Nt is batch for spatial  
Nh.nw is batch for temporal

- This operation can be performed efficiently by reshaping the tokens  $z$  from  $nt \cdot nh \cdot nw \times d$  to  $nt \times nh \cdot nw \cdot d$  to compute spatial self-attention. Similarly, the input to temporal self-attention, is reshaped to  $nh \cdot nw \times nt \cdot d$ .
- Computing self-attention separately on spatial and temporal dimensions reduce the quadratic complexity typically associated with attention mechanisms from  $O((nt \cdot nh \cdot nw)^2)$  to  $O(nt \cdot (nh \cdot nw)^2) + O(nh \cdot nw \cdot nt^2)$



Divided Space-Time  
Attention (T+S)

# Training

- ViT has been shown to only be effective when trained on large-scale datasets, as transformers lack some of the inductive biases of convolutional networks.
- However, even the largest video datasets such as Kinetics, have several orders of magnitude less labelled examples when compared to their image counterparts.
- As a result, training large models from scratch to high accuracy is extremely challenging.
- To sidestep this issue, and enable more efficient training we initialise our video models from pretrained image models.

# Positional Embeddings

- A positional embedding  $p$  is added to each input token. However, our video models have  $nt$  times more tokens than the pretrained image model.
- As a result, we initialise the positional embeddings by “repeating” them temporally from  $nw \cdot nh \times d$  to  $nt \cdot nh \cdot nw \times d$ .
- Therefore, at initialisation, all tokens with the same spatial index have the same embedding which is then fine-tuned.

# Embedding weights

- When using the “tubelet embedding” tokenisation method, the embedding filter  $E$  is a 3D tensor, compared to the 2D tensor in the pre-trained model,  $E_{\text{image}}$ .
- A common approach for initialising 3D convolutional filters from 2D filters for video classification is to “inflate” them by replicating the filters along the temporal dimension and averaging them
- $E = [E_{\text{image}}, \dots, E_{\text{image}}, \dots, E_{\text{image}}]/t$ .
- We consider an additional strategy, which we denote as “central frame initialisation”, where  $E$  is initialised with zeroes along all temporal positions, except at the centre  $[t/2]$ ,
- $E = [0, \dots, E_{\text{image}}, \dots, 0]$ .

# Transformer weights for Model 3

- The transformer block in Model 3 (Fig. 5) differs from the pretrained ViT model, in that it contains two multi-headed self attention (MSA) modules.
- In this case, we initialise the spatial MSA module from the pretrained module, and initialise all weights of the temporal MSA with zeroes.

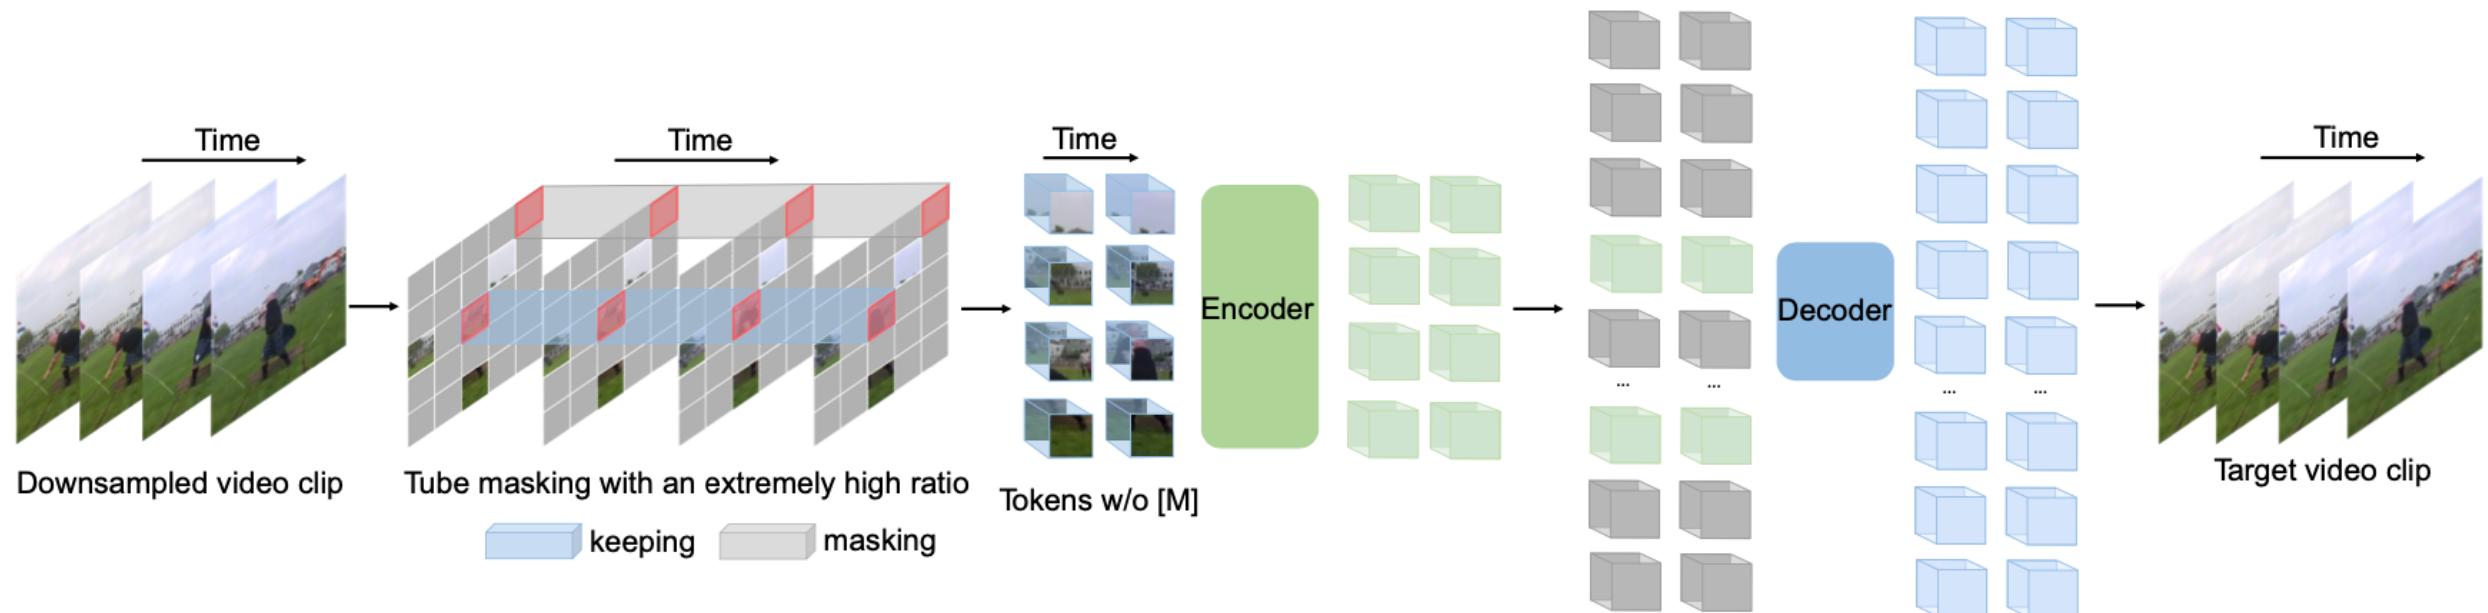
- ViT-Base (ViT-B, L=12, NH =12, d=768),
- L is the number of trans- former layers, each with a self-attention block of NH heads and hidden dimension d
- ViViT-B/16x2 denotes a ViT- Base backbone with a tubelet size of h×w×t = 16×16×2

Table 1: Comparison of input encoding methods using ViViT-B and spatio-temporal attention on Kinetics. Further details in text.

Top-1 accuracy	
Uniform frame sampling	78.5
<i>Tubelet embedding</i>	
Random initialisation [25]	73.2
Filter inflation [8]	77.6
Central frame	79.2

# Video MAE – NeurIPS22

- VideoMAE takes the *downsampled frames* as inputs and uses the *cube embedding* to obtain video tokens. Then, we propose a simple design of *tube masking with high ratio* to perform MAE pre-training with an asymmetric encoder- decoder architecture. Our backbone uses the vanilla ViT with *joint space-time attention*.



# Contd.

- one video clip consisting of  $t$  consecutive frames is first randomly sampled from the original video  $V$ .
- We then use temporal sampling to compress the clip to  $T$  frames, each of which contains  $H \times W \times 3$  pixels.
- In experiments, the stride  $\tau$  is set to 4 and 2 on Kinetics and Something-Something, respectively.

# Cube embedding

- We adopt the joint space-time cube embedding in our VideoMAE, where we treat each cube of size  $2 \times 16 \times 16$  as one token embedding.
- Thus, the cube embedding layer obtains  $T/2 \times H/16 \times W/16$  3D tokens and maps each token to the channel dimension D.
- This design  $2 \times 16 \times 16$  can decrease the spatial and temporal dimension of input, which helps to alleviate the spatiotemporal redundancy in videos.
- Mask ratio 90-95%.

Method	Backbone	Extra data	Ex. labels	Frames	GFLOPs	Param	Top-1	Top-5
NL I3D [78]	ResNet101	ImageNet-1K	✓	128	359×10×3	62	77.3	93.3
TANet [41]	ResNet152		✓	16	242×4×3	59	79.3	94.1
TDN <sub>En</sub> [75]	ResNet101		✓	8+16	198×10×3	88	79.4	94.4
TimeSformer [6]	ViT-L	ImageNet-21K	✓	96	8353×1×3	430	80.7	94.7
ViViT FE [3]	ViT-L		✓	128	3980×1×3	N/A	81.7	93.8
Motionformer [51]	ViT-L		✓	32	1185×10×3	382	80.2	94.8
Video Swin [39]	Swin-L		✓	32	604×4×3	197	83.1	95.9
ViViT FE [3]	ViT-L	JFT-300M	✓	128	3980×1×3	N/A	83.5	94.3
ViViT [3]	ViT-H	JFT-300M	✓	32	3981×4×3	N/A	84.9	95.8
VIMPAC [65]	ViT-L	HowTo100M+DALLE	✗	10	N/A×10×3	307	77.4	N/A
BEVT [77]	Swin-B	IN-1K+DALLE	✗	32	282×4×3	88	80.6	N/A
MaskFeat↑352 [80]	MViT-L	Kinetics-600	✗	40	3790×4×3	218	87.0	97.4
ip-CSN [69]	ResNet152	<i>no external data</i>	✗	32	109×10×3	33	77.8	92.8
SlowFast [23]	R101+NL		✗	16+64	234×10×3	60	79.8	93.9
MViTv1 [22]	MViTv1-B		✗	32	170×5×1	37	80.2	94.4
MaskFeat [80]	MViT-L		✗	16	377×10×1	218	84.3	96.3
<b>VideoMAE</b>	ViT-S	<i>no external data</i>	✗	16	57×5×3	22	79.0	93.8
<b>VideoMAE</b>	ViT-B		✗	16	180×5×3	87	81.5	95.1
<b>VideoMAE</b>	ViT-L		✗	16	597×5×3	305	85.2	96.8
<b>VideoMAE</b>	ViT-H		✗	16	1192×5×3	633	<b>86.6</b>	<b>97.1</b>
<b>VideoMAE</b> ↑320	ViT-L	<i>no external data</i>	✗	32	3958×4×3	305	86.1	97.3
<b>VideoMAE</b> ↑320	ViT-H		✗	32	7397×4×3	633	<b>87.4</b>	<b>97.6</b>

Comparison with the state-of-the-art methods on Kinetics-400

# Further reading

- An Empirical Study of End-to-End Video-Language Transformers with Masked Visual Modeling, CVPR2023
- Masked Video Distillation: Rethinking Masked Feature Modeling for Self-supervised Video Representation Learning, CVPR2023
- VideoMAE V2: Scaling Video Masked Autoencoders with Dual Masking, CVPR2023
-